

CS229T/STAT231: Statistical Learning Theory (Winter 2015)

Percy Liang

Last updated Sat Mar 14 2015 13:19

These lecture notes will be updated periodically as the course goes on. Please let us know if you find typos. Section [A.1](#) describes the basic notation and definitions; Section [A.2](#) describes some basic theorems.

Contents

1	Overview	2
1.1	What is this course about? (Lecture 1)	2
1.2	Online learning (Lecture 1)	4
1.3	Uniform convergence (Lecture 1)	5
1.4	Direct analysis (Lecture 1)	6
1.5	Kernel methods (Lecture 1)	7
1.6	Latent-variable models (Lecture 1)	8
2	Online learning	10
2.1	Introduction (Lecture 1)	10
2.2	Warm-up (Lecture 1)	12
2.3	Online convex optimization (Lecture 2)	14
2.4	Follow the leader (FTL) (Lecture 2)	19
2.5	Follow the regularized leader (FTRL) (Lecture 3)	23
2.6	Online subgradient descent (OGD) (Lecture 3)	26
2.7	Online mirror descent (OMD) (Lecture 4)	30
2.8	Regret bounds with Bregman divergences (Lecture 4)	32
2.9	Strong convexity and smoothness (Lecture 4)	34
2.10	Local norms (Lecture 5)	39
2.11	Multi-armed bandits (Lecture 5)	42
2.12	Online-to-batch conversion (Lecture 5)	44
2.13	Summary (Lecture 5)	47
2.14	References	48
3	Uniform convergence	49
3.1	Motivation (Lecture 6)	49
3.2	Formal setup (Lecture 6)	50
3.3	Realizable finite hypothesis classes (Lecture 6)	53
3.4	Generalization bounds via uniform convergence (Lecture 6)	55
3.5	Concentration inequalities (Lecture 7)	58

3.6	Finite hypothesis classes (Lecture 8)	67
3.7	Rademacher complexity (Lecture 8)	68
3.8	Finite-dimensional hypothesis classes (Lecture 9)	74
3.9	Shattering coefficient (Lecture 9)	75
3.10	VC dimension (Lecture 9)	77
3.11	Norm-constrained hypothesis classes (Lecture 10)	83
3.12	Covering numbers (metric entropy) (Lecture 10)	87
3.13	Algorithmic stability (Lecture 11)	95
3.14	PAC-Bayesian bounds (Lecture 11)	99
3.15	Interpretation of bounds (Lecture 11)	102
3.16	Summary (Lecture 11)	103
3.17	References	105
4	Direct analysis	106
4.1	Introduction (Lecture 12)	106
4.2	Fixed design linear regression (Lecture 12)	106
4.3	Finite-dimensional asymptotics (Lecture 13)	112
4.4	References	122
5	Kernel methods	123
5.1	Motivation (Lecture 14)	123
5.2	Kernels: definition and examples (Lecture 14)	125
5.3	Three views of kernel methods (Lecture 14)	128
5.4	Reproducing kernel Hilbert spaces (RKHS) (Lecture 14)	130
5.5	Learning using kernels (Lecture 15)	135
5.6	Fourier properties of shift-invariant kernels (Lecture 15)	139
5.7	Efficient computation (Lecture 16)	144
5.8	Universality (skipped in class)	150
5.9	RKHS embedding of probability distributions (skipped in class)	151
5.10	Summary (Lecture 16)	154
5.11	References	154
6	Neural networks	155
6.1	Motivation (Lecture 17)	155
6.2	Setup (Lecture 17)	156
6.3	Approximation error (universality) (Lecture 17)	157
6.4	Generalization bounds (Lecture 17)	159
6.5	Approximation error for polynomials (Lecture 17)	160
6.6	References	164
7	Method of moments	164
7.1	Motivation (Lecture 17)	164
7.2	Method of moments framework (Lecture 17)	166

7.3	Method of moments for latent-variable models (Lecture 17)	169
7.4	Summary (Lecture 17)	180
7.5	References	180
8	Conclusions and outlook	181
8.1	Review (Lecture 18)	181
8.2	Changes at test time (Lecture 18)	183
8.3	Alternative forms of supervision (Lecture 18)	185
8.4	Interaction between computation and statistics (Lecture 18)	186
A	Appendix	188
A.1	Notation	188
A.2	Basic theorems	190

1 Overview

1.1 What is this course about? (Lecture 1)

- Machine learning has become an indispensable part of many application areas, ranging from science (biology, neuroscience, psychology, astronomy, etc.) to engineering (natural language processing, computer vision, robotics, etc.). But machine learning is not one coherent approach, but rather consists of a dazzling array of seemingly disparate frameworks and paradigms spanning classification, regression, clustering, matrix factorization, graphical models, probabilistic programming, etc.
- This course tries to uncover the common **statistical principles** underlying this diverse array of techniques, using theoretical analysis as the main tool for understanding. We will try to answer two questions:
 - When do learning algorithms work (or fail)? Anyone with hands-on machine learning experience will probably know that learning algorithms rarely work the first time. When they don't, a deeper theoretical understanding can offer a new perspective and can aid in troubleshooting.
 - How do we make learning algorithms better? Theoretical analyses often point to weaknesses in current algorithms, and suggest new algorithms to address the limitations of existing ones. For example, is it possible to combine weak classifiers into a strong classifier? This theoretical challenge eventually led to the development of AdaBoost in the mid 1990s, a beautifully simple and practical algorithm with strong theoretical guarantees.
 - In a more recent example, Google's latest **22-layer convolutional neural network** that won the 2014 ImageNet Visual Recognition Challenge was in part inspired by **a theoretically-motivated algorithm for learning deep neural networks with sparsity structure**. Although there is very often still a gap between theory and practice, the transfer of ideas can be nonetheless fruitful.
- Example: text classification
 - Suppose we want to build a classifier to predict the topic of a document (e.g., sports, politics, technology, etc.).
 - We train an SVM on bag-of-words features and obtain 8% training error on 1000 training documents, test error is 13% on 1000 documents. Now what? Here are several questions which could be useful in figuring out what's going on.
 - * How reliable are these numbers? If we reshuffled the data, would we get the same answer?

- * How much should we expect the test error to change if we double the number of examples?
 - * What if we double the number of features? What if our features or parameters are sparse?
 - * What if we double the regularization?
 - * Should we change the learning algorithm altogether?
- In this class, we develop tools to tackle some of these questions. Our goal isn't to give precise quantitative answers (just like analyses of algorithms doesn't tell you how exactly many hours a particular algorithm will run). Rather, the analyses will reveal the relevant quantities (e.g., dimension, regularization strength, number of training examples), and reveal how they influence the final test error.
 - Also, we will use prediction as the primary motivation, but much of the techniques we will cover in this class can be applied to estimation, hypothesis testing, etc.
- Error decomposition
 - How do we even begin to analyze or even formalize the machine learning enterprise, which seems to have so many moving parts?
 - It will be conceptually useful to break the error of the classifier (more generally, predictor¹) returned by the learning algorithm into two parts: **approximation error** and **estimation error**: A learning algorithm can be viewed as operating




Figure 1: Cartoon showing error decomposition into approximation and estimation errors.

- on a **hypothesis class** \mathcal{H} , which governs what kind of predictions it can make.
- Approximation error is error made by the best predictor in \mathcal{H} . If the approximation error is large, then the hypothesis class is probably too small.
- Estimation error is the difference between the error of the learner and the error of the best predictor. If the estimation error is large, then the hypothesis class is probably too large for the given amount of data.
- We should strive to balance these two errors so that the sum is minimized. Of course, this is not a one-dimensional problem, for there is a lot of wiggle room in choosing \mathcal{H} based on knowledge about the learning task at hand.

¹Note that in statistics, *predictor* is synonymous with *feature*. In this class, we use *predictor* in the machine learning sense.

- In the next sections, we give a brief overview of the various topics in this course: online learning, uniform convergence, direct analysis, kernel methods, and latent-variable models. Very roughly speaking, the first half of the class will focus mostly on estimation error, and the second half focuses on approximation error. The class will also tend to proceed from learning settings with very few assumptions (data can be generated completely adversarially) to ones with more (data can be generated i.i.d. from a known parametric family). As we make more assumptions, we have more to work with and can prove stronger results.

1.2 Online learning (Lecture 1)

- Online learning is a nice place to start since it leads directly to many simple and efficient algorithms which are used heavily in practice, while at the same time offering a remarkably clean theory and sharp results.
- The goal of online learning is to solve a prediction task: given input $x \in \mathcal{X}$ (e.g., $\{0,1\}^d$), predict output $y \in \mathcal{Y}$ (e.g., $\{+1, -1\}$ in classification or \mathbb{R} in regression). Data arrives at the learner's doorstep in real-time and the learner must also make predictions on-the-fly.
- We will cast online learning as the following game between a learner and nature:




Figure 2: Online learning game.

- Iterate $t = 1, \dots, T$:
 - * Learner receives input $x_t \in \mathcal{X}$
 - * Learner outputs prediction $p_t \in \mathcal{Y}$
 - * Learner receives true label $y_t \in \mathcal{Y}$
 - * (Learner updates model parameters)
- How do we evaluate?
 - Loss function: $\ell(y_t, p_t) \in \mathbb{R}$

- Let \mathcal{H} be a set of *fixed* expert predictors.
- **Regret**: cumulative difference between learner’s loss and best expert’s loss (low regret is good).
- We will prove statements of the following form:

$$\text{Regret} \leq \text{SomeFunction}(\mathcal{H}, T). \quad (1)$$

For example, for finite \mathcal{H} , we will be able to prove:

$$\text{Regret} \leq O(\sqrt{T \log |\mathcal{H}|}). \quad (2)$$

Things to note:

- The average regret goes to zero (in the long run, we’re doing as well as the best expert).
- The bound only depends logarithmically on the number of experts, which means that the number of experts $|\mathcal{H}|$ can be exponentially larger than the number of data points T .
- We make zero assumptions about the examples—could be generated by an **adversary**!
- Regret gets at estimation error. We are only comparing with the fixed set of experts; it says nothing about how well the learner is doing in absolute terms.
- The algorithms that achieve these results, as we’ll see later, are dirt simple and as efficient as one could hope for.
- Main mathematical tool: **convexity**

1.3 Uniform convergence (Lecture 1)

- Next, we turn to the batch setting, where a learner gets a set of training examples, does some learning, and then is evaluated on generalization (test) error.
- We will study the **empirical risk minimizer** (known as an M-estimator in statistics):

$$\hat{h}_{\text{ERM}} \in \arg \min_{h \in \mathcal{H}} \hat{L}(h), \quad (3)$$

where $\hat{L}(h)$ is the empirical risk of h , the average loss over n i.i.d. training examples; in other words, the training error. Empirical risk minimization includes maximum likelihood for the special case where the loss is the negative log-likelihood. We won’t worry about how the ERM is computed; if the loss is convex, then this can be readily done via standard convex optimization.

- We assume all examples are drawn i.i.d. from some unknown distribution $p^*(x, y)$ but make no further assumptions about p^* .
- We will derive two types of **generalization bounds**. The first compares generalization and training error:

$$\underbrace{L(\hat{h}_{\text{ERM}})}_{\text{generalization error}} \leq \underbrace{\hat{L}(\hat{h}_{\text{ERM}})}_{\text{training error}} + O_p \left(\sqrt{\frac{\text{Complexity}(\mathcal{H})}{n}} \right). \quad (4)$$

The second relates the empirical risk minimizer to the best $h^* \in \mathcal{H}$ (this bounds the estimation error):

$$L(\hat{h}_{\text{ERM}}) \leq L(h^*) + O_p \left(\sqrt{\frac{\text{Complexity}(\mathcal{H})}{n}} \right). \quad (5)$$

- For a fixed $h \in \mathcal{H}$, the training error $\hat{L}(h)$ converges to generalization error $L(h)$ by Hoeffding's inequality or the central limit theorem.
- However, the key point is that \hat{h}_{ERM} is random, so plain convergence isn't enough. We need to get **uniform convergence** over all functions $h \in \mathcal{H}$, which ensures convergence for $L(\hat{h}_{\text{ERM}})$ as well.
- The rate of convergence is governed by the complexity of \mathcal{H} , so we will devote a good deal of time computing the complexity of various function classes \mathcal{H} . VC dimension, covering numbers, and Rademacher complexity are different ways of measuring how big a set of functions is. For example, for finite \mathcal{H} , the right measure of complexity is $\log |\mathcal{H}|$, which agrees with the online learning bound we saw earlier.
- Generalization bounds are in some sense the heart of statistical learning theory. But along the way, we will develop generally useful **concentration inequalities** whose applicability extends beyond machine learning (e.g., showing convergence of eigenvalues in random matrix theory).
- Main mathematical tool: **probability**

1.4 Direct analysis (Lecture 1)

- Uniform convergence only gives us upper bounds, so we can't directly compare the generalization error of two algorithms. Also, it is worst case over all distributions p^* , so it lacks sensitivity to the exact problem structure.
- In simplified settings, we can actually perform a direct analysis to get sharper results which are exact. Here, the goal will not be to obtain the most general results, but to obtain intuition about problem structure for certain cases.

- For example, for linear regression, we can see that the expected estimation error is exactly:

$$\mathbb{E}[L(\hat{h}_{\text{ERM}})] = L(h^*) + \frac{d\sigma^2}{n - d - 1}, \quad (6)$$

where d is the dimensionality, n is the number of examples, and σ^2 is the variance of the noise. Note that the expectation is taken over random training sets.

- For non-linear loss functions, we can't compute the error directly. But we can appeal to asymptotics, a common tool used in statistics. Assuming twice differentiability, we can perform a **Taylor expansion** of the expected risk to obtain:

$$L(\hat{h}_{\text{ERM}}) = L(h^*) + \frac{\text{SomeFunction}(p^*, h^*)}{n} + \dots \quad (7)$$

The asymptotics approach is fairly lightweight and gives equality up to the first-order, which allows us to compare different estimators, at least in the asymptotic sense.

- Main mathematical tool: **Taylor expansion**

1.5 Kernel methods (Lecture 1)

- Now we turn our attention from estimation error to approximation error. Real-world data is complex, so we need expressive models. Kernels provide a rigorous mathematical framework to build complex, non-linear models based on the machinery of linear models.
- For concreteness, suppose we're trying to solve a regression task: predicting $y \in \mathbb{R}$ from $x \in \mathcal{X}$.
- There are three ways to think about kernels:




Figure 3: The three key mathematical concepts in kernel methods.

- Feature view: define functions via features $\phi(x)$:

– $\phi(x) = (1, x)$ yields linear functions

- $\phi(x) = (1, x, x^2)$ yields quadratic functions
- This is the usual way of approaching machine learning.
- Kernel view: define **positive semidefinite kernel** $k(x, x')$, which captures “similarity” between x and x' .
 - Kernels allow us to construct complex non-linear functions (e.g., Gaussian kernel $k(x, x') = \exp(\frac{-\|x-x'\|^2}{2\sigma^2})$) that are **universal**, in that it can represent *any* continuous function (and with enough data, it will learn it).
 - For strings, trees, graphs, etc., can define kernels that exploit dynamic programming for efficient computation.
- Function view: directly study the class of functions, known as an **reproducing kernel Hilbert space** (RKHS). This provides a mathematically more elegant way at getting hold of the object we actually care about.
- One point is that kernels and RKHSes are in one-to-one correspondence, and are at the core of linear models. On the other hand, there could be many feature functions that yield the same kernel, so one can choose the one based on computational convenience.
- Kernel methods generally require $O(n^2)$ time (n is the number of training points) to even compute the kernel matrix. Approximations based on sampling offer efficient solutions. For example, by generating lots of random features of the form $\cos(\omega \cdot x + b)$, we can approximate any shift-invariant kernel.
- Main mathematical tool: **functional analysis**

1.6 Latent-variable models (Lecture 1)

- So far, we have focused mostly on prediction with fixed data representations (fixed features, kernels). This only works well when the representations are good. The natural question to ask is whether these representations can be learned automatically. We will consider several types of methods that learn representations in an unsupervised way by factorizing the data.
- One main challenge with such endeavors is that it is easy to run into non-convex optimization problems. We will show that convex relaxations, while optimizing a different objective, can produce the same solutions statistically. In practice, alternating minimization is used, which can get stuck in local minima. However, with the right statistical assumptions, we can show that these approaches can provably work too for problems such as matrix completion.
- Next, we consider parameter estimation in probabilistic latent-variable models such as mixtures of Gaussians, hidden Markov models, and latent Dirichlet allocation. Again, these models are difficult to estimate, and traditional methods such as EM get stuck

in local optima. We can sidestep this issue via the **method of moments**, which leads to simple algorithms based on tensor decomposition which (i) do not have problems with local optima and (ii) reveal additional insight into the model.

- Main mathematical tools: **linear algebra**

2 Online learning

We will first discuss the online learning framework, focusing on prediction. Then, we will cast online learning as online convex optimization and develop several algorithms and prove regret bounds for these algorithms. Finally, we will look at multi-armed bandit problems, where the learner obtains partial feedback. Throughout this section, there will be very little probability (since we will be working mostly in the adversarial setting), but we will draw quite a bit from convex analysis.

2.1 Introduction (Lecture 1)

- Framework
 - Prediction task: we want to map inputs $x \in \mathcal{X}$ to outputs $y \in \mathcal{Y}$.
 - The online learning setting can be thought of as the following game between a learner and nature:




Figure 4: Online learning game.

- * Iterate $t = 1, \dots, T$:
 - Learner receives input $x_t \in \mathcal{X}$
 - Learner outputs prediction $p_t \in \mathcal{Y}$
 - Learner receives true label $y_t \in \mathcal{Y}$
 - (Learner updates model parameters)
- **Example 1 (online binary classification for spam filtering)**
 - Inputs: $\mathcal{X} = \{0, 1\}^d$ are boolean feature vectors (presence or absence of a word).
 - Outputs: $\mathcal{Y} = \{+1, -1\}$: whether a document is spam or not spam.
 - Zero-one loss: $\ell(y_t, p_t) = \mathbb{I}[y_t \neq p_t]$ is whether the prediction was incorrect.
- Remarks

- The typical training phase (setting model parameters) and testing phase (making predictions for evaluation) are **interleaved** in online learning.
 - Note that the the online learning setting leaves completely open the time and memory usage of the algorithms that operate in this setting. Technically, we could just train an SVM on all the data that we’ve seen so far, and predict on the next example. However, the spirit of online learning suggests that the amount of work an algorithm does per iteration should not grow with t . In practice, online learning algorithms update parameters after each example, and hence tend to be **faster** than traditional batch optimization algorithms such as Newton’s method.
 - The real world is complex and constantly-changing, but online learning algorithms have the potential to **adapt** (although we will not analyze their adaptive capabilities in this course).
 - In some applications such as spam filtering, the inputs could be generated by an **adversary**. In our analyses, we will make no assumptions about the input/output sequence.
- Evaluation: how good is a learner?
 - Let \mathcal{H} be a set of experts, where each **expert** is a function $h : \mathcal{X} \rightarrow \mathcal{Y}$ that predicts $h(x)$ on input $x \in \mathcal{X}$. Note that the learner can adapt whereas the experts are fixed.
 - **Definition 1 (regret)**
 - * The regret of a learner with respect to an expert h is the cumulative difference between the loss incurred by the learner and the loss incurred by expert h :

$$\text{Regret}(h) = \sum_{t=1}^T [\ell(y_t, p_t) - \ell(y_t, h(x_t))]. \quad (8)$$
 - * The regret with respect to a class of experts \mathcal{H} is the maximum regret (attained by best expert):

$$\text{Regret} = \max_{h \in \mathcal{H}} \text{Regret}(h). \quad (9)$$
 - Note that Regret depends on \mathcal{H} , T , the sequences $x_{1:T}, y_{1:T}$, and of course the algorithm itself; but we’re eliding the dependence in the notation.
 - Intuition: having fixed experts makes it possible to compete in an adversarial setting. If the adversary tries to screw over the learner, it will probably screw over the experts too.
 - Objective: minimize regret. We want to obtain a result of the form: for all $\mathcal{H}, T, x_{1:T}, y_{1:T}$, we have:

$$\text{Regret} \leq \text{SomeFunction}(\mathcal{H}, T). \quad (10)$$

- Low regret is good. Usually, we want the regret to be sublinear in T , which means that the average regret goes to zero.
- Aside: note that regret can be negative in online learning (since the learner can adapt but the experts cannot).

2.2 Warm-up (Lecture 1)

We will give two examples, one where online learning is impossible and one where it is possible.

• Example 2 (negative result)

- Assume binary classification: $y \in \{-1, +1\}$
- Assume zero-one loss: $\ell(y_t, p_t) = \mathbb{I}[y_t \neq p_t]$
- Assume the learner is fully **deterministic**.
- Claim: there exists an \mathcal{H} and input/output sequence $x_{1:T}, y_{1:T}$ such that:

$$\boxed{\text{Regret} \geq T/2} \quad [\text{awful!}] \quad (11)$$

- **Key point:** adversary (having full knowledge of learner) can choose y_t to be always different from p_t .
- Learner's cumulative loss: T (make mistake on every example).
- We're not done yet, because remember regret is the difference between learner's cumulative loss and the best expert's, so we have to check how well the experts do in this case.
- Consider two experts, $\mathcal{H} = \{h_{-1}, h_{+1}\}$, where h_y always predicts y .
- The sum of the cumulative losses of the experts equals T because $\ell(y_t, h_{-1}(x_t)) + \ell(y_t, h_{+1}(x_t)) = 1$, so one of the experts must achieve loss $\leq T/2$.
- Therefore, the difference between T and $\leq T/2$ is $\geq T/2$.
- It is perhaps not surprising that no learner can do well because nature is too powerful here (it can choose any sequence with full knowledge of what the learner will do). So we will need assumptions to get positive results.

• Example 3 (positive result (learning with expert advice))

- Assume zero-one loss: $\ell(y_t, p_t) = \mathbb{I}[y_t \neq p_t]$.
- Clearly, we need to make *some* assumptions to make progress. Here, we will make a fairly strong assumption just to get some intuition.

– **Assumption 1 (realizable)**

Assume the best expert $h^* \in \mathcal{H}$ obtains zero cumulative loss ($\ell(y_t, h^*(x_t)) = 0$ for all $t = 1, \dots, T$); or equivalently, since we're dealing with zero-one loss, $y_t = h^*(x_t)$ for all t . This assumption places a joint restriction on $\mathcal{H}, x_{1:T}, y_{1:T}$.

– We will design an algorithm that queries the experts on each example x_t , and try to combine their predictions in some way; this setting is called **learning with expert advice**.

– **Algorithm 1 (majority algorithm)**

- * Maintain a set $V_t \subseteq \mathcal{H}$ of valid experts (those compatible with the first $t - 1$ examples).
- * On each iteration t , predict p_t to be the majority vote over predictions of valid experts $\{h(x_t) : h \in V_t\}$.
- * Keep experts which were correct: $V_{t+1} = \{h \in V_t : y_t = h(x_t)\}$.

– Analysis:

- * On each mistake, at least half of the experts are eliminated.
- * So $1 \leq |V_{T+1}| \leq |\mathcal{H}|2^{-M}$, where M is the number of mistakes (equal to Regret since the best expert has zero loss).
- * Note that the lower bound is due to the realizability assumption.
- * M is the exactly the regret here.
- * Take logs and rearrange:

$$\boxed{\text{Regret} \leq \log_2 |\mathcal{H}|.} \tag{12}$$

– Notes:

- * This is a really strong result: note that the regret is constant; after a finite number of iterations, we cease to make any more mistakes forever.
- * However, realizability (that some expert is perfect) is too strong of an assumption.
- * Also, the regret bound is useless here if there are an infinite number of experts ($|\mathcal{H}| = \infty$).

2.3 Online convex optimization (Lecture 2)

- In order to obtain more general results, we move to a framework called online convex optimization. We will see that **convexity** will give us considerable leverage. Later, we'll connect online convex optimization with online learning.
- Let $S \subseteq \mathbb{R}^d$ be a convex set (e.g., representing set of allowed weight vectors).

– Example: $S = \{u : \|u\|_2 \leq B\}$.

- **Definition 2 (convexity)**

A function $f : S \rightarrow \mathbb{R}$ is convex iff for all points $w \in S$, there is some vector $z \in \mathbb{R}^d$ such that

$$\boxed{f(u) \geq f(w) + z \cdot (u - w) \quad \text{for all } u \in S.} \quad (13)$$

This says that at any point $w \in S$, we can find a linear approximation (RHS of (13)) that lower bounds the function f (LHS of (13)).

- **Definition 3 (subgradient)**

For each $w \in S$, the set of all z satisfying (13) are known as the subgradients at w :

$$\boxed{\partial f(w) \stackrel{\text{def}}{=} \{z : f(u) \geq f(w) + z \cdot (u - w) \text{ for all } u \in S\}.} \quad (14)$$

If f is differentiable at w , then there is one subgradient equal to the gradient: $\partial f(w) = \{\nabla f(w)\}$.

- Convexity is remarkable because it allows you to say something (in particular, lower bound) the global behavior (for all $u \in S$) of a function f by something local and simple (a linear function). An important consequence of f being convex is that if $0 \in \partial f(w)$, then w is a global minimum of f .
- Checking convexity
 - How do you know if a function is convex? You can try to work it out from the definition or if the function is twice-differentiable, compute the Hessian and show that it's positive semidefinite.
 - But often, we can check convexity by decomposing the function using a few rules. This is by no means an exhaustive list, but these are essentially all the important cases that we'll need in this class.
 - * Linear functions: $f(w) = w \cdot z$ for any $z \in \mathbb{R}^d$

- * Quadratic functions: $f(w) = w^\top A w$ for positive semidefinite $A \in \mathbb{R}^{d \times d}$
- * Negative entropy: $f(w) = \sum_{i=1}^d w_i \log w_i$ on $w \in \Delta_d$.
- * Sum: $f + g$ if f and g are convex
- * Scale: cf where $c \geq 0$ and f is convex
- * Supremum: $\sup_{f \in \mathcal{F}} f$, where \mathcal{F} is a family of convex functions
- Example: hinge loss (skipped)
 - * Function: $f(w) = \max\{0, 1 - y(w \cdot x)\}$.
 - * Subgradient:
 - $\partial f(w) = \{0\}$ if $y(w \cdot x) > 1$ (w achieves margin at least 1)
 - $\partial f(w) = \{-yx\}$ if $y(w \cdot x) < 1$
 - $\partial f(w) = \{-yx\alpha : \alpha \in [0, 1]\}$ if $y(w \cdot x) = 1$
- The setup for online convex optimization is similar to online learning:
 - Iterate $t = 1, \dots, T$:
 - * Learner chooses $w_t \in S$
 - * Nature chooses convex loss function $f_t : S \rightarrow \mathbb{R}$
 - Regret is defined in the way you would expect (cumulative difference of losses):

$$\text{Regret}(u) \stackrel{\text{def}}{=} \sum_{t=1}^T [f_t(w_t) - f_t(u)]. \quad (15)$$

$$\text{Regret} \stackrel{\text{def}}{=} \max_{u \in S} \text{Regret}(u). \quad (16)$$

- The set S plays two roles: it is the set of experts with which we define our regret, and it is also the set of parameters that our learner is going to consider. For simplicity, we assume these two are the same, although in general, they do not have to be.
- We now turn to our original goal of doing online learning. We will show some examples of reducing online learning to online convex optimization. In particular, given an input/output sequence input $x_{1:T}, y_{1:T}$, we will construct a sequence of convex functions $f_{1:T}$ such that the low regret incurred by a learner on the online convex optimization problem implies low regret on the online learning problem. Let OL be the online learner who has access to OCO, a online convex optimizer.

– **Example 4 (linear regression)**

- * Assume we are doing linear regression with the squared loss, $\ell(y_t, p_t) = (p_t - y_t)^2$.
- * On each iteration $t = 1, \dots, T$:
 - OL receives an input $x_t \in \mathbb{R}^d$ from nature.

- OL asks OCO for a weight vector $w_t \in \mathbb{R}^d$.
- OL sends the prediction $p_t = w_t \cdot x_t$ to nature.
- OL receives the true output y_t from nature.
- OL relays the feedback to OCO via the loss function

$$f_t(w) = (w \cdot x_t - y_t)^2. \quad (17)$$

Note that (x_t, y_t) is baked into f_t , which changes each iteration. Since the squared loss is convex, f_t is convex. One can check easily based on matching definitions that the regret of OCO is exactly the same as regret of OL.

– **Example 5 (learning with expert advice)**

- * Now let's turn to a problem where the convexity structure isn't as apparent.
- * Assume we have a finite number (this is important) of experts \mathcal{H} , which the learner can query. Let

$$\mathcal{H} = \{h_1, \dots, h_d\}. \quad (18)$$

- * Assume we are doing binary classification with the zero-one loss (this is not important—any bounded loss function will do), $\ell(y_t, p_t) = \mathbb{I}[y_t \neq p_t]$.
- * Note that neither the loss function nor the set of experts \mathcal{H} is convex; in fact, this doesn't really make sense, since the domains are discrete.
- * Nonetheless, we can convexify the problem using **randomization**. Specifically, we allow the learner to produce a probability distribution $w_t \in \Delta_d$ ($\Delta_d \subseteq \mathbb{R}^d$ is the $(d-1)$ -dimensional simplex).² over the d experts \mathcal{H} , sample an expert from this distribution, and predict according to that expert. The expected zero-one loss is then a convex (in fact, linear) function of w_t .
- * Formally, the reduction is as follows:
 - OL receives an input x_t from nature.
 - OL asks OCO for a weight vector $w_t \in \Delta_d$, which represents a distribution over d experts.
 - OL samples an expert $j_t \sim w_t$ and send prediction $p_t = h_{j_t}(x_t)$ to nature.
 - OL receives the true output y_t from nature.
 - OL relays the feedback to OCO via the loss function

$$f_t(w) = w \cdot z_t, \quad (19)$$

where z_t is the vector of losses incurred by each expert:

$$z_t = [\ell(y_t, h_1(x_t)), \dots, \ell(y_t, h_d(x_t))] \in \{0, 1\}^d. \quad (20)$$

² Formally, think of a probability distribution over a random variable $J \in \{1, \dots, d\}$, with $\mathbb{P}[J = j] = w_j$.

Again, f_t bakes the loss and the data into the same function, and one can check that the expected regret of OL is the same as the regret of OCO. Note that we assume no structure on x_t and y_t —they could be arbitrarily complicated; we are only exposed to them via the experts and the loss function.

- * This convexify-by-randomization trick applies more generally to any loss function and any output space \mathcal{Y} . The key is that the set of experts is finite, and the learner is just choosing a convex combination of those experts.
- * Yet another way to convexify non-convex losses without randomization is to use an upper bound (e.g., hinge loss or logistic loss upper bounds the zero-one loss). However, minimizing the convex upper bound does not guarantee minimizing the zero-one loss.

- Relationship to convex optimization

- Before we talk about how to solve online convex optimization, let's see what the implications are for ordinary convex optimization, which is what one would typically turn to in the batch learning setting, where one gets all the data up front.
- Recall that in convex optimization, we are given a single convex function f and asked to minimize it, that is, compute:

$$w^* \in \arg \min_{w \in S} f(w). \quad (21)$$

Often in machine learning, the function f we're minimizing is a sum of convex functions, one for each of n data points:

$$f(w) = \frac{1}{n} \sum_{i=1}^n g_i(w), \quad (22)$$

where each g_i is convex. For example, g_i could be the hinge loss on training example (x_i, y_i) : $g_i(w) = \max\{0, 1 - y_i(w \cdot x_i)\}$.

- We can use an online convex optimization algorithm to minimize f as follows. At each time step t , choose a random g_i and feed it to the learner. Formally, let $f_t = g_{i_t}$, where i_t is drawn uniformly from $\{1, \dots, n\}$. An important property of this construction is that f_t is equal to f in expectation:

$$f(w) = \mathbb{E}[f_t(w)]. \quad (23)$$

- After T iterations, take the average over all the weight vectors that the learner predicted:

$$\bar{w} = \frac{1}{T} \sum_{t=1}^T w_t. \quad (24)$$

- We claim now that \bar{w} is good with respect to w^* . By the definition of regret:

$$\text{Regret}(w^*) = \sum_{t=1}^T [f_t(w_t) - f_t(w^*)]. \quad (25)$$

After dividing by T , rearranging, and taking expectations:

$$\mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T f_t(w_t) \right] = \mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T f_t(w^*) \right] + \frac{\mathbb{E}[\text{Regret}(w^*)]}{T}. \quad (26)$$

- * Let's analyze the LHS of (26). First, by definition of \bar{w} and Jensen's inequality (uniform distribution over $\{1, \dots, T\}$), we have:

$$f(\bar{w}) \leq \frac{1}{T} \sum_{t=1}^T f(w_t). \quad (27)$$

Taking expectations, and using linearity of expectation, and the fact that $f = \mathbb{E}[f_t]$:

$$\mathbb{E}[f(\bar{w})] \leq \mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T f(w_t) \right] = \mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T f_t(w_t) \right]. \quad (28)$$

- * Let's now analyze the first term of the RHS of (26). Since $f = \mathbb{E}[f_t]$,

$$f(w^*) = \mathbb{E} \left[\frac{1}{T} \sum_{t=1}^T f_t(w^*) \right]. \quad (29)$$

- * Therefore, we can conclude

$$\boxed{\mathbb{E}[f(\bar{w})] \leq f(w^*) + \frac{\mathbb{E}[\text{Regret}(w^*)]}{T}.} \quad (30)$$

- You should think of $\mathbb{E}[\text{Regret}(w^*)]$ as either $O(1)$, $O(\log T)$, or $O(\sqrt{T})$, so that $\mathbb{E}[\text{Regret}(w^*)]$ going to zero means convergence to the global minimum of f (solving the original convex optimization problem).
- By optimization standards, $O(1/\sqrt{T})$ or even $O(1/T)$ is very slow. Standard (batch) optimization algorithms such as gradient descent (with the appropriate step size) can, though not in all cases, obtain linear (confusingly, also called exponential or geometric) convergence ($O(e^{-c(T/n)})$) or even superlinear convergence ($O(e^{-c(T/n)^2})$).
- There are two reasons why you might deliberately not choose an algorithm with these rates:

- * First, note the difference in the dependence on n in the batch and online settings. Each iteration of an **online** algorithm processes only one example, whereas each iteration of an **batch** algorithm processes all n examples to compute $\nabla f(w)$, which is n times more expensive. When we have large datasets, then online algorithms can be faster.
- * Second, for machine learning applications, there is no need to optimize f to death since there is noise in the data anyway, and we'll show later that the statistical error is $O(1/\sqrt{T})$, so one could even argue that running an online algorithm over the data once ($T = n$) is theoretically good enough.

2.4 Follow the leader (FTL) (Lecture 2)

We first start out with a natural algorithm called follow the leader (FTL), which in some sense is the analog of the majority algorithm for online learning. We'll see that it works for quadratic functions but fails for linear functions. This will give us intuition about how to fix up our algorithm.

- **Algorithm 2 (follow the leader (FTL))**

- Let f_1, \dots, f_T be the sequence of loss functions played by nature.
- The learner chooses the weight vector $w_t \in S$ that minimizes the cumulative loss so far on the previous $t - 1$ iterations:

$$w_t \in \arg \min_{w \in S} \sum_{i=1}^{t-1} f_i(w). \quad (31)$$

(If there are multiple minima, choose any one of them. This is not important.)

- Aside: solving this optimization problem at each iteration is expensive in general (which would seem to destroy the spirit of online learning), but we'll consider special cases with analytic solutions.
- Note: We can think of FTL as an empirical risk minimizer (we will study this in batch learning), where the training set is the first $t - 1$ examples.
- We want to now study the regret of FTL. Regret compares the learner (FTL) with any expert $u \in S$, but this difference can be hard to reason about. So to get started, we will use the following result (Lemma 1) to replace u in the bound by (i) something easier to compare w_t to and (ii) at least as good as any $u \in S$.
- **Lemma 1 (compare FTL with one-step lookahead cheater)**
 - Let f_1, \dots, f_T be any sequence of loss functions.

- Let w_1, \dots, w_T be produced by FTL according to (31). For any $u \in S$:

$$\boxed{\text{Regret}(u) \stackrel{\text{def}}{=} \sum_{t=1}^T [f_t(w_t) - f_t(u)] \leq \sum_{t=1}^T [f_t(w_t) - f_t(w_{t+1})]}. \quad (32)$$

- Note: This is saying our regret against the best fixed u is no worse than comparing against the one-step lookahead w_{t+1} that peeks at the current function f_t (which is cheating!).
- Note: The RHS terms $f_t(w_t) - f_t(w_{t+1})$ measure how *stable* the algorithm is; smaller is better. Stability is an important intuition to develop.

• Proof of Lemma 1:

- Subtracting $\sum_t f_t(w_t)$ from both sides, it suffices to show

$$\sum_{t=1}^T f_t(w_{t+1}) \leq \sum_{t=1}^T f_t(u) \quad (33)$$

for all $u \in S$. Intuitively, this says that w_{t+1} (which takes the minimum over the first t functions) is better than using a fixed u for all time.

- Proof by induction:

- * Assume the inductive hypothesis on $T - 1$:

$$\sum_{t=1}^{T-1} f_t(w_{t+1}) \leq \sum_{t=1}^{T-1} f_t(u) \quad \text{for all } u \in S. \quad (34)$$

- * Add $f_T(w_{T+1})$ to both sides:

$$\sum_{t=1}^T f_t(w_{t+1}) \leq \sum_{t=1}^{T-1} f_t(u) + f_T(w_{T+1}) \quad \text{for all } u \in S. \quad (35)$$

- * In particular, this holds for $u = w_{T+1}$, so we have:

$$\sum_{t=1}^T f_t(w_{t+1}) \leq \sum_{t=1}^T f_t(w_{T+1}). \quad (36)$$

- * Since $w_{T+1} \in \arg \min_u \sum_{t=1}^T f_t(u)$ by definition of FTL, we have:

$$\sum_{t=1}^T f_t(w_{t+1}) \leq \sum_{t=1}^T f_t(u) \quad \text{for all } u \in S, \quad (37)$$

which is the inductive hypothesis for T .

- Note that Lemma 1 doesn't actually require on convexity of f_t , but rather only stability of the iterates $\{w_t\}$ as measured by f_t . As we'll see later, strong convexity is the main way we will achieve this stability. For now, let's consider two examples to gain some intuition: one where the $\{w_t\}$ are stable (Example 6) and one where they are not (Example 7).

- **Example 6 (quadratic optimization: FTL works)**

- Assume nature always chooses quadratic functions:

$$f_t(w) = \frac{1}{2} \|w - z_t\|_2^2, \quad (38)$$

where the points are bounded: $\|z_t\|_2 \leq L$ for all $t = 1, \dots, T$.

- FTL (minimizing over $S = \mathbb{R}^d$) has a closed form solution, which is just the average of the previous points:

$$w_t = \frac{1}{t-1} \sum_{i=1}^{t-1} z_i. \quad (39)$$

- Bound one term of the RHS of Lemma 1 (intuitively the difference is only one term):

$$f_t(w_t) - f_t(w_{t+1}) = \frac{1}{2} \|w_t - z_t\|_2^2 - \frac{1}{2} \|(1 - 1/t)w_t + (1/t)z_t - z_t\|_2^2 \quad (40)$$

$$= \frac{1}{2} (1 - (1 - 1/t)^2) \|w_t - z_t\|_2^2 \quad (41)$$

$$\leq (1/t) \|w_t - z_t\|_2^2 \quad (42)$$

$$\leq (1/t) 4L^2. \quad (43)$$

- Side calculation: summing $1/t$ yields $\log T$:

$$\sum_{t=1}^T (1/t) \leq 1 + \int_1^T (1/t) dt = \log(T) + 1. \quad (44)$$

- Summing over t yields:

$$\boxed{\text{Regret} \leq \sum_{t=1}^T [f_t(w_t) - f_t(w_{t+1})] \leq 4L^2(\log(T) + 1).} \quad (45)$$

- The important thing here is that the difference between w_t and w_{t+1} (measured in terms of loss) is really small (only $1/t$), which means that FTL for quadratic functions is really stable. This makes sense because averages are stable: adding an extra data point should only affect the running average by $O(1/t)$.

- **Example 7 (linear optimization: FTL fails)**

- We will construct an evil example to make FTL fail.
- Let $S = [-1, 1]$ be FTL's possible predictions. This is a nice bounded one-dimensional convex set, so we're not even trying hard to be pathological.
- Consider linear functions $f_t(w) = wz_t$ in $d = 1$ dimension, where

$$(z_1, z_2, \dots) = (-0.5, 1, -1, 1, -1, 1, -1, \dots). \quad (46)$$

- The minimizer computed by FTL will be attained at an extreme point, causing oscillating behavior.

$$(w_1, w_2, \dots) = (0, 1, -1, 1, -1, 1, -1, \dots). \quad (47)$$

- FTL obtains $T - 1$ cumulative loss (get loss 1 on every single example except the first).
- Expert $u = 0$ obtains 0 cumulative loss (not necessarily even the best).
- Therefore, the regret is pretty depressing:

$$\boxed{\text{Regret} \geq T - 1}. \quad (48)$$

- What's the lesson?

- For these quadratic functions, w_t and w_{t+1} must get closer (low regret).
- For these linear functions, w_t and w_{t+1} do not get closer (high regret).
- It seems then that FTL works when functions f_t offer “stability” (e.g., quadratic) but fail when they do not (e.g., linear).
- We will reveal the more general principle (strong convexity) later work.

[begin lecture 3] (3)

2.5 Follow the regularized leader (FTRL) (Lecture 3)

- It would be nice if nature just handed us quadratic-looking f_t 's, but in general, we're not the ones calling the shots there. But we do control the learner, so the key idea is to *add some regularization* of our own to stabilize the learner.
- **Algorithm 3 (follow the regularized leader (FTRL))**
 - Let $\psi : S \rightarrow \mathbb{R}$ be a function called a **regularizer** (this defines the learner).
 - Let f_1, \dots, f_T be the sequence of loss functions played by nature.
 - On iteration t , the learner chooses the weight vector that minimizes the regularizer plus the losses on the first $t - 1$ examples:

$$w_t \in \arg \min_{w \in S} \psi(w) + \sum_{i=1}^{t-1} f_i(w). \quad (49)$$

- Note: FTL is just FTRL with $\psi = 0$.
- Quadratic ψ , linear f_t
 - For the remainder of the section, just to build the right intuition in a transparent way, let's specialize to quadratic regularizers ψ and linear loss functions f_t :

$$\psi(w) = \frac{1}{2\eta} \|w\|_2^2, \quad f_t(w) = w \cdot z_t. \quad (50)$$

- Then the FTRL optimization problem (49) is:

$$w_t = \arg \min_{w \in S} \left\{ \frac{1}{2\eta} \|w\|_2^2 - w \cdot \theta_t \right\}, \quad (51)$$

where

$$\theta_t = - \sum_{i=1}^{t-1} z_i \quad (52)$$

is the negative sum of the gradients z_t . Interpret θ_t as the direction that we want to move in to reduce the loss, but now, unlike in FTL, we're held back by the quadratic regularizer.

- If $S = \mathbb{R}^d$, then FTRL has a closed form solution:

$$w_t = \eta\theta_t, \quad (53)$$

a scaled down version of θ_t . We can write $w_t = -\eta z_1 - \eta z_2 - \dots - \eta z_{t-1}$ and equivalently think of the weights as being updated incrementally according to the following recurrence:

$$\boxed{w_{t+1} = w_t - \eta z_t.} \quad (54)$$

From this perspective, the recurrence in (54) looks awfully like an online subgradient update where z_t is the gradient and η is the step size.

- If $S \neq \mathbb{R}^d$, then FTRL requires a projection onto S . We rewrite (51) by completing the square (add $\frac{\eta}{2}\|\theta_t\|_2^2$):

$$\boxed{w_t \in \arg \min_{w \in S} \frac{1}{2\eta} \|w - \eta\theta_t\|_2^2 = \Pi_S(\eta\theta_t),} \quad (55)$$

which is a Euclidean projection of $\eta\theta_t$ onto set S . This is called a **lazy projection** since θ_t still accumulates unprojected gradients, and we only project when we need to obtain a weight vector w_t for prediction. This is also known as Nesterov's **dual averaging** algorithm.

- Regularizers in online and batch learning
 - It's worth pausing to examine the difference in the way regularization enters online learning versus batch learning, with which you're probably more familiar.
 - In batch learning (e.g., in training an SVM or ridge regression), one typically seeks to optimize a function which is the sum of the training loss plus the regularizer. Notably, the regularizer is part of the objective function.
 - In online learning here, our objective in some sense is the regret, which makes no mention of the regularizer. The regularizer lies purely inside the learner's head, and is used to define the updates. In our example so far, the regularizer (in the context of FTRL) gives birth to the online gradient algorithm in (54).
- Now we will analyze the regret FTRL for quadratic regularizers and linear losses.
- **Theorem 1 (regret of FTRL)**
 - Let $S \subseteq \mathbb{R}^d$ be a convex set (of weight vectors).
 - Let f_1, \dots, f_T be any sequence of linear loss functions: $f_t(w) = w \cdot z_t$ for some $z_t \in \mathbb{R}^d$.
 - Let $\psi(w) = \frac{1}{2\eta}\|w\|_2^2$ be a quadratic regularizer for any $\eta > 0$.

- Then the regret of FTRL (as defined in Algorithm 3) with respect to any $u \in S$ is as follows:

$$\boxed{\text{Regret}(u) \leq \frac{1}{2\eta} \|u\|_2^2 + \frac{\eta}{2} \sum_{t=1}^T \|z_t\|_2^2.} \quad (56)$$

- Interpretation: the step size η allows us to trade off two terms:
 - First term: “bias” due to regularization. To compete with u , the learner’s iterates w_t must somehow get close to u . Smaller η means more regularization (remember $w_t = \eta\theta_t$), which means it’s harder to get there.
 - Second term: “variance” due to changing z_t . A smaller η means that successive weight vectors are closer and thus stabler (recall $w_{t+1} - w_t = -\eta z_t$). With a small η , we can hope to avoid the bad scenario in Example 7.
- FIGURE: [ellipse with u at edge, draw iterates w_t trying to reach u]
- Corollary:
 - To make the bound look cleaner:
 - * Let $\|z_t\|_2 \leq L$.
 - * Let $\|u\|_2 \leq B$ for all experts $u \in S$.
 Now (56) can be rewritten as:

$$\text{Regret}(u) \leq \frac{B^2}{2\eta} + \frac{\eta TL^2}{2}. \quad (57)$$

- Side calculation:
 - * Suppose we want to minimize some function with the following form: $C(\eta) = a/\eta + b\eta$.
 - * Take the derivative: $-a/\eta^2 + b = 0$, resulting in $\eta = \sqrt{a/b}$ and $C(\eta) = 2\sqrt{ab}$, which is just twice the geometric average of a and b .
- Letting $a = B^2/2$ and $b = TL^2/2$, we get $\eta = \frac{B}{L\sqrt{T}}$ and

$$\boxed{\text{Regret} \leq BL\sqrt{T}.} \quad (58)$$

Note that the average regret goes to zero as desired, even though not as fast as for quadratic functions ($\log T$).

- Proof of weakened version of Theorem 1

- We will prove Theorem 1 later using Bregman divergences, but just to give some intuition without requiring too much technical machinery, we will instead prove a slightly weaker result (note that the second term is looser by a factor of 2):

$$\boxed{\text{Regret}(u) \leq \frac{1}{2\eta} \|u\|_2^2 + \eta \sum_{t=1}^T \|z_t\|_2^2.} \quad (59)$$

- The key idea is to reduce FTRL to FTL. Observe that FTRL is the same as FTL where the first function is the regularizer.
- Let us then apply Lemma 1 to the sequence of functions ψ, f_1, \dots, f_T (when applying the theorem, note that the indices are shifted by 1). This results in the bound:

$$[\psi(w_0) - \psi(u)] + \sum_{t=1}^T [f_t(w_t) - f_t(u)] \leq [\psi(w_0) - \psi(w_1)] + \sum_{t=1}^T [f_t(w_t) - f_t(w_{t+1})]. \quad (60)$$

- Canceling $\psi(w_0)$, noting that $\psi(w_1) \geq 0$, and rearranging, we get:

$$\text{Regret}(u) \stackrel{\text{def}}{=} \sum_{t=1}^T [f_t(w_t) - f_t(u)] \leq \frac{1}{2\eta} \|u\|_2^2 + \sum_{t=1}^T [f_t(w_t) - f_t(w_{t+1})]. \quad (61)$$

- Now let's bound one term of the RHS sum:

$$f_t(w_t) - f_t(w_{t+1}) = z_t \cdot (w_t - w_{t+1}) \quad [\text{since } f_t(w) = w \cdot z_t] \quad (62)$$

$$\leq \|z_t\|_2 \|w_t - w_{t+1}\|_2 \quad [\text{Cauchy-Schwartz}] \quad (63)$$

$$= \|z_t\|_2 \|\Pi_S(\eta\theta_t) - \Pi_S(\eta\theta_{t+1})\|_2 \quad [\text{since } w_t = \Pi_S(\eta\theta_t)] \quad (64)$$

$$\leq \|z_t\|_2 \|\eta\theta_t - \eta\theta_{t+1}\|_2 \quad [\text{projection decreases distance}] \quad (65)$$

$$= \eta \|z_t\|_2^2 \quad [\text{since } \theta_{t+1} = \theta_t - z_t]. \quad (66)$$

Plugging this bound back into (61) completes the proof.

2.6 Online subgradient descent (OGD) (Lecture 3)

- So far, we have proven regret bounds for FTRL with linear loss functions. However, many loss functions we care about in practice (e.g., squared loss, hinge loss) are not linear.
- Even if did derive a regret for FTRL with more general losses, there would still be a computational problem: FTRL (see (49)) requires minimizing over all the loss functions seen so far, which in general is computationally impractical, especially for an online learning algorithm. For linear loss functions, on the other hand, we could optimize w_t easily by maintaining θ_t as a “sufficient statistics” (by linearity).

- Our strategy to handling general losses efficiently is by appealing to the linear machinery we already have. The key idea is to *run FTRL on a linear approximation of f_t* .
- What linear approximation $w \mapsto w \cdot z_t$ should we use? Let's use the subgradient of f_t at the current weights w_t : take any $z_t \in \partial f_t(w_t)$. Just to highlight the simplicity of the algorithm, here it is:
- **Algorithm 4 (Online subgradient descent (OGD))**

- Let $w_1 = 0$.
- For iteration $t = 1, \dots, T$:
 - * Predict w_t and receive f_t .
 - * Take any subgradient $z_t \in \partial f_t(w_t)$.
 - * If $S = \mathbb{R}^d$, perform the update:

$$w_{t+1} = w_t - \eta z_t. \quad (67)$$

- * If $S \subseteq \mathbb{R}^d$, project cumulative gradients onto S :

$$w_{t+1} = \Pi_S(\eta \theta_{t+1}), \quad \theta_{t+1} = \theta_t - z_t. \quad (68)$$

- To emphasize: OGD on f_t is nothing more than FTRL on quadratic regularizers and linear subgradient approximations of f_t .
- Analyzing regret
 - From our earlier analysis of FTRL (Theorem 1), we already have a bound on the regret on the linearized losses:

$$\sum_{t=1}^T [w_t \cdot z_t - u \cdot z_t]. \quad (69)$$

- We are interested on controlling the actual regret with respect to f_t :

$$\sum_{t=1}^T [f_t(w_t) - f(u)]. \quad (70)$$

- How do we relate these two? Here's where *convexity* comes in a crucial way.
- Since $z_t \in \partial f_t(w_t)$ is a subgradient, we have by the definition of subgradient:

$$f_t(u) \geq f_t(w_t) + z_t \cdot (u - w_t). \quad (71)$$

Rearranging, we get a direct comparison for each term of the regret:

$$\boxed{f_t(w_t) - f_t(u) \leq (w_t \cdot z_t) - (u \cdot z_t)}. \quad (72)$$

- The upshot is that the bounds we got for linear losses apply without modification to general losses! Intuitively, linear functions are the hardest to optimize using online convex optimization.
- FIGURE: [draw convex f_t with linearized]

- Remarks:

- OGD works for any convex loss function f_t (so does FTRL, but we only analyzed it for linear losses).
- OGD is in some sense the first practical, non-toy algorithm we’ve developed.
- Gradient-based methods are most commonly thought of as a procedure for optimizing a global objective, but this analysis provides a different perspective: that of doing full minimization of linearized losses with quadratic regularization.
- The minimization viewpoint opens way to many other algorithms, all of which can be thought of as using different regularizers or using better approximations of the loss functions, while maintaining efficient parameter updates.

- **Example 8 (Online SVM)**

- Let us use our result on OGD to derive a regret bound for learning SVMs in an online manner.
- We will just apply OGD on the hinge loss for classification ($x_t \in \mathbb{R}^d, y_t \in \{+1, -1\}$):

$$f_t(w) = \max\{0, 1 - y_t(w \cdot x_t)\}. \quad (73)$$

The algorithm (assume $S = \mathbb{R}^d$, so we don’t need to project):

- * If $y_t(w_t \cdot x_t) \geq 1$ (classify correctly with margin 1): do nothing.³
- * Else: $w_{t+1} = w_t + \eta y_t x_t$.
- Analysis:
 - * Assume the data points are bounded: $\|x_t\|_2 \leq L$. Then $z_t \in \partial f_t(w_t)$ also satisfies that bound $\|z_t\|_2 \leq L$.
 - * Assume that expert weights are bounded: $\|u\|_2 \leq B$.
 - * The regret bound from Theorem 1 is as follows:

$$\boxed{\text{Regret} \leq BL\sqrt{T}.} \quad (74)$$

- **Example 9 (Learning with expert advice)**

³This algorithm is very similar to the Perceptron algorithm; the only difference is that Perceptron just requires any positive margin, not necessarily 1.

- Now let us consider learning with expert advice.
 - * We maintain a distribution $w_t \in \Delta_d$ over d experts and predict by sampling an expert from that distribution.
 - * Assume the zero-one loss: $\ell(y_t, p_t) = \mathbb{I}[y_t \neq p_t]$.
 - * The loss function is linear: $f_t(w_t) = w_t \cdot z_t$, where

$$z_t = [\ell(y_t, h_1(x_t)), \dots, \ell(y_t, h_d(x_t))] \in \{0, 1\}^d \quad (75)$$

is the loss vector.

- Bound on set of experts (B): the experts live in the simplex $S = \Delta_d$, which has its 2-norm bounded by $B = 1$ (attained at a corner).
- Bound on loss gradient (L):
 - * The Lipschitz constant is bounded by the norm of the gradient z_t , which is at most \sqrt{d} .
 - * Therefore, the regret bound we get is

$$\boxed{\text{Regret} \leq BL\sqrt{T} = \sqrt{dT}.} \quad (76)$$

- Note that we are depending on the square root of the number of experts d rather than the log of the number of experts in our first online learning bound for learning with expert advice. Can we obtain a $\log d$ dependence without assuming realizability? We will find out in the next section.

2.7 Online mirror descent (OMD) (Lecture 4)

So far, we have analyzed FTRL for quadratic regularizers, which leads to (lazy projected) gradient-based algorithms. Quadratic regularization is imposing a certain prior knowledge, namely that there is a good parameter vector w in a small L_2 ball. But perhaps we know some dimensions to be more important than others. Then we might want to use a non-spherical regularizer. Or in the case of learning with expert advice, we know that $w \in \Delta_d$ (a probability distribution), so negative entropy might be more appropriate. In this section, we will develop a general way of obtaining regret bounds for general regularizers, and make explicit the glorious role that strong convexity plays. We make make extensive use of **Fenchel duality** and **Bregman divergences**.

- The goal for this lecture is to analyze FTRL (Algorithm 3) for arbitrary convex loss functions and regularizers. The resulting algorithm is this:
- **Algorithm 5 (online mirror descent (OMD))**

- Let $\psi : \mathbb{R}^d \rightarrow \mathbb{R}$ be the regularizer (this defines the learner).
- Let f_1, \dots, f_T be the sequence of loss functions played by nature.
- On each iteration $t = 1, \dots, T$, the learner chooses weights w_t to minimize the regularized (linearized) loss:

$$w_t \in \arg \min_{w \in \mathbb{R}^d} \{\psi(w) - w \cdot \theta_t\}, \quad (77)$$

where $z_t \in \partial f_t(w_t)$ is the t -th subgradient, and $\theta_t = -\sum_{i=1}^{t-1} z_i$ is the negative sum of the first $t-1$ subgradients.

- Technical note: we will let the domain of weight vectors be unconstrained ($S = \mathbb{R}^d$). We can always fold a constraint into the regularizer by setting $\psi(w) = \infty$ if w violates the constraint.
- To recap the terminology:
 - OMD on f_t is equivalent to FTRL on the linearizations $w \mapsto w \cdot z_t$.
 - OGD is OMD with the quadratic regularizer $\psi(w) = \frac{1}{2\eta} \|w\|_2^2$
- Examples of regularizers:
 - Quadratic regularizer: $\psi(w) = \frac{1}{2\eta} \|w\|_2^2$.
 - Non-spherical quadratic regularizer: $\psi(w) = \frac{1}{2\eta} w^\top A w$ for $A \succeq 0$.

- Entropic regularizer: $\psi(w) = \frac{1}{\eta} \sum_{j=1}^d w_j \log w_j$ if $w \in \Delta_d$ (this is the negative entropy defined on the probability simplex), and ∞ otherwise.
- Note: the difference between the two regularizers is that the entropic regularizer slopes up violently when w approaches the boundary of the simplex Δ_d (the function value is finite but the gradient goes to infinity).
- We now need to build up some tools to help us analyze OMD. First, we introduce Fenchel duality, an important tool in optimization:

– **Definition 4 (Fenchel conjugate)**

- * The **Fenchel conjugate**⁴ of a function (not necessarily convex) $\psi : \mathbb{R}^d \rightarrow \mathbb{R}$ is

$$\psi^*(\theta) \stackrel{\text{def}}{=} \sup_{w \in \mathbb{R}^d} \{w \cdot \theta - \psi(w)\}. \quad (78)$$

– Intuition

- * For scalars $w, \theta \in \mathbb{R}$, given a fixed θ (interpreted as a slope), $-\psi^*(\theta)$ is the position where the supporting hyperplane of ψ with slope θ hits the vertical axis.
- FIGURE: [draw ψ]
- * One can think of sweeping θ and reading out $\psi^*(\theta)$; this information in some sense encodes the epigraph of ψ if ψ is convex.

– Useful facts:

- * ψ^* is always convex (even if ψ is not), since it's just a supremum over a collection of linear functions $\theta \mapsto w \cdot \theta - \psi(w)$.
- * $\psi^*(\theta) \geq w \cdot \theta - \psi(w)$ for all $w \in \mathbb{R}^d$ (**Fenchel-Young inequality**). This follows directly from the definition of ψ^* .
- * If $r(w) = a\psi(w)$ with $a > 0$, then $r^*(\theta) = a\psi^*(\theta/a)$. This is useful because once we've computed the Fenchel conjugate for one function, we know it for different scalings. In fact, Fenchel duality has a very nice algebra that makes computing Fenchel conjugates modular.
- * $\psi^{**} = \psi$ iff ψ is convex (and lower semi-continuous). In this case, we have

$$\psi(w) = \psi^{**}(w) = \sup_{\theta \in \mathbb{R}^d} \{w \cdot \theta - \psi^*(\theta)\}. \quad (79)$$

- * If ψ is differentiable, then

$$\nabla \psi^*(\theta) = \arg \max_{w \in \mathbb{R}^d} \{w \cdot \theta - \psi(w)\}. \quad (80)$$

This is because the gradient of the supremum of a collection of functions at θ is the gradient of the function that attains the max at θ .

⁴Also known as the convex conjugate or Legendre-Fenchel transformation.

– Mirroring

- * Comparing this with the the OMD update (77), we see that $w_t = \nabla\psi^*(\theta_t)$, and $-\psi^*(\theta_t)$ is the corresponding value of the regularized loss.
- * Since w_t attains the supremum of the Fenchel conjugate ψ^* , the optimality conditions (differentiate with respect to w) tell us that $\theta_t = \nabla\psi(w_t)$.
- * We have a one-to-one mapping between weights w and negative cumulative subgradients θ , linked via the gradients of ψ and ψ^* , which are inverses of one another:

$$\boxed{w_t = \nabla\psi^*(\theta_t), \quad \theta_t = \nabla\psi(w_t).} \quad (81)$$

- * This provides a very elegant view of what online mirror descent is doing. OMD takes a series of gradient updates $\theta_{t+1} = \theta_t - z_t$, generating $\theta_1, \theta_2, \dots, \theta_T$. These steps are *mirrored* via Fenchel duality in the sequence w_1, w_2, \dots, w_T .
- * FIGURE: [mapping]

– **Example 10 (Quadratic regularizer)**

- * Let $\psi(w) = \frac{1}{2\eta}\|w\|_2^2$.
- * Then $\psi^*(\theta) = \frac{\eta}{2}\|\theta\|_2^2$, attained by $w = \nabla\psi^*(\theta) = \eta\theta$.
- * In this case, w and θ are simple rescalings of each other.

– **Example 11 (Entropic regularizer)**

- * Let $\psi(w) = \frac{1}{\eta} \sum_{j=1}^d w_j \log w_j$ for $w \in \Delta_d$ (negative entropy).
- * Then $\psi^*(\theta) = \frac{1}{\eta} \log \left(\sum_{j=1}^d e^{\eta\theta_j} \right)$ (log partition function), attained by $w_j = \frac{e^{\eta\theta_j}}{\sum_{k=1}^d e^{\eta\theta_k}}$.
- * Aside: this is maximum entropy duality in exponential families, where w and θ represent the mean and canonical parametrization of a multinomial distribution.

2.8 Regret bounds with Bregman divergences (Lecture 4)

• Motivation

- Having reinterpreted OMD as a mapping using a conjugate function, let's turn to proving a regret bound. Recall that in order to prove bounds, we needed to ensure that w_t and w_{t+1} don't change too much according to some criteria.
- For quadratic regularization, this criteria was based on Euclidean distance.
- We now generalize this by using Bregman divergences, which generalizes the notion of distance based on the regularizer.
- **Definition 5 (Bregman divergence)**

- * Let f be a continuously-differentiable convex function.
- * The **Bregman divergence** between w and u is the difference at w between f and a linear approximation around u :

$$D_f(w||u) \stackrel{\text{def}}{=} f(w) - f(u) - \nabla f(u) \cdot (w - u). \quad (82)$$

- * Intuitively, the divergence captures the the error of the linear approximation of f based on the gradient $\nabla f(u)$.
- * FIGURE: [show gap between f and its linear approximation]
- Property: $D_f(w||u) \geq 0$ (by definition of convexity).
- Note: Bregman divergences are not symmetric and therefore not a distance metric.
- **Example 12 (Quadratic regularizer)**
 - * Let $f(w) = \frac{1}{2}\|w\|_2^2$.
 - * Then $D_f(w||u) = \frac{1}{2}\|w - u\|_2^2$ (squared Euclidean distance).
- **Example 13 (Entropic regularizer)**
 - * Let $f(w) = \sum_{j=1}^d w_j \log w_j$ for $w \in \Delta_d$ (negative entropy).
 - * Then $D_f(w||u) = \text{KL}(w||u) = \sum_j w_j \log(w_j/u_j)$ for $w \in \Delta_d$ (KL divergence).
- Property (scaling): $D_{af}(w||u) = aD_f(w||u)$.

• **Theorem 2 (regret of OMD using Bregman divergences)**

- OMD (Algorithm 5) obtains the following regret bound:

$$\text{Regret}(u) \leq [\psi(u) - \psi(w_1)] + \sum_{t=1}^T D_{\psi^*}(\theta_{t+1}||\theta_t). \quad (83)$$

- Furthermore, if all f_t 's are linear, the inequality is actually an equality if u is the best expert. So this bound is pretty air tight.

• **Proof of Theorem 2:**

- We assume all the loss functions are linear; recall in our analysis of OGD that linear functions are the worst case.
- The key step is to find a potential function which allows us to monitor progress. The pivoting quantity is $\psi^*(\theta_{T+1})$, the negative regularized loss of the best fixed expert. This will allow us to relate the learner (w_t) to the expert (u).
- Recall:
 - * Learner's loss is $\sum_{t=1}^T w_t \cdot z_t$.

- * Expert's loss is $\sum_{t=1}^T u \cdot z_t$.
- * The regret is the difference between the two.
- The expert:

$$\psi^*(\theta_{T+1}) \geq u \cdot \theta_{T+1} - \psi(u) \quad (84)$$

$$= \sum_{t=1}^T [-u \cdot z_t] - \psi(u) \quad (85)$$

by the Fenchel-Young inequality. Note that we have equality if u is the best expert, by definition of ψ^* .

- The learner:

$$\psi^*(\theta_{T+1}) = \psi^*(\theta_1) + \sum_{t=1}^T [\psi^*(\theta_{t+1}) - \psi^*(\theta_t)] \quad [\text{telescoping sums}] \quad (86)$$

$$= \psi^*(\theta_1) + \sum_{t=1}^T [\nabla \psi^*(\theta_t) \cdot (\theta_{t+1} - \theta_t) + D_{\psi^*}(\theta_{t+1} \| \theta_t)] \quad [\text{Bregman definition}] \quad (87)$$

$$= \psi^*(\theta_1) + \sum_{t=1}^T [-w_t \cdot z_t + D_{\psi^*}(\theta_{t+1} \| \theta_t)] \quad [\text{definition of OMD (81) and } \theta_t]. \quad (88)$$

Note that $\psi^*(\theta_1) = -\psi(w_1)$ since $\theta_1 = 0$.

- Combining last equations for the expert and learner and rearranging yields the result.
- The regret bound (83) is a generalization of (56), where $\psi(w) = \frac{1}{2\eta} \|w\|_2^2$ is quadratic regularizer, and $D_{\psi^*}(\theta_{t+1} \| \theta_t) = \frac{\eta}{2} \|z_t\|_2^2$.
- However, a general Bregman divergence usually doesn't have such a nice form, and thus it is useful to bound it using a nicer quantity: a norm of some kind.

2.9 Strong convexity and smoothness (Lecture 4)

- First, let's review some intuition behind norms.

- L_p norms decrease as p increases:

$$\|w\|_1 \geq \|w\|_2 \geq \|w\|_\infty. \quad (89)$$

- The difference between the norms can be huge. For example, take $w = (1, \dots, 1) \in \mathbb{R}^d$. Then $\|w\|_1 = d$, $\|w\|_2 = \sqrt{d}$, $\|w\|_\infty = 1$.

- Recall that the dual norm of a norm $\|\cdot\|$ is

$$\|x\|_* = \sup_{\|y\| \leq 1} (x \cdot y). \quad (90)$$

Thinking of y as a linear operator on x , the dual norm measures the gain when we measure perturbations in x using $\|\cdot\|$.

- The norms $\|\cdot\|_p$ and $\|\cdot\|_q$ are dual to each other when $\frac{1}{p} + \frac{1}{q} = 1$. Two common examples are:

$$\|\cdot\|_1 \text{ is dual to } \|\cdot\|_\infty. \quad (91)$$

$$\|\cdot\|_2 \text{ is dual to } \|\cdot\|_2. \quad (92)$$

- We will now use these squared norms to control Bregman divergences provided the Bregman divergences have an important special structure called *strong convexity*:

- **Definition 6 (strong convexity/smoothness)**

- A function f is α -strongly convex with respect to a norm $\|\cdot\|$ iff for all w, u :

$$D_f(w\|u) \geq \frac{\alpha}{2} \|w - u\|^2. \quad (93)$$

- A function f is α -strongly smooth with respect to a norm $\|\cdot\|$ iff for all w, u :

$$D_f(w\|u) \leq \frac{\alpha}{2} \|w - u\|^2. \quad (94)$$

- **Intuitions**

- Strong convexity means that f must be growing at least quadratically.
- Strong smoothness means that f must be growing slower than some quadratic function.
- Example: the quadratic regularizer $\psi(w) = \frac{1}{2} \|w\|_2^2$ is both 1-strongly convex and 1-strongly smooth with respect to the L_2 norm, since $D_\psi(w\|u) = \frac{1}{2} \|w - u\|_2^2$.

- Duality links strong convexity and strong smoothness, as the following result shows.

- **Lemma 2 (strong convexity and strong smoothness)**

- The following two statements are equivalent:
 - * $\psi(w)$ is $1/\eta$ -strongly convex with respect to a norm $\|\cdot\|$.
 - * $\psi^*(\theta)$ is η -strongly smooth with respect to the dual norm $\|\cdot\|_*$.
- Sanity check the quadratic regularizer: $\psi(w) = \frac{1}{2\eta} \|w\|_2^2$ and $\psi^*(\theta) = \frac{\eta}{2} \|\theta\|_2^2$.

- With these tools, we can finally make some progress on our regret bound from Theorem 2, rewriting the Bregman divergences in terms of norms.

- **Theorem 3 (regret of OMD using norms)**

- Suppose ψ is a $\frac{1}{\eta}$ -strongly convex regularizer.
- Then the regret of online mirror descent is

$$\boxed{\text{Regret}(u) \leq [\psi(u) - \psi(w_1)] + \frac{\eta}{2} \sum_{t=1}^T \|z_t\|_*^2.} \quad (95)$$

- **Proof of Theorem 3**

- By Lemma 2, since ψ is $1/\eta$ -strongly convex, the Fenchel conjugate ψ^* is η -strongly smooth.
- By definition of strong smoothness and the fact that $\theta_{t+1} = \theta_t - z_t$,

$$\boxed{D_{\psi^*}(\theta_{t+1} \| \theta_t) \leq \frac{\eta}{2} \|z_t\|_*^2.} \quad (96)$$

- Plugging this bound into (83) gives us the result.

- **Remarks:**

- If we plug in the quadratic regularizer $\psi(w) = \frac{1}{2\eta} \|w\|_2^2$, then we get back the original result (56).
- However, (95) generalizes to other regularizers.
- We get to measure the size of z_t using the dual norm $\|\cdot\|_*$ rather than the default L_2 norm, which will help us get tighter bounds for the learning from expert advice problem.

- **Learning from expert advice**

- Let's now use our tools to improve the regret bound that we got for learning from expert advice.
- Recall that when we used the quadratic regularizer, we got a regret bound of \sqrt{dT} because $\|z_t\|_2$ could be as big as \sqrt{d} .
- To reduce this, let's just use another norm: $\|z_t\|_\infty \leq 1$; no dependence on d !
- But this means that the regularizer ψ has to be strongly convex with respect to the L_1 norm, but this is harder because $\|\cdot\|_1 \geq \|\cdot\|_2$.

- Failure: the quadratic regularizer $\psi(w) = \frac{1}{2}\|w\|_2^2$ is only $\frac{1}{d}$ -strongly convex with respect to the L_1 norm:

$$D_\psi(w\|u) \geq \frac{1}{2}\|w - u\|_2^2 \geq \frac{1}{2d}\|w - u\|_1^2. \quad (97)$$

Sanity check: $\|(1, \dots, 1)\|_2^2 = d$ but $\|(1, \dots, 1)\|_1^2 = d^2$. So if we try to use this in the regret bound, we get $\frac{1}{2\eta} + \frac{T\eta d}{2}$, which is still \sqrt{dT} (for optimal η).

- Failure: the L_1 regularizer $\psi(w) = \|w\|_1$ is neither strongly convex nor strongly smooth with respect to the any norm for any α : it doesn't grow fast enough for large w and grows too fast for small w .
- Success: entropic regularizer $\psi(w) = \frac{1}{\eta} \sum_j w_j \log w_j$ for $w \in \Delta_d$ is $1/\eta$ -strongly convex with respect to the L_1 norm. This requires some algebra and an application of Cauchy-Schwartz (see Example 2.5 in Shai Shalev-Shwartz's online learning tutorial).
- FIGURE: [compare quadratic and entropic regularizer for $d = 2$]

• **Example 14 (exponentiated gradient (EG))**

- Entropic regularizer: $\psi(w) = \frac{1}{\eta} \sum_{j=1}^d w_j \log w_j$ for $w \in \Delta_d$.
- Recall $\psi^*(\theta) = \frac{1}{\eta} \log \sum_{j=1}^d e^{\eta\theta_j}$ and $\nabla\psi_j^*(\theta) = \frac{e^{\eta\theta_j}}{\sum_{k=1}^d e^{\eta\theta_k}}$.
- OMD updates:

$$\boxed{w_{t,j} \propto e^{\eta\theta_{t,j}}.} \quad (98)$$

The equivalent recursive formula:

$$\boxed{w_{t+1,j} \propto w_{t,j} e^{-\eta z_{t,j}}.} \quad (99)$$

Interpretation: we maintain a distribution over the d experts. We use the gradient z_t to reweight the experts, normalizing afterwards to ensure a proper distribution.

- Recap: the exponentiated gradient (EG) algorithm is just online mirror descent using the entropic regularizer.

• **Example 15 (EG for learning with expert advice)**

- Consider learning with expert advice (Example 5).
- We use the expected zero-one loss: $f_t(w) = w \cdot z_t$, where z_t is the loss vector.
- We have that the dual norm of the gradients are bounded $\|z_t\|_\infty \leq 1$.
- The minimum and maximum values of the regularizer:
 - * $\max_w \psi(w) = 0$ (minimum entropy)

$$* \min_w \psi(w) = \frac{\log(1/d)}{\eta} = \frac{-\log d}{\eta} \text{ (maximum entropy)}$$

– Then

$$\text{Regret} = \frac{\log(d)}{\eta} + \frac{\eta T}{2}. \quad (100)$$

– Setting $\eta = \sqrt{\frac{2\log d}{T}}$ yields:

$$\boxed{\text{Regret} = \sqrt{2\log(d)T}.} \quad (101)$$

- To compare with quadratic regularization (OGD):

– Quadratic: $[\max_u \psi(u) - \min_u \psi(u)] \leq \frac{1}{2\eta}, \|z_t\|_2 \leq \sqrt{d}$

– Entropic: $[\max_u \psi(u) - \min_u \psi(u)] \leq \frac{\log d}{\eta}, \|z_t\|_\infty \leq 1$

- Discussion

- Online mirror descent (OMD) allows us to use different regularizers based on our prior knowledge about the expert vector u and the data z_t . As we see with EG, tailoring the regularizer can lead to better bounds. (this is too loose by a factor of \sqrt{d} because we know u lives in the simplex, which is than the L_2 ball).
- Using the L_2 norm means that we use the bound $\|z_t\|_2 \leq \sqrt{d}$. To get rid of the dependence on d here, we use the L_∞ norm with $\|z_t\|_\infty \leq 1$.
- However, this means that ψ must be strongly convex with respect to the L_1 norm. The quadratic regularizer isn't strong enough (only $\frac{1}{d\eta}$ -strongly convex with respect to L_2), so we need the entropic regularizer, which is 1-strongly convex with respect to L_1 .
- Curvature hurts us because $[\psi(u) - \psi(w_1)]$ is now larger, but the simplex is small, so $[\psi(u) - \psi(w_1)]$ only grows from 1 (with the quadratic regularizer) to $\log d$ (with the entropic regularizer).
- So the tradeoff was definitely to our advantage.

2.10 Local norms (Lecture 5)

- Recall that the general online mirror descent (OMD) analysis (Theorem 2) yields:

$$\text{Regret}(u) \leq [\psi(u) - \psi(w_1)] + \sum_{t=1}^T D_{\psi^*}(\theta_{t+1} \| \theta_t). \quad (102)$$

Using the fact that ψ is $1/\eta$ -strongly convex with respect to some norm $\|\cdot\|$, we can upper bound the Bregman divergence by the following (95):

$$D_{\psi^*}(\theta_{t+1} \| \theta_t) \leq \frac{\eta}{2} \|z_t\|_*^2. \quad (103)$$

- If we use the entropic regularizer for ψ with norm $\|\cdot\| = \|\cdot\|_1$ and dual norm $\|\cdot\|_* = \|\cdot\|_\infty$, we get our key $\sqrt{2 \log(d) T}$ regret bound for EG.
- In this section, we will use the local norms technique to improve (103). This will allow us to do two things:

- Recover the strong $O(\log d)$ bound for the realizable setting.
- Allow us to analyze the multi-armed bandit setting.

- Why we should do better:

- Consider an example where there are two experts: one which is perfect ($z_{t,1} = 0$ for all t) and one which is horrible ($z_{t,2} = 1$ for all t).
- The EG algorithm will quickly downweight the bad expert exponentially fast:

$$w_{t,1} \propto 1 \quad w_{t,2} \propto e^{-\eta(t-1)}. \quad (104)$$

- So as $t \rightarrow \infty$, we have basically put all our weight on the first expert, and should be suffering no loss.
- But $\|z_t\|_\infty^2 = 1$, so in the regret bound we still pay $\frac{\eta T}{2}$, which is simply a travesty.
- We would hope that $\|z_t\|_\infty^2 = \max_{1 \leq j \leq d} z_{t,j}^2$ should be replaced with something that is sensitive to how much weight we're placing on it, which is $w_{t,j}$. Indeed the following theorem fulfills this dream:

- Theorem 4 (exponentiated gradient (analysis using local norms))**

- Assume nature plays a sequence of linear loss functions $f_t(w) = w \cdot z_t$, where $z_{t,j} \geq 0$ for all $t = 1, \dots, T$ and $j = 1, \dots, d$.

- Then the exponentiated gradient (EG) algorithm (Example 15) achieves the following regret bound:

$$\boxed{\text{Regret}(u) \leq [\psi(u) - \psi(w_1)] + \eta \sum_{t=1}^T \sum_{j=1}^d w_{t,j} z_{t,j}^2.} \quad (105)$$

- This bound (105) is better than the bound in Theorem 3 because we are taking an average (with respect to the distribution $w_t \in \Delta_d$) instead of a max:

$$\sum_{j=1}^d w_{t,j} z_{t,j}^2 \leq \max_{1 \leq j \leq d} z_{t,j}^2 \stackrel{\text{def}}{=} \|z_t\|_\infty^2. \quad (106)$$

This allows some components of the loss subgradients $z_{t,j}$ to be large provided that the corresponding weights $w_{t,j}$ are small.

- **Example 16 (exponentiated gradient in the realizable setting)**

- Assume the loss vector is bounded: $z_t \in [0, 1]^d$.
- Assume there exists an expert $u \in \Delta_d$ with zero cumulative loss (realizability).
- Recall the regret bound from using local norms:

$$\text{Regret}(u) \stackrel{\text{def}}{=} \sum_{t=1}^T (w_t \cdot z_t) - \underbrace{\sum_{t=1}^T (u \cdot z_t)}_{=0} \leq \frac{\log d}{\eta} + \eta \sum_{t=1}^T \underbrace{\sum_{j=1}^d w_{t,j} z_{t,j}^2}_{\leq w_t \cdot z_t} \quad (107)$$

- Note that $\sum_{j=1}^d w_{t,j} z_{t,j}^2 \leq w_t \cdot z_t$, because all quantities are non-negative and $a^2 \leq a$ for $a \in [0, 1]$.
- Rearranging, we get:

$$\text{Regret}(u) \leq \frac{\log d}{\eta(1 - \eta)}. \quad (108)$$

- Minimize the bound with $\eta = \frac{1}{2}$ to obtain:

$$\boxed{\text{Regret}(u) \leq 4 \log d.} \quad (109)$$

- Recall that the majority algorithm (which aggressively zeros out the weights of components as soon as they err) also obtained the very low $O(\log d)$ regret (see Example 3), so it's really nice to see that EG obtains the same regret guarantee.

- If the problem isn't realizable, the majority algorithm isn't even correct (it will eliminate all the experts), whereas EG will gracefully fall back to $O(\sqrt{\log(d)T})$ regret.
- Now that you are hopefully somewhat convinced that the theorem is useful, let's prove it.
- Proof of Theorem 4:
 - This proof mostly involves starting with the Bregman divergence, and performing some low-level algebraic manipulation. Perhaps the most useful high-level take-away is whenever we're trying to massage some expression with log's and exp's, it's useful to try to approximate the functions using linear and quadratic approximations (think Taylor approximations).
 - Specifically, we will use the following two facts:
 - * Fact 1: $e^{-a} \leq 1 - a + a^2$ for $a \geq 0$ (this actually holds for smaller a , but let's keep it simple)
 - * Fact 2: $\log(1 - a) \leq -a$
 - Recall the Fenchel conjugate of the entropic regularizer is the log-partition function:

$$\psi^*(\theta) = \frac{1}{\eta} \log \sum_{j=1}^d e^{\eta \theta_j}. \quad (110)$$

- By the definition of Bregman divergences (this was used in the proof of Theorem 2), we have:

$$D_{\psi^*}(\theta_{t+1} || \theta_t) = \psi^*(\theta_{t+1}) - \psi^*(\theta_t) - \underbrace{\nabla \psi^*(\theta_t)}_{w_t} \cdot \underbrace{(\theta_{t+1} - \theta_t)}_{-z_t}. \quad (111)$$

- The rest is just applying the two facts and watching stuff cancel:

$$D_{\psi^*}(\theta_{t+1} || \theta_t) = \psi^*(\theta_{t+1}) - \psi^*(\theta_t) + w_t \cdot z_t \quad (112)$$

$$= \frac{1}{\eta} \log \left(\frac{\sum_{j=1}^d e^{\eta \theta_{t+1,j}}}{\sum_{j=1}^d e^{\eta \theta_{t,j}}} \right) + w_t \cdot z_t \quad [\text{definition of } \psi^*] \quad (113)$$

$$= \frac{1}{\eta} \log \left(\sum_{j=1}^d w_{t,j} e^{-\eta z_{t,j}} \right) + w_t \cdot z_t \quad [\text{definition of } \theta_t] \quad (114)$$

$$\leq \frac{1}{\eta} \log \left(\sum_{j=1}^d w_{t,j} [1 - (\eta z_{t,j} - \eta^2 z_{t,j}^2)] \right) + w_t \cdot z_t \quad [\text{fact 1}] \quad (115)$$

$$= \frac{1}{\eta} \log \left(1 - \sum_{j=1}^d w_{t,j} (\eta z_{t,j} - \eta^2 z_{t,j}^2) \right) + w_t \cdot z_t \quad [w_t \in \Delta_d] \quad (116)$$

$$\leq \frac{1}{\eta} \sum_{j=1}^d w_{t,j} (-\eta z_{t,j} + \eta^2 z_{t,j}^2) + w_t \cdot z_t \quad [\text{fact 2}] \quad (117)$$

$$= \eta \sum_{j=1}^d w_{t,j} z_{t,j}^2 \quad [\text{algebra}]. \quad (118)$$

2.11 Multi-armed bandits (Lecture 5)

- In online learning with expert advice, on each iteration, after we choose one of the d experts/actions, nature reveals the loss vector z_t for every single action.
- However, in many applications such as clinical trials or advertisement placement, packet routing, we only get to observe the loss of the action we took, not the losses of the actions we didn't take.
- This setting is known as the multi-armed bandit problem, which is a type of online learning problem with **partial feedback**.
- This problem is much more difficult. Intuitively, the learner should choose actions that we expect to yield low loss, but it must choose which actions to explore to get more information about the losses. Thus, the multi-armed bandit problem exposes the challenges of the classic exploration/exploitation tradeoff.⁵
- Setup
 - FIGURE: [matrix with loss functions]
 - There are d possible actions (corresponding to the arms of a row of slot machines).

⁵ Reinforcement learning requires managing the exploration/exploitation tradeoff, but is even more difficult because actions take the learner between different states in a way that is unknown to the learner.

- Let $z_t \in [0, 1]^d$ denote the vector of losses of the d actions at iteration t .
- For each iteration $t = 1, \dots, T$:
 - * Learner chooses a distribution $w_t \in \Delta_d$ over actions, and samples an action $a_t \sim w_t$.
 - * Learner observes the loss of *only that particular action* z_{t,a_t} and no others.
- The expected regret with respect to an expert $u \in \Delta_d$ is defined in the same way as it was for online learning with expert advice:

$$\mathbb{E}[\text{Regret}(u)] = \sum_{t=1}^T [w_t \cdot z_t - u \cdot z_t]. \quad (119)$$

Note that taking the max over randomized experts $u \in \Delta_d$ is equivalent to taking the max over single actions $a \in [d]$, since the maximum over a linear function is attained at one of the vertices.

- **Estimating the loss vector**

- Recall that in online learning, we would observe the entire loss vector z_t . In that case, we could use the exponentiated gradient (EG) algorithm ($w_{t+1,j} \propto w_{t,j} e^{-\eta z_{t,j}}$) to obtain a regret bound of $\text{Regret} \leq \sqrt{2 \log(d) T}$ (see Example 15).
- In the bandit setting, we don't observe z_t , so what can we do? Let's try to estimate it with some \hat{z}_t that (i) we can observe and (ii) is equal to z_t in expectation (unbiased) in that for all $a = 1, \dots, d$:

$$\mathbb{E}_{a_t \sim w_t}[\hat{z}_{t,a} \mid w_t] = z_{t,a}. \quad (120)$$

- Given these two constraints, it's not hard to see that the only choice for \hat{z}_t is:

$$\hat{z}_{t,a} = \begin{cases} \frac{z_{t,a}}{w_{t,a}} & \text{if } a = a_t \\ 0 & \text{otherwise.} \end{cases} \quad (121)$$

Note that dividing $w_{t,a}$ compensates for sampling $a_t \sim w_t$.

- **Algorithm 6 (Bandit-EG)**

- Run EG with the unbiased estimate \hat{z}_t rather than z_t . Really, that's it.

- **Theorem 5 (Bandit-EG analysis)**

- Bandit-EG obtains expected regret (expectation taken over learner's randomness) of

$$\mathbb{E}[\text{Regret}(u)] \leq 2\sqrt{d \log(d) T}. \quad (122)$$

- Comparison of Bandit-EG (bandit setting) and EG (online learning setting):
 - Compared with the bound for EG in the full information online learning setting (101), we see that this bound (122) is worse by a factor of \sqrt{d} .
 - In other words, the number of iterations T for Bandit-EG needs to be d times larger in order to obtain the same average regret as EG.
 - This is intuitive since in the bandit setting, each iteration reveals $(1/d)$ -th the amount of information compared with the online learning setting.
- Proof of Theorem 5:
 - If EG is run with the unbiased estimate \hat{z}_t ($z_t = \mathbb{E}[\hat{z}_t \mid w_t]$), then the expected regret is simply:

$$\mathbb{E}[\text{Regret}(u)] \leq \frac{\log d}{\eta} + \eta \sum_{t=1}^T \mathbb{E}_{a_t \sim w_t} \left[\sum_{a=1}^d w_{t,a} \hat{z}_{t,a}^2 \mid w_t \right]. \quad (123)$$

Note that the random variable \hat{z}_t only depends on the previous actions through the random variable w_t .

- So now, we just have to compute the expected local norm:

$$\mathbb{E}_{a_t \sim w_t} \left[\sum_{a=1}^d w_{t,a} \hat{z}_{t,a}^2 \mid w_t \right] = \sum_{a=1}^d w_{t,a}^2 \left(\frac{z_{t,a}^2}{w_{t,a}^2} \right) \leq d, \quad (124)$$

since $z_t \in [0, 1]^d$.

- Note that it is crucial that we use local norms here; the generic bound (95) of $\|z_t\|_\infty^2$ is not good enough because $\|z_t\|_\infty$ is unbounded.
- Minimizing the regret bound with respect to η , we get $2\sqrt{d \log(d)T}$ with $\eta = \sqrt{\log(d)/(dT)}$.
- Note that we have only gotten results in expectation (over randomness of the learner $a_t \sim w_t$). To get high probability results, we also need to control the variance of \hat{z}_t . To do this, we modify the EG algorithm by smoothing w_t with the uniform distribution over actions. This results in the standard Exp3 algorithm.

2.12 Online-to-batch conversion (Lecture 5)

- So far, we have been focusing on the **online** setting, where the learner receives one example at a time and is asked to make a prediction on each one. Good online learners have low **regret**.
- Sometimes it is more natural to operate in the **batch** setting, where the learner gets a set of training examples, learns a model, and then is asked to predict on new test examples. Good batch learners have low **generalization error**.

- In this section, we will show that low regret implies low generalization error by explicitly reducing the online setting to the batch setting.
- Batch setting
 - Assume we have a unknown data-generating distribution $p^*(z)$, where we use $z = (x, y) \in \mathcal{Z}$ to denote an input-output pair.
 - Assume our hypothesis class is a convex set of weight vectors $S \subseteq \mathbb{R}^d$.
 - As in online learning, we define a convex loss function $\ell(z, w) \in \mathbb{R}$; for example, the squared loss for linear regression would be $\ell((x, y), w) = (w \cdot x - y)^2$. Assume $\ell(z, w)$ is convex in w for each $z \in \mathcal{Z}$.
 - The **generalization error** of a weight vector $w \in S$ is defined as follows:

$$L(w) = \mathbb{E}_{z \sim p^*}[\ell(z, w)]. \quad (125)$$

- Define the weight vector that minimizes the generalization error:

$$w^* \in \arg \min_{w \in S} L(w). \quad (126)$$

- The batch learner gets T i.i.d. training examples (z_1, \dots, z_T) , where each $z_t \sim p^*$. The goal is to output some estimate $w \in S$ that hopefully has low $L(w)$.

Assuming we have an online learner as a black box, we will construct a batch learner as follows:

- **Algorithm 7 (online-to-batch conversion)**

- Input: T training examples z_1, \dots, z_T .
- Iterate $t = 1, \dots, T$:
 - * Receive w_t from the online learner.
 - * Send loss function $f_t(w) = \ell(z_t, w)$ to the online learner.
- Return the average of the weight vectors: $\bar{w} = \frac{1}{T} \sum_{t=1}^T w_t$.

- Remarks about randomness

- In contrast to the online learning (adversarial) setting, many of the quantities we are working with now have distributions associated with them.
- For example, f_t is a random function that depends on z_t .
- Each w_t is a random variable which depends on (i.e., is in the sigma-algebra of) $z_{1:t-1}$.
- Note that $L(w^*)$ is not random.

- **Theorem 6 (Online-to-batch conversion)**

- Recall the usual definition of regret (which depends on $z_{1:T}$):

$$\text{Regret}(u) = \sum_{t=1}^T [f_t(w_t) - f_t(u)]. \quad (127)$$

- Online-to-batch conversion obtains the following expected generalization error:

$$\boxed{\mathbb{E}[L(\bar{w})] \leq L(w^*) + \frac{\mathbb{E}[\text{Regret}(w^*)]}{T}}. \quad (128)$$

- Interpretation: the expected generalization error of the online-to-batch conversion $L(\bar{w})$ is bounded by the best possible generalization error (attained by $w^* \in S$) plus the online regret.
- Note that $\mathbb{E}[L(\bar{w})]$ has two expectations here: the $\mathbb{E}[\cdot]$ is an expectation over possible training datasets, and L contains an expectation over test examples.
- Proof:
 - The first key insight is that $f_t(w_t)$ provides an unbiased estimate of the generalization error of w_t .

$$\mathbb{E}[f_t(w_t) \mid w_t] = L(w_t). \quad (129)$$

This is true because all the examples are independent, so w_t (deterministic function of $z_{1:t-1}$) is independent of f_t (deterministic function of z_t).

- The second key insight is that averaging can only reduce loss. Since $\ell(z, \cdot)$ is convex, and an average of convex functions is convex, L is convex. By Jensen's inequality:

$$L(\bar{w}) \leq \frac{1}{T} \sum_{t=1}^T L(w_t). \quad (130)$$

- Now, the rest is just putting the pieces together. Putting (130) and (129) together:

$$L(\bar{w}) \leq \frac{1}{T} \sum_{t=1}^T \mathbb{E}[f_t(w_t) \mid w_t]. \quad (131)$$

Adding and subtracting $L(w^*)$ to the RHS, noting that $L(w^*) = \mathbb{E}[f_t(w^*)]$:

$$L(\bar{w}) \leq L(w^*) + \frac{1}{T} \sum_{t=1}^T (\mathbb{E}[f_t(w_t) \mid w_t] - \mathbb{E}[f_t(w^*)]). \quad (132)$$

- Taking expectations on both sides:

$$\mathbb{E}[L(\bar{w})] \leq L(w^*) + \frac{1}{T} \sum_{t=1}^T (\mathbb{E}[f_t(w_t) - f_t(w^*)]). \quad (133)$$

- The second term of the RHS is upper bounded by the regret, so we have:

$$\mathbb{E}[L(\bar{w})] \leq L(w^*) + \frac{\mathbb{E}[\text{Regret}(w^*)]}{T}. \quad (134)$$

- Remarks

- If you run online subgradient descent once over your training data,⁶ you should expect $O(\sqrt{T})$ regret by our previous analyses.
- The resulting average weight vector will, *in expectation*,⁷ have a generalization error which is within $O(1/\sqrt{T})$ of the best weight vector.
- If the batch learner returns the last weight vector w_{T+1} , then our analysis above doesn't apply.
- In general, averaging is a useful concept for stabilizing learning algorithms.

2.13 Summary (Lecture 5)

- This concludes our tour of online learning.
- Our measure of success is getting low **regret**, the difference between the learner's cumulative losses and the best *fixed* expert's cumulative losses. In particular, we hope for sublinear regret: $\text{Regret} = o(T)$.
- Without no additional assumptions (even with two experts), any deterministic algorithm must have $\text{Regret} = \Omega(T)$ (fail).
- In the realizable setting (some expert is perfect), the majority algorithm achieves $O(\log d)$ regret (constant in T).
- We started with the **follow the leader (FTL)**, and saw that it worked for quadratic loss functions ($\text{Regret} = O(\log T)$), but failed for linear loss functions ($\text{Regret} = \Omega(T)$).
- Inspecting the regret bounds reveals that in order to get low regret, we need to have a *stable* learner, in the sense that w_t and w_{t+1} should be close (according to some notion of proximity).

⁶In practice, it helps to run the online learning algorithm multiple times over the training data.

⁷You can get high probability bounds too, but we won't discuss those here.

- This motivated us to look at **follow the regularized leader (FTRL)**, where we add a quadratic regularizer, which gave us a regret bound with two terms: (i) a bias-like term (value of regularizer applied to the best expert) and (ii) a variance-like term (sum of the norm of the subgradients). Balancing the two by controlling the amount of regularization (inverse step size) yields $\text{Regret} = O(\sqrt{T})$.
- If our loss functions were non-linear, we could linearize using the subgradient. Coupled with a quadratic regularizer, we get the **online subgradient descent (OGD)** algorithm. We also showed that it suffices to analyze linear functions, since they result in the most regret.
- We looked at learning with expert advice, and got regret bounds of $O(\sqrt{dT})$, where d is the number of experts. Inspired by the logarithmic dependence on d in the majority algorithm, this motivated us to consider different regularizers.
- The general algorithm that deals with different regularizers is **online mirror descent (OMD)**. We analyzed this algorithm using Fenchel duality and Bregman divergences, which allowed us to look at arbitrary regularizers through the lens of the mapping between w_t and θ_t . Specializing to learning with expert advice, we get a $O(\sqrt{\log(d)T})$ regret bound with the entropic regularizer, resulting in the **exponentiated gradient (EG)** algorithm.
- By exploiting properties of exponentiated gradient, we can perform a refined analysis again using **local norms** to achieve stronger results, matching the $O(\log d)$ regret in the realizable setting.
- In the bandit setting with partial feedback, we get $O(\sqrt{d \log(d)T})$ regret again using local norms to control the size of the now unbounded subgradients.
- Finally, we showed that learners with low regret in the online setting directly lead to learners with low **generalization error** in the batch setting via an online-to-batch conversion.

2.14 References

- [Shalev-Shwartz, 2012: Online Learning and Online Convex Optimization \(survey paper\)](#)
 - This is a nice introduction to online learning, on which much of these online learning notes are based.

3 Uniform convergence

3.1 Motivation (Lecture 6)

- In online learning, training and testing were intertwined: the learner receives and is evaluated on a fresh example, but then is able to immediately train on that new example before proceeding to the next example.
- In the batch setting (which is the more standard one in machine learning), the training and testing phases are distinct: the learner gets a training set and outputs a hypothesis (from some fixed hypothesis class) to be judged on how well it generalizes to fresh test examples.
- As a learning algorithm gets increasingly more training data, one would expect it to perform better, and maybe eventually even do as well the best hypothesis in our hypothesis class. How do we formalize this?
- Our main object of study is the **empirical risk minimizer**, which corresponds to the hypothesis from a class with the lowest training error. This is the batch learning analogue of the follow the leader (FTL) algorithm from online learning.
- We measure the quality of a hypothesis by its **generalization error**, the expected loss of a hypothesis on a new test example. In online learning, we already saw that we could bound the *expected* generalization error of online gradient descent using online-to-batch conversion. This section will allow us to derive high probability bounds as well as more general results for hypothesis classes without relying on convexity.
- The generalization error of the empirical risk minimizer can be quite a complicated random variable depending on the training data. As we'll see later, our approach for studying generalization error is based on **uniform convergence**. In the process, we will develop some fairly general machinery from probability theory; these tools are more broadly applicable outside machine learning. These techniques are mathematically quite involved, so make sure you have a good understanding of probability!
- In summary, the central question is:

Why does minimizing training error reduce test error?

The answer is not obvious for the training error and test error are two separate quantities which can in general be arbitrarily far apart. This deep question is at the core of statistical learning theory, and answering it reveals what it means to learn.

3.2 Formal setup (Lecture 6)

- In this section, we formalize the (batch) supervised learning setting. Much of what we will do also works for unsupervised learning, but we will describe it in the context of supervised learning to provide intuition.
- Consider the problem of predicting an output $y \in \mathcal{Y}$ given an input $x \in \mathcal{X}$. Example: $\mathcal{X} = \mathbb{R}^d$, $\mathcal{Y} = \{-1, +1\}$.
- Let \mathcal{H} be a set of **hypotheses** (i.e. experts). Usually, each $h \in \mathcal{H}$ maps \mathcal{X} to \mathcal{Y} . Example: $\mathcal{H} = \{x \mapsto \text{sign}(w \cdot x) : w \in \mathbb{R}^d\}$ is all thresholded linear functions.
- Let $\ell : (\mathcal{X} \times \mathcal{Y}) \times \mathcal{H} \rightarrow \mathbb{R}$ be a **loss function**. Example: $\ell((x, y), h) = \mathbb{I}[y \neq h(x)]$ is the zero-one loss.
- Let p^* denote the true underlying data-generating distribution over input-output pairs $\mathcal{X} \times \mathcal{Y}$. This underlying distribution is the crucial piece that distinguishes batch learning from the online learning, in which we assumed that data could be generated adversarially.
- **Definition 7 (generalization error (expected risk))**

- Let $L(h)$ be the **generalization error** (also known as the **expected risk**) of a hypothesis $h \in \mathcal{H}$, which is the loss that h incurs on a new test example (x, y) in expectation:

$$L(h) \stackrel{\text{def}}{=} \mathbb{E}_{(x,y) \sim p^*} [\ell((x, y), h)]. \quad (135)$$

Getting low generalization error is in some sense the definition of successful learning.

- Define an **expected risk minimizer** h^* to be any hypothesis that minimizes the expected risk:

$$h^* \in \arg \min_{h \in \mathcal{H}} L(h). \quad (136)$$

This is the thing that we can only aspire to. $L(h^*)$ is the lowest possible generalization error (which might be large if the learning problem is noisy or your hypothesis class is too small).

- To do learning, we are given n **training examples**, which are a set of input-output pairs:

$$(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)}), \quad (137)$$

where each $(x^{(i)}, y^{(i)})$ is drawn **i.i.d.** from p^* .

- Note: the training and test distributions are the same. While this assumption often doesn't hold exactly in practice, the training and test distributions morally have to be related somehow, or else there's no hope that what we learned from training would be useful at test time.⁸
- Note: the i.i.d. assumption, which also doesn't hold exactly in practice, ensures that more training data gives us more information (or else we would get the same training example over and over again, which would be useless).⁹

• **Definition 8 (training error (empirical risk))**

- Let $\hat{L}(h)$ be the **training error** (also known as **empirical risk**) of a hypothesis $h \in \mathcal{H}$ as the average loss over the training examples:

$$\hat{L}(h) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n \ell((x^{(i)}, y^{(i)}), h). \quad (138)$$

Note that for a fixed h , $\hat{L}(h)$ is just an empirical average with mean $L(h)$. This is key.

- Define an **empirical risk minimizer** (ERM) be any hypothesis that minimizes the empirical risk:

$$\hat{h} \in \arg \min_{h \in \mathcal{H}} \hat{L}(h). \quad (139)$$

- Let us pause a moment to remember what are the random variables and what the independence assumptions are: \hat{h} is a random variable that depends the training examples (in a rather complicated way), but h^* is non-random. The training error $\hat{L}(h)$ is a random variable for each h , but the generalization error $L(h)$ is non-random.
- Recall that we are interested in the generalization error of the ERM:

$$L(\hat{h}). \quad (140)$$

As in the online learning setting, we will not study this quantity directly, but rather look at its difference with a baseline. There are two questions we can ask:

- How does training and generalization error relate?

$$\underbrace{L(\hat{h})}_{\text{generalization error of ERM}} - \underbrace{\hat{L}(\hat{h})}_{\text{training error of ERM}}. \quad (141)$$

⁸ If the two distributions are different but still related somehow, not all hope is lost; dealing with this discrepancy is called domain adaptation.

⁹ Pure i.i.d. is not necessary, but certainly some amount of independence is necessary.

- How well is ERM doing with respect to the best in the hypothesis class?

$$\underbrace{L(\hat{h})}_{\text{generalization error of ERM}} - \underbrace{L(h^*)}_{\text{lowest generalization error}}. \quad (142)$$

This is known as the **excess risk**, which is the batch analogue of the regret in the online learning setting.

- How do we analyze the excess risk? The excess risk is a random quantity that depends on the training examples (through \hat{h}). It is possible that the excess risk is high even for large n (for instance, if we just happened to see the same example over and over again). So we can't deterministically bound the excess risk.
- Fortunately, we can show that bad outcomes are not too likely. We will prove bounds of the following flavor: With probability at least $1 - \delta$, the excess risk is upper bounded by some ϵ ($L(\hat{h}) - L(h^*) \leq \epsilon$), where ϵ is generally a function that depends on δ (and other aspects of the learning problem). More formally, the types of statements we'd like to show can be written compactly as:

$$\boxed{\mathbb{P}[L(\hat{h}) - L(h^*) > \epsilon] \leq \delta.} \quad (143)$$

Note that the randomness in the probability is over draws of the n training examples: $(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)}) \sim p^*$.

- It is important to note that there are two sources of randomness at play here:
 - Generalization error (upper bounded by ϵ) is defined with respect to randomness over the test example.
 - Confidence (lower bounded by $1 - \delta$) is defined with respect to randomness over the training examples.
- Probably Approximately Correct (PAC) framework [Leslie Valiant, 1984]
 - A learning algorithm \mathcal{A} PAC learns \mathcal{H} if for any distribution p^* over $\mathcal{X} \times \mathcal{Y}$, $\epsilon > 0$, $\delta > 0$, \mathcal{A} (which takes as input n training examples along with ϵ and δ), returns $\hat{h} \in \mathcal{H}$ such that with probability at least $1 - \delta$, $L(\hat{h}) - L(h^*) \leq \epsilon$, and \mathcal{A} runs in $\text{poly}(n, \text{size}(x), 1/\epsilon, 1/\delta)$ time.
 - Remark: time complexity upper bounds sample complexity, because you have to go through all the data points. In this class, we will not focus so much on the computational aspect in our presentation, but just work with the ERM, and assume that it can be optimized efficiently (even though that's not true for general non-convex losses).
 - The ideas in PAC learning actually predate Valiant and were studied in the field of empirical process theory, but Valiant and others focused more on more combinatorial problems with computation in mind.

3.3 Realizable finite hypothesis classes (Lecture 6)

- So far, we have presented a framework that is very general, and in fact, so general, that we can't say anything. So we need to make some assumptions. Ideally, the assumptions would be both realistic (not too strong) but still allow us to prove something interesting (not too weak).
- In this section, we will start with an easy case, where we have a finite number of hypotheses, at least one of which has zero generalization error (the outcome will be analogous to the analysis of the majority algorithm in the online learning setting). These assumptions are as formalized as follows:

- **Assumption 2 (finite hypothesis space)**

- Assume \mathcal{H} is finite.

- **Assumption 3 (realizable)**

- Assume there exists a hypothesis $h^* \in \mathcal{H}$ that obtains zero generalization error, that is:

$$L(h^*) = \mathbb{E}_{(x,y) \sim p^*} [\ell((x,y), h^*)] = 0. \quad (144)$$

- **Theorem 7 (realizable finite hypothesis class)**

- Let \mathcal{H} be a hypothesis class, where each hypothesis $h \in \mathcal{H}$ maps some \mathcal{X} to \mathcal{Y} .
- Let ℓ be the zero-one loss: $\ell((x,y), h) = \mathbb{I}[y \neq h(x)]$.
- Let p^* be any distribution over $\mathcal{X} \times \mathcal{Y}$.
- Assume Assumptions 2 and 3 hold.
- Let \hat{h} be the empirical risk minimizer.
- Then the following two equivalent statements hold (each has a different interpretation depending on what we're interested in):
 - * Interpretation 1: what is the error after training on n examples (**generalization error**)? Answer: with probability at least $1 - \delta$,

$$\boxed{L(\hat{h}) \leq \frac{\log |\mathcal{H}| + \log(1/\delta)}{n}} \quad (145)$$

Usually, think of $\log(1/\delta)$ as a constant (e.g., $\delta = 0.01$, then $\log(1/\delta) \cong 4.6$), so the

$$\underbrace{L(\hat{h})}_{\text{generalization error}} = O \left(\frac{\overbrace{\log |\mathcal{H}|}^{\text{complexity}}}{\underbrace{n}_{\text{number of training examples}}} \right) \quad (146)$$

- * Interpretation 2: how many examples n (**sample complexity**) do I need to obtain generalization error at most ϵ with confidence at least $1 - \delta$? Answer: With probability at least $1 - \delta$:

$$\boxed{n \geq \frac{\log |\mathcal{H}| + \log(1/\delta)}{\epsilon} \Rightarrow L(\hat{h}) \leq \epsilon.} \quad (147)$$

- Remarks

- Statisticians generally talk about error, and computer scientists like to talk about sample complexity. But they’re just two sides of the same coin.
- In this case, the excess risk behaves as $O(1/n)$, which is known as a “fast” rate (analogous to $O(1/T)$ average regret). This is because we’ve assumed realizability: as soon as h makes even a single mistake on a training example, we can throw it away.
- In online-to-batch conversions, we only proved generalization bounds for convex losses in expectation; here, we are getting a **high probability** bound, which is much stronger.
- The excess risk only grows logarithmically with $|\mathcal{H}|$, so we can use pretty big hypothesis classes.
- Note that our result is independent of $p^*(x, y)$. This is known as a **distribution-free** result, which is great, because typically we don’t know what p^* is.

- Proof of Theorem 7

- FIGURE: [a row of hypotheses, sorted by increasing $L(h)$]
- We’d like to upper bound the probability of the bad event that $L(\hat{h}) > \epsilon$.
- Let $B \subseteq \mathcal{H}$ be the set of bad hypotheses h : $B = \{h \in \mathcal{H} : L(h) > \epsilon\}$. We can rewrite our goal as upper bounding the probability of selecting a bad hypothesis:

$$\mathbb{P}[L(\hat{h}) > \epsilon] = \mathbb{P}[\hat{h} \in B]. \quad (148)$$

- Recall that the empirical risk of the ERM is always zero ($\hat{L}(\hat{h}) = 0$) because at least $\hat{L}(h^*) = L(h^*) = 0$. So if we selected a bad hypothesis ($\hat{h} \in B$), then some bad hypothesis must have zero empirical risk:

$$\mathbb{P}[\hat{h} \in B] \leq \mathbb{P}[\exists h \in B : \hat{L}(h) = 0]. \quad (149)$$

- We now get to the heart of the argument, which consists of two steps.
- Step 1: bound $\mathbb{P}[\hat{L}(h) = 0]$ for a **fixed** $h \in B$.
 - * On each example, hypothesis h does not err with probability $1 - L(h)$.

- * Since the training examples are i.i.d. and the fact that $L(h) > \epsilon$ for $h \in B$:

$$\mathbb{P}[\hat{L}(h) = 0] = (1 - L(h))^n \leq (1 - \epsilon)^n \leq e^{-\epsilon n}, \quad (150)$$

where the last step follows since $1 - a \leq e^{-a}$.

- * Remark: this probability **decreases exponentially** with n , which is important.
- Step 2: show that step 1 holds simultaneously for all $h \in B$:
 - * We apply the **union bound** to bound the probability of the event for any $h \in B$:

$$\mathbb{P}[\exists h \in B : \hat{L}(h) = 0] \leq \sum_{h \in B} \mathbb{P}[\hat{L}(h) = 0]. \quad (151)$$

- * The rest is straightforward:

$$\sum_{h \in B} \mathbb{P}[\hat{L}(h) = 0] \leq |B|e^{-\epsilon n} \quad (152)$$

$$\leq |\mathcal{H}|e^{-\epsilon n} \quad (153)$$

$$\stackrel{\text{def}}{=} \delta. \quad (154)$$

- Taking logs of the last equality and rearranging:

$$\epsilon = \frac{\log |\mathcal{H}| + \log(1/\delta)}{n}. \quad (155)$$

The theorem follows by substituting this expression for δ .

3.4 Generalization bounds via uniform convergence (Lecture 6)

- The proof of Theorem 7 is elementary but illustrates an important pattern that will recur again in more complex scenarios. At a high level, we are interested in expected risk L , but only have access to empirical risk \hat{L} to choose the ERM \hat{h} . In the proof, we saw two steps:
 - Step 1 (convergence): For a **fixed** h , show that $\hat{L}(h)$ is close to $L(h)$ with high probability. In the above proof, this meant showing that if $L(h) > \epsilon$, then $\hat{L}(h) = 0$ is unlikely.
 - Step 2 (uniform convergence): Show that the above holds simultaneously for **all** hypotheses $h \in \mathcal{H}$. In the above proof, this meant using a union bound.

The difference between convergence and uniform convergence is absolutely crucial. It is important to note that \hat{h} is a random variable that depends on the training examples, so $\hat{L}(\hat{h})$ is not just a sum of i.i.d. variables, so step 1 does not apply directly.

- Theorem 7 also made two restrictive assumptions.
 - First, there exists a perfect hypothesis (realizability). What happens when the problem is not realizable (all hypotheses make some error)? To answer this, we consider the general problem of convergence of random variables using **concentration inequalities**.
 - Second, the hypothesis class is finite. What happens when the number of hypotheses is infinite? We can't just apply a union bound any more. To answer this, we need to have more suitable ways of measuring the “size” of a set other than cardinality. This leads to **Rademacher complexity**, **VC dimension**, and **covering numbers** as ways for measuring the “size” of an infinite set from the point of view of the loss function.
- Breaking free of these restrictive assumptions, we will show how bounding generalization error can be reduced to one of uniform convergence. Recall that our goal is to bound the excess risk, the amount by which ERM's generalization error exceeds the lowest possible generalization error:

$$\mathbb{P}[L(\hat{h}) - L(h^*) \geq \epsilon] \leq \delta. \quad (156)$$

Note that the difference between \geq and $>$ isn't important here, and it will be more convenient to use \geq .

- FIGURE: [plot \hat{L} and L]
- Let us first relate generalization errors to training errors, since that's what the ERM is defined in terms of:

$$L(\hat{h}) - L(h^*) = \underbrace{[L(\hat{h}) - \hat{L}(\hat{h})]}_{\text{concentration}} + \underbrace{[\hat{L}(\hat{h}) - \hat{L}(h^*)]}_{\leq 0} + \underbrace{[\hat{L}(h^*) - L(h^*)]}_{\text{concentration}}. \quad (157)$$

- The second term is non-positive by definition of the empirical risk minimizer.
- The third term involves a comparison of $\hat{L}(h^*)$ and $L(h^*)$. If we expand things, we realize that this is just a question of the difference between an average of n i.i.d. random variables and its mean:

$$\hat{L}(h^*) = \frac{1}{n} \sum_{i=1}^n \ell((x^{(i)}, y^{(i)}), h^*), \quad L(h^*) = \mathbb{E}_{(x,y) \sim p^*} [\ell((x, y), h^*)]. \quad (158)$$

We'll see how concentration inequalities can be used to control this difference.

- What about the first term? The same reasoning doesn't apply because \hat{h} depends on the training examples, and so $\hat{L}(\hat{h})$ is not a sum of i.i.d. random variables! We'll need something something more sophisticated: uniform convergence. Suppose we could ensure that $\hat{L}(h)$ and $L(h)$ were close (say within $\frac{\epsilon}{2}$) for all $h \in \mathcal{H}$. Then, we could be ensure that $\hat{L}(\hat{h})$ and $L(\hat{h})$ were within $\frac{\epsilon}{2}$, as well as $\hat{L}(h^*)$ and $L(h^*)$.

- The contrapositive can be written formally as:

$$\mathbb{P}[L(\hat{h}) - L(h^*) \geq \epsilon] \leq \mathbb{P}\left[\sup_{h \in \mathcal{H}} |L(h) - \hat{L}(h)| \geq \frac{\epsilon}{2}\right]. \quad (159)$$

On the LHS is a statement about excess risk, and on the RHS is a statement about uniform convergence. The RHS is the probability of the event that the largest difference between the empirical and expected risk is at least $\frac{\epsilon}{2}$, or equivalently, the event that this difference exceeds $\frac{\epsilon}{2}$ for at least one $h \in \mathcal{H}$.

- Note: the classic example of uniform convergence is the Glivenko-Cantelli theorem (also called the uniform law of large numbers), for estimating distribution functions. Given x_1, \dots, x_n drawn i.i.d. from some distribution with cumulative distribution function (CDF) $F(x)$, we can form the empirical CDF $F_n(x) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}[x \leq x_i]$. One can ask for the convergence of the CDF function in the uniform norm:

$$\|F_n - F\|_\infty \stackrel{\text{def}}{=} \sup_{x \in \mathbb{R}} |F_n(x) - F(x)| \xrightarrow{P} 0. \quad (160)$$

What we will be studying is a generalization of (160), where we have arbitrary hypotheses $h \in \mathcal{H}$ rather than $x \in \mathbb{R}$.

- Note: if we look at the difference between \hat{L} and L , we can construct something called an **empirical process**:

$$\{G_n(h)\}_{h \in \mathcal{H}}, \quad G_n \stackrel{\text{def}}{=} \sqrt{n}(\hat{L}(h) - L(h)), \quad (161)$$

which is a stochastic process (collection of random variables indexed by $h \in \mathcal{H}$). Empirical process theory focuses on studying empirical processes. We know that for a given $h \in \mathcal{H}$, $G_n(h)$ converges to a normal distribution by the central limit theorem. We can think about G_n converging to a Gaussian process G with covariance function

$$\text{Cov}(G(h), G(h')) = \mathbb{E}[\ell(z, h)\ell(z, h')]. \quad (162)$$

This stochastic process viewpoint allows one to talk not just about the supremum $\sup_{h \in \mathcal{H}} G_n(h)$, but also get distributional results, which is useful for computing confidence intervals. This is outside the scope of the class; for additional information, Pollard has an excellent book on this.

3.5 Concentration inequalities (Lecture 7)

- Concentration inequalities are a very powerful set of techniques from probability theory that shows that an appropriate combination of independent random variables will *concentrate* around its expectation. From the point of view of learning theory, the random variables of interest are the losses of hypotheses on training examples.

- Mean estimation

- Let X_1, \dots, X_n be i.i.d. real-valued random variables with mean $\mu \stackrel{\text{def}}{=} \mathbb{E}[X_1]$
- Define the empirical mean as follows:

$$\hat{\mu}_n \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n X_i \quad (163)$$

- Question: how does $\hat{\mu}_n$ relate to μ ?
- Examples
 - * X_i is the height of the i -th person sampled from some population.
 - * X_i is the loss of a *fixed* hypothesis $h \in \mathcal{H}$ on the i -th example.

- FIGURE: $[\mu \text{ with distribution over } \hat{\mu}_n]$

- Types of statements

- **Consistency**: by the law of large numbers,

$$\hat{\mu}_n - \mu \xrightarrow{P} 0, \quad (164)$$

where \xrightarrow{P} denotes convergence in probability.¹⁰ Consistency assures us that as we get more data ($n \rightarrow \infty$), we will approach the correct answer, but it doesn't tell us how quickly.

- **Asymptotic normality**: Letting $\text{Var}[X_1] = \sigma^2$, by the central limit theorem,

$$\sqrt{n}(\hat{\mu}_n - \mu) \xrightarrow{d} \mathcal{N}(0, \sigma^2), \quad (165)$$

where \xrightarrow{d} denotes convergence in distribution.¹¹ Asymptotic normality says that if n is large enough, then $\hat{\mu}_n - \mu$ behaves as $\frac{\sigma}{\sqrt{n}}$, where the variance is decreasing at a rate of $1/n$. But this result is only asymptotic, it doesn't tell us anything precise for a particular value of n (say, $n = 10$).

¹⁰Convergence in probability: For each $\epsilon > 0$, $\mathbb{P}[|\hat{\mu}_n - \mu| \geq \epsilon] \rightarrow 0$ as $n \rightarrow \infty$.

¹¹Convergence in distribution: For each t , $\mathbb{P}[\frac{\sqrt{n}(\hat{\mu}_n - \mu)}{\sigma} \leq t] \rightarrow \Phi(t)$ as $n \rightarrow \infty$, where Φ is the cumulative distribution of the standard Gaussian distribution.

- **Tail bounds:** Ideally, we want a statement of the following form:

$$\mathbb{P}[|\hat{\mu}_n - \mu| \geq \epsilon] \leq \text{SomeFunction}(n, \epsilon) = \delta. \quad (166)$$

Based on the Gaussian approximation, we expect that the bounding function on the RHS would decay double exponentially in ϵ and exponentially in n . We shall see shortly that this intuition is indeed true. In the context of learning theory, ϵ would be the bound on the difference between empirical and expected risks, and $1 - \delta$ would be the confidence.

- Note: of course, as we sweep ϵ from 0 to ∞ and look at how much probability mass is past ϵ , we get a complete picture of the full distribution. However, typically tail bounds are simple upper bounds which are often loose and only become more reliable for small ϵ .
- Our starting point is Markov's inequality, a very simple tool that allows us to control the deviation of a non-negative random variable from its mean using the expectation of that random variable. In short, it turns expectations (which are easier to work with) into tail probabilities (what we want).
- **Theorem 8 (Markov's inequality)**

- Let $Z \geq 0$ be a random variable.
- Then

$$\mathbb{P}[Z \geq t] \leq \frac{\mathbb{E}[Z]}{t}. \quad (167)$$

- Proof:
 - Since Z is non-negative, we have $t\mathbb{I}[Z \geq t] \leq Z$.
 - Taking expectations on both sides and rearranging completes the proof.
- Remarks
 - We can apply $Z = (X - \mu)^2$ (second moment) and $t = \epsilon^2$ to obtain **Chebyshev's inequality**:

$$\mathbb{P}[|X - \mu| \geq \epsilon] \leq \frac{\text{Var}(X)}{\epsilon^2}. \quad (168)$$

Applying the inequality to the average over i.i.d. variables ($\hat{\mu}_n = \frac{1}{n} \sum_{i=1}^n X_i$), then $\text{Var}(\hat{\mu}_n) = \frac{\text{Var}(X_1)}{n}$. This is a very weak result, because the tail probability is decaying only at polynomial rate ($1/n$).

- To get stronger bounds, we need to apply Markov's inequality on higher order moments. In particular, we will look at all moments by considering $Z = e^{tX}$, where t is a free parameter we will use later to optimize the bound. "All the moments" is captured by the moment generating function:

- **Definition 9 (moment generating function)**

– For a random variable X , the moment generating function (MGF) of X is:

$$\boxed{M_X(t) \stackrel{\text{def}}{=} \mathbb{E}[e^{tX}].} \quad (169)$$

- One useful way to think about the MGF is in terms of its Taylor expansion:

$$M_X(t) = 1 + t \mathbb{E}[X] + \frac{t^2}{2} \mathbb{E}[X^2] + \frac{t^3}{6} \mathbb{E}[X^3] + \dots \quad (170)$$

- The moment generating function receives its name because the k -th derivative evaluated at $t = 0$ yield the k -th moment (assuming we can swap integration and differentiation):

$$\left. \frac{d^k M_X(t)}{dt^k} \right|_{t=0} = \mathbb{E}[X^k]. \quad (171)$$

- One important property is that the MGF of a sum of independent random variables is simply the product of the MGFs. If X_1 and X_2 are independent random variables, then we have:

$$M_{X_1+X_2}(t) = M_{X_1}(t)M_{X_2}(t). \quad (172)$$

The distribution over $X_1 + X_2$ can be computed using a convolution, which is typically cumbersome, but MGFs (like Fourier transforms) turn convolutions into products.

- Applying Markov's inequality to $Z = e^{tX}$, we get that

$$\mathbb{P}[X \geq \epsilon] \leq \frac{M_X(t)}{e^{t\epsilon}} \text{ for all } t > 0. \quad (173)$$

We can apply this to the case of sample means ($X = \hat{\mu}_n$) by computing $\mathbb{P}[\hat{\mu}_n \geq \epsilon] = \mathbb{P}[X_1 + \dots + X_n \geq n\epsilon]$. We get that all $t > 0$:

$$\mathbb{P}[\hat{\mu}_n \geq \epsilon] \leq \left(\frac{M_{X_1}(t)}{e^{t\epsilon}} \right)^n. \quad (174)$$

- Provided that $\frac{M_{X_1}(t)}{e^{t\epsilon}} < 1$, getting n independent samples means that our tail probability will decrease exponentially. This is key.
- Why should we expect this to happen? If $\mathbb{E}[X_1] = 0$, then the dominant term in the Taylor expansion of $M_{X_1}(t)$ (170) is $O(t^2)$, which grows slower than $O(t)$ (what is in the denominator of (174)), so we can always choose t small enough such that the ratio is strictly less than 1.

- Note that $M_{X_1}(t)$ could be infinite for some t , in which case the bounds are vacuous. In this class, we will work with distributions X such that $M_{X_1}(t) < \infty$ for all $t > 0$.
- Now let's actually compute the MGF of some probability distributions of interest. The Gaussian distribution is a natural place to start, as we will see.
- **Example 17 (MGF of Gaussian variables)**
 - Let $X \sim \mathcal{N}(0, \sigma^2)$.
 - Then $M_X(t) = e^{\sigma^2 t^2 / 2}$.
 - Derivation (by completing the square):

$$M_X(t) = \mathbb{E}[e^{tX}] = \int (2\pi\sigma^2)^{-\frac{1}{2}} \exp\left(\frac{x^2 - 2\sigma^2 tx}{-2\sigma^2}\right) dx \quad (175)$$

$$= \int (2\pi\sigma^2)^{-\frac{1}{2}} \exp\left(\frac{(x - \sigma^2 t)^2 - \sigma^4 t^2}{-2\sigma^2}\right) dx \quad (176)$$

$$= \exp\left(\frac{\sigma^2 t^2}{2}\right). \quad (177)$$

- **Lemma 3 (Tail bound for Gaussian variables)**
 - Having control on the Gaussian MGF, we can now derive a tail bound by plugging the form of the MGF into (173). This yields:

$$\mathbb{P}[X \geq \epsilon] \leq \inf_t \exp\left(\frac{\sigma^2 t^2}{2} - t\epsilon\right). \quad (178)$$

The infimum on the RHS is attained by setting $t = \epsilon/\sigma^2$, yielding:

$$\mathbb{P}[X \geq \epsilon] \leq \exp\left(\frac{-\epsilon^2}{2\sigma^2}\right). \quad (179)$$

- What about non-Gaussian variables? Note that the bounds would still hold if we replaced $M_X(t)$ with an upper bound. This motivates the following definition:
- **Definition 10 (sub-Gaussian)**

- A mean-zero random variable X is **sub-Gaussian** with parameter σ^2 if its moment generating function is bounded as follows:

$$M_X(t) \leq \exp\left(\frac{\sigma^2 t^2}{2}\right). \quad (180)$$

- It follows immediately by analogy with (179) that:

$$\mathbb{P}[X \geq \epsilon] \leq \exp\left(\frac{-\epsilon^2}{2\sigma^2}\right). \quad (181)$$

- Examples

- Gaussian random variables: If $X \sim \mathcal{N}(0, \sigma^2)$, then X is sub-Gaussian with parameter σ^2 (trivial). Note that the sub-Gaussian parameter and the variance coincide in this case, but this is not true in general.
- Bounded random variables (**Hoeffding's lemma**): If $a \leq X \leq b$ with probability 1 and $\mathbb{E}[X] = 0$, then X is sub-Gaussian with parameter $(b - a)^2/4$.
 - * Intuition (not a proof): in the Gaussian case, the sub-Gaussian parameter was the variance. Suppose $a = -b$. Then the variance of X is $b^2 = (b - a)^2/4$, the sub-Gaussian parameter given by the lemma. In general, the variance is at most the sub-Gaussian parameter.
- Non-examples: exponential and Gamma variables, since these distributions have tails which are too fat, decaying exponentially rather than double exponentially.¹²

- Properties

- Sum: If X_1 and X_2 are independent sub-Gaussian variables with parameters σ_1^2 and σ_2^2 , respectively, then $X_1 + X_2$ is sub-Gaussian with parameter $\sigma_1^2 + \sigma_2^2$.
- Multiplication by a constant: if X is sub-Gaussian with parameter σ^2 , then for any $c > 0$, cX is sub-Gaussian with parameter $c^2\sigma^2$.
- Not surprisingly, these properties coincide with those of Gaussians.
- Note that the product of two sub-Gaussian variables is in general not sub-Gaussian.
- Given the machinery thus far, we can easily obtain the following classic tail bound for bounded random variables:
- **Theorem 9 (Hoeffding's inequality)**
 - Let X_1, \dots, X_n be independent random variables.
 - Assume each X_i is bounded: $a_i \leq X_i \leq b_i$.
 - Let $\hat{\mu}_n = \frac{1}{n} \sum_{i=1}^n X_i$ be the sample mean.
 - Then

$$\mathbb{P}[\hat{\mu}_n \geq \mathbb{E}[\hat{\mu}_n] + \epsilon] \leq \exp\left(\frac{-2n^2\epsilon^2}{\sum_{i=1}^n (b_i - a_i)^2}\right). \quad (182)$$

¹²These variables have moment generating functions which are defined not for all t but only small t . These are known as sub-exponential variables.

- Special case ($a_i = -B, b_i = +B$):

$$\mathbb{P}[\hat{\mu}_n \geq \mathbb{E}[\hat{\mu}_n] + \epsilon] \leq \exp\left(\frac{-n\epsilon^2}{2B^2}\right). \quad (183)$$

- Proof of Theorem 9:

- Using compositional properties of sub-Gaussian, $\hat{\mu}_n - \mathbb{E}[\hat{\mu}_n]$ is sub-Gaussian with parameter $\frac{1}{n^2} \sum_{i=1}^n \frac{(b_i - a_i)^2}{4}$.
- Apply (181).

- Summary so far

- We’ve shown that Gaussian and bounded random variables are sub-Gaussian and enjoy sharp tail bounds.
- Furthermore, if we have an average of n independent sub-Gaussian variables, then the bound simply gets powered up by n .

- Next, we will show a generalization of Hoeffding’s inequality, where we want to bound not the average of X_1, \dots, X_n , but any function on X_1, \dots, X_n satisfying an appropriate bounded differences condition. After all, learning algorithms do more complex things than taking averages.

- Theorem 10 (**McDiarmid’s inequality** (bounded differences inequality))

- Let f be a function satisfying the following bounded differences condition:

$$|f(x_1, \dots, x_i, \dots, x_n) - f(x_1, \dots, x'_i, \dots, x_n)| \leq c_i \quad (184)$$

for all $i = 1, \dots, n$ and all x_1, \dots, x_n, x'_i . Intuition: modifying one coordinate doesn’t change the function value too much.

- Let X_1, \dots, X_n be independent random variables.
- Then

$$\mathbb{P}[f(X_1, \dots, X_n) - \mathbb{E}[f(X_1, \dots, X_n)] \geq \epsilon] \leq \exp\left(\frac{-2\epsilon^2}{\sum_{i=1}^n c_i^2}\right). \quad (185)$$

- Remarks

- McDiarmid’s inequality generalizes Hoeffding’s inequality as follows. Define the function $f(x_1, \dots, x_n) = \frac{1}{n} \sum_{i=1}^n x_i$, where each x_i satisfies $a_i \leq x_i \leq b_i$. Then f satisfies the bounded differences condition with $c_i = \frac{1}{n}(b_i - a_i)$. Substituting this value of c_i recovers Hoeffding’s inequality exactly.

- This result is quite powerful, as it holds for any independent random variables, and f could be quite complex (e.g., fitting a neural network). As long as the function isn't too sensitive to perturbations in one of its arguments, we get good concentration.
- The proof is not difficult, but it requires introducing martingales, which generalize partial sums of independent random variables.

- **Definition 11 (martingale)**

- A sequence of random variables Z_0, Z_1, \dots, Z_n is a **martingale sequence** with respect to another sequence of random variables X_1, \dots, X_n iff Z_i is a function of $X_{1:i}$, $\mathbb{E}[|Z_i|] < \infty$, and

$$\mathbb{E}[Z_i \mid X_{1:i-1}] = Z_{i-1}. \quad (186)$$

- Define $D_i \stackrel{\text{def}}{=} Z_i - Z_{i-1}$. We call $D_{1:n}$ a **martingale difference sequence** with respect to $X_{1:n}$. Another way to write (186) is

$$\mathbb{E}[D_i \mid X_{1:i-1}] = 0. \quad (187)$$

- **Examples**

- Random walk: $Z_0 = 0$, $Z_i = Z_{i-1} + 1$ with probability $\frac{1}{2}$ and $Z_i = Z_{i-1} - 1$ with probability $\frac{1}{2}$. This is just a sum of i.i.d. random variables.
- Random walk with absorbing state: same as above but $Z_i = Z_{i-1}$ if $Z_{i-1} = 42$.

- **Intuitions**

- Given $X_{1:i-1}$ (the past), the expected value of Z_i is the same as Z_{i-1} (i.e., can't predict the future).
- Think of $D_{1:n}$ as a generalization of an i.i.d. sequence.

- Recall that we showed that sums of independent sub-Gaussian variables are sub-Gaussian. The exact same result holds for martingales, as shown in the following lemma:

- **Lemma 4 (sub-Gaussian martingales)**

- Let Z_0, Z_1, \dots, Z_n be a martingale with respect to X_1, \dots, X_n .
- Suppose that each difference $D_i = Z_i - Z_{i-1}$ is *conditionally* sub-Gaussian with parameter σ_i^2 , that is:

$$\mathbb{E}[e^{tD_i} \mid X_{1:i-1}] \leq \exp(\sigma_i^2 t^2 / 2). \quad (188)$$

– Then $Z_n - Z_0 = \sum_{i=1}^n D_i$ is sub-Gaussian with parameter $\sigma^2 \stackrel{\text{def}}{=} \sum_{i=1}^n \sigma_i^2$.

• Proof of Lemma 4

– This proof is straightforward by induction on n :

$$\mathbb{E}[\exp(t(Z_n - Z_0))] = \mathbb{E}[\exp(tD_n) \exp(t(Z_{n-1} - Z_0))] \quad (189)$$

$$= \mathbb{E}[\mathbb{E}[\exp(tD_n) \exp(t(Z_{n-1} - Z_0)) \mid X_{1:n-1}]] \quad (190)$$

$$\leq \exp(\sigma_n^2 t^2 / 2) \mathbb{E}[\exp(t(Z_{n-1} - Z_0))] \quad (191)$$

$$\leq \exp\left(\sum_{i=1}^n \sigma_i^2 t^2 / 2\right). \quad (192)$$

– The key is that conditioned on $X_{1:i-1}$, D_i is just a sub-Gaussian variable and $Z_{i-1} - Z_0$ is just a constant. The last inequality follows by repeating the argument recursively.

• Proof of Theorem 10 (McDiarmid's inequality)

– We construct a particular type of martingale called **Doob martingale**:

$$Z_i = \mathbb{E}[f(X_1, \dots, X_n) \mid X_{1:i}]. \quad (193)$$

Note the extremes: $Z_0 = \mathbb{E}[f(X_1, \dots, X_n)]$ and $Z_n = f(X_1, \dots, X_n)$. Using this terminology, we are interested in bounding $\mathbb{P}[Z_n - Z_0 \geq \epsilon]$.

– Now let's study $D_i = Z_i - Z_{i-1}$. Both Z_i and Z_{i-1} condition on $X_{1:i-1}$ and take expectations over $X_{i+1:n}$. The only difference is in the treatment of X_i . Therefore we would D_i is contained in some interval of length c_i

– To make this precise, define the lower and upper bounds:

$$L_i = \inf_x \mathbb{E}[f(X_{1:n}) \mid X_{1:i-1}, X_i = x] - \mathbb{E}[f(X_{1:n}) \mid X_{1:i-1}], \quad (194)$$

$$U_i = \sup_x \mathbb{E}[f(X_{1:n}) \mid X_{1:i-1}, X_i = x] - \mathbb{E}[f(X_{1:n}) \mid X_{1:i-1}]. \quad (195)$$

Note that $L_i \leq D_i \leq U_i$.

– Let x_L and x_U correspond to the x 's achieving L_i and U_i , respectively. By the bounded differences assumption,

$$f(X_{1:i-1}, x_U, X_{i+1}) - f(X_{1:i-1}, x_L, X_{i+1}) \leq c_i. \quad (196)$$

By independence of the X_i 's, the distribution over $X_{i+1:n}$ is the same no matter on whether we condition on $X_i = x_L$ or $X_i = x_U$. Therefore, we can take an expectation over $X_{i+1:n}$ to get that

$$U_i - L_i = \mathbb{E}[f(X_{1:i-1}, x_U, X_{i+1}) \mid X_{1:i-1}, X_i = x_U] \quad (197)$$

$$- \mathbb{E}[f(X_{1:i-1}, x_L, X_{i+1}) \mid X_{1:i-1}, X_i = x_L] \leq c_i. \quad (198)$$

This means that $D_i \in [L_i, U_i]$ is sub-Gaussian with parameter $c_i^2/4$ conditioned on $X_{1:i-1}$.

- Applying Lemma 4, we have that $Z_n - Z_0$ is sub-Gaussian with parameter $\sum_{i=1}^n c_i^2/4$. Finally, apply the sub-Gaussian tail inequality (181).

3.6 Finite hypothesis classes (Lecture 8)

- **Theorem 11 (finite hypothesis class)**

- Let \mathcal{H} be a hypothesis class, where each hypothesis $h \in \mathcal{H}$ maps \mathcal{X} to \mathcal{Y} .
- Let ℓ be the zero-one loss: $\ell((x, y), h) = \mathbb{I}[y \neq h(x)]$.
- Assume Assumption 2 holds.
- Let \hat{h} be the empirical risk minimizer.
- Then with probability at least $1 - \delta$,

$$L(\hat{h}) - L(h^*) \leq \sqrt{\frac{2(\log |\mathcal{H}| + \log(2/\delta))}{n}} = O\left(\sqrt{\frac{\log |\mathcal{H}|}{n}}\right). \quad (199)$$

- **Proof of Theorem 11:**

- The high-level strategy is to use Hoeffding's inequality to show convergence of a fixed hypothesis $h \in \mathcal{H}$ and then a union bound to get uniform convergence.
- Step 1 (convergence)
 - * For a fixed $h \in \mathcal{H}$, note that $\hat{L}(h)$ is an empirical average over n i.i.d. loss terms (bounded in $[0, 1]$) with expectation $L(h)$. Therefore, by Hoeffding's inequality,

$$\mathbb{P}[\hat{L}(h) - L(h) \geq \epsilon] \leq \exp(-2n\epsilon^2). \quad (200)$$

- * To bound the absolute value, we apply the bound again on the negative loss and combine using the union bound:

$$\mathbb{P}[|\hat{L}(h) - L(h)| \geq \epsilon] \leq 2 \exp(-2n\epsilon^2). \quad (201)$$

- Step 2 (uniform convergence).
 - * Since \mathcal{H} is finite, we can apply the union bound over $|\mathcal{H}|$ hypotheses, obtaining:

$$\mathbb{P}\left[\sup_{h \in \mathcal{H}} |\hat{L}(h) - L(h)| \geq \frac{\epsilon}{2}\right] \leq |\mathcal{H}| \cdot 2 \exp\left(-2n\left(\frac{\epsilon}{2}\right)^2\right) \stackrel{\text{def}}{=} \delta. \quad (202)$$

Note that we substituted $\frac{\epsilon}{2}$ for ϵ , which is needed by the generalization bound (159).

- * Rearranging completes the proof.

- Comparison with the realizable case (145):
 - We still have logarithmic dependence on $|\mathcal{H}|$ and $1/\delta$.
 - The main difference is that the realizable bound had a $\frac{1}{n}$ rate, whereas when there is noise, we now get a $\frac{1}{\sqrt{n}}$ rate. This rate difference should be reminiscent of the situation in online learning, where we got $O(\frac{\log T}{T})$ and $O(\frac{1}{\sqrt{T}})$ average regret with and without a perfect expert, respectively. Of course the batch setting is different from the online setting: we have i.i.d. assumptions on the data but no assumptions of convexity on the loss.

3.7 Rademacher complexity (Lecture 8)

- So far, we've relied on concentration inequalities (Hoeffding's inequality) along with the union bound to derive generalization bounds via uniform convergence. However, we can't directly apply the union bound to infinite hypothesis classes (set of all linear classifiers). We need a more sophisticated way to measure the complexity of a hypothesis class. In this section, we will develop such a framework known as Rademacher complexity, which has many nice properties and connects to other measures of complexity such as VC dimension and covering numbers.
- We will arrive at Rademacher complexity by deriving generalization bounds. Specifically, we carry out the following steps:
 - Step 1: Define a random variable G_n (maximum difference between expected and empirical risk).
 - Step 2: Show that it concentrates to $\mathbb{E}[G_n]$ using McDiarmid's inequality.
 - Step 3: Use a technique called **symmetrization** to bound the expectation using a quantity known as the Rademacher complexity.
- Step 1 (setup)
 - Consider the largest difference between the expected and the empirical risk over all possible hypotheses:

$$G_n \stackrel{\text{def}}{=} \sup_{h \in \mathcal{H}} L(h) - \hat{L}(h). \quad (203)$$

Here, G_n is a random variable that depends on the data points Z_1, \dots, Z_n . An upper bound on G_n would ensure that if you observed empirical risk $\hat{L}(\hat{h})$, the expected risk $L(\hat{h})$ will not be much higher.

- In order to bound the excess risk (159), we actually need to also bound G'_n defined analogously for the negative loss $\ell'(z, h) = -\ell(z, h)$, so that

$$\mathbb{P}[L(\hat{h}) - L(h^*) \geq \epsilon] \leq \mathbb{P}\left[G_n \geq \frac{\epsilon}{2}\right] + \mathbb{P}\left[G'_n \geq \frac{\epsilon}{2}\right]. \quad (204)$$

Usually, the tail bound for G_n is the same as for G'_n , as we'll see later.

- Step 2 (concentration): convert tail bound into an expectation

- Let g be the deterministic function such that $G_n = g(Z_1, \dots, Z_n)$.
- Then g satisfies the following bounded differences condition:

$$|g(Z_1, \dots, Z_i, \dots, Z_n) - g(Z_1, \dots, Z'_i, \dots, Z_n)| \leq \frac{1}{n}. \quad (205)$$

- Proof:

- * Recall $\hat{L}(h) = \frac{1}{n} \sum_{i=1}^n \ell(Z_i, h)$.
- * We have:

$$\left| \underbrace{\sup_{h \in \mathcal{H}} [L(h) - \hat{L}(h)]}_{g(Z_1, \dots, Z_i, \dots, Z_n)} - \underbrace{\sup_{h \in \mathcal{H}} [L(h) - \hat{L}(h) + \frac{1}{n}(\ell(Z_i, h) - \ell(Z'_i, h))]}_{g(Z_1, \dots, Z'_i, \dots, Z_n)} \right| \leq \frac{1}{n}. \quad (206)$$

- * For each $h \in \mathcal{H}$, the difference between the G_n and the perturbed G_n is at most $\frac{1}{n}$ since the loss is bounded: $\ell(z, h) \in [0, 1]$. Taking the supremum (which is a contraction) cannot increase the difference.
- Now we can apply McDiarmid's inequality (Theorem 10) to get that:

$$\mathbb{P}[G_n \geq \mathbb{E}[G_n] + \epsilon] \leq \exp(-2n\epsilon^2). \quad (207)$$

- Remark: Note that g is quite a non-trivial function (involving a sup over a huge hypothesis class), but because it satisfies the bounded differences condition, we can simply apply McDiarmid's inequality.

- Step 3 (symmetrization)

- Now we need to bound $\mathbb{E}[G_n]$. This quantity is quite difficult since it depends on $L(h)$, an expectation over the unknown distribution p^* . The goal of symmetrization is to remove this strong dependence on p^* and replace it with a quantity that only depends on p^* through the data points Z_1, \dots, Z_n .
- The key idea of symmetrization is to introduce a ghost dataset Z'_1, \dots, Z'_n , drawn i.i.d. from p^* . Let $\hat{L}'(h) = \frac{1}{n} \sum_{i=1}^n \ell(Z'_i, h)$ be the empirical risk with respect to this ghost dataset.
- Rewriting $L(h)$ in terms of the ghost dataset:

$$\mathbb{E}[G_n] = \mathbb{E}[\sup_{h \in \mathcal{H}} \mathbb{E}[\hat{L}'(h)] - \hat{L}(h)]. \quad (208)$$

- Bring $\hat{L}(h)$ into the inner expectation by conditioning on the original dataset Z_1, \dots, Z_n (the two datasets are independent):

$$\mathbb{E}[G_n] = \mathbb{E}[\sup_{h \in \mathcal{H}} \mathbb{E}[\hat{L}'(h) - \hat{L}(h) \mid Z_1, \dots, Z_n]]. \quad (209)$$

- Pushing the sup inside the expectation can only increase:

$$\mathbb{E}[G_n] \leq \mathbb{E}[\mathbb{E}[\sup_{h \in \mathcal{H}} \hat{L}'(h) - \hat{L}(h) \mid Z_1, \dots, Z_n]]. \quad (210)$$

- Apply law of iterated conditional expectation:

$$\mathbb{E}[G_n] \leq \mathbb{E}[\sup_{h \in \mathcal{H}} \hat{L}'(h) - \hat{L}(h)]. \quad (211)$$

- Introduce i.i.d. **Rademacher variables** $\sigma_1, \dots, \sigma_n$, where σ_i is uniform over $\{-1, +1\}$. Since $\ell(Z'_i, h) - \ell(Z_i, h)$ is symmetric around 0, multiplying by σ_i doesn't change its distribution. Now expanding the definition of the empirical risk:

$$\mathbb{E}[G_n] \leq \mathbb{E} \left[\sup_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \sigma_i [\ell(Z'_i, h) - \ell(Z_i, h)] \right]. \quad (212)$$

- Pushing the sup inside $\sup_h [a_h - b_h] \leq \sup_h [a_h] + \sup_h [-b_h]$:

$$\mathbb{E}[G_n] \leq \mathbb{E} \left[\sup_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \sigma_i \ell(Z'_i, h) + \sup_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (-\sigma_i) \ell(Z_i, h) \right]. \quad (213)$$

- By linearity of expectation and the fact that Z'_i has the same distribution as Z_i , and σ_i and $-\sigma_i$ have the same distribution:

$$\boxed{\mathbb{E}[G_n] \leq 2 \mathbb{E} \left[\sup_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \sigma_i \ell(Z_i, h) \right]}. \quad (214)$$

The RHS motivates the following general definition of the Rademacher complexity:

• **Definition 12 (Rademacher complexity)**

- Let \mathcal{F} be a class of real-valued functions $f : \mathcal{Z} \rightarrow \mathbb{R}$. Example: $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ consists of input-output pairs.
- Define the **Rademacher complexity** (or Rademacher average) of \mathcal{F} to be

$$\boxed{R_n(\mathcal{F}) \stackrel{\text{def}}{=} \mathbb{E} \left[\sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i) \right]}, \quad (215)$$

where

- * Z_1, \dots, Z_n are drawn i.i.d. from p^* (data points); and
- * $\sigma_1, \dots, \sigma_n$ are drawn i.i.d. from the uniform distribution over $\{-1, +1\}$ (Rademacher variables).
- Interpretation of Rademacher complexity:
 - * Consider a binary classification problem with inputs Z_1, \dots, Z_n and *random labels* $\sigma_1, \dots, \sigma_n$. Clearly, this is a meaningless learning problem.
 - * The Rademacher complexity captures (in expectation) how well the best function from the function class \mathcal{F} can align with these random labels. In other words, how well \mathcal{F} can fit noise?
 - * We'd like $R_n(\mathcal{F})$ to go to zero as n increases.
- Define the **empirical Rademacher complexity** of \mathcal{F} to be:

$$\hat{R}_n(\mathcal{F}) \stackrel{\text{def}}{=} \mathbb{E} \left[\sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i) \mid Z_1, \dots, Z_n \right], \quad (216)$$

which is a random variable depending on the data. Note that $R_n(\mathcal{F}) = \mathbb{E}[\hat{R}_n(\mathcal{F})]$, where the expectation is taken over the n training examples.

• **Theorem 12 (generalization bound based on Rademacher complexity)**

- Define $\mathcal{A} = \{z \mapsto \ell(z, h) : h \in \mathcal{H}\}$ to be the **loss class**, the composition of the loss function with each of the hypotheses. With probability at least $1 - \delta$,

$$L(\hat{h}) - L(h^*) \leq 4R_n(\mathcal{A}) + \sqrt{\frac{2 \log(2/\delta)}{n}}. \quad (217)$$

• Proof of Theorem 12:

- Note that $\mathbb{E}[G_n] \leq 2R_n(\mathcal{A})$ by definition of Rademacher complexity.
- Since negation doesn't change the Rademacher complexity, $R_n(\mathcal{A}) = R_n(-\mathcal{A})$, so $\mathbb{E}[G'_n] \leq 2R_n(-\mathcal{A})$.
- Let's apply the tail bound (207) on both G_n and G'_n :

$$\mathbb{P} \left[G_n \geq \frac{\epsilon}{2} \right] \leq \exp \left(-2n \left(\frac{\epsilon}{2} - \mathbb{E}[G_n] \right)^2 \right) \quad (218)$$

$$\leq \exp \left(-2n \left(\frac{\epsilon}{2} - 2R_n(\mathcal{A}) \right)^2 \right) \quad [\text{for } \epsilon \geq 4R_n(\mathcal{A})] \quad (219)$$

$$\stackrel{\text{def}}{=} \frac{\delta}{2}. \quad (220)$$

We have an analogous inequality for G'_n . Adding them and solving for ϵ yields the result.

- Remark: We see that the essence of generalization is captured in the Rademacher complexity of the loss class $R_n(\mathcal{A})$.
- Summary
 - We have reduced the problem of bounding G_n , which involved a complex comparison between $L(h)$ and $\hat{L}(h)$, to something that only depends on samples. In fact, the empirical Rademacher complexity can even be computed from data. We'll see that in the analysis to follow, we will often condition on $Z_{1:n}$; the points take a backseat as it is the randomness of σ_i and the supremum over \mathcal{F} that really determine the Rademacher complexity.
 - We will study the Rademacher complexity $R_n(\mathcal{F})$ of various function classes \mathcal{F} . First, let us discuss some basic compositional properties that Rademacher complexity enjoys, mostly due to linearity of expectation:
- **Basic properties of Rademacher complexity**
 - Boundedness
 - * $R_n(\mathcal{F}) \leq \max_{f \in \mathcal{F}} \max_z f(z)$.
 - * This is a trivial bound, which means that it's not impressive to say that the Rademacher complexity is a constant. We'd ideally like it to go to zero as $n \rightarrow \infty$.
 - Singleton
 - * $R_n(\{f\}) = 0$
 - * Proof: σ_i has zero mean and is independent of everything else, so $\mathbb{E}[\sigma_i f(Z_i)] = 0$
 - Monotonicity
 - * $R_n(\mathcal{F}_1) \leq R_n(\mathcal{F}_2)$ if $\mathcal{F}_1 \subseteq \mathcal{F}_2$
 - * Proof: $\sup_{f \in \mathcal{F}_2}$ ranges over at least as many functions as $\sup_{f \in \mathcal{F}_1}$, which makes it at least as large.
 - Linear combination
 - * $R_n(\mathcal{F}_1 + \mathcal{F}_2) = R_n(\mathcal{F}_1) + R_n(\mathcal{F}_2)$ for $\mathcal{F}_1 + \mathcal{F}_2 = \{f_1 + f_2 : f_1 \in \mathcal{F}_1, f_2 \in \mathcal{F}_2\}$
 - * Proof: linearity of expectation
 - Scaling
 - * $R_n(c\mathcal{F}) = |c|R_n(\mathcal{F})$
 - * Proof: for $c > 0$, this is trivial; if $c < 0$, then we can absorb the negation into σ_i 's, which are symmetric around zero.
 - Lipschitz composition (a generalization of scaling):
 - * $R_n(\phi \circ \mathcal{F}) \leq c_\phi R_n(\mathcal{F})$, where $\phi \circ \mathcal{F} = \{z \mapsto \phi(f(z)) : f \in \mathcal{F}\}$ and c_ϕ is the Lipschitz constant of ϕ : $|\phi(z) - \phi(z')| \leq c_\phi |z - z'|$.

- * Proof: see Corollary 3.17 of Ledoux and Talagrand 1991
 - * If ϕ is differentiable, then c_ϕ is just a bound on the L_2 norm of the gradient.
 - * This property is useful because we can analyze the complexity of our hypothesis class (linear functions), and compose with ϕ to get the loss class.
- Convex hull
- * $R_n(\text{convex-hull}(\mathcal{F})) = R_n(\mathcal{F})$ for finite \mathcal{F}
 - * Proof: supremum over the convex hull is attained at one of its vertices
 - * This property is useful because if we want to compute the Rademacher of a polytope (an infinite set), it suffices to compute the Rademacher complexity of its vertices (a finite set). We will use this property when we look at L_1 regularization.

3.8 Finite-dimensional hypothesis classes (Lecture 9)

- So far, we've set up Rademacher complexity as a means of bounding the complexity of a function class (specifically, the loss class associated with a hypothesis class). Now, let's instantiate Rademacher complexity for the case where the function class is finite. This may initially seem like overkill (since we already analyzed the finite hypothesis case), but doing this analysis in the Rademacher complexity is a good sanity check, and it will reveal an unexpected but important aspect of finiteness.

- **Lemma 5 (Massart's finite lemma)**

- Throughout this lemma, condition on Z_1, \dots, Z_n (treat them like constants).
- Let \mathcal{F} be a finite class of functions.
- Let M be a bound satisfying:

$$\sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n f(Z_i)^2 \leq M^2. \quad (221)$$

- Then

$$\boxed{\hat{R}_n(\mathcal{F}) \leq \sqrt{\frac{2M^2 \log |\mathcal{F}|}{n}}.} \quad (222)$$

- Proof of Lemma 5:

- Let $W_f = \frac{1}{n} \sum_{i=1}^n \sigma_i f(Z_i)$.
- We are interested in bounding $\hat{R}_n(\mathcal{F}) = \mathbb{E}[\sup_{f \in \mathcal{F}} W_f \mid Z_{1:n}]$.
- Exponentiate and use convexity of $x \mapsto \exp(tx)$ for $t \geq 0$ to push the exp inside the expectation:

$$\exp(t \mathbb{E}[\sup_{f \in \mathcal{F}} W_f \mid Z_{1:n}]) \leq \mathbb{E}[\exp(t \sup_{f \in \mathcal{F}} W_f) \mid Z_{1:n}]. \quad (223)$$

- By monotonicity, we can pull the sup out of the exponent:

$$= \mathbb{E}[\sup_{f \in \mathcal{F}} \exp(t W_f) \mid Z_{1:n}]. \quad (224)$$

- The sup over non-negative terms can be upper bounded by the sum:

$$\leq \sum_{f \in \mathcal{F}} \mathbb{E}[\exp(t W_f) \mid Z_{1:n}]. \quad (225)$$

- Now we have to bound $\mathbb{E}[\exp(tW_f)]$, which is the moment generating function of W_f .
 - * By Hoeffding's lemma (after Lemma 3.5), σ_i is a bounded random variable with sub-Gaussian parameter $2^2/4 = 1$.
 - * By scaling and independence of σ_i , we get that W_f is sub-Gaussian with parameter $\frac{1}{n^2} \sum_{i=1}^n f(Z_i)^2 \leq \frac{M^2}{n}$ (remember that we're conditioning on $Z_{1:n}$, which are just treated like constants).
 - * By the definition of sub-Gaussian, we have

$$\exp(t\hat{R}_n(\mathcal{F})) \leq \sum_{f \in \mathcal{F}} \mathbb{E}[\exp(tW_f) \mid Z_{1:n}] \leq |\mathcal{F}| \exp\left(\frac{t^2 M^2}{2n}\right). \quad (226)$$

- Taking logs and dividing by t :

$$\hat{R}_n(\mathcal{F}) \leq \frac{\log |\mathcal{F}|}{t} + \frac{tM^2}{2n}. \quad (227)$$

- Optimizing over t yields the result as desired (by just taking the twice the geometric average of $\log |\mathcal{F}|$ and $M^2/(2n)$):

$$\hat{R}_n(\mathcal{F}) \leq \sqrt{\frac{2M^2 \log |\mathcal{F}|}{n}}. \quad (228)$$

- We can immediately apply Massart's finite lemma to any finite loss class \mathcal{A} . For example, for the zero-one loss, we have each $f(Z_i) = \ell(Z_i, h) \in [0, 1]$ (for $h \in \mathcal{H}$ corresponding to $f \in \mathcal{F}$) so we can take $M = 1$. Combined with (217), this gives us

$$L(\hat{h}) - L(h^*) \leq \sqrt{\frac{32 \log |\mathcal{H}|}{n}} + \sqrt{\frac{2 \log(2/\delta)}{n}} = O\left(\sqrt{\frac{\log |\mathcal{H}|}{n}}\right), \quad (229)$$

which gives the same result (with worse constants) compared with Theorem 11.

3.9 Shattering coefficient (Lecture 9)

- But we can do more than that. Notice that Massart's finite lemma places a bound on the empirical Rademacher complexity \hat{R}_n , which only depends on the n data points. If we had an infinite function class \mathcal{F} , but acts the same on the n points as a finite function class \mathcal{F}' , then the two function classes have the same empirical Rademacher complexity. More formally, if

$$\{[f(Z_1), \dots, f(Z_n)] : f \in \mathcal{F}\} = \{[f'(Z_1), \dots, f'(Z_n)] : f' \in \mathcal{F}'\}, \quad (230)$$

then $\hat{R}_n(\mathcal{F}) = \hat{R}_n(\mathcal{F}')$. Therefore, all that matters as far as empirical Rademacher complexity is concerned is **the behavior of a function class on n data points**.

- This key observation is useful when we are trying to analyze **boolean functions**, those that return either 0 or 1. An important example is the loss class $\mathcal{A} = \{z \mapsto \ell(z, h) : h \in \mathcal{H}\}$ where ℓ is the zero-one loss. Then the number of behaviors on n points is certainly finite (and perhaps even small). The behavior of \mathcal{F} on n data points is captured in the shattering coefficient:

- **Definition 13 (shattering coefficient (growth function))**

- Let \mathcal{F} be a family of functions that map \mathcal{Z} to a finite set (usually $\{0, 1\}$).
- The **shattering coefficient** of \mathcal{F} is the maximum number of behaviors over n points:

$$s(\mathcal{F}, n) \stackrel{\text{def}}{=} \max_{z_1, \dots, z_n \in \mathcal{Z}} |\{[f(z_1), \dots, f(z_n)] : f \in \mathcal{F}\}|. \quad (231)$$

- We can use Massart's finite lemma to upper bound the Rademacher complexity by the shattering coefficient. If \mathcal{F} contains boolean functions, we have the bound $M = 1$, so that:

$$\boxed{R_n(\mathcal{F}) \leq \sqrt{\frac{2 \log s(\mathcal{F}, n)}{n}}}. \quad (232)$$

The significance is that we can apply Massart's finite lemma to **infinite function classes with finite shattering coefficient**. In order to have applied this bound, it is important to condition on the data (note that the distribution over the data has disappeared completely!) After applying Massart's finite lemma, we can take expectations over the data without changing the bound. We have turned an impossible expectation over p^* into a purely combinatorial notion of complexity.

- Intuition:

- For boolean functions, if $s(\mathcal{F}, n) = 2^n$ (we obtain all possible labelings), then we say \mathcal{F} **shatters** the n points that achieve the maximum of (231).
- To get meaningful bounds, we want $s(\mathcal{F}, n)$ to grow sub-exponentially with n ; otherwise the Rademacher complexity will not go to zero, and we will not obtain uniform convergence. This is expected since if \mathcal{F} can really hit all labelings for all n , then we would be able to fit any labeling of the data, leading to massive overfitting.

- Shattering coefficient of hypotheses class and loss class

- The bounds we derive depend on the shattering coefficient $s(\mathcal{A}, n)$ of the loss class \mathcal{A} (via (232) and (217)). However, it will be more intuitive to talk about the shattering coefficient hypothesis class \mathcal{H} . We now show that the two are the same for zero-one loss and binary classifiers.

- Let \mathcal{H} be a set of binary classifiers $h : \mathcal{X} \mapsto \{0, 1\}$ and zero-one loss $\ell((x, y), h) = \mathbb{I}[y \neq h(x)]$.
- Note that the hypothesis class \mathcal{H} contains functions on \mathcal{X} , and the loss class \mathcal{A} contains functions on $\mathcal{X} \times \{0, 1\}$: $\mathcal{A} = \{(x, y) \mapsto \mathbb{I}[y \neq h(x)] : h \in \mathcal{H}\}$.
- The key point is that

$$\boxed{s(\mathcal{H}, n) = s(\mathcal{A}, n).} \quad (233)$$

- The reason is as follows: for any $(x_1, y_1), \dots, (x_n, y_n)$, there is a bijection between the behavior of the losses and the hypotheses:

$$[\ell((x_1, y_1), h), \dots, \ell((x_n, y_n), h)] \Leftrightarrow [h(x_1), \dots, h(x_n)], \quad (234)$$

where the translation between the two sets of vectors is obtained by taking the XOR with $[y_1, \dots, y_n]$. Therefore, as we range over $h \in \mathcal{H}$, we obtain the same number of behaviors from ℓ as we do from h .

3.10 VC dimension (Lecture 9)

- In the previous section, we saw that the shattering coefficient can be used to upper bound the Rademacher complexity of a family of boolean functions (232), which in turn can be used to upper bound the excess risk (Theorem 12). Although the shattering coefficient nicely captures the behavior of an infinite \mathcal{H} , it is not necessarily the most convenient quantity to get a handle on. In this section, we will use a concept called VC dimension to gain more intuition about the shattering coefficient.

- **Definition 14 (VC dimension)**

- The **VC dimension** of a family of functions \mathcal{H} with boolean outputs is the maximum number of points that can be shattered by \mathcal{H} :

$$\boxed{\text{VC}(\mathcal{H}) = \sup\{n : s(\mathcal{H}, n) = 2^n\}.} \quad (235)$$

- Intuition: the VC dimension of \mathcal{H} is the maximum number of points whose labels can be memorized perfectly by choosing some $h \in \mathcal{H}$.
- Overloading notation, we can think of each $h \in \mathcal{H}$ (a boolean function on \mathcal{X}) as a subset of \mathcal{X} for which $h(x)$ returns 1.
- Example (intervals)
 - Let $\mathcal{H} = \{[a, b] : a, b \in \mathbb{R}\}$.
 - $s(\mathcal{H}, 1) = 2 = 2^1$ (can shatter)

- $s(\mathcal{H}, 2) = 4 = 2^2$ (can shatter)
 - $s(\mathcal{H}, 3) = 7 < 2^3$ (can't shatter, because can't isolate the middle point)
 - $s(\mathcal{H}, n) = \binom{n+1}{2} + 1$
 - Therefore, $\text{VC}(\mathcal{H}) = 2$.
 - Important note: to show that a class \mathcal{H} has VC dimension d , one needs to
 - Derive an upper bound: show that no $d + 1$ points can be shattered. This is done by showing that for any $d + 1$ points, there is some labeling that cannot be achieved by \mathcal{H} .
 - Derive a lower bound: show that there exists d points can be shattered. This is done by showing that there exists d points such that all 2^d labelings can be achieved by \mathcal{H} .
 - One can verify that the VC dimension of rectangles in two dimensions is 4. Based on this, one might be tempted to conclude that the VC dimension of a hypothesis class with d parameters is d , but the following example shows that this is in general not the case (there are pathological examples).
 - Example (infinite VC dimension with one parameter)
 - Let $\mathcal{H} = \{\{x : \sin(\theta x) \geq 0\} : \theta \in \mathbb{R}\}$.
 - We have that $\text{VC}(\mathcal{H}) = \infty$.
 - It turns out that it's not the number of parameters that matter, but the dimension of the function space:
 - **Theorem 13 (finite-dimensional function class)**
 - Let \mathcal{F} be a function class containing functions $f : \mathcal{X} \rightarrow \mathbb{R}$.
 - Recall that the dimension of \mathcal{F} is the number of elements in a basis of \mathcal{F} .
 - Let $\mathcal{H} = \{\{x : f(x) \geq 0\} : f \in \mathcal{F}\}$ (for example, \mathcal{F} could be linear functions, but not necessarily).
 - Then we have
- $$\boxed{\text{VC}(\mathcal{H}) \leq \dim(\mathcal{F})}. \tag{236}$$
- Remark: this allows us to connect the linear algebraic properties of \mathcal{F} (dimension) with the combinatorial properties of \mathcal{H} (VC dimension).
 - Proof of Theorem 13:
 - Take any n points x_1, \dots, x_n with $n > \dim(\mathcal{F})$. We will show that these n points cannot be shattered.

- Consider the linear map $H(f) \stackrel{\text{def}}{=} (f(x_1), \dots, f(x_n)) \in \mathbb{R}^n$ which maps each function $f \in \mathcal{F}$ to the function values on the n points (function evaluation is linear).
- The vector space $\{H(f) : f \in \mathcal{F}\} \in \mathbb{R}^n$ has dimension at most $\dim(\mathcal{F})$ (applying linear maps can't increase dimension).
- Since $n > \dim(\mathcal{F})$, there is some non-zero vector $c \in \mathbb{R}^n$ such that $H(f) \cdot c = 0$ for all $f \in \mathcal{F}$.
- Without loss of generality, some component of c is negative (otherwise, just take $-c$, which satisfies $H(f) \cdot c = 0$ too).
- So we have for all $f \in \mathcal{F}$:

$$\sum_{i:c_i \geq 0} c_i f(x_i) + \sum_{i:c_i < 0} c_i f(x_i) = 0. \quad (237)$$

- For the purposes of contradiction, suppose \mathcal{H} shatters $\{x_1, \dots, x_n\}$.
 - * Then we could find an $A = \{x : f(x) \geq 0\} \in \mathcal{H}$ (equivalently some $f \in \mathcal{F}$) such that
 - $x_i \in A$ ($f(x_i) \geq 0$) whenever $c_i \geq 0$ (so the first term of (237) is non-negative)
 - $x_i \notin A$ ($f(x_i) < 0$) whenever $c_i < 0$ (so the second term of (237) is strictly positive).
 - * The sum of the two terms couldn't equal zero, which contradicts (237).
- Therefore, \mathcal{H} can't shatter $\{x_1, \dots, x_n\}$ for any choice of x_1, \dots, x_n , so $\text{VC}(\mathcal{H}) \leq \dim(\mathcal{F})$.

• **Example 18 (Half-spaces passing through the origin)**

- Let $\mathcal{H} = \{x : w \cdot x \geq 0\} : w \in \mathbb{R}^d\}$.
- By Theorem 13, the VC dimension of \mathcal{H} is at most d (this is a linear function class of dimension d). It remains to show that the VC dimension is at least d .
- In general, upper bounds suffice for obtaining generalization bounds, but just for fun, let's get a lower bound on the VC dimension.
- Showing that the VC dimension of some \mathcal{H} is at least d is easier because we just have to construct some set of d points that can be shattered rather than showing there does not exist $d + 1$ points that can be shattered.
- Create d points:

$$x_1 = (1, 0, \dots, 0, 0) \quad (238)$$

$$\dots \quad (239)$$

$$x_d = (0, 0, \dots, 0, 1) \quad (240)$$

- Given any subset $I \subseteq \{1, \dots, d\}$ (labeling of x_1, \dots, x_d), we can construct a w as follows to obtain that labeling:
 - * Set $w_i = 1$ for $i \in I$
 - * Set $w_i = -1$ for $i \notin I$
- This establishes that $\text{VC}(\mathcal{H}) \geq d$.
- Putting the upper and lower bounds, we get that $\boxed{\text{VC}(\mathcal{H}) = d}$.
- We have defined VC dimension in terms of shattering coefficient (236), and have given some examples of VC dimension. Let us now upper bound the shattering coefficient in terms of the VC dimension.
- **Lemma 6 (Sauer's lemma)**

- For a class \mathcal{H} be a class with VC dimension d .
- Then

$$\boxed{s(\mathcal{H}, n) \leq \sum_{i=0}^d \binom{n}{i} \leq \begin{cases} 2^n & \text{if } n \leq d \\ \left(\frac{en}{d}\right)^d & \text{if } n > d. \end{cases}} \quad (241)$$

- Intuition
 - For $n \leq d$, the shattering coefficient grows exponentially in n .
 - For $n > d$, the shattering coefficient grows only polynomially in n .
 - In some sense, \mathcal{H} can only represent any subset of up to d of the n points.
- The ramification of Sauer's lemma is that now we have a new upper upper bound on the Rademacher complexity (and thus on uniform convergence):

$$\hat{R}_n(\mathcal{A}) \leq \sqrt{\frac{2 \log s(\mathcal{A}, n)}{n}} = \sqrt{\frac{2 \log s(\mathcal{H}, n)}{n}} \leq \sqrt{\frac{2 \text{VC}(\mathcal{H})(\log n + 1)}{n}}, \quad (242)$$

where the first inequality follows from (232), the second equality follows from (233), and the final inequality follows from (241) and holds for $n \geq d$.

- Proof of Lemma 6 (somewhat technical and specific, skipped in class)
 - The idea is to take \mathcal{H} and transform it into \mathcal{H}' with the same shattering coefficient ($s(\mathcal{H}, n) = s(\mathcal{H}', n)$) but where $s(\mathcal{H}', n)$ will be easier to bound.
 - Key: all that matters is the action of \mathcal{H} and \mathcal{H}' on a finite set of n points, which we will represent as a table.
 - Draw a table whose columns are the n points and rows are the $s(\mathcal{H}, n)$ possible labelings, and each entry is either a 0 or 1. The question is how many rows there are.

- Here's an example table T :

x_1	x_2	x_3	x_4
0	1	0	1
0	0	0	1
1	1	1	0
1	0	1	0

- We will transform the table to a canonical form as follows:

- * Pick a column j .
- * For each row r with $r_j = 1$, set $r_j = 0$ if the resulting r doesn't exist in the table.
- * Repeat until no more changes are possible.

- Here is the resulting table T' (corresponding to some \mathcal{H}'):

x_1	x_2	x_3	x_4
0	1	0	1
0	0	0	1
0	1	0	0
0	0	0	0

- Step 1: Note that the number of rows is still the same and all the rows are all distinct, so $s(\mathcal{H}, n) = s(\mathcal{H}', n)$. So we just have to compute $s(\mathcal{H}', n)$, which should be easier.
- Step 2: We show that the VC dimension doesn't increase by transformation ($\text{VC}(\mathcal{H}') \leq \text{VC}(\mathcal{H})$)

- * The transformations proceed one column at a time:

$$T \rightarrow T_1 \rightarrow \dots T_k \xrightarrow{\text{transform column } j} T_{k+1} \rightarrow \dots \rightarrow T'. \quad (243)$$

- * Claim: After transforming any column j , if some subset $S \subseteq \{1, \dots, n\}$ of points is shattered (all $2^{|S|}$ labelings exist on those columns) after transformation (in T_{k+1}), then S was also shattered before transformation (in T_k).
 - * Case 1: trivially true for all subsets S that don't contain j .
 - * Case 2: take any subset S that contains j .
 - For any row i with 1 in column j , there is a row i' with 0 in column j and agrees with r on all columns except j : $T_{k+1}(i, j') = T_{k+1}(i', j')$ for all $j' \in \{1, \dots, n\} \setminus \{j\}$, but $T_{k+1}(i, j) = 1$ and $T_{k+1}(i', j) = 0$.
 - Note that $T_k(i, j) = 1$ (because we never turn zeros into ones).
 - Note that $T_k(i', j) = 0$ because if it had been a 1, then rows i and i' would have been identical, and we maintain the invariant that there are no duplicate rows.
 - * So all $2^{|S|}$ labelings on S existed before transformation in T_k .
- Step 3: Each row of T' must contain at most d ones.
- * Suppose if T' has a row with k ones in columns $S \subseteq \{1, \dots, n\}$.

- * Then for each $j \in S$, there must be another row with a labeling that assigns ones to exactly $S \subseteq \{j\}$ (otherwise we would have been able to transform column j by changing the 1 to a 0).
 - * Reasoning recursively, all 2^k subsets must exist.
 - * Since T' has VC dimension at most d , $k \leq d$.
 - * Based on simple counting, we find that number of rows (remember, they're all distinct!) is upper bounded by $\sum_{i=0}^d \binom{n}{i}$, completing the first inequality of (241).
- Finishing the second part of the inequality of (241) is just algebra. Observe that for $n \geq d$,

$$\sum_{i=0}^d \binom{n}{i} \leq \left(\frac{n}{d}\right)^d \sum_{i=0}^d \binom{n}{i} \left(\frac{d}{n}\right)^i \quad (244)$$

$$\leq \left(\frac{n}{d}\right)^d \sum_{i=0}^n \binom{n}{i} \left(\frac{d}{n}\right)^i \quad (245)$$

$$= \left(\frac{n}{d}\right)^d \left(1 + \frac{d}{n}\right)^n \quad (246)$$

$$\leq \left(\frac{n}{d}\right)^d e^d. \quad (247)$$

3.11 Norm-constrained hypothesis classes (Lecture 10)

- We started by establishing Rademacher complexity $R_n(\mathcal{F})$ as a measure of complexity of a function class that yielded generalization bounds when applied to the loss class \mathcal{A} . Then we specialized to boolean functions and zero-one losses, which led to notions of shattering coefficients and VC dimension as combinatorial means of bounding the Rademacher complexity.
- However, combinatorial notions are not the most appropriate way to think about complexity. For example, it is common in machine learning to have a huge number of features (large d) in a linear classifier, but where we regularize or constrain the norm of the weights. In this case, d is the VC dimension doesn't really capture the true complexity of the hypothesis class.
- Deriving Rademacher complexity will actually be easier for the norm-constrained setting. We will study the Rademacher complexity of three such examples:
 - Linear functions with L_2 norm constraints (suitable for kernel methods)
 - Linear functions with L_1 norm constraints (suitable for sparsity)
- Rademacher complexity of linear functions with weights bounded in an L_2 ball
 - **Theorem 14 (Rademacher complexity of L_2 ball)**
 - * Let $\mathcal{F} = \{z \mapsto w \cdot z : \|w\|_2 \leq B_2\}$ (bound on weight vectors).
 - * Assume $\mathbb{E}_{Z \sim p^*}[\|Z\|_2^2] \leq C_2^2$ (bound on spread of data points).
 - * Then

$$\boxed{R_n(\mathcal{F}) \leq \frac{B_2 C_2}{\sqrt{n}}.} \tag{248}$$

- Proof of Theorem 14:
 - * Expand the definition:

$$R_n(\mathcal{F}) = \frac{1}{n} \mathbb{E} \left[\sup_{\|w\|_2 \leq B_2} \sum_{i=1}^n \sigma_i(w \cdot Z_i) \right]. \tag{249}$$

- * By Cauchy-Schwartz applied to w and $\sum_{i=1}^n \sigma_i Z_i$:

$$\leq \frac{B_2}{n} \mathbb{E} \left[\left\| \sum_{i=1}^n \sigma_i Z_i \right\|_2 \right]. \tag{250}$$

* By concavity of $\sqrt{\cdot}$, we can push it outside the expectation:

$$\leq \frac{B_2}{n} \sqrt{\mathbb{E} \left[\left\| \sum_{i=1}^n \sigma_i Z_i \right\|_2^2 \right]}. \quad (251)$$

* Distribute the sum; **expectation of cross terms is zero** by independence of σ_i (this is the key point, which turns n^2 terms into n terms):

$$= \frac{B_2}{n} \sqrt{\mathbb{E} \left[\sum_{i=1}^n \|\sigma_i Z_i\|_2^2 \right]}. \quad (252)$$

* We can drop σ_i because it changes sign, not magnitude:

$$= \frac{B_2}{n} \sqrt{\mathbb{E} \left[\sum_{i=1}^n \|Z_i\|_2^2 \right]}. \quad (253)$$

* Use the bound on Z_i :

$$\leq \frac{B_2}{n} \sqrt{n C_2^2}. \quad (254)$$

Simple algebra completes the proof.

- Rademacher complexity of linear functions with weights bounded in an L_1 ball

– Motivation

- * Working with L_2 regularization has the advantage that we can use kernel methods, and the dimensionality could be infinite as long the norm is bounded.
- * In some applications, we have a finite but large set of features, and we believe that there are a relatively small subset that are relevant to our task (i.e., we believe in **parameter sparsity**). It is common to use L_1 regularization, or similarly, assume that the weights satisfy $\|w\|_1 \leq B_1$.

Let us compute the Rademacher complexity of $\mathcal{F} = \{z \mapsto w \cdot z : \|w\|_1 \leq B_1\}$.

– **Theorem 15 (Rademacher complexity of L_1 ball)**

- * Assume that the coordinates are bounded: $\|Z_i\|_\infty \leq C_\infty$ with probability 1 for all data points $i = 1, \dots, n$.
- * Then

$$\boxed{R_n(\mathcal{F}) \leq \frac{B_1 C_\infty \sqrt{2 \log(2d)}}{\sqrt{n}}}. \quad (255)$$

– Proof of Theorem 15

- * The key step is to realize that the L_1 ball ($\{w : \|w\|_1 \leq B_1\}$) is the convex hull of the following $2d$ weight vectors:

$$W = \cup_{j=1}^d \{B_1 e_j, -B_1 e_j\}. \quad (256)$$

Since the Rademacher complexity of a class is the same as the Rademacher complexity of its convex hull, we just need to look at the finite class:

$$R_n(\mathcal{F}) = \mathbb{E} \left[\sup_{w \in W} \frac{1}{n} \sum_{i=1}^n \sigma_i(w \cdot Z_i) \right]. \quad (257)$$

- * Apply Hölder's inequality, we have $w \cdot Z_i \leq \|w\|_1 \|Z_i\|_\infty \leq B_1 C_\infty$.
- * Applying Massart's finite lemma (Lemma 5) to the function class specified by W (of size $2d$) with $M^2 = B_1^2 C_\infty^2$, we get that

$$R_n(\mathcal{F}) \leq \sqrt{\frac{2B_1^2 C_\infty^2 \log(2d)}{n}}. \quad (258)$$

– Remarks

- * It is useful to recall that as p increases,
 - p -norms decrease: $\|w\|_p \geq \|w\|_q$
 - Size of balls increase: $\{w : \|w\|_p \leq B\} \subseteq \{w : \|w\|_q \leq B\}$
- * Note that a L_1 bound on the parameters is placing a much stronger constraint than the L_2 norm, which allows us to measure the L_∞ norm of the data (which is much better) rather than the L_2 norm at the expense of a logarithmic dependence on d .
- * This tradeoff came up earlier when we looked at exponentiated gradient (EG) for online learning, where the experts resided in an L_1 ball, which allowed us to get a bound that depended logarithmically on the number of experts.

– Ramifications under **sparsity**

- * L_1 regularization is often used when we believe that most features are irrelevant; formally, that the desired weight vector has $s \ll d$ non-zero entries.
- * You might have seen the intuition that L_1 regularization has sharp corners which encourage weights to be identically zero, but that doesn't directly tell us anything about generalization. We would like a stronger justification.
- * For convenience, assume that all entries of w and x have magnitude at most 1 ($\|w\|_\infty \leq 1, \|x\|_\infty \leq 1$).
- * It suffices to consider the hypothesis class $\|w\|_1 \leq B_1 = s$.
- * Then the Rademacher complexity (and thus the generalization error) is $O\left(\frac{s\sqrt{\log d}}{\sqrt{n}}\right)$.

- * Interpretation: essentially the number of relevant features (s) controls the complexity, and we can have a ton of irrelevant features (an exponentially large number).
 - * In contrast, if we use L_2 regularization, we would have $B_2 = \sqrt{s}$ and $C_2 = \sqrt{d}$. The Rademacher complexity of $\{w : \|w\|_2 \leq B_2\}$ is $O\left(\frac{s\sqrt{d/s}}{\sqrt{n}}\right)$.
 - * When $s \ll d$, then L_1 regularization is desirable, but if $s = d$, then L_1 regularization is worse by a factor of $\sqrt{\log d}$.
 - * Note that these are heuristic arguments since we are comparing upper bounds. In this case, the heuristics are accurate, but in general, one should exercise caution.
- From hypothesis class to loss class (for binary classification)
 - So far, we have bounded the Rademacher complexity of various hypothesis classes \mathcal{F} . We still need to turn that into a bound on the loss class \mathcal{A} .
 - * Recall that the zero-one loss on a data point is $\ell((x, y), w) = \mathbb{I}[yx \cdot w \leq 0]$, where $x \in \mathbb{R}^d$, $y \in \{-1, +1\}$.
 - * Recall the hypothesis class that we've bounded: $\mathcal{F} = \{z \mapsto w \cdot z : \|w\|_2 \leq B_2\}$.
 - * To leverage this result, think of each data point as $z = yx$, so that each function in \mathcal{F} maps $z = xy$ to the margin that we get on that data point $m = yx \cdot w$.
 - * Then the loss class can simply be expressed as function composition: $\mathcal{A} = \phi \circ \mathcal{F}$, where $\phi(m) = \mathbb{I}[m \leq 0]$ maps the margin m to the zero-one loss. However, the problem is that we can't apply the composition directly because ϕ is not Lipschitz due to the discontinuity at 0.
 - Our strategy will be to instead introduce a surrogate loss function, the **truncated hinge loss**, which upper bounds the zero-one loss:

$$\phi_{\text{surrogate}}(m) = \min(1, \max(0, 1 - m)). \quad (259)$$

Similarly, define the loss class $\mathcal{A}_{\text{surrogate}} = \phi_{\text{surrogate}} \circ \mathcal{F}$ and expected risk $L_{\text{surrogate}}(h)$ in terms of the surrogate loss. Note that the Lipschitz constant of $\phi_{\text{surrogate}}$ is 1, so we can apply the Lipschitz composition rule:

- By applying the Lipschitz composition rule:

$$R_n(\mathcal{A}_{\text{surrogate}}) \leq R_n(\mathcal{F}). \quad (260)$$

- Using the fact that the surrogate is an upper bound and plugging into (217), we obtain:

$$L(\hat{h}) \leq L_{\text{surrogate}}(\hat{h}) \leq L_{\text{surrogate}}(h^*) + 4R_n(\mathcal{F}) + \sqrt{\frac{2 \log(2/\delta)}{n}}. \quad (261)$$

Note that on the LHS, we have the zero-one loss of the ERM and on the RHS, we have the surrogate loss of the minimizer of expected risk. This says that if there exists a classifier that obtains low expected surrogate risk (by classifying correctly with large margin), then one should expect low expected risk $L(\hat{h})$ with a norm-constrained set of functions.

– Remarks

- * These bounds either do not depend (L_2 constrained) or only depend weakly (L_1 constrained) on the dimensionality d . This supports the prevailing wisdom that it doesn't matter how many features you have (how big d is). As long as you regularize properly (constrain the norm of the weights), then you will still have good generalization.
- * Compare this bound to the regret bound (58) we got in online learning. Note that this essentially gives the same result, but with an extra term for high probability guarantees, despite the fact that the analysis techniques were quite different.
- * We make no assumptions about convexity here (only Lipschitz is needed). This allows us to use the truncated hinge-loss rather than the zero-one loss, which is a tighter approximation to the zero-one loss (of course, computation with non-convex loss functions is hard).

3.12 Covering numbers (metric entropy) (Lecture 10)

- Motivation

- We can measure the complexity of a finite hypothesis class simply by computing its cardinality. For infinite hypothesis classes, we observed that shattering coefficient was an appropriate measure since (thanks to symmetrization) all that mattered was the behavior of a function class on a finite set of points. However, shattering coefficient only works for functions that return a finite number of values. Can we retain the combinatorial nature of shattering coefficients, but allow for real-valued functions (for example, to handle regression problems)?
- The key measure we will explore in this section is covering numbers, which counts the number of balls of size ϵ one needs to cover the hypothesis class.
- We can use the Massart's finite lemma to control the representatives; the behavior of the other functions is controlled by virtue of being close to some representative. In essence, covering numbers allows us to discretize the problem.

To talk about closeness of functions, we introduce the general notion of a metric space:

- **Definition 15 (metric space)**

- A metric space (\mathcal{F}, ρ) is defined over a set \mathcal{F} with equipped with a metric ρ .

- A metric $\rho : \mathcal{F} \times \mathcal{F} \rightarrow \mathbb{R}$ must be non-negative, symmetric, satisfy the triangle inequality, and evaluate to 0 iff its arguments are equal.
- If $\rho(f, f') = 0$ is possible for $f \neq f'$, then we say ρ is a pseudometric. In this section, technically we will be working with the pseudometric.
- Examples of metrics ρ
 - Euclidean distance
 - * For real vectors $\mathcal{F} = \mathbb{R}^d$.
 - * The metric is $\rho(f, f') = \|f - f'\|_2$.
 - $L_2(P_n)$
 - * This is the L_2 distance with respect to the empirical distribution over n data points: $P_n = \frac{1}{n} \sum_{i=1}^n \delta_{z_i}$.
 - * Let \mathcal{F} be a family of functions mapping \mathcal{Z} to \mathbb{R} .
 - * The metric is $\rho(f, f') = \left(\frac{1}{n} \sum_{i=1}^n (f(z_i) - f'(z_i))^2 \right)^{\frac{1}{2}}$.
 - * Remark: this metric can be thought of as computing an empirical standard deviation over differences between f and f' .
 - * Remark: since we are restricting the functions to evaluations at n points, we can really think of these functions as n -dimensional vectors, with a metric which is Euclidean distance scaled down by \sqrt{n} .
 - * FIGURE: [two functions and n points]

• **Definition 16 (ball)**

- Let (\mathcal{F}, ρ) be a metric space.
- Then the ball with radius $\epsilon > 0$ centered at $f \in \mathcal{F}$ is defined as:

$$B_\epsilon(f) \stackrel{\text{def}}{=} \{f' \in \mathcal{F} : \rho(f, f') \leq \epsilon\}. \quad (262)$$

• **Definition 17 (covering number)**

- An **ϵ -cover** of a set \mathcal{F} with respect to a metric ρ is a finite subset $C = \{f_1, \dots, f_m\} \subseteq \mathcal{F}$ such that if we put a ball of radius ϵ at each f_j , the resulting union is a superset of \mathcal{F} ; that is $\mathcal{F} \subseteq \cup_{j=1}^m B_\epsilon(f_j)$.
- Equivalently, every point in \mathcal{F} is at most distance ϵ away (under the metric ρ) from some point f_j in the cover C .
- Define the **ϵ -covering number** of \mathcal{F} with respect to ρ to be the size of the smallest cover:

$$N(\epsilon, \mathcal{F}, \rho) \stackrel{\text{def}}{=} \min\{m : \exists \{f_1, \dots, f_m\} \subseteq \mathcal{F} \subseteq \cup_{j=1}^m B_\epsilon(f_j)\}. \quad (263)$$

- The **metric entropy** of \mathcal{F} is $\log N(\epsilon, \mathcal{F}, \rho)$.
- FIGURE: [covering a set with balls]
- As ϵ decreases, the points f_j in the cover are a better approximation of the points in $B_\epsilon(f_j)$, but $N(\epsilon, \mathcal{F}, \rho)$ also increases. What is this tradeoff?
- **Example 19 (all functions)**
 - Let \mathcal{F} be all functions from $\mathcal{Z} = \mathbb{R}$ to $[0, 1]$.
 - Recall that using the metric $\rho = L_2(P_n)$ means that only function evaluations on the points z_1, \dots, z_n matter.
 - In order to cover \mathcal{F} , fix any $f \in \mathcal{F}$. Our strategy is to construct some $g \in \mathcal{F}$ such that $\rho(f, g) \leq \epsilon$ and count the number of g 's we obtain as we sweep f across \mathcal{F} .
 - For each point z_i , we cover the range $[0, 1]$ with the set of discrete points $Y = \{2\epsilon, 4\epsilon, \dots, 1\}$. For any value of $f(z_i) \in [0, 1]$, we can pick $g(z_i) \in Y$ so that $|f(z_i) - g(z_i)| \leq \epsilon$. The value of $g(z)$ for z not equal to any of the n points can be chosen arbitrarily. Summing over all z_i , we get that $\rho(f, g) \leq \epsilon$.
 - Furthermore, $|Y| = \frac{1}{2\epsilon}$, so

$$N(\epsilon, \mathcal{F}, L_2(P_n)) \leq \left(\frac{1}{2\epsilon}\right)^n. \quad (264)$$

The metric entropy is thus $O(n \log(1/\epsilon))$, which intuitively is too large. To see this, even if we ignored the discretization error ϵ for the moment, we would see that the empirical Rademacher complexity of the cover, by Massart's finite lemma is $O(\sqrt{\frac{n \log(1/\epsilon)}{n}}) = O(1)$, which does not go to zero. Clearly this class of functions is too large. Let's consider something smaller but still pretty rich:

- **Example 20 (non-decreasing functions)**
 - Let \mathcal{F} be all non-decreasing functions from $\mathcal{Z} = \mathbb{R}$ to $[0, 1]$.
 - Recall that z_1, \dots, z_n are the n fixed point, assumed to be sorted in increasing order.
 - Similar to before, we break up the range $[0, 1]$ into $\frac{1}{\epsilon}$ levels $Y = \{\epsilon, 2\epsilon, \dots, 1\}$.
 - Fix any function $f \in \mathcal{F}$. We will construct a piecewise constant approximation g for which $\rho(f, g) \leq \epsilon$.
 - For each level $y \in Y$, consider the points z_i for which $f(z_i) \in [y - \epsilon, y]$. Set $g(z_i) = y$ for these points. By construction, $\rho(f, g) \leq \epsilon$.

- Now let's count the number of possible g 's there are. For each of the $|Y| = 1/\epsilon$ levels, we choose one of n points to serve as the leftmost z_i for which $f(z_i) \in [y - \epsilon, y]$. Therefore

$$N(\epsilon, \mathcal{F}, L_2(P_n)) = O(n^{1/\epsilon}). \quad (265)$$

The metric entropy is $O((\log n)/\epsilon)$, which is much better than that of all functions. The main difference is that we're choosing a point for each level rather than a level for each point.

Having set up covering numbers with some examples, we are ready to state our main theorems. Specifically, we will show that the metric entropy leads to an upper bound on the Rademacher complexity. This can be accomplished two ways:

- First, we will show a simple discretization argument that applies covering numbers at a single resolution ϵ .
- Next, we provide a more sophisticated argument called **chaining** that adaptively chooses the resolution, yielding tighter bounds.

- **Theorem 16 (simple discretization)**

- Let \mathcal{F} be a family of functions mapping \mathcal{Z} to $[-1, 1]$.
- The empirical Rademacher complexity of a function class \mathcal{F} can be upper bounded using its covering number:

$$\hat{R}_n(\mathcal{F}) \leq \inf_{\epsilon > 0} \left(\sqrt{\frac{2 \log N(\epsilon, \mathcal{F}, L_2(P_n))}{n}} + \epsilon \right). \quad (266)$$

- Remark: the RHS involves two terms
 - * The first is the covering number, which increases as ϵ decreases. This comes from Massart's finite lemma.
 - * The second is ϵ (which decreases as ϵ decreases). This is the penalty we pay for having a discretization.

- **Preparation**

- We will also assume $Z_{1:n}$ are constant and suppress the explicit conditioning, writing $\mathbb{E}[A]$ instead of $\mathbb{E}[A \mid Z_{1:n}]$.
- To simplify notation, we write $\|f\|$ for $\|f\|_{L_2(P_n)}$ and $\langle f, g \rangle$ for $\langle f, g \rangle_{L_2(P_n)}$.
- Overloading notation, let $\sigma : \mathcal{Z} \rightarrow \{-1, +1\}$ be defined as a function which when evaluated on z_i returns the random sign σ_i . This allows us to write the empirical Rademacher complexity succinctly as:

$$\hat{R}_n(\mathcal{F}) = \mathbb{E} \left[\sup_{f \in \mathcal{F}} \langle \sigma, f \rangle \right]. \quad (267)$$

– Note that $\|\sigma\| = 1$ since it's just a vector of +1s and -1s.

• **Proof of Theorem 16:**

– Fix $\epsilon > 0$ and let C be an ϵ -cover of \mathcal{F} .

– The proof follows from manipulating the empirical Rademacher complexity:

$$\hat{R}_n(\mathcal{F}) = \mathbb{E} \left[\sup_{f \in \mathcal{F}} \langle \sigma, f \rangle \right] \quad (268)$$

$$= \mathbb{E} \left[\sup_{g \in C} \sup_{f \in \mathcal{F} \cap B_\epsilon(g)} \langle \sigma, g \rangle + \langle \sigma, f - g \rangle \right] \quad (269)$$

$$= \mathbb{E} \left[\sup_{g \in C} \frac{1}{n} \langle \sigma, g \rangle + \epsilon \right] \quad [\text{Cauchy-Schwartz}] \quad (270)$$

$$= \sqrt{\frac{2N(\epsilon, \mathcal{F}, L_2(P_n))}{n}} + \epsilon \quad [\text{Massart's finite lemma}]. \quad (271)$$

• **Example 21 (non-decreasing functions (with simple discretization))**

– Let \mathcal{F} be all non-decreasing functions from $\mathcal{Z} = \mathbb{R}$ to $[0, 1]$.

– Plugging the covering number of \mathcal{F} into Theorem 16:

$$\hat{R}_n(\mathcal{F}) \leq \inf_{\epsilon > 0} \left(\sqrt{\frac{2 \cdot O\left(\frac{\log n}{\epsilon}\right)}{n}} + \epsilon \right). \quad (272)$$

The RHS is minimized when the two terms are equal. Solving for ϵ and substituting it back yields:

$$\hat{R}_n(\mathcal{F}) = O \left(\left(\frac{\log n}{n} \right)^{\frac{1}{3}} \right). \quad (273)$$

– Note that this bound provides a rate which is worse than the usual $\frac{1}{\sqrt{n}}$ rates. Is it because non-decreasing functions are just more complex than parametric function classes? Not in this case. We'll see shortly that this is just an artifact of the analysis being too weak, because it is only able to work at one level of resolution ϵ . We will now introduce a clever technique called chaining which fixes this problem:

• **Theorem 17 (Chaining (Dudley's theorem))**

– Let \mathcal{F} be a family of functions mapping \mathcal{Z} to $[-1, 1]$.

- The empirical Rademacher complexity of a function class \mathcal{F} can be upper bounded using an integral over its covering number:

$$\hat{R}_n(\mathcal{F}) \leq 12 \int_0^\infty \sqrt{\frac{\log N(\epsilon, \mathcal{F}, L_2(P_n))}{n}} d\epsilon. \quad (274)$$

- Remark: note that compared with Theorem 16, this bound involves an integral that sweeps across different resolutions ϵ , and importantly removes the additive ϵ penalty.

- Proof of Theorem 17

- FIGURE: [lines corresponding to levels, where each level is discretized more finely than the previous]
- FIGURE: [plot showing $N(\epsilon, \mathcal{F}, \rho)$ as a function of ϵ , with locations of $\epsilon_0, \epsilon_1, \dots$]
- The key is to work at multiple levels of resolution.
- Let $\epsilon_0 = \sup_{f \in \mathcal{F}} \|f\|$ be the maximum norm of a function $f \in \mathcal{F}$, which is the coarsest resolution, and let $\epsilon_j = 2^{-j}\epsilon_0$ for $j = 1, \dots, m$ be successively finer resolutions.
- For each $j = 0, \dots, m$, let C_j be an ϵ_j -cover of \mathcal{F} .
- Fix any $f \in \mathcal{F}$.
- Let $g_j \in C_j$ be such that $\|f - g_j\| \leq \epsilon_j$; take $g_0 = 0$. Note that g_j 's depend on f .
- Let us decompose f as follows:

$$f = f - g_m + \underbrace{g_0}_{=0} + \sum_{j=1}^m (g_j - g_{j-1}). \quad (275)$$

- Restating Massart's finite lemma, if we have $\|f\| \leq M$ for all $f \in \mathcal{F}$, then $\hat{R}_n(\mathcal{F}) \leq M \sqrt{\frac{2 \log |\mathcal{F}|}{n}}$.
- Let us bound some norms:
 - * $\|f - g_m\| \leq \epsilon_m$
 - * $\|g_j - g_{j-1}\| \leq \|g_j - f\| + \|f - g_{j-1}\| \leq \epsilon_j + \epsilon_{j-1} = 3\epsilon_j$ (since $2\epsilon_j = \epsilon_{j-1}$)

– Now compute the empirical Rademacher complexity:

$$\hat{R}_n(\mathcal{F}) = \mathbb{E} \left[\sup_{f \in \mathcal{F}} \langle \sigma, f \rangle \right] \quad [\text{definition}] \quad (276)$$

$$= \mathbb{E} \left[\sup_{f \in \mathcal{F}} \langle \sigma, f - g_m \rangle + \sum_{j=1}^m \langle \sigma, g_j - g_{j-1} \rangle \right] \quad [\text{decompose } f] \quad (277)$$

$$\leq \epsilon_m + \mathbb{E} \left[\sup_{f \in \mathcal{F}} \sum_{j=1}^m \langle \sigma, g_j - g_{j-1} \rangle \right] \quad [\text{Cauchy-Schwartz}] \quad (278)$$

$$\leq \epsilon_m + \sum_{j=1}^m \mathbb{E} \left[\sup_{f \in \mathcal{F}} \langle \sigma, g_j - g_{j-1} \rangle \right] \quad [\text{push sup inside}] \quad (279)$$

$$\leq \epsilon_m + \sum_{j=1}^m \mathbb{E} \left[\sup_{g_j \in C_j, g_{j-1} \in C_{j-1}} \langle \sigma, g_j - g_{j-1} \rangle \right] \quad [\text{refine dependence}] \quad (280)$$

$$\leq \epsilon_m + \sum_{j=1}^m (3\epsilon_j) \sqrt{\frac{2 \log(|C_j| |C_{j-1}|)}{n}} \quad [\text{Massart's finite lemma}] \quad (281)$$

$$\leq \epsilon_m + \sum_{j=1}^m (6\epsilon_j) \sqrt{\frac{\log |C_j|}{n}} \quad [\text{since } |C_j| \geq |C_{j-1}|] \quad (282)$$

$$= \epsilon_m + \sum_{j=1}^m 12(\epsilon_j - \epsilon_{j+1}) \sqrt{\frac{\log |C_j|}{n}} \quad [\text{since } \epsilon_j = 2(\epsilon_j - \epsilon_{j+1})] \quad (283)$$

$$\leq 12 \int_0^\infty \sqrt{\frac{\log N(\epsilon, \mathcal{F}, L_2(P_n))}{n}} d\epsilon \quad [\text{bound sum with integral}] \quad (284)$$

In the last step, we took $m \rightarrow \infty$, which makes the additive penalty $\epsilon_m \rightarrow 0$.

- Remark: The technical reason why chaining provides better results is the following. In simple discretization, we apply Massart's finite lemma on functions whose range is $[-1, 1]$, whereas in chaining, we apply Massart's finite lemma on functions C_j whose range has size $O(\epsilon_j)$; this leads to a scaling of ϵ_j .
- Remark: when j increases by one, ϵ_j is multiplied by $\frac{1}{2}$, but for some examples (such as the next one), the integrand, because it's under the square root, will only be multiplied by $\sqrt{2}$, not 2. So we win.

• **Example 22 (non-decreasing functions (with chaining))**

- Let \mathcal{F} be all non-decreasing functions from \mathcal{Z} to $[0, 1]$.
- Note that $\|f\| \leq 1$ for all $f \in \mathcal{F}$, so the coarsest resolution is $\epsilon_0 = 1$, so we only have to integrate ϵ up to 1, not ∞ , since $\log N(\epsilon, \mathcal{F}, L_2(P_n)) = 0$ for $\epsilon \geq \epsilon_0$.

- Plugging the covering number of \mathcal{F} into Theorem 17:

$$\hat{R}_n(\mathcal{F}) \leq 12 \int_0^1 \left(\sqrt{\frac{O(\frac{\log n}{\epsilon})}{n}} \right) d\epsilon \quad (285)$$

$$= O \left(\sqrt{\frac{\log n}{n}} \right) \int_0^1 \sqrt{\frac{1}{\epsilon}} d\epsilon \quad (286)$$

$$= O \left(\sqrt{\frac{\log n}{n}} \right). \quad (287)$$

- Remark: compared with the bound using the simple discretization, we get a better bound ($n^{-\frac{1}{2}}$ versus $n^{-\frac{1}{3}}$).

- Summary

- Covering numbers is a powerful technique that allows us to handle rich function classes such as all non-decreasing functions.
- The essence is discretization of a function class, so that we can use Massart's finite lemma.
- The discretization resolution ϵ can be set via chaining to yield tight bounds.
- Covering numbers give you a lot of freedom for customization: choice of metric, choice of discretization.
- Covering numbers can sometimes be more convenient; for example, the covering number of $\mathcal{F}_1 \cup \mathcal{F}_2$ is at most that of \mathcal{F}_1 plus that of \mathcal{F}_2 , a property not shared by Rademacher complexity.

3.13 Algorithmic stability (Lecture 11)

- Motivation

- Let us shelve excess risk $L(\hat{h}) - L(h^*)$ for the moment, and instead consider the gap between the expected and empirical risk. It is easy to see that this quantity can be upper bounded using uniform convergence:

$$\mathbb{P}[L(\hat{h}) - \hat{L}(\hat{h}) \geq \epsilon] \leq \mathbb{P}[\sup_{h \in H} L(h) - \hat{L}(h) \geq \epsilon]. \quad (288)$$

Colloquially, this answers the following question: if I get 10% error on my training set, what should I expect my test error to be? (The answer is no more than $10\% + \epsilon$.) Analyzing the excess risk has relied on \hat{h} being the ERM, but (288) actually holds for *any* estimator \hat{h} . It is useful to think about \hat{h} as a function of the training examples: $\hat{h} = A(z_1, \dots, z_n)$, where A is an *algorithm*.

- Uniform convergence hones in on studying the “size” of \mathcal{H} (using Rademacher complexity, VC dimension, etc.). However, what if a learning algorithm A does not “use” all of \mathcal{H} ? For example, what if you use naive Bayes, regularization via a penalty term, early stopping, or dropout training? All of these could in principle can return any hypothesis from \mathcal{H} , but are somehow constrained by their objective function or algorithm, so we might expect better generalization than simply looking at the complexity of \mathcal{H} in an algorithm-independent way.
- To be more concrete, let us analyze $\mathcal{H} = \{x \mapsto w \cdot x : w \in \mathbb{R}^d\}$, the class of linear functions (with no bounds on the norm). Define the norm $\|h\|_{\mathcal{H}}$ to be $\|w\|_2^2$ of the associated weight vector. Define the **regularized empirical risk minimizer** as follows:

$$\hat{h} = \arg \min_{h \in \mathcal{H}} \hat{L}(h) + \frac{\lambda}{2} \|h\|_{\mathcal{H}}^2. \quad (289)$$

- Of course, for a given realization of the data $z_{1:n}$, the solution to (289) can be gotten by finding the ERM on the constrained set $\{h \in \mathcal{H} : \|h\|_{\mathcal{H}} \leq B\}$ for some bound B . But this connection is weak in that B would need to depend on the data to obtain an exact equivalence. How can we analyze (289) directly?
- In this section, we will introduce bounds based on a notion of **algorithmic stability**, where we view a learning algorithm as a function A that take as input the data $z_{1:n}$ and outputs some hypothesis \hat{h} . We will study the generalization properties of \hat{h} as a function of how stable the algorithm is. Notice the focus on algorithms (similar to in online learning) rather than a focus on the hypothesis class.

- Stability will be measured with respect to the difference in behavior of an algorithm on a training set S and a perturbed version S^i :

- $S = (z_1, \dots, z_n)$: training data, drawn i.i.d. from p^*
- $S^i = (z_1, \dots, z'_i, \dots, z_n)$: training data with an i.i.d. copy of the i -th example
- z_0 is a new test example

We start with a definition of stability:

– **Definition 18 (uniform stability)**

- * We say that an algorithm $A : \mathcal{Z}^n \rightarrow \mathcal{H}$ has **uniform stability** β (or in the context of these notes, simply β -stable) with respect to a loss function ℓ if for all training sets $S \in \mathcal{Z}^n$, perturbations $S^i \in \mathcal{Z}^n$, and test example $z_0 \in \mathcal{Z}$,

$$|\ell(z_0, A(S)) - \ell(z_0, A(S^i))| \leq \beta. \quad (290)$$

- * Note that this is a very strong condition in that the bound must hold uniformly for all z_0, S, S^i and is not reliant on the probability distribution.

– **Example 23 (stability of mean estimation)**

- * Assume all points $z \in \mathbb{R}^d$ are bounded $\|z\|_2 \leq B$.
- * Define the squared loss: $\ell(z, h) = \frac{1}{2}\|z - h\|_2^2$.
- * Define an algorithm that computes the regularized empirical risk minimizer:

$$A(S) \stackrel{\text{def}}{=} \arg \min_{h \in \mathbb{R}^d} \hat{L}(h) + \frac{\lambda}{2}\|h\|_2^2 \quad (291)$$

$$= \frac{1}{(1 + \lambda)n} \sum_{i=1}^n z_i. \quad (292)$$

- * Then A has uniform stability $\beta = \frac{6B^2}{(1 + \lambda)n}$.

- * To derive this, define $v_1 = A(S) - z_0$ and $v_2 = \frac{1}{(1 + \lambda)n}[z'_i - z_i]$.

$$|\ell(z_0, A(S)) - \ell(z_0, A(S^i))| \leq \left| \frac{1}{2}\|v_1\|_2^2 - \frac{1}{2}\|v_1 + v_2\|_2^2 \right| \quad (293)$$

$$\leq \|v_2\|_2(\|v_1\|_2 + \frac{1}{2}\|v_2\|_2) \quad (294)$$

$$\leq \frac{2B}{(1 + \lambda)n}(2B + B). \quad (295)$$

- * We can see that as we increase regularization (larger λ) or amount of data (larger n), we have more stability (smaller β).

– **Example 24 (stability of linear predictors)**

- * Let $\mathcal{H} = \{x \mapsto w \cdot x : w \in \mathbb{R}^d\}$.
- * Assume that $\|x\|_2 \leq C_2$ with probability 1 (according to the data-generating distribution p^*).
- * Assume the loss ℓ is 1-Lipschitz: for all $z_0 \in \mathcal{Z}$ and $h, h' \in \mathcal{H}$:

$$|\ell(z_0, h) - \ell(z_0, h')| \leq \|h - h'\|_\infty \stackrel{\text{def}}{=} \sup_{x \in \mathbb{R}^d} |h(x) - h'(x)|. \quad (296)$$

For example, for classification ($y \in \{-1, +1\}$), this holds for the hinge loss $\ell((x, y), h) = \max\{1 - yh(x), 0\}$.

- * Define the regularized empirical risk minimizer:

$$A(S) \stackrel{\text{def}}{=} \arg \min_{h \in \mathcal{H}} \hat{L}(h) + \frac{\lambda}{2} \|h\|_{\mathcal{H}}^2 \quad (297)$$

- * Then A has uniform stability $\boxed{\beta = \frac{C_2^2}{\lambda n}}$ with respect to ℓ . See the Bousquet/Elisseeff paper on stability for the proof.

And now for the main theorem:

– **Theorem 18 (generalization under uniform stability)**

- * Let A be an algorithm with uniform stability β .
- * Assume the loss is bounded: $\sup_{z, h} |\ell(z, h)| \leq M$.
- * Then with probability at least $1 - \delta$,

$$\boxed{L(A(S)) \leq \hat{L}(A(S)) + \beta + (\beta n + M) \sqrt{\frac{2 \log(1/\delta)}{n}}}. \quad (298)$$

- * Remark: Due to the presence of βn , for this bound to be not vacuous, we must have $\beta = o\left(\frac{1}{\sqrt{n}}\right)$. Generally, we will have $\beta = O\left(\frac{1}{n}\right)$, which guarantees that $L(A(S)) - \hat{L}(A(S)) = O\left(\frac{1}{\sqrt{n}}\right)$.

– Proof of Theorem 18:

- * Our goal is to bound the difference between the expected and empirical risks:

$$D(S) \stackrel{\text{def}}{=} L(A(S)) - \hat{L}(A(S)). \quad (299)$$

- * Step 1: Bound the expectation of $D(S)$.
 - Recall that $S = (z_1, \dots, z_n), (z'_1, \dots, z'_n), z_0$ are all independent and therefore exchangeable. The basic trick is just to rename variables that preserve the dependency structure and therefore the expectation $\mathbb{E}[D(S)]$,

and get it into a form where we can apply the definition of uniform stability:

$$\mathbb{E}[D(S)] = \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n [\ell(z_0, A(S)) - \ell(z_i, A(S))] \right] \quad [\text{definition}] \quad (300)$$

$$= \mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n [\ell(z'_i, A(S)) - \ell(z'_i, A(S^i))] \right] \quad [\text{renaming}] \quad (301)$$

$$\leq \beta \quad [\text{definition of uniform stability}] \quad (302)$$

The point is that in the first term z_0 is not in S and z_i is in S . This logic is preserved: z'_i is not in S and z'_i is in S^i .

- * Step 2: show that $D(S)$ satisfies the bounded differences property.
 - We should expect such a property to hold given the definition of uniform stability. But we are not directly applying the bounded differences to the loss $\ell(z_0, A(S))$, but rather to $D(S)$. So there is slightly more work. The trick here is to break down differences using the triangle inequality into a chain of comparisons, each involving a single perturbation.
 - Let \hat{L}^i denote the empirical expectation with respect to S^i .
 - We have

$$|D(S) - D(S^i)| \quad (303)$$

$$\leq |L(A(S)) - \hat{L}(A(S)) - L(A(S^i)) + \hat{L}^i(A(S^i))| \quad (304)$$

$$\leq |L(A(S)) - L(A(S^i))| + |\hat{L}(A(S)) - \hat{L}^i(A(S^i))| \quad [\text{triangle inequality}] \quad (305)$$

$$\leq \underbrace{|L(A(S)) - L(A(S^i))|}_{\leq \beta} + \underbrace{|\hat{L}(A(S)) - \hat{L}(A(S^i))|}_{\leq \beta} + \underbrace{|\hat{L}(A(S^i)) - \hat{L}^i(A(S^i))|}_{\leq \frac{2M}{n}} \quad (306)$$

$$\leq 2\beta + \frac{2M}{n}. \quad (307)$$

In the first two cases, we just used the fact that A is β -stable; recall that $\hat{L}(h) = \frac{1}{n} \sum_{i=1}^n \ell(z_i, h)$ and $L(h) = \mathbb{E}_{z_0 \sim p^*} [\ell(z_0, h)]$; here it's important that β -stability holds uniformly for any first argument of ℓ , so that we can upper bound by β regardless of the dependence structure between the two arguments. In the third case, \hat{L} and \hat{L}^i just differ by one term which can differ by at most $2M$ and is scaled by $\frac{1}{n}$.

- * Step 3: apply McDiarmid's inequality to bound $D(S)$ in high probability.

• This is a straightforward application. We have

$$\mathbb{P}[D(S)) - \mathbb{E}[D(S)] \geq \epsilon] \leq \exp \left(\frac{-2\epsilon^2}{n(2\beta + \frac{2M}{n})^2} \right) \quad (308)$$

$$= \exp \left(\frac{-n\epsilon^2}{2(\beta n + M)^2} \right) \stackrel{\text{def}}{=} \delta. \quad (309)$$

Rewriting the bound and using the fact that $\mathbb{E}[D(S)] \leq \beta$, we have that with probability at least $1 - \delta$,

$$D(S) \leq \beta + (\beta n + M) \sqrt{\frac{2 \log(1/\delta)}{n}}. \quad (310)$$

– Application to linear functions

- * Recall that for regularized ERM on linear functions with 1-Lipschitz losses (like hinge loss), $\beta = \frac{C_2^2}{\lambda n}$, where λ is the regularization strength and $\|x\|_2 \leq C_2$. Plugging this value of β into Theorem 18, we get that with probability at least $1 - \delta$,

$$L(A(S)) \leq \hat{L}(A(S)) + O \left(\frac{C_2^2 + M}{\lambda \sqrt{n}} \right). \quad (311)$$

Note that this bound has the right dependence on n , but has a worse dependence on C_2 compared to the Rademacher bounds.

3.14 PAC-Bayesian bounds (Lecture 11)

- Motivation

- Everything we’ve been doing is under a frequentist paradigm, where we assume an unknown $h^* \in \mathcal{H}$ against which we assess our performance. We assume the worst case over h^* (defined via p^*).
- A Bayesian would instead start with a prior $P(h)$, observe the training data z_1, \dots, z_n , and produce a posterior $Q(h)$. Bayesian procedures are optimal assuming the prior (and the model) is “correct”, but this is in practice too strong of an assumption.
- From one perspective, a Bayesian procedure is just an algorithm that returns some hypothesis \hat{h} (perhaps stochastically) given z_1, \dots, z_n , and we already have analyses that work on any algorithm (via uniform convergence or algorithmic stability). However, the analysis itself is still worst case over h^* . Can we *incorporate the prior into the analysis itself while not assuming that the prior is correct*?
- PAC-Bayesian bounds provide exactly this. In this class, we will consider two types:

- * Occam bound: assume a countable hypothesis class, algorithm outputs a single hypothesis
- * PAC-Bayesian theorem: assume an infinite hypothesis class, algorithm outputs a posterior

Let us start with Occam bound, which captures the key intuitions:

- **Theorem 19 (Occam bound)**

- Let \mathcal{H} be a countable hypothesis class.
- Let the loss function be bounded: $\ell(z, h) \in [0, 1]$.
- Let P be any “prior” distribution over \mathcal{H} .
- Then with probability at least $1 - \delta$,

$$\boxed{\forall h \in \mathcal{H} : L(h) \leq \hat{L}(h) + \sqrt{\frac{\log(1/P(h)) + \log(1/\delta)}{2n}}.} \quad (312)$$

- Let’s try to understanding this bound with a few special cases.
 - If the prior puts all its mass on one hypothesis h_0 ($P(h) = \mathbb{I}[h = h_0]$), then the bound is just the standard Hoeffding bound you would get for a single hypothesis.
 - If we have a uniform distribution over some subset of hypotheses $S \subseteq \mathcal{H}$ ($P(h) = \mathbb{I}[h \in S]/|S|$), then we recover the standard result for finite hypotheses (similar to Theorem 11, which is for excess risk).
 - This reveals how PAC-Bayes is a generalization: rather than committing to prior distributions which are uniform over some support $S \subseteq \mathcal{H}$, we can have prior distributions which place different probability mass on different hypotheses. We can let \mathcal{H} be as large as we want as long as we still have probabilities ($\sum_{h \in \mathcal{H}} P(h) = 1$). In some sense, $P(h)$ defines a fuzzy hypothesis class.
- The bound also suggests an algorithm.
 - First, note that the penalty term $\log(1/P(h))$ is tailored towards the particular hypothesis h , whereas in our previous bounds, we were looking at the complexity of all of \mathcal{H} . This dependence on h presents a new possibility: that of treating the RHS bound as an objective to be minimized by an algorithm. Recall that this bound holds simultaneously for all $h \in \mathcal{H}$, which means that it will hold for the output of an algorithm $A(S)$.
 - Motivated by the bound, we can define an *algorithm that actually uses the bound* by minimizing the RHS:

$$A(S) \stackrel{\text{def}}{=} \arg \min_{h \in \mathcal{H}} \hat{L}(h) + R(h), \quad R(h) = \sqrt{\frac{\log(1/P(h)) + \log(1/\delta)}{2n}}. \quad (313)$$

From this perspective, the bound provides a regularizer $R(h)$ which penalizes h more if it has small prior probability $P(h)$. The algorithm $A(S)$ thus balances the empirical risk $\hat{L}(h)$ with the regularizer $R(h)$.

- As we get more data ($n \rightarrow \infty$), the regularizer also goes to zero, meaning that we will trust the empirical risk more, allowing us to consider more complex hypotheses in \mathcal{H} . This is a pretty nice behavior to have.

- **Proof of Theorem 19:**

- The proof is very simple. The key idea is to allocate our confidence parameter δ across different hypotheses proportionally based on the prior $P(h)$.
- By Hoeffding’s inequality, we have that for any fixed $h \in \mathcal{H}$:

$$\mathbb{P}[L(h) \geq \hat{L}(h) + \epsilon] \leq \exp(-2n\epsilon^2). \quad (314)$$

- With probability at most $\delta P(h)$,

$$L(h) \geq \hat{L}(h) + \sqrt{\frac{\log(1/P(h)) + \log(1/\delta)}{2n}}. \quad (315)$$

- Applying the union bound across all $h \in \mathcal{H}$, we have that with probability at most δ ,

$$\exists h \in \mathcal{H} : L(h) \geq \hat{L}(h) + \sqrt{\frac{\log(1/P(h)) + \log(1/\delta)}{2n}}. \quad (316)$$

- There are two things lacking about the Occam bound:

- It only applies to countable \mathcal{H} , which does not include the set of all weight vectors, for example.
- It only embraces half of the Bayesian story: while we have a prior $P(h)$, only a single $h \in \mathcal{H}$ is returned rather than a full posterior $Q(h)$.

The following theorem generalizes the Occam bound:

- **Theorem 20 (PAC-Bayesian theorem)**

- Let the loss function be bounded: $\ell(z, h) \in [0, 1]$.
- Let P be any “prior” distribution over \mathcal{H} .
- Let Q_S be any “posterior” distribution over \mathcal{H} , which is a function of the training data S .
- Then with probability at least $1 - \delta$,

$$\mathbb{E}_{h \sim Q_S}[L(h)] \leq \mathbb{E}_{h \sim Q_S}[\hat{L}(h)] + \sqrt{\frac{\text{KL}(Q_S \| P) + \log(4n/\delta)}{2n - 1}}. \quad (317)$$

- See the McAllester paper for the proof.
- To recover the Occam bound (up to constants), we simply set Q_S to be a delta function at some h .

3.15 Interpretation of bounds (Lecture 11)

- Now that we've derived a whole host of generalization bounds, let us take a step back and ask the question: how should we think about these bounds?
- Properties
 - One could evaluate these bounds numerically, but they will probably be too loose to use directly. The primary purpose of these bounds is to instead formalize the relevant properties of a learning problem and characterize their relationship to the generalization error, the quantity of interest.
 - The relationships solidify intuitions about learning. Here are some examples:
 - * If we have d features, $n \sim d$ training examples suffices to learn. If the number of features increase by 2, we need to increase n by 2 as well to maintain the same estimation error.
 - * We can actually have as many features d as we want (even infinite), so long as we regularize properly using L_2 regularization: bounds depend on norm B not dimension d .
 - * If there are many irrelevant features use L_1 regularization: the L_1 ball is just much smaller than the L_2 ball. Here, exploiting the structure of the problem leads to better bounds (and algorithms).
 - * If there is low noise in the problem (in the realizable setting, some predictor obtains zero generalization error), then estimation error is smaller ($O(1/n)$ versus $O(1/\sqrt{n})$ convergence).
- Focus on estimation error
 - It is important to note that generalization bounds focus on addressing the estimation error (excess risk) $L(\hat{h}) - L(h^*)$, not the approximation error $L(h^*)$.
 - For example, if d is the number of features, then $L(\hat{h}) - L(h^*) = O(\frac{d}{n})$ shows that by adding more features, the estimation error will worsen. However, we hope that $L(h^*)$ will improve.
 - In practice, one can still hope to reduce $L(\hat{h})$ by adding additional features.
 - The technical core of this section is **concentration of measure**: as you aggregate over increasing amounts of **independent** data, many of the relevant quantities converge. Insight is obtained by closely observing how fast these quantities converge.

- Loss function
 - Is the bound on the right measure of error?
 - The bounds derived using finite hypothesis classes or finite VC dimension operated on the zero-one loss (supposing for the moment that's our desired loss function). However, the empirical risk minimizer in this case is NP hard to compute.
 - However, the norm-based bounds using Rademacher complexity required a Lipschitz loss function such as the truncated hinge loss or the hinge loss, which is a surrogate loss. This gives us results on an empirical risk minimizer which we can actually evaluate in practice. One can say is that the zero-one loss is upper bounded by the hinge loss, but this is relatively weak, and in particular, minimizing the hinge loss even in the limit of infinite data will not give you something that minimizes the zero-one loss. A special case is that for universal kernels, minimizing the hinge loss (among others) does correspond to minimizing the zero-one loss in the limit of infinite data (Bartlett/Jordan/McAuliffe, 2005).
 - Note that this distinction mirrors our online learning regret bounds as well: in the finite experts case, we obtained bounds on expected loss for arbitrary (bounded) loss functions, whereas in general, we needed convexity (which is even stronger than what we require here).

3.16 Summary (Lecture 11)

- The main focus of this section was to study the **excess risk** $L(\hat{h}) - L(h^*)$ of the **empirical risk minimizer** \hat{h} . In particular, we wanted that with probability at least $1 - \delta$, the excess risk $L(\hat{h}) - L(h^*)$ is upper bounded by something that depends on the complexity of the learning problem and n , the number of i.i.d. training examples. (Another goal is to bound the difference $L(\hat{h}) - \hat{L}(\hat{h})$.)
- The excess risk is often within a factor of two of the difference between empirical and expected risk: $\sup_{h \in \mathcal{H}} |L(h) - \hat{L}(h)|$, which has a form which suggests using uniform convergence tools to bound it.
- Results on excess risk $L(\hat{h}) - L(h^*)$:
 - Realizable, finite hypothesis class: $O\left(\frac{\log |\mathcal{H}|}{n}\right)$
 - Finite hypothesis class: $O\left(\sqrt{\frac{\log |\mathcal{H}|}{n}}\right)$
 - Shattering coefficient / VC dimension: $O\left(\sqrt{\frac{\log s(\mathcal{H}, n)}{n}}\right) = O\left(\sqrt{\frac{VC(\mathcal{H}) \log n}{n}}\right)$
 - L_2 norm constrained—kernels ($\|w\|_2 \leq B_2, \|x\|_2 \leq C_2$): $O\left(\frac{B_2 C_2}{\sqrt{n}}\right)$

- L_1 norm constrained—sparsity ($\|w\|_1 \leq B_1, \|x\|_\infty \leq C_\infty$): $O\left(\frac{B_1 C_\infty \sqrt{\log d}}{\sqrt{n}}\right)$
- Non-decreasing functions (via covering numbers and chaining): $O\left(\int_0^\infty \sqrt{\frac{\log N(\epsilon, \mathcal{H}, L_2(P_n))}{n}} d\epsilon\right) = O\left(\sqrt{\frac{\log n}{n}}\right)$

- Technical tools

- Tail bounds
 - * How much do random variables deviate from their mean? We generally look for sharp concentration bounds, which means that the probability of deviation by a constant decays **exponentially** fast as a function of n : $\mathbb{P}[G_n - \mathbb{E}[G_n] \geq \epsilon] \leq c^{-n\epsilon^2}$.
 - * When G_n is an average of i.i.d. **sub-Gaussian** variables Z_i we can bound the moment generating function and get the desired bound (**Hoeffding's inequality** for bounded random variables). Sub-Gaussian random variables include Gaussian and bounded random variables (it is convenient to assume our loss or data is bounded), but not Laplace distribution, which has heavier tails.
 - * When G_n is the result of applying a function with bounded differences to i.i.d. variables, then **McDiarmid's inequality** gives us the same bound.
- Complexity control
 - * For a single hypothesis, we can directly apply the tail bound to control the difference $L(h) - \hat{L}(h)$. However, we seek **uniform convergence** over all $h \in \mathcal{H}$.
 - * Finite hypothesis classes: $\log |\mathcal{H}|$ (use simple union bound)
 - * For infinite hypothesis classes, the key intuition is that the complexity of \mathcal{H} is described by **how it acts on n data points**. Formally, **symmetrization** (introduce ghost dataset and Rademacher variables) reveals this.
 - * The complexity is the **shattering coefficient** $s(\mathcal{H}, n)$ (technically of the loss class \mathcal{A}). By **Sauer's lemma**, the shattering coefficient can be upper bounded using the **VC dimension** $VC(\mathcal{H})$.
 - * Rademacher complexity $R_n(\mathcal{H})$ measures how well \mathcal{H} can fit random binary labelings (noise). Rademacher complexity is nice because of the numerous compositional properties (convex hull, Lipschitz composition, etc.)
 - By **Massart's finite lemma**, we can relate $R_n(\mathcal{H})$ to the shattering coefficient.
 - For L_2 norm constrained linear functions, we used linearity and Cauchy-Schwartz, which enables us to analyze kernels.
 - For L_1 norm constrained linear functions, we used the fact that the L_1 polytope is really simple as it only has $2d$ vertices.

- For neural networks, we leveraged the compositionality of Rademacher complexity.
- Other paradigms
 - We also studied two alternative paradigms for analyzing learning, both of which were motivated by the need for greater nuance. Let A be any learning algorithm that maps training data S to a hypothesis $\hat{h} \in \mathcal{H}$. The algorithm A could be the ERM, but it need not be.
 - Typical uniform convergence bounds yield results that depend on the complexity of the entire hypothesis class \mathcal{H} :

$$L(A(S)) - \hat{L}(A(S)) \leq \text{SomeFunction}(\mathcal{H}). \quad (318)$$

- Algorithmic stability allows us to obtain a bound that depends on properties of the algorithm A (i.e., its stability β) rather than \mathcal{H} . We obtained bounds of the form:

$$L(A(S)) - \hat{L}(A(S)) \leq \text{SomeFunction}(A). \quad (319)$$

- PAC-Bayesian bounds allow us to incorporate the prior into the analysis without sacrificing objective rigor. We obtained bounds of the form:

$$L(A(S)) - \hat{L}(A(S)) \leq \text{SomeFunction}(A(S)). \quad (320)$$

Note that the bound depends on the output of the algorithm.

3.17 References

- [Bousquet/Boucheron/Lugosi, 2008: Introduction to Statistical Learning Theory](#)
- [Martin Wainwright's lecture notes](#)
- [Peter Bartlett's lecture notes](#)
- [Sham Kakade's lecture notes](#)
- [Tomaso Poggio and Lorenzo Rosasco's lecture notes](#)
- [Bartlett/Mendelson, 2002: Rademacher and Gaussian Complexities: Risk Bounds and Structural Results](#)
- [Kakade/Sridharan/Tewari, 2008: On the Complexity of Linear Prediction: Risk Bounds, Margin Bounds, and Regularization](#)
- [Bosquet/Elisseeff, 2002: Stability and Generalization](#)
- [David McAllester's notes on PAC-Bayesian bounds](#)

4 Direct analysis

4.1 Introduction (Lecture 12)

- In the last unit, we used uniform convergence to study the generalization ability of learning algorithms, in particular, the empirical risk minimizer (ERM). Uniform convergence is like a tank. While we were able to obtain quite general results, this required a fairly heavy piece of machinery, and there was a sense in which we were losing touch with the particular properties of the data.
- In this unit, we are still in pursuit of strong generalization bounds, but we will do it using a completely different approach that is more tailored to the problem. First, we will start with **fixed design linear regression**, which should be a breath of fresh air. This problem is simple enough that we can do everything in closed form. As a result, we will obtain new intuitions such as bias-variance tradeoffs.
- Of course, for most problems other than linear regression, we won't have nice closed form expressions. We will turn to **asymptotic analysis**, which by only track first-order terms enables us to yield simple closed form expressions which are exact in the limit.

4.2 Fixed design linear regression (Lecture 12)

- Setup
 - Our goal is to predict a real-valued output (response) $y \in \mathbb{R}$ given an input (covariates/features) $x \in \mathbb{R}^d$.
 - The **fixed design** setting means that we have a fixed set of n inputs x_1, \dots, x_n , which are treated as constants.
 - We assume that there is a true underlying parameter vector $\theta^* \in \mathbb{R}^d$, which governs a set of outputs:

$$y_i = x_i \cdot \theta^* + \epsilon_i, \tag{321}$$

where we assume that the noise terms ϵ_i are i.i.d. with mean $\mathbb{E}[\epsilon_i] = 0$ and variance $\text{Var}[\epsilon_i] = \sigma^2$.

- At training time, we observe one realization of y_1, \dots, y_n . For convenience, let's put the data into matrices:
 - * $X = [x_1, \dots, x_n]^\top \in \mathbb{R}^{n \times d}$

- * $\epsilon = [\epsilon_1, \dots, \epsilon_n]^\top \in \mathbb{R}^d$
- * $Y = [y_1, \dots, y_n]^\top \in \mathbb{R}^d$
- * $\Sigma = \frac{1}{n} X^\top X \in \mathbb{R}^{d \times d}$ (second moment matrix)

- FIGURE: [linear regression line over x_1, \dots, x_n]
- Our goal is to minimize the expected risk as defined by the squared loss:

$$L(\theta) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n \mathbb{E}[(x_i \cdot \theta - y_i)^2] = \frac{1}{n} \mathbb{E}[\|X\theta - Y\|_2^2]. \quad (322)$$

Note that the expectation is over the randomness in Y (recall that X is fixed).

- One can think of the fixed design setting as performing signal reconstruction, where the x_i 's correspond to fixed locations, and y_i is a noisy sensor reading.
- The regularized least squares estimator is as follows:

$$\hat{\theta} \stackrel{\text{def}}{=} \arg \min_{\theta \in \mathbb{R}^d} \frac{1}{n} \|X\theta - Y\|_2^2 + \lambda \|\theta\|_2^2, \quad (323)$$

where $\lambda \geq 0$ is the **regularization strength**. By taking derivatives and setting the result to zero, we obtained the following closed form solution:

$$\hat{\theta} = (X^\top X + \lambda I)^{-1} X^\top Y = \frac{1}{n} \Sigma_\lambda^{-1} X^\top Y, \quad (324)$$

where we define $\Sigma_\lambda = \Sigma + \lambda I$.

- Our goal is to study the expected risk of the regularized least squares estimator (ridge regression):

$$L(\hat{\theta}). \quad (325)$$

More specifically, we will study the expectation $\mathbb{E}[L(\hat{\theta})]$, where the expectation is taken over the training data. We could obtain high probability statements as we did for all our generalization bounds using uniform convergence, but we will not do that here.

- Studying the expected risk

- As a first step, let us understand the expected risk $L(\theta)$. The goal will be to understand this quantity geometrically. The basic idea is to expand $Y = X\theta^* + \epsilon$;

the rest is just algebra. We have:

$$L(\theta) = \frac{1}{n} \mathbb{E}[\|X\theta - Y\|_2^2] \quad (326)$$

$$= \frac{1}{n} \mathbb{E}[\|X\theta - X\theta^* + \epsilon\|_2^2] \quad [\text{by definition of } Y \text{ (321)}] \quad (327)$$

$$= \frac{1}{n} \mathbb{E}[\|X\theta - X\theta^*\|_2^2 + \|\epsilon\|_2^2] \quad [\text{cross terms vanish}] \quad (328)$$

$$= \frac{1}{n} (\theta - \theta^*)^\top (X^\top X) (\theta - \theta^*) + \sigma^2 \quad [\text{algebra, definition of } \epsilon] \quad (329)$$

$$= \|\theta - \theta^*\|_\Sigma^2 + \sigma^2. \quad (330)$$

- Intuitively, the first term of the expected risk is the squared distance between the estimate θ and the true parameters θ^* as measured by the shape of the data. If the data does not vary much in one direction, then the discrepancy between θ and θ^* will be downweighted in that direction.
- The second term is the unavoidable variance term coming from the noise, which is present even with the optimal parameters θ^* . Note that $L(\theta^*) = \sigma^2$.
- In conclusion, the excess risk is:

$$\boxed{L(\theta) - L(\theta^*) = \|\theta - \theta^*\|_\Sigma^2.} \quad (331)$$

- Unregularized case

- Now, let us analyze the expected risk of the ERM $L(\hat{\theta})$ for the case where we don't regularize ($\lambda = 0$). Assume that $X^\top X \succ 0$ (which means necessarily that $n > d$). The key is to expand $\hat{\theta}$ and Y based on their definitions and perform algebra. The expectation is over the test data. The main term of the expected risk is:

$$\|\hat{\theta} - \theta^*\|_\Sigma^2 = \frac{1}{n} \|X\hat{\theta} - X\theta^*\|_2^2 \quad (332)$$

$$= \frac{1}{n} \|X(X^\top X)^{-1}X^\top(X\theta^* + \epsilon) - X\theta^*\|_2^2 \quad (333)$$

$$= \frac{1}{n} \|\Pi X\theta^* + \Pi\epsilon - X\theta^*\|_2^2 \quad [\text{projection } \Pi \stackrel{\text{def}}{=} X(X^\top X)^{-1}X^\top] \quad (334)$$

$$= \frac{1}{n} \|\Pi\epsilon\|_2^2 \quad [\text{projection doesn't change } X\theta^*, \text{ cancel}] \quad (335)$$

$$= \frac{1}{n} \text{tr}(\Pi\epsilon\epsilon^\top) \quad [\text{projection is idempotent and symmetric}]. \quad (336)$$

- Taking expectations (over the training data), and using the fact that $\mathbb{E}[\epsilon\epsilon^\top] = \sigma^2 I$, and subtracting off $L(\theta^*)$, we get:

$$\boxed{\mathbb{E}[L(\hat{\theta}) - L(\theta^*)] = \frac{d\sigma^2}{n}.} \quad (337)$$

Note that the complexity of the problem is completely determined by the dimensionality d and the variance of the noise σ^2 .

- Intuitively, the noise ϵ is an n -dimensional vector gets projected onto d dimensions by virtue of having to fit the data using a linear function with d degrees of freedom.

- Regularized case

- The above shows that when $d \ll n$, regularization isn't important. But what happens when $d \gg n$? Intuitively, we should regularize (take $\lambda > 0$), but how do we justify this in terms of the expected risk? And by how much?
- The first main insight is the **bias-variance tradeoff**, whose balance is determined by λ . Let us decompose the excess risk:

$$\mathbb{E}[\|\hat{\theta} - \theta^*\|_{\Sigma}^2] = \mathbb{E}[\|\hat{\theta} - \mathbb{E}[\hat{\theta}] + \mathbb{E}[\hat{\theta}] - \theta^*\|_{\Sigma}^2] \quad (338)$$

$$= \underbrace{\mathbb{E}[\|\hat{\theta} - \mathbb{E}[\hat{\theta}]\|_{\Sigma}^2]}_{\stackrel{\text{def}}{=} \text{Var}} + \underbrace{\|\mathbb{E}[\hat{\theta}] - \theta^*\|_{\Sigma}^2}_{\stackrel{\text{def}}{=} \text{Bias}}, \quad (339)$$

where the cross terms are designed to cancel out. Note that in the unregularized case ($\lambda = 0$), the bias is zero since $\mathbb{E}[\hat{\theta}] = (X^{\top}X)^{-1}X^{\top}(X\theta^* + \mathbb{E}[\epsilon]) = \theta^*$, but when $\lambda > 0$, the bias will be non-zero.

- The second main insight is that the risk of the regularized least squares estimate on the original problem is the same as the risk of an equivalent problem which has been rotated into a basis that is easier to analyze.
 - * Suppose we rotate each data point x_i by an orthogonal R ($x_i \mapsto R^{\top}x_i$), so that $X \mapsto XR$ for some orthogonal matrix R . Correspondingly, we must rotate the parameters $\theta^* \mapsto R^{\top}\theta^*$. Then the claim is that the excess risk does not change.
 - * The excess risk of the modified problem is:

$$\mathbb{E}[\|XR(R^{\top}X^{\top}XR + \lambda I)^{-1}R^{\top}X^{\top}(XRR^{\top}\theta^* + \epsilon) - XRR^{\top}\theta^*\|_2^2]. \quad (340)$$

Simplification reveals that we get back exactly the original excess risk:

$$\mathbb{E}[\|X(X^{\top}X + \lambda I)^{-1}X^{\top}(X\theta^* + \epsilon) - X\theta^*\|_2^2]. \quad (341)$$

- * If we take the SVD $X = USV^{\top}$ and rotate by $R = V$, then we can see that $X^{\top}X \mapsto (V^{\top}VSU^{\top})(USV^{\top}V) = S^2$, which is diagonal. Therefore, for the purposes of analysis, we can assume that Σ is diagonal without loss of generality:

$$\Sigma = \text{diag}(\tau_1, \dots, \tau_d). \quad (342)$$

- Let us compute the mean estimator:

$$\bar{\theta}_j \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}_j] \quad (343)$$

$$= \mathbb{E}[\Sigma_\lambda^{-1} n^{-1} X^\top (X\theta^* + \epsilon)]_j \quad (344)$$

$$= \mathbb{E}[\Sigma_\lambda^{-1} \Sigma \theta^* + \Sigma_\lambda^{-1} X^\top \epsilon]_j \quad (345)$$

$$= \frac{\tau_j}{\tau_j + \lambda} \theta_j^*. \quad (346)$$

Thus, the expected value of the estimator is the true parameter value θ^* shrunk towards zero by a strength that depends on λ .

- Compute the bias:

$$\text{Bias} = \|\bar{\theta} - \theta^*\|_\Sigma^2 \quad (347)$$

$$= \sum_{j=1}^d \left(\frac{\tau_j}{\tau_j + \lambda} \theta_j^* - \theta_j^* \right)^2 \quad (348)$$

$$= \sum_{j=1}^d \frac{\tau_j \lambda^2 (\theta_j^*)^2}{(\tau_j + \lambda)^2}. \quad (349)$$

If $\lambda = 0$, then the bias is zero as expected.

- Compute the variance:

$$\text{Var} = \mathbb{E}[\|\hat{\theta} - \bar{\theta}\|_\Sigma^2] \quad (350)$$

$$= \mathbb{E}[\|\Sigma_\lambda^{-1} n^{-1} X^\top \epsilon\|_\Sigma^2] \quad [\text{recall definition of } \hat{\theta}] \quad (351)$$

$$= \frac{1}{n^2} \mathbb{E}[\epsilon^\top X \Sigma_\lambda^{-1} \Sigma \Sigma_\lambda^{-1} X^\top \epsilon] \quad (352)$$

$$= \frac{1}{n^2} \text{tr}(\Sigma_\lambda^{-1} \Sigma \Sigma_\lambda^{-1} X^\top \mathbb{E}[\epsilon \epsilon^\top] X) \quad (353)$$

$$= \frac{\sigma^2}{n} \text{tr}(\Sigma_\lambda^{-1} \Sigma \Sigma_\lambda^{-1} \Sigma) \quad (354)$$

$$= \frac{\sigma^2}{n} \sum_{j=1}^d \frac{\tau_j^2}{(\tau_j + \lambda)^2}. \quad (355)$$

If we didn't regularize, the variance would be $\frac{d\sigma^2}{n}$. Regularization clearly reduces the variance.

- Now let us balance the two terms. Fact: $(a + b)^2 \geq 2ab$.
- Bound the bias:

$$\text{Bias} \leq \sum_{j=1}^d \frac{\lambda (\theta_j^*)^2}{2} = \frac{\lambda \|\theta^*\|_2^2}{2} \quad (356)$$

$$\text{Var} \leq \sum_{j=1}^d \frac{\tau_j \frac{\sigma^2}{n}}{2\lambda} = \frac{\text{tr}(\Sigma) \sigma^2}{2n\lambda}. \quad (357)$$

- Optimizing yields an bound on the excess risk:

$$\mathbb{E}[L(\hat{\theta}) - L(\theta^*)] \leq \sqrt{\frac{\|\theta^*\|_2^2 \text{tr}(\Sigma) \sigma^2}{n}}, \quad (358)$$

with the regularization parameter set to:

$$\lambda = \sqrt{\frac{\text{tr}(\Sigma) \sigma^2}{\|\theta^*\|_2^2 n}}. \quad (359)$$

- Remarks

- * As $\lambda \rightarrow 0$, the bias vanishes and we are left with $\frac{d\sigma^2}{n}$.
- * We can do better in cases where $d > n$. In particular, d could be essentially infinite. The true dimensionality is the sum of the eigenvalues ($\text{tr}(\Sigma)$). So if the data can be described (via PCA) by a few dimensions, then we don't need to regularize much to get a dependence on that effective dimension.
- * However, the excess risk is now decaying at a rate of $O(\sqrt{\frac{1}{n}})$ rather than $O(\frac{1}{n})$.
- * This bound is similar to the one we got for online learning. If $\|\theta^*\|_2 \leq B$ and $\|x_i\|_2 \leq L$ (which means $\text{tr}(\Sigma) \leq L$), then we get a excess risk (analogous to regret) of $\frac{BL}{\sqrt{n}}$.
- * In the random design where X is a random variable, things are more complicated. In that setting, $X^\top X$ would be random and different from the expected covariance $\mathbb{E}[X^\top X]$. We would need to argue that the two converge. See the Hsu/Kakade/Zhang paper for more details.

4.3 Finite-dimensional asymptotics (Lecture 13)

- Motivation

- In the previous section, the simplicity of the regularized least squares problem in the fixed design setting allowed us to compute everything exactly. What happens if we are in the random design setting or if we wanted to handle other loss functions (e.g., logistic loss)? We can't hope to compute the excess risk $L(\hat{\theta}) - L(\theta^*)$ exactly, and we might lose too much if we try to derive an upper bound.
- In this unit, we will show that the excess risk approaches a simple quantity as the number of data points n goes to infinity by performing **Taylor expansions** around θ^* . In brief, we will see that $\hat{\theta} - \theta^*$ is approximately Gaussian with some variance that is $O(\frac{1}{n})$, and assuming that L is continuous, we can convert this result into one about the expected risk.

* FIGURE: $[\hat{\theta}$ in a ellipsoid Gaussian ball around θ^*]

- How large does n have to be for the asymptotic results to hold? While our analysis doesn't give a quantitative handle on this, asymptotics can provide valuable insight into the problem. In practice for even small n , asymptotic approximations can be more accurate than upper bounds since we need not be strict about enforcing one-sided errors.

- Comparing two estimators

- One compelling use of asymptotics arises when we are comparing estimators. Suppose we have two algorithms (estimators) that return hypotheses $\hat{\theta}_1$ and $\hat{\theta}_2$. Which one is better?
- Let $L(\hat{\theta}_1; p^*)$ and $L(\hat{\theta}_2; p^*)$ denote the expected risks of the two, where we've made the dependence on the data generating distribution $p^* \in \mathcal{P}$ explicit (\mathcal{P} is the family of distributions we're considering).
- We can derive upper bounds on the generalization error, but it's invalid to just compare these two upper bounds, because discrepancies could reflect our ability to prove bounds rather than real phenomena.
- We could try to derive lower bounds (which usually analyze the worst case), which is harder but doable:

$$\text{Lower}_i \leq \sup_{p^* \in \mathcal{P}} L(\hat{\theta}_i; p^*) \leq \text{Upper}_i \quad \text{for } i \in \{1, 2\}. \quad (360)$$

To show $\hat{\theta}_1$ is better than $\hat{\theta}_2$, we would need to show that $\text{Upper}_1 \leq \text{Lower}_2$. However, these lower bounds are obtained on the worst case p^* , and moreover,

the p^* might be different between $\hat{\theta}_1$ and $\hat{\theta}_2$. So this comparison might be quite crude.

- What we really want is control

$$L(\hat{\theta}_1; p^*) - L(\hat{\theta}_2; p^*) \quad (361)$$

for various values of $p^* \in \mathcal{P}$.

- We will show that asymptotics can give us a handle on this. The main idea is to approximate $L(\hat{\theta}_i; p^*)$ directly in terms of some function of p^* plus some lower order terms in n .

$$L(\hat{\theta}_i; p^*) \cong \text{Approx}_i(p^*). \quad (362)$$

We can compute $\text{Approx}_1(p^*) - \text{Approx}_2(p^*)$, which converges to the correct difference as $n \rightarrow \infty$ for any $p^* \in \mathcal{P}$.

- Probability refresher

- Let X_1, \dots, X_n be real vectors drawn i.i.d. from some distribution with mean μ and covariance matrix Σ .
- Let $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$.
- Convergence in probability
 - * Example: $\hat{\mu} \xrightarrow{P} \mu$ (weak law of large numbers)
- Convergence in distribution
 - * Example: $\sqrt{n}(\hat{\mu} - \mu) \xrightarrow{d} \mathcal{N}(0, \Sigma)$ (central limit theorem)
- Continuous mapping theorem
 - * If f is continuous and $Y_n \xrightarrow{P} Y$, then $f(Y_n) \xrightarrow{P} f(Y)$
 - * Example: $\|\hat{\mu}\|_2 \xrightarrow{P} \|\mu\|_2$
- Slutsky's theorem
 - * If $Y_n \xrightarrow{P} c$ and $Z_n \xrightarrow{d} Z$ then $Y_n Z_n \xrightarrow{d} cZ$.
 - * Example: $\sqrt{n}\hat{\mu} \cdot (\hat{\mu} - \mu) \xrightarrow{d} \mathcal{N}(0, \mu^\top \Sigma \mu)$.
- Notation: $X_n = O_p(f(n))$ if $X_n/f(n)$ is bounded in probability, that is, for every $\epsilon > 0$, there exists M_ϵ such that for all n , $\mathbb{P}[|X_n/f(n)| > M_\epsilon] < \epsilon$.
 - * Example: $\mathcal{N}(0, \Sigma) = O_p(1)$
 - * Example: $\hat{\mu} - \mu = O_p\left(\frac{1}{\sqrt{n}}\right)$
- These results turn complicated things on the LHS into simple things on the RHS. The continuous mapping theorem and Slutsky's theorem allow you to compose these results in simple ways. Generally, just use your intuitions from real analysis.

- Setup

- $z = (x, y)$ is an example
- $\ell(z, \theta)$ is a loss function on example z with parameters $\theta \in \mathbb{R}^d$
 - * Example: $\ell((x, y), \theta) = \frac{1}{2}(\theta \cdot x - y)^2$
- Let p^* be the true probability distribution over examples $z \in \mathcal{Z}$.
- Let $\theta^* \in \mathbb{R}^d$ be the minimizer of the expected risk:

$$\theta^* \stackrel{\text{def}}{=} \arg \min_{\theta \in \mathbb{R}^d} L(\theta), \quad L(\theta) \stackrel{\text{def}}{=} \mathbb{E}_{z \sim p^*} [\ell(z, \theta)] \quad (363)$$

- Let $\hat{\theta} \in \mathbb{R}^d$ be the minimizer of the empirical risk:

$$\hat{\theta} \stackrel{\text{def}}{=} \arg \min_{\theta \in \mathbb{R}^d} \hat{L}(\theta), \quad \hat{L}(\theta) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n \ell(z^{(i)}, \theta), \quad (364)$$

where $z^{(1)}, \dots, z^{(n)}$ are drawn i.i.d. from p^* .

- Recall that we are interested in studying the excess risk $L(\hat{\theta}) - L(\theta^*)$.

- Assumptions on the loss

- Loss function $\ell(z, \theta)$ is twice differentiable in θ (works for squared and logistic loss, but not hinge)
- Let $\nabla \ell(z, \theta) \in \mathbb{R}^d$ be the gradient of the loss at θ .
 - * Example: $\nabla \ell(z, \theta) = (\theta \cdot x - y)x$
- Let $\nabla^2 \ell(z, \theta) \in \mathbb{R}^{d \times d}$ be the Hessian of the loss at θ .
 - * Example: $\nabla^2 \ell(z, \theta) = xx^\top$
- Assume that the expected loss Hessian $\mathbb{E}_{z \sim p^*} [\nabla^2 \ell(z, \theta)] \succ 0$ is positive definite for all $\theta \in \mathbb{R}^d$. This assumption is actually not needed, but it will make the math simpler.¹³

- Outline of analysis

- Step 0 (**consistency**): show that $\hat{\theta} \xrightarrow{P} \theta^*$. This is obtained by a uniform convergence argument to show that \hat{L} approximates L well. Then, since the Hessian $\mathbb{E}[\nabla^2 \ell(z, \theta)]$ is positive definite, minimizing \hat{L} will eventually minimize L . We will not dwell on this point.

¹³ In fact, if the expected loss is rank deficient, things are actually even better, since the complexity will depend on the rank of the Hessian rather than the dimensionality.

- Step 1: obtain an asymptotic expression for the **parameter error** by Taylor expanding the gradient of the empirical risk.

$$\boxed{\hat{\theta} - \theta^*} \quad (365)$$

- Step 2: obtain an asymptotic expression for the **excess risk** by Taylor expanding the expected risk.

$$\boxed{L(\hat{\theta}) - L(\theta^*)} \quad (366)$$

• **Definition 19 (well-specified model)**

- Assume that the loss function corresponds to the log-likelihood under a probabilistic model p_θ :

$$\ell((x, y); \theta) = -\log p_\theta(y | x), \quad (367)$$

so that $\hat{\theta}$ is the (conditional) maximum likelihood estimate under this model.

- We say that this model family $\{p_\theta\}_{\theta \in \mathbb{R}^d}$ is **conditionally well-specified** if $p^*(x, y) = p^*(x)p_{\theta^*}(y | x)$ for some parameter $\theta^* \in \mathbb{R}^d$.
- Suppose each model θ actually specifies a joint distribution over both x and y : $p_\theta(x, y)$. We say that this model family $\{p_\theta\}_{\theta \in \mathbb{R}^d}$ is **jointly well-specified** if $p^*(x, y) = p_{\theta^*}(x, y)$ for some parameter $\theta^* \in \mathbb{R}^d$. This places a much stronger assumption on the data generating distribution. If x is an image and y is a single binary label, this is much harder to satisfy.
- Of course, jointly well-specified implies conditionally well-specified.
- In the conditionally well-specified case, the Bartlett identity allows us to equate the variance of the risk gradient with the risk Hessian. This quantity is the **Fisher information** matrix (or rather, a generalization of it).

• **Theorem 21 (Bartlett identity)**

- In the well-specified case (conditionally, and thus also jointly), the following identity holds:

$$\boxed{\nabla^2 L(\theta^*) = \text{Var}(\nabla \ell(z, \theta^*))} \quad (368)$$

• **Proof of Theorem 21:**

- Recall that $z = (x, y)$.

- Using the fact that probability densities integrate to one:

$$\int p^*(x) \underbrace{e^{-\ell(z, \theta^*)}}_{p_{\theta^*}(y|x)} dz = 1. \quad (369)$$

- Assuming regularity conditions, differentiate with respect to θ^* :

$$\int p^*(x) e^{-\ell(z, \theta^*)} (-\nabla \ell(z, \theta^*)) dz = 0. \quad (370)$$

Note that this implies $\mathbb{E}[\nabla \ell(z, \theta^*)] = 0$, which shouldn't be surprising since θ^* minimizes $L(\theta) = \mathbb{E}[\ell(z, \theta^*)]$.

- Differentiating again, using the product rule:

$$\int p^*(x) [-e^{-\ell(z, \theta^*)} \nabla^2 \ell(z, \theta^*) + e^{-\ell(z, \theta^*)} \nabla \ell(z, \theta^*) \nabla \ell(z, \theta^*)^\top] dz = 0. \quad (371)$$

- Re-arranging:

$$\mathbb{E}[\nabla^2 \ell(z, \theta^*)] = \mathbb{E}[\nabla \ell(z, \theta^*) \nabla \ell(z, \theta^*)^\top]. \quad (372)$$

- Using the fact that $\mathbb{E}[\nabla \ell(z, \theta^*)] = 0$ and the definition of $L(\theta)$ yields the result.

- Remark: our general analysis does not assume the model is well-specified. We will only make this assumption for examples to get simple expressions for intuition.

- **Example 25 (well-specified linear regression)**

- Assume that the conditional model is as follows:

- * $x \sim p^*(x)$ for some arbitrary $p^*(x)$
- * $y = \theta^* \cdot x + \epsilon$, where $\epsilon \sim \mathcal{N}(0, 1)$

- Loss function: $\ell((x, y), \theta) = \frac{1}{2}(\theta \cdot x - y)^2$ (the log-likelihood up to constants)

- Check that the following two are equal:

- * Hessian: $\nabla^2 L(\theta) = \mathbb{E}[xx^\top]$ (covariance matrix of data)
- * Variance of gradient: $\text{Var}(\nabla \ell(z, \theta)) = \mathbb{E}[\epsilon x x^\top \epsilon] = \mathbb{E}[xx^\top]$, where the last equality follows from independence of x and ϵ

Now let us analyze the parameter error (step 1) and expected risk (step 2) in the general (not necessarily well-specified) case. In the following, pay attention how fast each of the terms is going to zero.

- Step 1: analysis of **parameter error**

- Since ℓ is twice differentiable, we can perform a Taylor expansion of the gradient of the empirical risk ($\nabla \hat{L}$) around θ^* :

$$\nabla \hat{L}(\hat{\theta}) = \nabla \hat{L}(\theta^*) + \nabla^2 \hat{L}(\theta^*)(\hat{\theta} - \theta^*) + O_p(\|\hat{\theta} - \theta^*\|_2^2). \quad (373)$$

- Using the fact that the LHS $\nabla \hat{L}(\hat{\theta}) = 0$ (by optimality conditions of the empirical risk minimizer) and rearranging:

$$\hat{\theta} - \theta^* = -\nabla^2 \hat{L}(\theta^*)^{-1} \left(\nabla \hat{L}(\theta^*) + O_p(\|\hat{\theta} - \theta^*\|_2^2) \right). \quad (374)$$

- As $n \rightarrow \infty$:

- * By the weak law of large numbers, we have

$$\nabla^2 \hat{L}(\theta^*) \xrightarrow{P} \nabla^2 L(\theta^*). \quad (375)$$

Since the inverse is a smooth function around θ^* (we assumed $\nabla^2 L(\theta^*) \succ 0$), we can apply the continuous mapping theorem:

$$\nabla^2 \hat{L}(\theta^*)^{-1} \xrightarrow{P} \nabla^2 L(\theta^*)^{-1}. \quad (376)$$

- * $\hat{L}(\theta^*)$ is a sum of mean zero i.i.d. variables, so by the central limit theorem, $\sqrt{n} \cdot \nabla \hat{L}(\theta^*)$ converges in distribution:

$$\sqrt{n} \cdot \nabla \hat{L}(\theta^*) \xrightarrow{d} \mathcal{N}(0, \text{Var}(\nabla \ell(z, \theta^*))). \quad (377)$$

An intuitive implication of this result is that $\hat{L}(\theta^*) = O_p\left(\frac{1}{\sqrt{n}}\right)$.

- * Suppose $\hat{\theta} - \theta^* = O_p(f(n))$. By (374), $f(n)$ goes to zero at a rate which is the maximum of $\frac{1}{\sqrt{n}}$ and $f(n)^2$. This implies that $f(n) = \frac{1}{\sqrt{n}}$, so we have:

$$\sqrt{n} \cdot O_p(\|\hat{\theta} - \theta^*\|_2^2) \xrightarrow{P} 0. \quad (378)$$

- By Slutsky's theorem, we can substitute the limits into (374) to obtain:

$$\boxed{\sqrt{n} \cdot \underbrace{(\hat{\theta} - \theta^*)}_{\text{parameter error}} \xrightarrow{d} \mathcal{N}\left(0, \nabla^2 L(\theta^*)^{-1} \text{Var}(\nabla \ell(z, \theta^*)) \nabla^2 L(\theta^*)^{-1}\right)}, \quad (379)$$

where we used the fact that if $x_n \xrightarrow{d} \mathcal{N}(0, \Sigma)$, then $Ax_n \xrightarrow{d} \mathcal{N}(0, A\Sigma A^\top)$. This also establishes that the parameter error behaves $\hat{\theta} - \theta^* = O_p\left(\frac{1}{\sqrt{n}}\right)$ as expected.

- * $\nabla^2 L(\theta^*)$: measures the amount of **curvature** in the loss function at θ^* . The more there is, the more stable the parameter estimates are.

- * $\text{Var}(\nabla \ell(z, \theta^*))$: measures the **variance** in the loss gradient. The less there is, the better.
- * When $\sqrt{n}(\hat{\theta} - \theta^*) \xrightarrow{d} \mathcal{N}(0, \Sigma)$, Σ is known as the **asymptotic variance** of the estimator $\hat{\theta}$.

– **Example 26 (well-specified linear regression)**

- * In this case, due to (368), we have $\text{Var}(\nabla \ell(z, \theta^*)) = \mathbb{E}[\nabla^2 \ell(z, \theta^*)] = \mathbb{E}[xx^\top]$, so the variance factor is canceled out by one of the curvature factors.

$$\boxed{\sqrt{n} \cdot (\hat{\theta} - \theta^*) \xrightarrow{d} \mathcal{N}(0, \mathbb{E}[xx^\top]^{-1})}. \quad (380)$$

Intuition: the larger x is, the more stable the parameter estimates; think about wiggling a pencil by either holding it with two hands out at the ends (large x) or near the center (small x).

• Step 2: analysis of excess risk

- Perform a Taylor expansion of the expected risk around θ^* :

$$L(\hat{\theta}) = L(\theta^*) + \nabla L(\theta^*)^\top (\hat{\theta} - \theta^*) + \frac{1}{2}(\hat{\theta} - \theta^*)^\top \nabla^2 L(\theta^*)(\hat{\theta} - \theta^*) + O_p(\|\hat{\theta} - \theta^*\|_2^3). \quad (381)$$

- By optimality conditions of the expected risk minimizer θ^* , we have $\nabla L(\theta^*) = 0$. This is the key to getting $O(\frac{1}{n})$ rates of convergence.
- Multiplying by n and rearranging:

$$n(L(\hat{\theta}) - L(\theta^*)) = \frac{1}{2}\sqrt{n}(\hat{\theta} - \theta^*)^\top \nabla^2 L(\theta^*)\sqrt{n}(\hat{\theta} - \theta^*) + O_p(n\|\hat{\theta} - \theta^*\|_2^3). \quad (382)$$

- Substituting in the parameter error:

* Facts

- If $x_n \xrightarrow{d} \mathcal{N}(0, \Sigma)$, then $x_n x_n^\top \xrightarrow{d} \mathcal{W}(\Sigma, 1)$, where $\mathcal{W}(\Sigma, 1)$ is a Wishart distribution with mean Σ and 1 degree of freedom.
- Taking the trace of both sides, we have that $x_n^\top x_n = \text{tr}(x_n x_n^\top) \xrightarrow{d} \text{tr}(\mathcal{W}(\Sigma, 1))$.¹⁴
- The distribution on the RHS is a weighted sum of d chi-squared distributed variables, whose distribution is the same as $\sum_{j=1}^d \Sigma_{jj} v_j^2$, where $v_j \sim \mathcal{N}(0, 1)$ is a standard Gaussian and $v_j^2 \sim \chi_1^2$ is a chi-squared.

¹⁴We are abusing notation slightly by writing the trace of a distribution D to mean the distribution of the trace of $x \sim D$.

* In our context, let us define

$$x_n = \sqrt{n}(\nabla^2 L(\theta^*))^{\frac{1}{2}}(\hat{\theta} - \theta^*). \quad (383)$$

Then

$$\Sigma = \nabla^2 L(\theta^*)^{-\frac{1}{2}} \text{Var}(\nabla \ell(z, \theta^*)) \nabla^2 L(\theta^*)^{-\frac{1}{2}}. \quad (384)$$

Therefore:

$$\boxed{n(L(\hat{\theta}) - L(\theta^*)) \xrightarrow{d} \frac{1}{2} \text{tr } \mathcal{W} \left(\nabla^2 L(\theta^*)^{-\frac{1}{2}} \text{Var}(\nabla \ell(z, \theta^*)) \nabla^2 L(\theta^*)^{-\frac{1}{2}}, 1 \right)}, \quad (385)$$

where $\mathcal{W}(V, n)$ is the Wishart distribution with scale matrix V and n degrees of freedom.

– **Example 27 (well-specified models)**

* Since the model is well-specified, everything cancels nicely, resulting in:

$$\boxed{n(L(\hat{\theta}) - L(\theta^*)) \xrightarrow{d} \frac{1}{2} \text{tr } \mathcal{W}(I_{d \times d}, 1)}. \quad (386)$$

- * The limiting distribution is half times a χ_d^2 distributed random variable, which has mean $\frac{d}{2}$ and variance d .
- * To get an idea of their behavior, we can compute the mean and variance:
 - Mean: $\mathbb{E}[n(L(\hat{\theta}) - L(\theta^*))] \rightarrow \frac{d}{2}$.
 - Variance: $\text{Var}(n(L(\hat{\theta}) - L(\theta^*))) \rightarrow d$.

In short,

$$\boxed{L(\hat{\theta}) - L(\theta^*) \sim \frac{d}{2n}}. \quad (387)$$

Note that this recovers the result from fixed-design linear regression (337) (the factor of $\frac{1}{2}$ is due to defining the loss function with an extra $\frac{1}{2}$).

- * Interestingly, in the well-specified case, the expected risk does not depend asymptotically on any properties of x .
 - For parameter estimation, the more x varies, the more accurate the parameter estimates.
 - For prediction, the more x varies, the harder the prediction problem.
 - The two forces cancel each other out exactly (asymptotically).

• **Remarks**

- For this brief section, suppose the model is well-specified.

- We have shown that regret $L(\hat{\theta}) - L(\theta^*)$ is exactly $\frac{d}{2n}$ asymptotically; we emphasize that there are no hidden constants and this is equality, not just a bound.
- Lower-order terms
 - * Of course there could be more error lurking in the lower order $(\frac{1}{n^2})$ terms.
 - * For linear regression, the low-order terms $O_p(\|\hat{\theta} - \theta^*\|_2^2)$ in the Taylor expansion are actually zero, and the only approximation comes from estimating the second moment matrix $\mathbb{E}[xx^\top]$.
 - * For the fixed design linear regression setting, $\nabla^2 \hat{L}(\theta^*) = \nabla^2 L(\theta^*)$, so all lower-order terms are identically zero, and so our asymptotic expressions are exact.
- Norm/regularization
 - * We only obtained results in the unregularized case. Asymptotically as $n \rightarrow \infty$ and the dimension d is held constant, the optimal thing to do (up to first order) is not use regularization.
 - * So these results are only meaningful when n is large compared to the complexity (dimensionality) of the problem. This is consistent with the fact that the type of analyses we do are fundamental local around the optimum.
 - * What about constraining the norm? This also doesn't do anything since if the norm constraint even slightly too big, then locally it is not tight.
- Let us compare the asymptotic expression (386) with results that we've derived previously in this class.
 - Using uniform convergence, we were only able to get a $\frac{1}{\sqrt{n}}$ convergence rate for the unrealizable setting. This was unavoidable using our techniques since we relied on concentration of empirical risk to expected risk, which even for a single hypothesis is already $\frac{1}{\sqrt{n}}$.
 - How did we get a faster rate? While asymptotics gives us the correct rate (as long as $n \rightarrow \infty$ while all other quantities such as dimensionality remain constant), there are non-asymptotic effects hidden away in the $O_p\left(n^{-\frac{3}{2}}\right)$ terms. capturing different aspects regarding the complexity of the learning problem. For instance, $\frac{1}{\sqrt{n}}$ is the rate associated with the norm, while $\frac{1}{n}$ is associated with the dimensionality.
 - In online learning, we were able to get a $\frac{\log n}{n}$ bound for strongly convex loss functions. However, strong convexity is too much to ask for, since any linear model will not satisfy this. In the asymptotic setting, we only need strong convexity in expectation (averaged over data points) at θ^* (remember, the entire analysis operates locally around θ^*). It is also possible to analyze online learning under statistical assumptions and obtain bounds that depend on strong convexity in expectation.

- Comparison of generative versus discriminative models (skipped in lecture)
 - Recall one of the main motivations of asymptotic analysis was that we could compare different estimators.
 - In this section, let's assume that the model is *jointly* well-specified and is an exponential family (includes logistic regression, conditional random fields, MRFs, etc):

$$p_{\theta}(x, y) = \exp(\phi(x, y)^{\top} \theta - A(\theta)), \quad (388)$$

where

- * $\phi(x, y) \in \mathbb{R}^d$ is the feature vector,
 - * $\theta \in \mathbb{R}^d$ are the parameters, and
 - * $A(\theta) = \log \int_{\mathcal{X} \times \mathcal{Y}} \exp(\phi(x, y)^{\top} \theta) dx dy$ is the joint log-partition function.
 - * $A(\theta; x) = \log \int_{\mathcal{Y}} \exp(\phi(x, y)^{\top} \theta) dy$ is the conditional log-partition function (useful later).
- We consider two estimators which are used to train:
 - * Generative: $\ell_{\text{gen}}((x, y), \theta) = -\log p_{\theta}(x, y)$ defines estimator $\hat{\theta}_{\text{gen}}$
 - * Discriminative: $\ell_{\text{dis}}((x, y), \theta) = -\log p_{\theta}(y | x)$ defines estimator $\hat{\theta}_{\text{dis}}$

Here, we are being careful to define the estimators with respect to the same model, but only changing the estimator as to pin down the underlying essence between generative and discriminative estimation.

- Important: note that we are using different loss functions at training time, although at test time, we still evaluate using the discriminative loss ℓ_{dis} .
- Recall that the asymptotic variance of the estimators $\hat{\theta} - \theta^*$ are functions of $\nabla^2 \ell(z; \theta)^{-1}$.
- For exponential families, the derivatives are simply the moments of the distributions:
 - * $L_{\text{gen}}(\theta^*) = \mathbb{E}[\nabla^2 \ell_{\text{gen}}((x, y), \theta^*)] = \nabla^2 A(\theta^*) = \text{Var}(\phi(x, y))$.
 - * $L_{\text{dis}}(\theta^*) = \mathbb{E}[\nabla^2 \ell_{\text{dis}}((x, y), \theta^*)] = \mathbb{E}[\nabla^2 A(\theta; x)] = \mathbb{E}[\text{Var}(\phi(x, y) | x)]$.
- Key variance decomposition identity:

$$\text{Var}(\phi(x, y)) = \mathbb{E}[\text{Var}(\phi(x, y) | x)] + \text{Var}(\mathbb{E}[\phi(x, y) | x]). \quad (389)$$

Since variance matrices are PSD, we have that

$$\text{Var}(\phi(x, y)) \succeq \mathbb{E}[\text{Var}(\phi(x, y) | x)]. \quad (390)$$

Inverting:

$$\underbrace{\text{Var}(\phi(x, y))}_{\text{asymptotic variance of } \hat{\theta}_{\text{gen}} - \theta^*} \preceq \underbrace{\mathbb{E}[\text{Var}(\phi(x, y) | x)]}_{\text{asymptotic variance of } \hat{\theta}_{\text{dis}} - \theta^*}. \quad (391)$$

This says that the asymptotic variance of generative estimator ($\hat{\theta}_{\text{gen}}$) to be at most the asymptotic variance of the discriminative estimator ($\hat{\theta}_{\text{dis}}$).

- To get the generalization error from the parameter error, we simply left and right multiply by the *same* matrix $\nabla^2 L_{\text{dis}}(\theta^*)^{-\frac{1}{2}}$. Therefore, the generalization error of the generative estimator is at most the generalization error of the discriminative estimator.
- However, if the model is not jointly well-specified, then the two estimators will not even converge to the same θ^* in general, and the discriminative estimator will clearly be better since it converges to the expected risk minimizer.

4.4 References

- [Sham Kakade's statistical learning theory course](#)
- [Hsu/Kakade/Zhang, 2014: Random design analysis of ridge regression](#)
- [van der Vaart, 2000: Asymptotic Statistics](#)

5 Kernel methods

5.1 Motivation (Lecture 14)

- So far in this class, we have studied excess risk $L(\hat{h}) - L(h^*)$, which measures how far our estimated predictor \hat{h} is away from the best possible predictor $h^* \in \mathcal{H}$ in terms of expected risk. But this is only half of the story, since the expected risk that we care about is the sum of the **estimation error** (excess risk) and the **approximation error**:

$$L(\hat{h}) = L(\hat{h}) - L(h^*) + L(h^*). \quad (392)$$

The approximation has to do with how expressive \mathcal{H} is.

- We have mainly focused on linear models, where the prediction is a function of the inner product $\langle w, x \rangle$ between a weight vector $w \in \mathbb{R}^d$ and an input $x \in \mathbb{R}^d$ (e.g., for regression, the prediction function is just $f(x) = \langle w, x \rangle$). However, real data often exhibit highly non-linear relationships which are important to model.
- Fortunately, all our algorithms so far (e.g., online gradient descent) actually only require convexity in w to ensure that the loss functions are convex, not x . So we can sneakily replace x with an arbitrary feature vector $\phi(x) \in \mathbb{R}^d$.
 - Example: $\phi(x) = (1, x, x^2)$ for $x \in \mathbb{R}$
 - Example: $\phi(x) = (\text{count of } a \text{ appearing in } x, \dots)$ for a string x .

Note that x does not even need to be a real vector. In general, we assume that $x \in \mathcal{X}$ for some set \mathcal{X} of all possible inputs (we won't assume any further structure on \mathcal{X} for now).

- We can actually represent very expressive **non-linear** functions by simply augmenting $\phi(x)$, but the problem is that $\phi(x)$ would have to be very high-dimensional in order to attain the desired degree of expressiveness, resulting in computationally expensive algorithms. A secondary consideration is that for some problems, it might be hard to design good features directly.
- Kernels address the two issues above by offering:
 - A *computationally* efficient way of working with high (and even infinite) dimensional $\phi(x)$ **implicitly**.
 - A different perspective on features, which can be more natural from a *modeling* perspective for certain applications.

- Before we define kernels formally, let's provide some intuition.

- Consider online gradient descent on, say, the squared loss: $\ell((x, y), w) = \frac{1}{2}(y - \langle w, \phi(x) \rangle)^2$.
- The weight update is ($w_1 = 0$):

$$w_{t+1} = w_t + \underbrace{\eta(y_t - \langle w_t, \phi(x_t) \rangle)}_{\stackrel{\text{def}}{=} \alpha_t} \phi(x_t). \quad (393)$$

- We see that the weight vector will always be a linear combination of the feature vectors (we will revisit this property in much greater generality when we study the representer theorem):

$$w_t = \sum_{i=1}^{t-1} \alpha_i \phi(x_i), \quad (394)$$

and the prediction is:

$$\langle w_t, \phi(x_t) \rangle = \sum_{i=1}^{t-1} \alpha_i \langle \phi(x_i), \phi(x_t) \rangle, \quad (395)$$

by linearity of inner products. Note that predictions only depend on the **inner product** between the feature vectors. This suggests we can work with in high (even infinite) dimensions as long as we can compute this inner product efficiently. The specific algorithmic tradeoff of expressing online gradient descent in this peculiar way is that we store the α_t 's (T numbers) rather than the w (d numbers). If d is much larger than T , then we win on space. On the other hand, if we have a lot of data (T), we must resort to approximations.

– **Example 28 (Computing with quadratic features)**

- * Let the raw input be $x \in \mathbb{R}^b$.
- * Feature map with all quadratic terms:

$$\phi(x) = (x_1^2, \dots, x_b^2, \sqrt{2}x_1x_2, \dots, \sqrt{2}x_1x_b, \sqrt{2}x_2x_3, \dots, \sqrt{2}x_2x_b, \dots, \sqrt{2}x_{b-1}x_b), \quad (396)$$

There are $O(b^2)$ dimensions.

- * Explicit computation: $\langle w, \phi(x) \rangle$ takes $O(b^2)$ time.
- * Implicit computation: $\langle \phi(x), \phi(x') \rangle = \langle x, x' \rangle^2$, which takes $O(b)$ time. $\langle w, \phi(x) \rangle$ requires doing this for each of T data points, which takes $O(bT)$ time.
- It's important to realize that mathematically, we're still running online gradient descent, and that all we've done is perform a computational sleight of hand, known as the **kernel trick**.

- Aside: sometimes, we can compute dot products efficiently in high (even infinite) dimensions without using the kernel trick. If $\phi(x)$ is sparse (as is often the case in natural language processing), then $\langle \phi(x), \phi(x') \rangle$ can be computed in $O(s)$ time rather than $O(d)$ time, where s is the number of nonzero entries in $\phi(x)$.
- Summary thus far
 - We start with a feature map $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$.
 - We can recast algorithms such as online gradient descent in a way so that they only depend on the inner product $\langle \phi(x), \phi(x') \rangle$.
 - We can sometimes compute this inner product efficiently.
- Since these algorithms only depend on the inner product, maybe we can just cut to the chase and directly write down functions k that correspond to inner products: $k(x, x') = \langle \phi(x), \phi(x') \rangle$ for some feature map ϕ .
- This is a key conceptual change: it shifts our perspective from thinking in terms of features of single inputs to thinking about a similarity $k(x, x')$ between two examples x and x' . Sometimes, similarities might be more convenient from a modeling point of view.

5.2 Kernels: definition and examples (Lecture 14)

- We will first define kernels as a standalone concept without reference to feature maps. Later, we'll establish the connection.
- **Definition 20 (kernel)**
 - A function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is a positive semidefinite kernel (or more simply, a kernel) iff for every finite set of points $x_1, \dots, x_n \in \mathcal{X}$, the **kernel matrix** $K \in \mathbb{R}^{n \times n}$ defined by $K_{ij} = k(x_i, x_j)$ is positive semidefinite.
- Let us now give some examples of kernels, and then prove that they are indeed valid according to Definition 20. Assume in the following that the input space is $\mathcal{X} = \mathbb{R}^b$. Note that we can visualize a kernel for $\mathcal{X} = \mathbb{R}$ by fixing x to some value (say 1) and plotting $k(x, x')$ against x' .

- Linear kernel:

$$k(x, x') = \langle x, x' \rangle. \quad (397)$$

- Polynomial kernel:

$$k(x, x') = (1 + \langle x, x' \rangle)^p. \quad (398)$$

- * Intuition: for boolean features ($x \in \{0, 1\}^b$), this corresponds to forming conjunctions of the original features.
- * Here, we can check that the corresponding dimensionality (number of features) is $O(b^p)$, which is exponential in p .
- Gaussian kernel:

$$k(x, x') = \exp\left(\frac{-\|x - x'\|_2^2}{2\sigma^2}\right). \quad (399)$$

- * A Gaussian kernel puts a smooth bump at x .
- * The bandwidth parameter σ^2 governs how smooth the functions should be: larger σ^2 yields more smoothness.
- * The corresponding dimensionality is infinite (as we'll see later), so computationally, we really have no choice but to work with kernels.
- * The Gaussian is in some sense the “go to” kernel in machine learning, because it defines a very expressive hypothesis class, as we will see.
- Sequence mis-match kernel:
 - * The above kernels have been defined for continuous input spaces $\mathcal{X} \subseteq \mathbb{R}^b$, but we can define kernels on any type of object.
 - * Suppose $\mathcal{X} = \Sigma^*$ is the set of all possible sequences (strings) over some alphabet Σ . We want to define a kernel between two strings which measures their similarity, a problem that arises in NLP and computational biology. For example, consider the strings *format* and *fmt*.
 - * A string $u \in \mathcal{X}$ is a subsequence of $x \in \mathcal{X}$ if there exists a sequence of indices $\mathbf{i} = (i_1, \dots, i_{|u|})$ such that $1 \leq i_1 < \dots < i_{|u|} \leq |x|$ such that $u_j = x_{i_j}$. In this case, we write $u = x(\mathbf{i})$.
 - * Note that x has exponentially many subsequences in general.
 - * Now define a kernel between two sequences x, x' to be a weighted number of common subsequences:

$$k(x, x') = \sum_{u \in \Sigma^*} \sum_{(\mathbf{i}, \mathbf{j}): x(\mathbf{i})=x'(\mathbf{j})=u} \lambda^{|\mathbf{i}|+|\mathbf{j}|}, \quad (400)$$

for some decay parameter $0 \leq \lambda \leq 1$. Smaller values of λ discount longer subsequences more.

- Non-example: $k(x, x') = \mathbb{I}[\|x - x'\|_2 \leq 1]$
 - Exercise: show that k is not a kernel function.
- Now we show the above kernels (linear, polynomial, Gaussian) are actually valid according to Definition 20. Let $x_1, \dots, x_n \in \mathcal{X}$ be any points. We have to check that the kernel matrix K formed by $K_{ij} = k(x_i, x_j)$ is positive semidefinite. But we first start with some general principles that will allow us to easily check whether a given function k is a kernel easily.

- General principles for checking kernels
 - Base case: for any function $f : \mathcal{X} \rightarrow \mathbb{R}$, $k(x, x') = f(x)f(x')$ is positive semidefinite.
 - * Proof: the kernel matrix can be written as $K = uu^\top \succeq 0$, where $u = (f(x_1), \dots, f(x_n))$.
 - Recursive case: given two kernels k_1, k_2 , we can create new kernels k . Note that to check that k is a kernel, it suffices to check that $K_1, K_2 \succeq 0 \Rightarrow K \succeq 0$, where K_1, K_2, K are the corresponding kernel matrices of k_1, k_2, k .
 - Sum (recursive case): $k(x, x') = k_1(x, x') + k_2(x, x')$
 - * Since positive semidefiniteness is closed under addition, $K = K_1 + K_2 \succeq 0$.
 - Product (recursive case): $k(x, x') = k_1(x, x')k_2(x, x')$
 - * $K = K_1 \circ K_2$ corresponds to elementwise product.
 - * Since K_1, K_2 are positive semidefinite, we can take their eigendecompositions:
 - $K_1 = \sum_{i=1}^n \lambda_i u_i u_i^\top$
 - $K_2 = \sum_{j=1}^n \tau_j z_j z_j^\top$
 - * Taking the elementwise product yields the following eigendecomposition, showing that K is also positive semidefinite:
 - $K = \sum_{i=1}^n \sum_{j=1}^n \lambda_i \tau_j (u_i \circ z_j)(u_i \circ z_j)^\top$, where \circ denotes elementwise products.
- Using these three principles, we can show that the linear, polynomial, and Gaussian kernels are valid.
 - Linear kernel: sum over kernels defined by functions of the form $f(x) = x_i$.
 - Polynomial kernel: given the linear kernel $\langle x, x' \rangle$, add 1 (which is a kernel); this shows that the sum $\langle x, x' \rangle + 1$ is a kernel. By the product property, $(\langle x, x' \rangle + 1)^2$ is also a kernel. Repeat $p - 1$ times to show that $(\langle x, x' \rangle + 1)^p$ is a kernel.
 - Gaussian kernel:
 - * Rewrite

$$k(x, x') = \exp\left(\frac{-\|x\|_2^2}{2\sigma^2}\right) \exp\left(\frac{-\|x'\|_2^2}{2\sigma^2}\right) \exp\left(\frac{\langle x, x' \rangle}{\sigma^2}\right). \quad (401)$$

- * The first two factors are handled by the base case.
- * For the third factor, take the Taylor expansion:

$$\exp\left(\frac{\langle x, x' \rangle}{\sigma^2}\right) = 1 + \frac{\langle x, x' \rangle}{\sigma^2} + \frac{1}{2} \frac{\langle x, x' \rangle^2}{\sigma^4} + \frac{1}{6} \frac{\langle x, x' \rangle^3}{\sigma^6} + \dots \quad (402)$$

Each term is just a homogenous polynomial kernel. Summing a finite number of terms yields a kernel. Kernels are closed under taking limits (since the set of positive semidefinite matrices is closed).

5.3 Three views of kernel methods (Lecture 14)

- We now start laying the mathematical foundation for kernel methods. Specifically, we will develop three views of kernel methods, as illustrated in Figure 5.
 - Feature map ϕ : maps from a data point $x \in \mathcal{X}$ to an element of an inner product space (the feature vector). This allows us to think about properties of single data points.
 - Kernel k : takes two data points $x, x' \in \mathcal{X}$ and returns a real number. This allows us to think about similarity between two data points.
 - RKHS \mathcal{H} : a set of functions $f : \mathcal{X} \rightarrow \mathbb{R}$ equipped a norm $\|\cdot\|_{\mathcal{H}}$ for measuring the complexity of functions. This allows us to think about the prediction function f itself.

We will define each of the three views separately, but eventually show that they are all in a sense equivalent.




Figure 5: The three key mathematical concepts in kernel methods.

- First, we need a formal notion of infinite feature vectors (the range of a feature map ϕ) that generalizes \mathbb{R}^d .
- **Definition 21 (Hilbert space)**

- A Hilbert space \mathcal{H} is an complete¹⁵ vector space with an inner product $\langle \cdot, \cdot \rangle : \mathcal{H} \times \mathcal{H} \rightarrow \mathbb{R}$ that satisfies the following properties:
 - * Symmetry: $\langle f, g \rangle = \langle g, f \rangle$
 - * Linearity: $\langle \alpha_1 f_1 + \alpha_2 f_2, g \rangle = \alpha_1 \langle f_1, g \rangle + \alpha_2 \langle f_2, g \rangle$
 - * Positive definiteness: $\langle f, f \rangle \geq 0$ with equality only if $f = 0$

The inner product gives us a norm: $\|f\|_{\mathcal{H}} \stackrel{\text{def}}{=} \sqrt{\langle f, f \rangle}$.

- Examples

¹⁵ Completeness means that all Cauchy sequences (in which elements get closer and closer to each other) converge to some element in the space (see Definition 30). Examples: set of real numbers is complete, set of rational numbers is not.

- * Euclidean space: \mathbb{R}^d , with $\langle u, v \rangle = \sum_{i=1}^d u_i v_i$
- * Square summable sequences: $\ell^2 = \{(u_i)_{i \geq 1} : \sum_{i=1}^{\infty} u_i^2 < \infty\}$, with $\langle u, v \rangle = \sum_{i=1}^{\infty} u_i v_i$.
- * Square integrable functions on $[0, 1]$: $L^2([0, 1]) = \{f : \int_0^1 f(x)^2 < \infty\}$, with $\langle f, g \rangle = \int_0^1 f(x)g(x)dx$.¹⁶

• **Definition 22 (feature map)**

- Given a Hilbert space \mathcal{H} , a **feature map** $\phi : \mathcal{X} \rightarrow \mathcal{H}$ takes inputs $x \in \mathcal{X}$ to infinite feature vectors $\phi(x) \in \mathcal{H}$.

• **Theorem 22 (feature map defines a kernel)**

- Let $\phi : \mathcal{X} \rightarrow \mathcal{H}$ be a feature mapping some input space \mathcal{X} to a Hilbert space \mathcal{H} .
- Then $k(x, x') \stackrel{\text{def}}{=} \langle \phi(x), \phi(x') \rangle$ is a kernel.

• **Proof:**

- The key idea is that the definition of the kernel only needs to look at n points, which reduces everything to a finite problem.
- Let x_1, \dots, x_n be a set of points, and let K be the kernel matrix where $K_{ij} = \langle \phi(x_i), \phi(x_j) \rangle$.
- To show that K is positive semidefinite, take any $\alpha \in \mathbb{R}^n$. We have

$$\alpha^\top K \alpha = \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j \langle \phi(x_i), \phi(x_j) \rangle \quad (403)$$

$$= \left\langle \sum_{i=1}^n \alpha_i \phi(x_i), \sum_{j=1}^n \alpha_j \phi(x_j) \right\rangle \quad (404)$$

$$\geq 0, \quad (405)$$

where we use linearity of the inner product.

• **Theorem 23 (kernel defines feature maps)**

- For every kernel k (Definition 20), there exists a Hilbert space \mathcal{H} and a feature map $\phi : \mathcal{X} \rightarrow \mathcal{H}$ such that $k(x, x') = \langle \phi(x), \phi(x') \rangle$.

¹⁶ Technically, $L^2([0, 1])$ is not a vector space since a function f which is non-zero on a measure zero set will still have $\|f\|_{\mathcal{H}} = 0$. But the quotient space (with respect to functions f with $\|f\|_{\mathcal{H}} = 0$) is a vector space.

- We will prove Theorem 23 later since it requires more sophisticated machinery (RKHSes). But to get some intuition, let's prove it for the case where the number of inputs \mathcal{X} is finite.
 - Let $\mathcal{X} = \{x_1, \dots, x_n\}$ and define the kernel matrix K (which defines the entire kernel).
 - Since $K \in \mathbb{R}^{n \times n}$ is positive semidefinite, we can write it as $K = \Phi\Phi^\top$. For example we can take an eigendecomposition $K = UDU^\top$ and let $\Phi = UD^{1/2}$ or the Cholesky decomposition $K = LL^\top$.
 - Let the feature vector $\phi(x_i) \in \mathbb{R}^n$ be the i -th row of Φ . We can verify that $K = \Phi\Phi^\top$ (equivalently, $k(x_i, x_j) = \langle \phi(x_i), \phi(x_j) \rangle$).
 - Note: the feature map is not unique, since we can also define an alternate feature matrix $\Phi' = \Phi Q$ for any orthogonal matrix Q . In this light, kernels are a more “pure” way of defining models, because feature vectors have this rotational indeterminacy.
 - If the input space \mathcal{X} is infinite, then we need to generalize our notion of feature vector from \mathbb{R}^n to an infinite dimensional space. What is that space?

5.4 Reproducing kernel Hilbert spaces (RKHS) (Lecture 14)

- We will now introduce RKHSes (a type of Hilbert space), the third view on kernel methods. Initially, we will use it as means to show that every kernel k actually corresponds to some feature map.
- But RKHSes stand on their own right as an object worth studying since they directly allow us to work on the prediction *function*. Concretely, in linear regression, we fit a weight vector w , constraining or regularizing the norm $\|w\|_2$, and we use w to predict on a new input x via $x \mapsto \langle w, \phi(x) \rangle$. But can we get a handle on this prediction function $x \mapsto f(x)$ directly as well as a norm $\|f\|_{\mathcal{H}}$ measuring the complexity of f ? Let \mathcal{H} be the space of all prediction functions f we are (implicitly) considering when using a kernel k . What is this space and how is it related to k ?
- We can consider Hilbert spaces over functions $f : \mathcal{X} \rightarrow \mathbb{R}$. But not all Hilbert spaces are suitable for machine learning.
 - For example, consider $\mathcal{H} = L^2([0, 1])$. Recall that every $f \in \mathcal{H}$ is actually an equivalence class over functions which differ on a measure zero set, which means pointwise evaluations $f(x)$ at individual x 's is not even defined.
 - This is highly distressing given that the whole point is to learn an f for the purpose of doing pointwise evaluations (a.k.a. prediction)!
 - RKHSes remedy this problem.
- **Definition 23 (bounded functional)**

- Given a Hilbert space \mathcal{H} , a functional $L : \mathcal{H} \rightarrow \mathbb{R}$ is bounded iff there exists an $M < \infty$ such that

$$|L(f)| \leq M\|f\|_{\mathcal{H}} \text{ for all } f \in \mathcal{H}. \quad (406)$$

- Example: $\mathcal{H} = \mathbb{R}^d$ with the usual inner product, $L(f) = \langle c, f \rangle$ is bounded (with $M = \|c\|_2$ by Cauchy-Schwartz)

• **Definition 24 (evaluation functional)**

- Let \mathcal{H} be a Hilbert space consisting of functions $f : \mathcal{X} \rightarrow \mathbb{R}$.
- For each $x \in \mathcal{X}$, define the evaluation functional $L_x : \mathcal{H} \rightarrow \mathbb{R}$ as

$$L_x(f) \stackrel{\text{def}}{=} f(x). \quad (407)$$

- Example: for $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{H} = \{f_c(x) = c \cdot x : c \in \mathbb{R}^d\}$ be linear functions, then the evaluation functional is $L_x(f_c) = \langle c, x \rangle$.
- Intuitively, the evaluation functional is a projection operator that turns a function f into one component $f(x)$. Hence, evaluation functionals are linear. This will be important later.

• **Definition 25 (reproducing kernel Hilbert space)**

- A reproducing kernel Hilbert space \mathcal{H} is a Hilbert space over functions $f : \mathcal{X} \rightarrow \mathbb{R}$ such that for each $x \in \mathcal{X}$, the **evaluation functional** L_x is bounded.

• Non-example

- Let \mathcal{H} be the set of all square integrable continuous functions from $[0, 1]$ to \mathbb{R} .
- Consider $f_{\epsilon}(x) = \max(0, 1 - \frac{|x - \frac{1}{2}|}{\epsilon})$, which is zero except for a small spike at $x = \frac{1}{2}$ up to $f(x) = 1$. Note that $\|f_{\epsilon}\|_{\mathcal{H}} \rightarrow 0$ as $\epsilon \rightarrow 0$.
- Consider the evaluation functional $L_{\frac{1}{2}}$. Note that $L_{\frac{1}{2}}(f_{\epsilon}) = f_{\epsilon}(\frac{1}{2}) = 1$. So there cannot exist an M such that $L_{\frac{1}{2}}(f_{\epsilon}) \leq M\|f_{\epsilon}\|_{\mathcal{H}}$ for all $\epsilon > 0$.
- So this \mathcal{H} is not a RKHS.

• **Theorem 24 (RKHS defines a kernel)**

- Every RKHS \mathcal{H} over functions $f : \mathcal{X} \rightarrow \mathbb{R}$ defines a unique kernel $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$.

• Proof (construction of the kernel)

- Note that $L_x(f)$ is *linear*: $L_x(cf) = cL_x(f)$ and $L_x(f + g) = f(x) + g(x)$.

- The **Riesz representation theorem** states that all *bounded linear functionals* L on a Hilbert space can be expressed as an inner product $L(f) = \langle R, f \rangle$ for a *unique* $R \in \mathcal{H}$.
- Applying this theorem to the evaluation functionals L_x , we can conclude that for each $x \in \mathcal{X}$, there exists a unique **representer** $R_x \in \mathcal{H}$ such that $L_x(f) = \langle R_x, f \rangle$. Recall that we also have $L_x(f) = f(x)$ by definition. Combining yields the **reproducing property**:

$$\boxed{f(x) = \langle R_x, f \rangle \text{ for all } f \in \mathcal{H}.} \quad (408)$$

This is the key property: *function evaluations can be expressed as inner products*.

- Now let's define a function k :

$$\boxed{k(x, x') \stackrel{\text{def}}{=} R_x(x').} \quad (409)$$

Applying the reproducing property one more time with $f = R_x$ yields

$$k(x, x') = R_x(x') = \langle R_x, R_{x'} \rangle. \quad (410)$$

If we define a feature map $\phi(x) \stackrel{\text{def}}{=} R_x$, we can invoke Theorem 22 to conclude that $k(x, x')$ is a valid kernel.

- In summary, any RKHS \mathcal{H} gives rise to a kernel k called the **reproducing kernel** of \mathcal{H} . The key is the Riesz representation, which turns function evaluations into inner products.

To complete the picture, we need to show that a kernel defines an unique RKHS.

- **Theorem 25 (Moore-Aronszajn theorem)**

- For every kernel k , there exists a unique RKHS \mathcal{H} with reproducing kernel k .

- **Proof sketch:**

- Let k be a kernel. We will construct a RKHS \mathcal{H} from the functions $\{k(x, \cdot) : x \in \mathcal{X}\}$.
- First, define \mathcal{H}_0 to contain all finite linear combinations of the form

$$f(x) = \sum_{i=1}^n \alpha_i k(x_i, x), \quad (411)$$

for all $n, \alpha_{1:n}, x_{1:n}$. By construction, \mathcal{H}_0 is a vector space (not necessarily complete though).

- Second, define the inner product between $f(x) = \sum_{i=1}^n \alpha_i k(x_i, x)$ and $g(x) = \sum_{j=1}^{n'} \beta_j k(x'_j, x)$ as follows:

$$\langle f, g \rangle \stackrel{\text{def}}{=} \sum_{i=1}^n \sum_{j=1}^{n'} \alpha_i \beta_j k(x_i, x'_j). \quad (412)$$

Let's check that our definition of $\langle \cdot, \cdot \rangle$ is an actual inner product:

- * Symmetry ($\langle f, g \rangle = \langle g, f \rangle$): by symmetry of k
- * Linearity ($\langle \alpha_1 f_1 + \alpha_2 f_2, g \rangle = \alpha_1 \langle f_1, g \rangle + \alpha_2 \langle f_2, g \rangle$): by definition of f and g (they are just a linear sum of terms).
- * Positive definiteness ($\langle f, f \rangle \geq 0$ with equality only if $f = 0$):
 - For any $f \in \mathcal{H}_0$, we have $\langle f, f \rangle = \alpha^\top K \alpha \geq 0$ by the positive semidefinite property of kernels.
 - Now we will show that $\langle f, f \rangle = 0$ implies $f = 0$. Let $f(x) = \sum_{i=1}^n \alpha_i k(x_i, x)$. Take any $x \in \mathcal{X}$ and define $c = [k(x_1, x), \dots, k(x_n, x)]^\top$. Since

$$\begin{pmatrix} K & c \\ c^\top & k(x, x) \end{pmatrix} \succeq 0, \quad (413)$$

we must have that

$$\alpha^\top K \alpha + 2bc^\top \alpha + b^2 k(x, x) \geq 0 \quad (414)$$

for all b . Note that $\langle f, f \rangle = \alpha^\top K \alpha = 0$. We now argue that $c^\top \alpha = 0$. If $c^\top \alpha > 0$, then as b approaches 0 from the negative side, we have that the LHS of (414) is strictly negative, which is a contradiction. If $c^\top \alpha < 0$, then as b approaches 0 from the positive side, we get a contradiction as well. Therefore, $f(x) = c^\top \alpha = 0$.

- So far we have a valid Hilbert space, but we need to still check that all evaluation functionals L_x are bounded to get an RKHS. Also, we should check that $R_x \stackrel{\text{def}}{=} k(x, \cdot)$ is indeed a representer of function evaluations. Take any $f \in \mathcal{H}_0$. Then:

$$f(x) = \sum_{i=1}^n \alpha_i k(x_i, x) \quad [\text{definition of } f] \quad (415)$$

$$= \langle f, k(x, \cdot) \rangle \quad [\text{definition of inner product}] \quad (416)$$

$$= \langle R_x, f \rangle \quad [\text{definition of } R_x]. \quad (417)$$

- Finally, let \mathcal{H} be the completion of \mathcal{H}_0 (by including all limit points of sequences in \mathcal{H}_0). For details, see the references at the end of this section.
- This proves Theorem 23 because the RKHS \mathcal{H} is an inner product space by construction.

- Summary

- A feature map $\phi : \mathcal{X} \rightarrow \mathcal{H}$: maps points in \mathcal{X} to some inner product space \mathcal{H} .
- A (positive semidefinite) kernel $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$: every derived kernel matrix K is positive semidefinite
- A reproducing kernel Hilbert Space (RKHS) \mathcal{H} containing functions $f : \mathcal{X} \rightarrow \mathbb{R}$ such that function evaluations are bounded linear operators.
- Equivalences
 - * $f(x) = \sum_{i=1}^{\infty} \alpha_i k(x_i, x)$, $R_x = k(x, \cdot)$: Moore-Aronszajn establishes connection between kernels and RKHSes
 - * $\phi(x) = R_x$: can set feature map (not unique) to map x to the representer of x in the RKHS
 - * $k(x, x') = \langle \phi(x), \phi(x') \rangle$: every kernel k corresponds to some inner product (via RKHS) and vice-versa (easy)

5.5 Learning using kernels (Lecture 15)

- We have established that kernels k provide a space of functions \mathcal{H} ; this is the hypothesis class.¹⁷ Now let's talk about learning, which is about combining the hypothesis class \mathcal{H} with actual data.
- Let's start with kernelized ridge regression, where we obtain examples $\{(x_i, y_i)\}_{i=1}^n$, and want to find a function $f \in \mathcal{H}$ that fits the data, where \mathcal{H} is an RKHS. A natural objective function is to penalize the squared loss plus a penalty for the complexity of f , where the complexity is the RKHS norm:

$$f^* \in \arg \min_{f \in \mathcal{H}} \sum_{i=1}^n \frac{1}{2} (f(x_i) - y_i)^2 + \frac{\lambda}{2} \|f\|_{\mathcal{H}}^2. \quad (418)$$

- More generally, a learning problem can be posed as the following optimization problem:

$$f^* \in \arg \min_{f \in \mathcal{H}} L(\{(x_i, y_i, f(x_i))\}_{i=1}^n) + Q(\|f\|_{\mathcal{H}}^2), \quad (419)$$

where

- $L : (\mathcal{X} \times \mathcal{Y} \times \mathbb{R})^n \rightarrow \mathbb{R}$ is an arbitrary loss function on n examples.
 - * Example (regression): $L(\{(x_i, y_i, f(x_i))\}_{i=1}^n) = \sum_{i=1}^n \frac{1}{2} (f(x_i) - y_i)^2$.
- $Q : [0, \infty) \rightarrow \mathbb{R}$ is a strictly increasing function (regularizer).
 - * Example (quadratic): $Q(\|f\|_{\mathcal{H}}^2) = \frac{\lambda}{2} \|f\|_{\mathcal{H}}^2$.

This optimization problem may seem daunting since it is optimizing over a potentially very large function space \mathcal{H} . But the following representer theorem reassures us that all minimizers can be written as a linear combination of the kernel functions evaluated at the training points.

- **Theorem 26 (representer theorem)**

- Let V denote the span of the representer of the training points:

$$V \stackrel{\text{def}}{=} \text{span}(\{k(x_i, \cdot) : i = 1, \dots, n\}) = \left\{ \sum_{i=1}^n \alpha_i k(x_i, \cdot) : \alpha \in \mathbb{R}^n \right\}. \quad (420)$$

- Then all minimizers f^* of (419) satisfy $f^* \in V$.

¹⁷For regression, our prediction at x is simply $f(x)$, where $f \in \mathcal{H}$. For classification and ranking problems, we need to pass the function values through some non-linear transformation.

- Proof

- FIGURE: [projection of f^* onto V]
- The key is to use the fact that an RKHS has an inner product structure, which allows us to use linear algebra.
- Define the orthogonal complement:

$$V_{\perp} = \{g \in \mathcal{H} : \langle f, g \rangle = 0 \text{ for all } f \in V\}. \quad (421)$$

- Any $f \in \mathcal{H}$ can be decomposed into a part in the span of the examples and an orthogonal part:

$$f^* = f + f_{\perp}, \quad (422)$$

where $f \in V$ and $f_{\perp} \in V_{\perp}$.

- The idea is that the loss is unchanged by f_{\perp} but the regularizer grows with non-zero f_{\perp} , so we must have $f_{\perp} = 0$.
- The loss depends on f^* only through $\{f^*(x_j) : j = 1, \dots, n\}$, which can be written as:

$$f^*(x_j) = f(x_j) + \langle f_{\perp}, k(x_j, \cdot) \rangle. \quad (423)$$

The second term is zero, so the loss doesn't depend on f_{\perp} .

- The regularizer:

$$Q(\|f^*\|_{\mathcal{H}}^2) = Q(\|f\|_{\mathcal{H}}^2 + \|f_{\perp}\|_{\mathcal{H}}^2). \quad (424)$$

Since Q is strictly monotonic and f^* is a minimizer, we must have $f_{\perp} = 0$.

- Therefore, $f^* \in V$.

- Remark: the representer theorem does not require the loss function L to be convex.
- The representer theorem tells us α 's exist, but how to find the α 's depends on the actual loss function and regularizer. Let's now look at some examples.

- **Example 29 (Kernelized ridge regression)**

- Recall the optimization problem for regression:

$$\min_{f \in \mathcal{H}} \sum_{i=1}^n \frac{1}{2} (f(x_i) - y_i)^2 + \frac{\lambda}{2} \|f\|_{\mathcal{H}}^2. \quad (425)$$

By the representer theorem, we have the equivalent optimization problem:

$$\min_{f \in \mathcal{H}} \sum_{i=1}^n \frac{1}{2} \left(\sum_{j=1}^n \alpha_j k(x_i, x_j) - y_i \right)^2 + \frac{\lambda}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j k(x_i, x_j). \quad (426)$$

- Letting $K \in \mathbb{R}^{n \times n}$ be the kernel matrix and $Y \in \mathbb{R}^n$ denote the vector of outputs, we have:

$$\min_{\alpha \in \mathbb{R}^n} \frac{1}{2} \|K\alpha - Y\|_2^2 + \frac{\lambda}{2} \alpha^\top K \alpha. \quad (427)$$

Differentiating with respect to α and setting to zero:

$$K(K\alpha - Y) + \lambda K\alpha = 0. \quad (428)$$

Rearranging:

$$K(K + \lambda I)\alpha = KY. \quad (429)$$

Solving yields a solution:

$$\boxed{\alpha = (K + \lambda I)^{-1} Y.} \quad (430)$$

Note that the solution is not necessarily unique, since we could add any vector in the null space of K , but there's no reason to consider them.

- To predict on a new example x , we form kernel evaluations $c \in \mathbb{R}^n$ where $c_i = k(x_i, x)$, and then predict

$$y = c^\top \alpha. \quad (431)$$

• Example 30 (SVM classification)

- This was done in CS229, so we won't go through it again.
- Primal ($y_i \in \{-1, +1\}$):

$$\min_{f \in \mathcal{H}} \sum_{i=1}^n \max\{0, 1 - y_i f(x_i)\} + \frac{\lambda}{2} \|f\|_{\mathcal{H}}^2. \quad (432)$$

- Dual (define $\tilde{K}_{ij} = y_i y_j K_{ij}$):

$$\min_{\alpha \in \mathbb{R}^n} -\mathbf{1}^\top \alpha + \alpha^\top \tilde{K} \alpha \quad \text{subject to} \quad 0 \leq \alpha_i \leq \frac{1}{\lambda}, Y^\top \alpha = 0. \quad (433)$$

The dual is computed by taking the Lagrange dual of the primal optimization problem.

• Example 31 (Kernel PCA)

- Review of PCA

- * Recall in (featurized) PCA, we want to find directions in our data with highest variance.
- * Suppose we have data points x_1, \dots, x_n and a feature map $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$.
- * Assume that the data points are centered at zero: $\sum_{i=1}^n \phi(x_i) = 0$.
- * Define the empirical covariance matrix as follows:

$$C \stackrel{\text{def}}{=} \frac{1}{n} \Phi^\top \Phi, \quad (434)$$

where the i -th row of Φ is $\phi(x_i)$.

- * Then PCA seeks to find a eigendecomposition of C :

$$Cv = \lambda v. \quad (435)$$

In practice, we will compute the eigendecomposition of C and take the first r eigenvectors v_1, \dots, v_r (principal components) as an approximation of the entire feature space. Note that the v_i 's form an orthonormal basis (have unit norm and are orthogonal).

- * The squared reconstruction error of a new point x is:

$$\left\| \sum_{i=1}^r \langle \phi(x), v_i \rangle v_i - \phi(x) \right\|_2^2. \quad (436)$$

For example, if we were doing anomaly detection, if a data point x has a large reconstruction error, then x is an anomaly.

– Heuristic derivation of kernel PCA

- * By the representer theorem (or even more simply, by inspecting the form of the covariance matrix), we have $v = \sum_{i=1}^n \alpha_i \phi(x_i) = \Phi^\top \alpha$, so an equivalent characterization is to project the vectors on the data points:

$$\Phi C v = \lambda \Phi v. \quad (437)$$

- * Using the definition of C and the fact that $\Phi v = K\alpha$, we have

$$\frac{1}{n} K^2 \alpha = \lambda K \alpha. \quad (438)$$

Again, any solution to the following is a valid solution to the above (but we can always add spurious vectors in the null space of K):

$$\boxed{\frac{1}{n} K \alpha = \lambda \alpha.} \quad (439)$$

- * In practice, we compute the eigendecomposition of K and take the first r eigenvectors as the approximation. For simplicity, let's assume we just take one principal component $v \in \mathcal{H}$.

- Computing in infinite dimensions
 - * Though the derivation assumed $\phi(x) \in \mathbb{R}^d$, the result is the same for $\phi(x) = k(x, \cdot) \in \mathcal{H}$ in general.
 - * We won't go through the details, but the idea is to define a covariance operator (rather than a matrix) $C : \mathcal{H} \rightarrow \mathcal{H}$. The intuition is the same.
 - * Recall the principal component is now a function $v \in \mathcal{H}$ with $\|v\|_{\mathcal{H}} = 1$,¹⁸ where $v(x) = \sum_{i=1}^n \alpha_i k(x_i, x)$.
 - * The squared reconstruction error of a new point x is:

$$\|v(x)v - \phi(x)\|_{\mathcal{H}}^2 = v(x)^2 - 2v(x)^2 + k(x, x) = k(x, x) - v(x)^2. \quad (440)$$

As expected, all we need are kernel evaluations.

- Interpretation
 - * The point of PCA is to reduce the dimensionality of the data, so it might seem strange at first we would want to use kernel PCA to first increase the number of dimensions (possibly to infinity).
 - * This is actually okay, because the point of kernels is to reshape the data (in non-linear ways). In doing so, we can expose better directions than those present in the original data.
 - * For example, if we use a quadratic kernel, we are saying that we believe the data lies close to a quadratic surface. Of course, with more dimensions, statistical error in the directions could increase.

5.6 Fourier properties of shift-invariant kernels (Lecture 15)

- Having explored how kernels can be used in practice, let us turn back to studying their theoretical properties. Specifically, we will use Fourier analysis to get a handle on what information about the data a kernel is capturing. We will also see that this leads to a new basis representation of kernels. In this section, we will focus on shift-invariant kernels, which are kernels that don't depend on the absolute position of the data points.

• Definition 26 (shift-invariant kernel)

- A kernel $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ where $\mathcal{X} \subseteq \mathbb{R}^b$ is **shift invariant** (a.k.a. stationary, translation invariant) iff k can be written as $k(x, x') = h(x - x')$ for some function h .
- Example: Gaussian kernel
- Non-example: linear kernel $((x + a)(x' + a) \neq xx')$

¹⁸To make v a unit vector, we just rescale v by $\|v(x)\|_{\mathcal{H}}^{-1}$.

- Our goal is to understand shift-invariant kernels in the frequency domain. In particular, it will shed insight into the smoothness properties of the kernel.
- First, let's warmup with some basic facts from Fourier analysis:
 - FIGURE: [complex unit circle]
 - $e^{i\omega t} = \cos(\omega t) + i \sin(\omega t)$
 - $e^{-i\omega t} = \cos(\omega t) - i \sin(\omega t)$
 - $\cos(\omega t) = \frac{1}{2}e^{i\omega t} + \frac{1}{2}e^{-i\omega t}$
 - Here $\omega \in \mathbb{R}$ is the frequency of the sinusoid.
 - Note that everything thus far generalizes to $\omega \in \mathbb{R}^b$ and $t \in \mathbb{R}^b$; just replace ωt with $\langle \omega, t \rangle$.
- Now let's try to construct a kernel using the Fourier basis:
 - First define a feature map for a fixed ω : $\phi(x) = e^{-i\langle \omega, x \rangle}$ (note that this is complex-valued, but that's okay).
 - Using this feature map, let's define a kernel:

$$k(x, x') = \phi(x) \overline{\phi(x')} = e^{-i\langle \omega, x - x' \rangle}, \quad (441)$$

where \bar{a} denotes the complex conjugate of a . This kernel deems two points to be similar if they are close modulo $2\pi/\omega$ (for some fixed scalar frequency ω); clearly this is a bit silly.

- To get more realistic kernels, we need to incorporate multiple frequencies. We can do this by simply averaging over multiple kernels (recall that the sum of kernels is a kernel). Specifically, let $\mu(\cdot)$ be a finite non-negative measure over frequencies. Then

$$k(x, x') = \int e^{-i\langle \omega, x - x' \rangle} \mu(d\omega). \quad (442)$$

is also a valid kernel. Or in terms of h (recall that $t = x - x'$):

$$h(t) = \int e^{-i\langle \omega, t \rangle} \mu(d\omega). \quad (443)$$

- The corresponding feature map consists of the following basis functions: $\{x \mapsto e^{-i\langle \omega, x \rangle} : \omega \in \mathbb{R}^b\}$
- Intuitively μ tells us how much focus to put on various frequencies.

- **Example 32 (constant)**

- Let the spectral measure $\mu = \delta_0$ place all its mass at 0.

- Then $k(x, x') = h(t) = 1$ is the constant kernel.

- **Example 33 (single frequency)**

- Suppose μ places mass only at $-\omega$ and ω :

$$\mu = \frac{1}{2}(\delta_{-\omega} + \delta_{\omega}). \quad (444)$$

- Then the resulting kernel is:

$$k(x, x') = h(t) = \cos(\omega t). \quad (445)$$

- In general, $h(t)$ defined via μ might be complex-valued, but if μ is **symmetric** (that is, $\mu(A) = \mu(-A)$ for all measurable sets A), then $h(t)$ will be **real-valued**.

- **Example 34 (sinc)**

- Let the spectral density s be the 1 over $[-a, a]$ (that is, we only keep frequencies below a):

$$s(\omega) = \mathbb{I}[-a \leq \omega \leq a]. \quad (446)$$

- Then the resulting kernel is:

$$h(t) = \frac{2 \sin(at)}{t}. \quad (447)$$

- Proof: just integrate:

$$h(t) = \int_{-a}^a e^{-i\omega t} d\omega = \frac{1}{-it}(e^{-iat} - e^{iat}) = \frac{1}{-it}(-2i \sin(at)). \quad (448)$$

- Note: as a increases, we cover more frequencies. If $a \rightarrow \infty$, then $h(t)$ converges to a delta function at 0, which corresponds to the degenerate kernel $k(x, x') = \mathbb{I}[x = x']$.

- At first, it might seem that this is a funny way of defining certain types of kernels, but what's remarkable is that *all* shift-invariant kernels can be written in the form (443) for some appropriate choice of μ . This statement is precisely Bochner's theorem:

- **Theorem 27 (Bochner's theorem)**

- Let $k(x, x') = h(x - x')$ be a continuous shift-invariant kernel ($x \in \mathbb{R}^b$).

- Then there exists a unique finite non-negative measure μ (called the **spectral measure**) on \mathbb{R}^b such that

$$h(t) = \int e^{-i\langle t, \omega \rangle} \mu(d\omega). \quad (449)$$

- Furthermore, if μ has a density s ,

$$\mu(d\omega) = s(\omega)d\omega, \quad (450)$$

then we call s the **spectral density**, and h is the Fourier transform of s .

- So far, we've defined kernels through various spectral measures μ . But we can also take a given kernel and compute its spectral measure to study the properties of the kernel. Given a candidate kernel function $k(x, x') = h(x - x')$, take the Fourier transform of h (which for symmetric functions is the inverse Fourier transform times $1/(2\pi)$) to get the spectral density s .

- **Example 35 (box is not a kernel)**

- Consider the following function:

$$h(t) = \mathbb{I}[-1 \leq t \leq 1]. \quad (451)$$

- The inverse Fourier transform times $1/(2\pi)$ is

$$s(\omega) = \frac{\sin(\omega)}{\pi\omega}, \quad (452)$$

reusing the result from above.

- But notice that $s(\omega)$ is negative in some places, which means, by Bochner's theorem, that $h(t)$ is not a valid (positive semidefinite) kernel! Now we have another way to check whether a shift-invariant function specifies a kernel—simply take the inverse Fourier transform and see whether it's non-negative everywhere.

- **Example 36 (Gaussian kernel)**

- Let the spectral density s be the density of the multivariate Gaussian distribution with variance $1/\sigma^2$:

$$s(\omega) = \left(\frac{2\pi}{\sigma^2}\right)^{-d/2} \exp\left(\frac{-\sigma^2 \|\omega\|_2^2}{2}\right). \quad (453)$$

- Then the resulting kernel is the Gaussian kernel with variance σ^2 (note the inverting of the variance):

$$h(t) = \exp\left(\frac{-\|t\|_2^2}{2\sigma^2}\right). \quad (454)$$

- Proof:

$$h(t) = \int \left(\frac{2\pi}{\sigma^2}\right)^{-d/2} \exp\left(\frac{(-\sigma^2\|\omega\|_2^2 - 2i\langle\omega, t\rangle - \sigma^{-2}i^2\|t\|_2^2) + \sigma^{-2}i^2\|t\|_2^2}{2}\right) d\omega. \quad (455)$$

Complete the square and note that the Gaussian distribution (with mean it/σ) integrates to 1.

- Intuition: the larger σ^2 is, one can see from $s(\omega)$ that high frequency components are dampened. Consequently, the smoother the kernel $h(t)$ is.

• **Example 37 (rational quadratic kernel)**

- Motivation: with Gaussian kernels, how do we set the variance σ^2 ?
- Putting on our Bayesian hats, let's define a prior over $\tau = \sigma^{-2}$.
- Let $k_\tau(x, x')$ be the Gaussian kernel with hyperparameter τ .
- Recalling that the sum of kernels is a kernel, we have that

$$\int k_\tau(x, x') p(\tau) d\tau \quad (456)$$

is also a kernel for any $p(\tau)$.

- For mathematical convenience, let's put a $\text{Gamma}(\alpha, \beta)$ prior on τ .
- By conjugacy, we can integrate a Gamma distribution against a Gaussian, which yields a student-t distribution.
- Ignoring normalization constants, the kernel is the **rational quadratic kernel** (derivation omitted):

$$h(t) = \left(1 + \frac{\beta t^2}{2\alpha}\right)^{-\alpha}. \quad (457)$$

- When $\alpha = 1$ and $\beta = 2$, we have:

$$h(t) = \frac{1}{1 + t^2}. \quad (458)$$

- Compared with the Gaussian kernel:
 - * The area near zero is steeper, so the function can change rapidly.
 - * The tails decay slower, function values can have longer range dependencies.

This flexibility comes from integrating over values of σ^2 .
- Note that as $\alpha \rightarrow \infty$, the rational quadratic approaches the Gaussian kernel.

5.7 Efficient computation (Lecture 16)

- We saw how kernel methods could be used for learning: Given n points x_1, \dots, x_n , we form the $n \times n$ kernel matrix K , and optimize an objective function whose variables are $\alpha \in \mathbb{R}^n$. For example, in kernel ridge regression, we have $\alpha = (K + \lambda I)^{-1}Y$, which requires $O(n^3)$ time. When we have large datasets (e.g., $n = 10^6$), this is prohibitively expensive. On the other hand, when the feature vector $\phi(x)$ is high-dimensional (especially infinite-dimensional), then scaling with n is the lesser of two evils. But can we do better? Note that even merely computing the full kernel matrix K takes $O(n^2)$ time, which is already too large, so we will probably have to cut some corners.
- We will introduce two types of kernel approximations:
 - Random features: We will write the kernel function as an integral, and using Monte Carlo approximations of this integral. These approximations are of the kernel function and are data-independent.
 - Nyström method: We will sample a subset of the n points and use these points to approximate the kernel matrix. These approximations are of the kernel matrix and are data-dependent.
- Gaussian kernel intuitions
 - Let's get some intuition about when approximations might be a sensible thing to do on a concrete scenario. Recall the Gaussian kernel:

$$k(x, x') = \exp\left(\frac{-\|x - x'\|_2^2}{2\sigma^2}\right). \quad (459)$$

- If the points x_1, \dots, x_n are far apart (relative to the bandwidth of σ^2), then the kernel matrix K will be roughly the identity matrix. In this case, we can't possibly hope for a low-rank approximation to capture everything.
 - On the other hand, if the points are tightly clustered into m clusters, then the kernel matrix (with sorted columns/rows) looks like a block diagonal matrix with m blocks, where each block is a rank 1 all-ones matrix. Here, you would expect a rank m approximation to be effective.
 - In reality, the situation is somewhere in between. Of course, kernel methods are exactly useful when the data are fairly complex, so we shouldn't expect these approximations to provide magical savings, unless the data is very redundant.
- Random Fourier features (Rahimi/Recht, 2008)

- Our starting point is Bochner’s theorem (Theorem 27), which allows us to write shift-invariant kernels in terms of an integral over some spectral measure μ :

$$k(x, x') = \int \phi_\omega(x) \overline{\phi_\omega(x')} \mu(d\omega), \quad (460)$$

where $\phi_\omega(x) = e^{-i\langle \omega, x \rangle}$ is a single (Fourier) feature.

- The key idea is to replace the integral with a finite sum over m elements. For simplicity, assume that μ is a probability distribution. If it is not, then we can normalize it and then multiply the result by $\mu(\mathbb{C}^b)$. Let $\omega_1, \dots, \omega_m$ be drawn i.i.d. from μ . Then, define the approximate kernel as:

$$\hat{k}(x, x') = \frac{1}{m} \sum_{i=1}^m \phi_{\omega_i}(x) \overline{\phi_{\omega_i}(x')}. \quad (461)$$

This kernel corresponds to having the following random feature map:

$$\hat{\phi}(x) = [\phi_{\omega_1}(x), \dots, \phi_{\omega_m}(x)] \in \mathbb{C}^b. \quad (462)$$

- As a concrete example, consider the Gaussian kernel, which has a Gaussian spectral density (recall $\mu(d\omega) = s(\omega)d\omega$) with the inverse variance:

$$k(x, x') = \exp\left(\frac{-\|x - x'\|_2^2}{2\sigma^2}\right), \quad (463)$$

$$s(\omega) = \left(\frac{2\pi}{\sigma^2}\right)^{-b/2} \exp\left(\frac{-\sigma^2\|\omega\|_2^2}{2}\right). \quad (464)$$

This means that each $\omega_i \sim \mathcal{N}(0, \sigma^{-2}I)$ is drawn from a Gaussian.

- Algorithm

- * The practical upshot of random Fourier features on the Gaussian kernel is that it is dirt simple.
- * Before you get data, draw $\omega_1, \dots, \omega_m \sim \mathcal{N}(0, \sigma^{-2}I)$, which defines the random feature map $\hat{\phi}$. This feature map is fixed once and for all.
- * In training/test, given a new data point x , we can apply the feature map $\hat{\phi}(x)$, which simply involves m Gaussian projections.
- Note that the approximate kernel is unbiased ($\mathbb{E}[\hat{k}(x, x')] = k(x, x')$), so as $m \rightarrow \infty$, we have that $\hat{k}(x, x')$ converges to $k(x, x')$ for a fixed x, x' . We want this to work well on average for all the data that we’re going to see, which smells almost like uniform convergence. The following theorem quantifies this:
- **Theorem 28 (random features (Rahimi/Recht, 2008))**
 - * Let k be a shift-invariant kernel on $x \in \mathbb{R}^b$.

* Let

$$\mathcal{F} \stackrel{\text{def}}{=} \left\{ x \mapsto \int \alpha(\omega) \phi_\omega(x) \mu(d\omega) : \forall \omega, |\alpha(\omega)| \leq C \right\} \quad (465)$$

be the subset of functions in the RKHS \mathcal{H} with bounded Fourier coefficients $\alpha(\omega)$.

* Let

$$\hat{\mathcal{F}} \stackrel{\text{def}}{=} \left\{ x \mapsto \frac{1}{m} \sum_{i=1}^m \alpha(\omega_i) \phi_{\omega_i}(x) : \forall \omega, |\alpha(\omega)| \leq C \right\} \quad (466)$$

be the subset that is spanned by the random feature functions, where $\omega_{1:k}$ be drawn i.i.d. from μ .

* Let p^* be any distribution over $\mathcal{X} = \mathbb{R}^b$.

* Define the inner product with respect to the data-generating distribution (this is not the RKHS norm):

$$\langle f, g \rangle \stackrel{\text{def}}{=} \mathbb{E}_{x \sim p^*} [f(x)g(x)]. \quad (467)$$

* Let $f^* \in \mathcal{F}$ be any true function.

* Then with probability at least $1 - \delta$, there exists $\hat{f} \in \hat{\mathcal{F}}$ that

$$\|\hat{f} - f^*\| \leq \frac{C}{\sqrt{m}} \left(1 + \sqrt{2 \log(1/\delta)} \right). \quad (468)$$

– Proof of Theorem 28:

* This proof uses fairly standard tools: McDiarmid's inequality and Jensen's inequality. The function we're applying involves taking a norm of a function, but we just need the bounded differences condition to hold.

* Fix $f^* \in \mathcal{F}$ with coefficients $\alpha(\omega)$.

* Construct \hat{f} with the same coefficients, and note that $\hat{f} \in \hat{\mathcal{F}}$ and $\mathbb{E}[\hat{f}] = f^*$.

* Define

$$D(\omega_{1:m}) = \|\hat{f} - f^*\|. \quad (469)$$

Note that D satisfies the bounded differences inequality: letting $\omega_{1:m}^i = \omega_{1:m}$ except on the i -th component, where it is ω'_i :

$$|D(\omega_{1:m}) - D(\omega_{1:m}^i)| \leq \|\hat{f} - f^*\| - \|\hat{f}^i - f^*\| \quad (470)$$

$$\leq \|\hat{f} - \hat{f}^i\| \quad [\text{triangle inequality}] \quad (471)$$

$$\leq \frac{1}{m} \|\alpha(\omega_i) \phi_{\omega_i} - \alpha(\omega'_i) \phi_{\omega'_i}\| \quad (472)$$

$$\leq \frac{2C}{m}. \quad (473)$$

Note that the last line follows because $|\alpha(\omega_i)| \leq C$ and $\phi_{\omega_i}(x) = e^{-i\langle \omega_i, x \rangle}$ and $|e^{-ia}| = 1$ for all a .

* We can bound the mean by passing to the variance:

$$\mathbb{E}[D(\omega_{1:m})] \leq \sqrt{\mathbb{E}[D(\omega_{1:m})^2]} \quad [\text{Jensen's inequality}] \quad (474)$$

$$= \sqrt{\mathbb{E} \left[\left\| \frac{1}{m} \sum_{i=1}^m (\alpha(\omega_i) \phi_{\omega_i} - f^*) \right\|^2 \right]} \quad [\text{expand}] \quad (475)$$

$$= \sqrt{\frac{1}{m^2} \sum_{i=1}^m \mathbb{E} [\|\alpha(\omega_i) \phi_{\omega_i} - f^*\|^2]} \quad [\text{variance of i.i.d. sum}] \quad (476)$$

$$\leq \frac{C}{\sqrt{m}} \quad [\text{use } |\alpha(\omega_i)| \leq C]. \quad (477)$$

* Applying McDiarmid's inequality (Theorem 10), we get that

$$\mathbb{P} \left[D(\omega_{1:m}) \geq \frac{C}{\sqrt{m}} + \epsilon \right] \leq \exp \left(\frac{-2\epsilon^2}{\sum_{i=1}^m (2C/m)^2} \right). \quad (478)$$

Rearranging yields the theorem.

– Remark: the definition of α here differs from the Rahimi/Recht paper.

– Corollary:

* Suppose we had a loss function $\ell(y, v)$ which is 1-Lipschitz in the second argument. (e.g., the hinge loss). Define the expected risk in the usual way:

$$L(f) \stackrel{\text{def}}{=} \mathbb{E}_{(x,y) \sim p^*} [\ell(y, f(x))]. \quad (479)$$

Then the approximation ratio is bounded:

$$L(\hat{f}) - L(f^*) \leq \mathbb{E} [|\ell(y, \hat{f}(x)) - \ell(y, f^*(x))|] \quad [\text{definition, add } |\cdot|] \quad (480)$$

$$\leq \mathbb{E} [|\hat{f}(x) - f^*(x)|] \quad [\text{fix } y, \ell \text{ is Lipschitz}] \quad (481)$$

$$\leq \|\hat{f} - f^*\| \quad [\text{concavity of } \sqrt{\cdot}]. \quad (482)$$

– So far, we have analyzed approximation error due to having a finite m , but assuming an infinite amount of data. Separately, there is the estimation error due to having n data points:

$$L(\hat{f}) - L(\hat{f}_{\text{ERM}}) \leq O_p \left(\frac{C}{\sqrt{n}} \right), \quad (483)$$

where \hat{f}_{ERM} minimizes the empirical risk over the random hypothesis class $\hat{\mathcal{F}}$. So, the total error, which includes approximation error and estimation error is

$$L(\hat{f}_{\text{ERM}}) - L(f^*) = O_p \left(\frac{C}{\sqrt{n}} + \frac{C}{\sqrt{m}} \right). \quad (484)$$

This bound suggests that the approximation and estimation errors are balanced when m and n are on the same order. One takeaway is that we shouldn't over-optimize one without the other. But one might also strongly object and say that if $m \cong n$, then we aren't really getting any savings! This is indeed a valid complaint, and in order to get stronger results, we would need to impose more structure on the problem.

- Dot product kernels (Kar/Karnick, 2012)

- We have seen that shift-invariant kernels admit an integral representation, which allows us to use Monte Carlo to approximate it. What about non-shift invariant kernels such as polynomial kernels, such as the following?

$$k(x, x') = \langle x, x' \rangle^p. \quad (485)$$

- Although random Fourier features will not work, we can still try to write the kernel as an expectation. The key is that if we draw a Rademacher variable $\omega \in \{-1, +1\}^b$ (uniform), randomly projecting x onto ω yields an unbiased estimate of the inner product:

$$\langle x, x' \rangle = \mathbb{E}[\langle \omega, x \rangle \langle \omega, x' \rangle]. \quad (486)$$

Of course, this isn't useful by itself, but it does reduce x to a scalar $\langle \omega, x \rangle$, which is useful.

- To generalize to polynomial kernels, we simply do the above construction p times and multiply it all together. For the quadratic kernel, let ω_1 and ω_2 be two independent Rademacher vectors. Then:

$$\langle x, x' \rangle^2 = \mathbb{E}[\langle \omega_1, x \rangle \langle \omega_1, x' \rangle] \mathbb{E}[\langle \omega_2, x \rangle \langle \omega_2, x' \rangle] \quad (487)$$

$$= \mathbb{E}[\langle \omega_1, x \rangle \langle \omega_2, x \rangle \langle \omega_1, x' \rangle \langle \omega_2, x' \rangle], \quad (488)$$

where the first line follows from the earlier calculation, and the second line follows from independence of ω_1 and ω_2 . Note that $\langle \omega_1, x \rangle \langle \omega_2, x \rangle$ is still just a number.

- More generally, if the kernel is an analytic function of the dot product, then it admits the following Taylor expansion around 0:

$$k(x, x') = f(\langle x, x' \rangle), \quad f(z) = \sum_{j=0}^{\infty} a_j z^j. \quad (489)$$

- To construct a random feature,
 - * Choose J with probability proportional to a_j (if we can't sample from a_j exactly, then we can use importance weighting).

* Choose $\omega_1, \dots, \omega_J$ Rademacher vectors, and let

$$\phi_{\omega_{1:J}}(x) = \prod_{j=1}^J \langle \omega_j, x \rangle. \quad (490)$$

- If we do this m times to form a m -dimensional feature vector, then we have a Monte Carlo estimate of the kernel k . Note that in the process, we have to draw an expected $m\mathbb{E}[J]$ Rademacher variables.
- At this point, we have only showed that we have an unbiased estimate of the kernel. We still need to show that the variance isn't too large. See the paper in the references below for that.

• Nyström method (Williams/Seeger, 2000)

- In the above, we have constructed random features, which were independent of the data. A technique that predates these, which can work better when the spectrum of the kernel matrix is to form a low-rank approximation. This method applies more generically to approximating large PSD matrices.
- Given a kernel matrix $K \in \mathbb{R}^{n \times n}$, we will sample a subset of the indices $I \subset \{1, \dots, n\}$ with $|I| = m$, and let $J = \{1, \dots, n\} \setminus I$ be the other indices. We then evaluate the kernel on points in I paired with all other points, for a total of $O(|I|n)$ evaluations. Then we can define the approximate kernel matrix:

$$K = \begin{pmatrix} K_{II} & K_{IJ} \\ K_{JI} & K_{JJ} \end{pmatrix} \cong \begin{pmatrix} K_{II} & K_{IJ} \\ K_{JI} & \textcolor{red}{K_{JI}K_{II}^\dagger K_{IJ}}, \end{pmatrix} = \tilde{K} \quad (491)$$

or more compactly:

$$\tilde{K} \stackrel{\text{def}}{=} K_{\cdot I} K_{II}^\dagger K_{I \cdot}. \quad (492)$$

Note that the difference $K_{JJ} - \textcolor{red}{K_{JI}K_{II}^\dagger K_{IJ}}$ is the Schur complement of K_{II} . If we interpret K as a covariance matrix of a Gaussian Z , then this is the conditional variance $\text{Var}(Z_J | Z_I)$.

- Note that if K is rank m and K_{II} also contains linearly independent columns (so that it captures the subspace of K), then the Schur complement is zero, and the Nyström method is exact. If not, then the error stems from simply not being able to capture the low rank solution by having a rank m matrix plus an error from doing column sampling (which doesn't yield the eigenvectors). We can think of this as projecting the kernel matrix K on to the subspace of the data points in I .
- How do we choose I ? Two popular choices are uniform sampling and sampling proportional to $K_{ii} = k(x_i, x_i)$, which corresponds to the squared magnitude of x_i . Intuitively, the weighted sampling focuses more energy on points which are more important.

- The following theorem formalizes the error bound:
- **Theorem 29 (Nyström with non-uniform sampling (Drineas/Mahoney, 2005))**
 - * Suppose we choose I by sampling (with replacement), choosing $i \in \{1, \dots, n\}$ with probability $K_{ii} / \sum_{j=1}^n K_{jj}$.
 - * Let \tilde{K}_m be the best rank m approximation of K .
 - * Let \tilde{K} be defined as in (491), but where we replace K_{II} with the best rank m approximation of K_{II} .
 - * Then with probability at least $1 - \delta$,

$$\|K - \tilde{K}\|_F^2 \leq \|K - \tilde{K}_m\|_F^2 + 4(1 + \sqrt{8 \log(1/\delta)}) \operatorname{tr}(K)^2 \sqrt{\frac{m}{|I|}}. \quad (493)$$

- Proof: follows from algebraic manipulation and concentration. Note that the theorem statement is a correct version of Drineas and Mahoney’s Theorem 3, where we just combined equation 31 with Lemma 9.
- The theorem suggests that we should take $|I| > m$, which gives us more opportunities to cover the column space of \tilde{K}_m .

5.8 Universality (skipped in class)

- We have explored several different kernels, and we can (and should) certainly choose one based on domain knowledge.
- But one can ask: is there a general purpose kernel k , in the sense that k can be used to solve *any* learning problem given sufficient data? The notion of general purpose is defined as follows.
- **Definition 27 (universal kernel)**
 - Let \mathcal{X} be a locally compact Hausdorff space (e.g., \mathbb{R}^b or any discrete set, but not infinite-dimensional spaces in general).
 - Let $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ be a kernel.
 - We say that k is a **universal kernel** (specifically, a c_0 -universal kernel) iff the RKHS \mathcal{H} with reproducing kernel k is dense in $C_0(\mathcal{X})$, the set of all continuous bounded functions on \mathcal{X} (with respect to the uniform norm). In other words, for any function $f \in C_0(\mathcal{X})$ and $\epsilon > 0$, there exists some $g \in \mathcal{H}$ such that $\sup_{x \in \mathcal{X}} \|f - g\| \leq \epsilon$.
- The premise is that the target function we want to learn is in $C_0(\mathcal{X})$, so by using a universal kernel, we are defining an RKHS which can approximate any function in $C_0(\mathcal{X})$ as well as we want.

- The following theorem characterizes universal kernels in terms of their Fourier properties:
- **Theorem 30 (Carmeli, 2010)**
 - Let k be a shift-invariant kernel with spectral measure μ on $\mathcal{X} = \mathbb{R}^d$.
 - If the support of μ is all of \mathbb{R}^b , then k is a universal kernel.

Intuition: in order to represent any $C_0(\mathcal{X})$ function, we must not have any gaps in our spectrum.

- Example: the Gaussian kernel is universal; the sinc kernel is not universal.
- The final piece of the puzzle is **universal consistency**, which means that that a learning algorithm will actually achieve the best possible error as the number of training examples tends to infinity. Steinwart showed that using SVMs with a universal kernel guarantees universal consistency. Of course, universality is only about how well we can represent the target function; it says nothing about readily we can actually estimate that function based on finite data.

5.9 RKHS embedding of probability distributions (skipped in class)

- So far, we've showed that kernels can be used for estimating functions for regression, classification, dimensionality reduction (PCA), etc. Now we will show how kernels can be used to represent and answer questions about probability distributions without having to explicitly estimate them.
- As a motivating example, consider the problem of testing whether two probability distributions P and Q are the same by only observing expectations under the distributions.
- Given a distribution P , we can look at various **moments** of the distribution $\mathbb{E}_{x \sim P}[f(x)]$ for various functions f . For example, if $f(x) = x$, then we get the mean. Such a f only gives us partial information about P : if two distributions differ in their mean, then we know they are different, but if they have the same mean, we cannot conclude that they are same.
- More generally, assume P and Q are defined on some locally compact Hausdorff space \mathcal{X} (e.g., \mathbb{R}^b). Define the **maximum mean discrepancy** (MMD) as follows:

$$D(P, Q, \mathcal{F}) \stackrel{\text{def}}{=} \sup_{f \in \mathcal{F}} (\mathbb{E}_{x \sim P}[f(x)] - \mathbb{E}_{x \sim Q}[f(x)]), \quad (494)$$

for some set of functions \mathcal{F} . Shorthand: $\mathbb{E}_P[f]$ means $\mathbb{E}_{x \sim P}[f(x)]$.

- Can we find \mathcal{F} so that

$$D(P, Q, \mathcal{F}) = 0 \Leftrightarrow P = Q? \quad (495)$$

Note that $P = Q$ always implies $D(P, Q, \mathcal{F}) = 0$, but the other direction requires some work.

- If we knew P and Q were Gaussian, then it suffices to take $\mathcal{F} = \{x \mapsto x, x \mapsto x^2\}$, since the first two moments define a Gaussian distribution. However, what about general P and Q ? We need a much larger class of functions \mathcal{F} :

- **Theorem 31 (Dudley, 1984)**

- If $\mathcal{F} = C_0(\mathcal{X})$ (all continuous bounded functions), then $D(P, Q, \mathcal{F}) = 0$ implies $P = Q$.

- However, $C_0(\mathcal{X})$ is a large and difficult set to work with. Fortunately, it suffices to take \mathcal{F} to be any set that is dense in $C_0(\mathcal{X})$, in particular an RKHS:

- **Theorem 32 (Steinwart, 2001)**

- Let $\mathcal{F} = \{f \in \mathcal{H} : \|f\|_{\mathcal{H}} \leq 1\}$, where \mathcal{H} is the RKHS defined by a universal kernel k .
- Then $D(P, Q, \mathcal{F}) = 0$ implies $P = Q$.

- **Proof:**

- Let $P \neq Q$ be two distinct distributions.
- Then there exists $f \in C_0(\mathcal{X})$ such that $|\mathbb{E}_P[f] - \mathbb{E}_Q[f]| = \epsilon > 0$.
- Since \mathcal{H} is universal (i.e., \mathcal{H} is dense in $C_0(\mathcal{X})$ with respect to the uniform norm), there exists $g \in \mathcal{H}$ with $g \neq 0$ such that

$$\|f - g\|_{\infty} \stackrel{\text{def}}{=} \sup_{x \in \mathcal{X}} |f(x) - g(x)| \leq \epsilon/3. \quad (496)$$

- This means $|\mathbb{E}_P[f] - \mathbb{E}_P[g]| \leq \epsilon/3$ and $|\mathbb{E}_Q[f] - \mathbb{E}_Q[g]| \leq \epsilon/3$.
- By the triangle inequality, $|\mathbb{E}_P[g] - \mathbb{E}_Q[g]| \geq \epsilon/3 > 0$.
- Rescale g : let $u = g/\|g\|_{\mathcal{H}} \in \mathcal{F}$.
- We still have $D(P, Q, \mathcal{F}) \geq |\mathbb{E}_P[u] - \mathbb{E}_Q[u]| > 0$.

- **Computing $D(P, Q, \mathcal{F})$**

- We've established that $D(P, Q, \mathcal{F})$ contains sufficient information for testing whether two distributions are equal. But how do we actually compute the max over \mathcal{F} ? This seems daunting at first sight. Fortunately, we can compute $D(P, Q, \mathcal{F})$ in closed form by exploiting properties of the RKHS.

- First, a general statement. By the reproducing property and linearity of the inner product, we can express expected function value as an inner product:

$$\mathbb{E}_{x \sim P}[f(x)] = \mathbb{E}_{x \sim P}[\langle k(\cdot, x), f \rangle] = \left\langle \underbrace{\mathbb{E}_{x \sim P}[k(x, \cdot)]}_{\stackrel{\text{def}}{=} \mu_P}, f \right\rangle. \quad (497)$$

Here, $\mu_P \in \mathcal{H}$ is the **RKHS embedding** of the probability distribution P .

- We can now write the MMD solution as follows:

$$D(P, Q, \mathcal{F}) = \sup_{f \in \mathcal{F}} \langle \mu_P - \mu_Q, f \rangle = \|\mu_P - \mu_Q\|_{\mathcal{H}}, \quad (498)$$

where the sup is obtained by setting f to be a unit vector in the direction of $\mu_P - \mu_Q$.

- Unpacking the square of the last expression and rewriting in terms of kernel evaluations:

$$\|\mu_P - \mu_Q\|_{\mathcal{H}}^2 = \mathbb{E}_{P \times P}[k(x, x')] - \mathbb{E}_{P \times Q}[k(x, y)] - \mathbb{E}_{Q \times P}[k(y, x)] + \mathbb{E}_{Q \times Q}[k(y, y')]. \quad (499)$$

- Of course, in practice, we only have samples from P, Q : let $x_1, \dots, x_n \sim P$ and $y_1, \dots, y_n \sim Q$ all be drawn independently.
- We can obtain an empirical estimate of $D(P, Q, \mathcal{F})$ as a U-statistic (a U-statistic is a function which is an average over some function applied to all pairs of points):

$$\hat{D}_n(P, Q, \mathcal{F}) = \frac{1}{\binom{n}{2}} \sum_{i < j} [k(x_i, x_j) - k(x_i, y_j) - k(y_i, x_j) + k(y_i, y_j)]. \quad (500)$$

This estimate is unbiased, since the expectation of each term is $D(P, Q, \mathcal{F})$.

- Let the null hypothesis be that $P = Q$. Under the null hypothesis, as $n \rightarrow \infty$, we know that $\hat{D}_n(P, Q, \mathcal{F}) \xrightarrow{P} 0$, but in order to use \hat{D}_n as a test statistic for hypothesis testing, we need to know its (approximate) distribution. (Recall that $\hat{D}_n(P, Q, \mathcal{F})$ is a random variable that is a function of the data points.)
- There are two ways to go about this:
 - We can derive finite sample complexity bounds to bound the deviation of $\hat{D}(P, Q, \mathcal{F})$ from its mean $D(P, Q, \mathcal{F})$.
 - We can show that $\hat{D}(P, Q, \mathcal{F})$ is asymptotically normal with some variance, and use the normal as an approximation of the distribution.
- In the next section, we will develop the tools to analyze random variables such as these.

5.10 Summary (Lecture 16)

- We began by noting that some algorithms (e.g., online gradient descent) do not require arbitrary inner products between weight vectors and feature vectors, but only **inner products between feature vectors**.
- This motivated the use of **kernels** (defined to be positive semidefinite functions), which can provide both computational (ability to implicitly compute inner products between infinite-dimensional feature vectors) and modeling advantages (thinking in terms of similarities between two inputs).
- Taking a step back, we saw that all that matters at the end of the day are functions evaluated at various inputs. This motivated the definition of **reproducing kernel Hilbert spaces** (RKHS), in which two important properties hold: (i) function evaluations were bounded (meaningful), and (ii) there is a nice inner product structure.
- We showed that the three distinct viewpoints above (features, kernels, functions) are actually all equivalent (due to the Moore-Aronszajn theorem).
- The **representer theorem** shows that the optimum over an appropriately regularized function space \mathcal{H} is attained by a function in the span of the training data. This allows us to derive kernelized SVMs, kernelized regression, kernel PCA, RKHS embeddings of probability distributions.
- **Bochner's theorem**, allows us to study the Fourier properties of shift-invariant kernels, relating universality and smoothness properties of a kernel to the frequencies that the kernel passes through.
- Bochner's theorem allowed us to obtain computationally efficient ways to approximate kernel methods by writing kernels as an integral over dot products of **random features**. This leads to efficient algorithms. but have uncertainty estimates over function values. Uncertainty estimates are critical for active learning and Bayesian optimization.

5.11 References

- Hofmann/Scholkopf/Smola, 2008: Kernel Methods in Machine Learning
- Drineas/Mahoney, 2005: On the Nystrom Method for Approximating a Gram Matrix for Improved Kernel-Based Learning
- Rahimi/Recht, 2008: Weighted Sums of Random Kitchen Sinks: Replacing minimization with randomization in learning
- Yang/Li/Mahdavi/Jin/Zhou, 2012: Nystrom Method vs Random Fourier Features: A Theoretical and Empirical Comparison
- Kar/Karnick, 2012: Random Feature Maps for Dot Product Kernels

6 Neural networks

6.1 Motivation (Lecture 17)

- One major motivating factor for studying neural networks is that they have had a lot of empirical success and attention in recent years. Here is the general recipe:
 - Train on large amounts of data (Krizhevsky, 2012: 1 million examples).
 - Use a very large neural network (Krizhevsky, 2012: 60 million parameters).
 - Running simple stochastic gradient descent (with some (important) tweaks such as step size control, sometimes momentum, dropout), and wait a moderately long period of time (a week).
 - Get state-of-the-art results across multiple domains.
 - * Object recognition (ImageNet): reduce error from 25.7% to 17.0% (Krizhevsky, 2012)
 - * Speech recognition (Switchboard): reduce error from 23% to 13% (Microsoft, 2009–2013)

Error reductions on these tasks/datasets are significant, since these are large realistic datasets (unlike MNIST) on which many serious people have tried to improve accuracy.

- However, in theory, neural networks are poorly understood:
 - The objective function is non-convex. SGD has no reason to work.
 - What about the particular hypothesis class of neural networks makes them perform well across so many tasks? What types of functions are they good at representing?

This makes neural networks an important but challenging area to do new theoretical research.

- Compared to the theory of online learning, uniform convergence, or kernel methods, the theory of neural networks (“why/when do they work?”) is spotty at best. In this lecture, we will attempt lay out the high-level important questions and provide preliminary thrusts towards answering them, which in turn produces a series of more concrete open questions.
- A caveat: there are many theorems that one could prove about neural networks. We will write some of these down. However, most of these theorems only nibble off the

corners of the problem, and do not really answer the hard questions. But that's where we are. We will try to point out the interpretation of these results, which is especially important to recognize in this prehistoric stage of understanding.

6.2 Setup (Lecture 17)

- Definition

- Let $x \in \mathbb{R}^d$ be an input.
- A neural network defines a non-linear function on x . For concreteness, let us focus on two-layer neural networks (which means it has one hidden layer).
- Let $\sigma : \mathbb{R} \mapsto \mathbb{R}$ be a non-linear function (called an activation or transfer function). Examples:

- * Logistic: $z \mapsto \frac{1}{1+e^{-z}}$ (maps real numbers monotonically to $[0, 1]$)
- * Hyperbolic tangent (tanh): $z \mapsto \frac{e^z - e^{-z}}{e^z + e^{-z}}$ (maps real numbers monotonically to $[-1, 1]$)
- * Rectified linear: $z \mapsto \max(0, z)$ (truncates negative numbers to zero)

We will extend σ to operate on vectors elementwise:

$$\sigma(x) \stackrel{\text{def}}{=} [\sigma(x_1), \dots, \sigma(x_d)]. \quad (501)$$

- An artificial **neural network** (not to be confused with and completely different from the thing that's in your brain) can be described by a function $f : \mathbb{R}^d \mapsto \mathbb{R}$ which has the following form:

$$f_\theta(x) = \sum_{i=1}^m \alpha_i \sigma(w_i \cdot x + b_i). \quad (502)$$

Here, the parameters are $\theta = (\alpha_{1:m}, w_{1:m}, b_{1:m})$. In matrix notation:

$$f_\theta(x) = \alpha \cdot \sigma(Wx + b), \quad (503)$$

where the i -th row of W is $w_i \in \mathbb{R}^d$.

- A useful way to think about a neural network is that the first layer $\sigma(Wx + b)$ computes some non-linear features of the input and the second layer simply performs a linear combination. So you can think of a neural network as parametrizing the features as well as the weights α .
- FIGURE: [draw neural network]
- In practice, people use many layers (for example, in speech recognition, seven is not uncommon). A three-layer neural network looks like this:

$$f_\theta(x) = \sigma(W_2 \sigma(W_1 x + b_1) + b_2). \quad (504)$$

People also use convolutional neural networks in which W has additional low-dimensional structure. Recurrent neural networks are good for representing time series. We will not discuss these here and focus on the two-layer neural networks in this lecture.

- Let $\mathcal{F} = \{f_\theta\}$ be the class of all neural networks as we range over values of θ .
- We can use \mathcal{F} for regression or classification in the usual way by minimizing the loss on the training set. For regression:

$$\hat{f}_{\text{ERM}} = \arg \min_{f_\theta} \hat{L}(\theta), \quad \hat{L}(\theta) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n (f_\theta(x^{(i)}) - y^{(i)})^2. \quad (505)$$

- The surprising thing about neural networks is that people use stochastic gradient descent (online gradient descent where we choose $i \in \{1, \dots, n\}$ at each step randomly), which converges to \hat{f}_{SGD} . Since the objective function \hat{L} is non-convex, \hat{f}_{SGD} will be different from \hat{f}_{ERM} . so there is an additional **optimization error** in addition to the usual approximation and estimation errors. The decomposition must be taken with a grain of salt, because having suboptimal optimization effectively restricts the hypothesis class, which actually improves estimation error (a simple example is early stopping). So approximation and optimization error are perhaps difficult to tease apart in methods whose success is tied to an algorithm not just an optimization problem.




Figure 6: Cartoon showing error decomposition into approximation, estimation, and optimization errors.

6.3 Approximation error (universality) (Lecture 17)

- Question: which functions can be represented by two-layer neural networks? Answer: basically all of them.
- We saw that Gaussian kernels were universal (Theorem 30) in the sense that they are dense in $C_0(\mathcal{X})$, the set of all continuous bounded functions on \mathcal{X} .
- Functions defined by (Gaussian) kernels can be represented in terms of the partial kernel evaluations

$$f(x) = \sum_{i=1}^n \alpha_i k(x_i, x). \quad (506)$$

or using Fourier features:

$$f(x) = \sum_{i=1}^m \alpha_i e^{-i\langle w_i, x \rangle}. \quad (507)$$

Note that these expressions resemble neural networks (502). All of these are linear combinations of some non-linear basis functions. So we might expect that neural networks would be universal. The following theorem answers affirmatively, which is not hard to show:

- **Theorem 33 (neural networks are universal)**

- Consider $\mathcal{X} = [0, 1]^d$.
- Then the class of two-layer neural networks \mathcal{F} is dense in $C_0(\mathcal{X})$ in the uniform metric: for every $f^* \in C_0(\mathcal{X})$, and $\epsilon > 0$, there exists an $f_\theta \in \mathcal{F}$ such that $\max_{x \in \mathcal{X}} |f^*(x) - f_\theta(x)| \leq \epsilon$.

- **Proof of Theorem 33:**

- Fix $f^* \in \mathcal{F}$. Because f^* is continuous over a compact set (\mathcal{X}), it is uniformly continuous, which means we can find a width δ such that $|x_1 - x_2| \leq \delta$ implies $|f(x_1) - f(x_2)| \leq \epsilon$.
- Grid \mathcal{X} into regions R_j which are of the form: $R_j = \{x : |x - x_j|_\infty \leq \delta\}$.
- FIGURE: [1D function]
- Note that $z \mapsto \sigma(Cz)$ converges to the step function as $C \rightarrow \infty$.
- Let us construct $3d$ hidden units for each R_j . For $d = 1$, let $R_j = [x_j - \delta, x_j + \delta] = [a_j, b_j]$.

$$f^*(x_j)\sigma(C(x - a_j)) + f^*(x_j)\sigma(C(b_j - x)) - f^*(x_j)\sigma(0). \quad (508)$$

This function is approximately $f^*(x_j)$ on R_j and zero elsewhere, So we we do this for all regions, then we can approximate f^* uniformly well.

- Note that the number of hidden units is exponential in d , since the number of regions is $O((\frac{1}{\delta})^d)$. This is just a glorified nearest neighbors. So while this theorem holds, it doesn't really provide much insight into why neural networks work well and generalize.
- Depth: One argument made in favor of deep neural networks is that they more compactly represent the data compared to a shallower one. To get intuition, let us think about networks that perform arithmetic operations on the raw inputs via addition and multiplication. Such networks define polynomial functions in a compact way, since internal nodes represent factors. For example, the polynomial

$$f(x) = x_1 x_2^2 x_3 + x_1 x_2 x_3 x_4 + x_2^2 x_3^3 + x_2 x_3^3 x_4 \quad (509)$$

$$= (a + b)(b + c), \quad (510)$$

where $a = x_1 x_2$, $b = x_2 x_3$, and $c = x_3 x_4$ correspond to internal nodes.

6.4 Generalization bounds (Lecture 17)

- The excess risk $L(\hat{h}) - L(h^*)$ captures the estimation error (the generalization ability) of an estimator \hat{h} . Recall that the excess risk is controlled by the Rademacher complexity (215) of the function class, so it suffices to study the Rademacher complexity. We will perform the analysis for two-layer neural networks.
- **Theorem 34 (Rademacher complexity of neural networks)**
 - Let $x \in \mathbb{R}^d$ be the input vector with $\|x\|_2 \leq C_2$.
 - Let $w_j \in \mathbb{R}^d$ be the weights connecting to the j -th hidden unit, $j = 1, \dots, m$
 - Let $\alpha \in \mathbb{R}^m$ be the weights connecting the hidden units to the output
 - Let $h : \mathbb{R} \rightarrow \mathbb{R}$ be a non-linear activation function with Lipschitz constant 1; examples include
 - * Logistic: $h(z) = (1 + e^{-z})^{-1}$
 - * Rectified linear: $h(z) = \max\{0, z\}$
 - For each set of weights (w, α) , define the predictor:

$$f_{w,\alpha}(x) = \sum_{j=1}^m v_j h(w_j \cdot x). \quad (511)$$

- Let \mathcal{F} be the class of prediction functions where the weights are bounded:

$$\mathcal{F} = \{f_{w,\alpha} : \|\alpha\|_2 \leq B'_2, \|w_j\|_2 \leq B_2 \text{ for } j = 1, \dots, m\}. \quad (512)$$

- Then

$$R_n(\mathcal{F}) \leq \frac{B_2 B'_2 C_2 \sqrt{m}}{\sqrt{n}}. \quad (513)$$

- Proof of Theorem 34:
 - The key is that the composition properties of Rademacher complexity aligns very nicely with the layer-by-layer compositionality of neural networks.
 - The function mapping the input layer to a hidden unit is just a linear function:

$$R_n(\{x \mapsto w_j \cdot x : \|w_j\|_2 \leq B_2\}) \leq \frac{B_2 C_2}{\sqrt{n}}. \quad (514)$$

- Sending each of these through the 1-Lipschitz non-linearity does not change the upper bound on the complexity:

$$R_n(\{x \mapsto h(w_j \cdot x) : \|w_j\|_2 \leq B_2\}) \leq \frac{B_2 C_2}{\sqrt{n}}. \quad (515)$$

For convenience, define \mathcal{G} as the set of vector-valued functions mapping the input to the hidden activations, as w ranges over different values:

$$g(x) = [h(w_1 \cdot x), \dots, h(w_m \cdot x)]. \quad (516)$$

– Now let us handle the second layer:

$$R_n(\mathcal{F}) \leq \mathbb{E} \left[\sup_{\|\alpha\|_2 \leq B'_2, g \in \mathcal{G}} \frac{1}{n} \sum_{i=1}^n \sigma_i(\alpha \cdot g(Z_i)) \right] \quad (517)$$

– Applying Cauchy-Schwartz:

$$R_n(\mathcal{F}) \leq B'_2 \mathbb{E} \left[\sup_{g \in \mathcal{G}} \left\| \frac{1}{n} \sum_{i=1}^n \sigma_i g(Z_i) \right\|_2 \right]. \quad (518)$$

– Using the fact that if $a \in \mathbb{R}^m$ then $\|a\|_2 \leq \sqrt{m} \max_{1 \leq j \leq m} |a_j|$:

$$R_n(\mathcal{F}) \leq B'_2 \sqrt{m} \mathbb{E} \left[\max_{1 \leq j \leq m} \sup_{\|w_j\|_2 \leq B_2} \left| \frac{1}{n} \sum_{i=1}^n \sigma_i h(w_j \cdot Z_i) \right| \right]. \quad (519)$$

– This is the same as taking the sup over a single generic w subject to $\|w\|_2 \leq B_2$, so by (515), we have:

$$R_n(\mathcal{F}) \leq B'_2 \sqrt{m} \left(\frac{B_2 C_2}{\sqrt{n}} \right). \quad (520)$$

Admittedly, there is a small sleight of hand in this last step because our definition of Rademacher complexity does not have an absolute value, whereas (519) does. In some parts of the literature, Rademacher complexity is defined with absolute values. For that definition, the Lipschitz composition property still holds and the Rademacher complexity of linear functions doesn't change, so our result holds anyway.

- Interpretation: the bottom line of this theorem is that neural networks with non-linearities which are smooth (Lipschitz) will generalize regardless of convexity.

6.5 Approximation error for polynomials (Lecture 17)

- There are two problems with the previous analyses:
 - First, we defined approximation error with respect to the set of all bounded continuous functions, which is clearly much too rich of a function class, and as a result we get exponential dependence on the dimension d .
 - Second, we have not yet said anything about optimization error.

- We will now present part of an ICML paper (Andoni/Panigrahy/Valiant/Zhang), which makes some progress on both points. We will focus on using neural networks to approximate bounded-degree polynomials, which is a reasonable class of functions which itself approximates Lipschitz functions. Further, we will show that choosing weights W *randomly* and optimizing α (which is convex) is sufficient.

- Setup

- Monomial $x^J = x^{J_1} \dots x^{J_d}$ with degrees $J = (J_1, \dots, J_d)$
 - * Example: $J = (1, 0, 7)$, $x^J = x_1 x_3^7$
- A polynomial is $f(x) = \sum_J b_J x^J$
 - * Example: $b_{(1,0,7)} = 3$ and $b_{(0,1,0)} = -5$, $f(x) = 3x_1 x_3^7 - 5x_2$
- Degree of polynomial: $\deg(f) = \max_{b_J \neq 0} |J|$, $|J| \stackrel{\text{def}}{=} \sum_{i=1}^d J_i$

- Inner product structure

- Let $p^*(x)$ is the uniform distribution over $\mathbb{C}(R)^d$, where $\mathbb{C}(R)$ is set of complex numbers c with norm at most R ($|c| = \sqrt{c\bar{c}} \leq R$).
- Define the inner product over functions and the associated norm:

$$\langle f, g \rangle_{p^*} \stackrel{\text{def}}{=} \mathbb{E}_{x \sim p^*} [f(x) \overline{g(x)}] \quad (521)$$

$$\|f\|_{p^*} \stackrel{\text{def}}{=} \sqrt{\langle f, f \rangle_{p^*}}. \quad (522)$$

- Neural network as an infinite polynomial

- Let $\sigma(z) = \sum_{j \geq 0} a_j z^j$
 - * Example: for $\sigma(z) = e^z$, $a_j = \frac{1}{j!}$
- For $w \in \mathbb{C}^d$, we define the basis function, which can be written as a weighted combination of monomials x^J :

$$\phi^w(x) \stackrel{\text{def}}{=} \sigma(w \cdot x) = \sum_{j \geq 0} a_j \left(\sum_{k=1}^d w_k x_k \right)^j \stackrel{\text{def}}{=} \sum_J a_J w^J x^J. \quad (523)$$

Here, $w^J = \prod_{k=1}^d w^{J_k}$ is the product of the base weights.

- The neural network is

$$f(x) = \sum_{i=1}^m \alpha_i \phi^{w_i}(x). \quad (524)$$

- Note that weights $\{w^J\}$ are orthogonal in the following sense:

$$\mathbb{E}_{w \sim \mathbb{C}(R)^d} [w^J \overline{w^{J'}}] = \begin{cases} R^{2|J|} & \text{if } J = J' \\ 0 & \text{otherwise.} \end{cases} \quad (525)$$

This stems from the fact that each $w_j \sim \mathbb{C}(R)$ means $w_j = e^{it}$ where $t \sim \text{Uniform}([0, 2\pi])$; and $\mathbb{E}_{t \sim \text{Uniform}([0, 2\pi])} [e^{iat} \overline{e^{ibt}}] = \mathbb{I}[a = b]$.

• **Theorem 35 (neural networks approximate bounded-degree polynomials)**

- Choose w_1, \dots, w_m i.i.d. from $\mathbb{C}(1/\sqrt{d})^d$.
- For any polynomial f^* of degree q and norm $\|f^*\| = 1$, exists $\alpha_{1:m}$ with $\|\alpha_{1:m}\|_2 = \sum_{i=1}^m |\alpha_i|^2 = O(d^{2q}/m)$ such that

$$\left\| \sum_{i=1}^m \alpha_i \phi^{w_i} - f^* \right\|_{p^*} = O_p \left(\sqrt{\frac{d^{2q}}{m}} \right). \quad (526)$$

• Proof of Theorem 35:

- Let the target function be:

$$f^*(x) = \sum_J b_J x^J. \quad (527)$$

- Construct an unbiased estimate of f^* :

- * For any $x \in \mathbb{C}^d$ and $J \in \mathbb{N}^d$,

$$\mathbb{E}_{w \sim \mathbb{C}(R)^d} [\phi^w(x) \overline{w^J}] = a_J R^{2|J|} x^J, \quad (528)$$

which follows from applying (525) to the expansion (523) and noting that all random coefficients except w^J drop out.

- * Define the coefficients

$$T(w) = \sum_J \frac{b_J \overline{w^J}}{a_J R^{2|J|}}. \quad (529)$$

- * Then by construction, we have an unbiased estimate of the target polynomial:

$$\mathbb{E}_{w \sim \mathbb{C}(R)^d} [T(w) \phi^w(x)] = f^*(x). \quad (530)$$

- Now we would like to control the variance of the estimator.

- * First, define the error:

$$\eta(x) \stackrel{\text{def}}{=} \frac{1}{m} \sum_{i=1}^m T(w_i) \phi^{w_i}(x) - f^*(x). \quad (531)$$

* We can show that the variance falls as $1/m$ due to independence of w_1, \dots, w_m :

$$\mathbb{E}_{w_1, \dots, w_m \sim \mathbb{C}(R)^d} [\|\eta\|_{p^*}^2] \leq \frac{1}{m} \mathbb{E}_{w \sim \mathbb{C}(R)^d} [|T(w)|^2 \|\phi^w\|_{p^*}^2]. \quad (532)$$

* To bound the coefficients:

- Let $a(q) = \min_{|J| \leq q} |a_J|$ (property of the non-linear activation function)
- Let $\|f^*\|_1 = \sum_J |b_J|$ (property of the target function)
- Let $q = \deg(f^*)$.
- Assume $R \leq 1$.
- Then

$$|T(w)| \leq \sum_J \left| \frac{b_J \overline{w^J}}{a_J R^{2|J|}} \right| \quad (533)$$

$$\leq \sum_J \left| \frac{b_J}{a_J R^{|J|}} \right| \quad [\text{since } |\overline{w^J}| \leq R^{|J|}] \quad (534)$$

$$\leq \frac{\sum_J |b_J|}{a(q) R^q} \quad [\text{since } |J| \leq q, a(q) \leq |a_J|] \quad (535)$$

$$= \frac{\|f^*\|_1}{a(q) R^q}. \quad (536)$$

* Define the following bounds on the variance of the basis functions ϕ^w and the 1-norm of the target function f^* , respectively:

$$\beta(q, R) \stackrel{\text{def}}{=} \mathbb{E}_{w \sim \mathbb{C}(R)^d} \left[\frac{\|\phi^w\|_{p^*}^2}{R^{2q}} \right], \quad (537)$$

$$\gamma(q) \stackrel{\text{def}}{=} \max_{\|f^*\|_{p^*}=1} \|f^*\|_1. \quad (538)$$

Then we have the following result (by plugging quantities in and applying concavity of $\sqrt{\cdot}$):

$$\mathbb{E}_{w_1, \dots, w_m \sim \mathbb{C}(R)^d} [\|\eta\|_{p^*}] \leq \sqrt{\frac{\gamma(q)^2 \beta(q, R)}{a(q)^2 m}}. \quad (539)$$

– Finally, we specialize to $\sigma(z) = e^z$ and $p^* = \mathbb{C}(1)^d$.

* Claim: $a(q) \geq 1/q!$

- We have $a_j = 1/j!$ and a_J is just a_j times a positive integer stemming from the multiplicities in J (for example, if $J = (2, 0, 3)$, then $a_J = 2!3!a_j \geq a_j$).

* Claim: $\beta(q, R) = O(d^q)$ if $R = O(1/\sqrt{d})$.

· We have that

$$\mathbb{E}_{w \sim \mathbb{C}(R)^d, x \sim \mathbb{C}(1)^d} [\phi^w(x) \overline{\phi^w(x)}] = \sum_J a_J^2 R^{2|J|} \quad (540)$$

$$= O(e^{2\sqrt{d}R}). \quad (541)$$

· Finally, $\beta(q, R) = O(e^{2\sqrt{d}R}/R^{2q}) = O(d^q)$ by definition of R . Note that we need R to be small to fight the explosion stemming from $e^{2\sqrt{d}}$.

* Claim: $\gamma(q) = O(\sqrt{d^q})$

* Claim: the coefficients are bounded:

$$\sum_{i=1}^m |\alpha_i|^2 = \frac{1}{m^2} \sum_{i=1}^m |T(w_i)|^2 \leq \frac{\|f^*\|_1^2}{ma(q)^2 R^{2q}} = O(d^{2q}/m). \quad (542)$$

See the paper for more details.

6.6 References

- Telgarsky, 2012: Representation Power of Feedforward Neural Networks (slides)
- Andoni/Panigrahy/Valiant/Zhang, 2014: Learning Polynomials with Neural Networks
- Livni/Shalev-Shwartz/Shamir, 2014: On the Computational Efficiency of Training Neural Networks
- Livni/Shalev-Shwartz/Shamir, 2014: An Algorithm for Training Polynomial Networks
- Schmidhuber, 2014: Deep Learning in Neural Networks: An Overview

7 Method of moments

Note: only a small part of this unit was covered in class.

7.1 Motivation (Lecture 17)

- In the last section, with maximum likelihood estimation in exponential families, we shifted from minimizing expected risk (where only the predictions of the model matter) to parameter estimation (where the parameters themselves matter). Parameter estimation is a stronger requirement: if you can estimate the parameters accurately, then you usually can predict accurately.
- However, obtaining accurate parameter estimates is not needed for predicting accurately. For example, consider a linear regression task with two identical features ($x \in \mathbb{R}^2$ and $x_1 = x_2$ always). In this case, if we have two distinct parameters θ and θ' with $\theta_1 + \theta_2 = \theta'_1 + \theta'_2$, then the predictions are the same ($\theta \cdot x = \theta' \cdot x$).

- For exponential families where all the variables are observed, parameter estimation via maximum likelihood or maximum pseudolikelihood results in convex optimization problems. However, **latent-variable models** such as Gaussian mixture models (GMMs), Hidden Markov Models (HMMs), and Latent Dirichlet Allocation (LDA) are quite useful in practice, because latent-variables often provide more compact representations of the data which are learned automatically.
- The main downside to latent-variable models is that standard estimation procedures such as maximum likelihood result in **non-convex** optimization problems. In practice, people use Expectation Maximization (EM) to optimize these objective functions,¹⁹ but EM is only guaranteed to converge to a local optimal setting of the parameters which can be arbitrarily far away from the global optimum.
- In this section, we will explore alternative techniques for parameter estimation which do not rely on likelihood-based estimation, but instead rely on **method of moments**. The method of moments dates back to Karl Pearson from 1894, which was before the maximum likelihood principle was fully developed by Fisher in the 1920s. Since then, maximum likelihood has been the dominant paradigm for parameter estimation, mainly due to its statistical efficiency and naturalness. There is more ad-hocness in the method of moments, but it is exactly this ad-hocness that allows us to get computationally efficient algorithms that converges in a sense to the **globally optimal solution** at the price of reduced statistical efficiency, a theme that we saw already with maximum likelihood and maximum pseudolikelihood.
- The abstract unsupervised learning problem is as follows:
 - Let x be the observed variables and h be the hidden variable.
 - Define model: $p_\theta(x, h)$ for $\theta \in \Theta \subseteq \mathbb{R}^d$
 - Input: examples $x^{(1)}, \dots, x^{(n)}$ drawn i.i.d. from p_{θ^*}
 - Output: parameter estimates $\hat{\theta}$
 - Throughout this section, we will assume that our models are well-specified, that is, data is actually generated from our model.

As a concrete example, consider the following latent-variable model:

- **Example 38 (Naive Bayes clustering model)**
 - Let k be the number of possible document clusters.
 - Let b be the number of possible word types in a vocabulary.
 - Model parameters $\theta = (\pi, B)$

¹⁹ One can use other optimization methods such as gradient descent, L-BFGS. Bayesian formulations are tackled by either variational inference, which also converge to only local optima, or MCMC methods are guaranteed to converge in the limit, but can take exponentially long.

- * $\pi \in \Delta_k$: prior distribution over clusters.
- * $B = (\beta_1, \dots, \beta_k) \in (\Delta_b)^k$: distributions over words given cluster

Let Θ denote the set of valid parameters.

- The generative model is as follows: For each document $i = 1, \dots, n$:
 - * Sample the cluster: $h^{(i)} \sim \pi$
 - * For each word $j = 1, \dots, L$:
 - Sample a word: $x_j^{(i)} \sim \beta_{h^{(i)}}$.
- Maximum (marginal) likelihood is the standard approach to parameter estimation, which is non-convex:

$$\min_{\theta \in \Theta} \sum_{i=1}^n -\log \sum_{h=1}^k p_{\theta}(h, x^{(i)}). \quad (543)$$

A popular optimization algorithm is to use the EM algorithm, which can be shown to be either a bound optimization algorithm (which repeatedly constructs a lower bound and optimize) or coordinate-wise ascent on a related objective.

- E-step: for each example i , compute the posterior $q_i(h) = p_{\theta}(h^{(i)} = h \mid x^{(i)})$.
- M-step: optimize the expected log-likelihood: $\max_{\theta} \sum_{i=1}^n \sum_{h=1}^k q_i(h) \log p_{\theta}(h, x^{(i)})$.

The EM algorithm is widely used and can get excellent empirical results, although there are no theoretical guarantees in general that EM will converge to a global optimum, and in practice, it can get stuck in bad local optima.

7.2 Method of moments framework (Lecture 17)

- The method of moments involves two steps:
 - Step 1: define a **moment mapping** relating the model parameters to properties (i.e., moments) of the data distribution specified by those parameters.
 - Step 2: **plug in** the empirical moments and invert the mapping to get parameter estimates.
- Moment mapping
 - Let $\phi(x) \in \mathbb{R}^p$ be an observation function which only depends on the observed variables x .
 - * Example: $\phi(x) = (x, x^2)$.




Figure 7: Schema for method of moments: We define a moment mapping M from parameters to moments. We estimate the moments (always easy) and invert the mapping to recover parameter estimates (sometimes easy).

- Define the **moment mapping**

$$M(\theta) \stackrel{\text{def}}{=} \mathbb{E}_{x \sim p_\theta}[\phi(x)], \quad (544)$$

which maps each parameter vector $\theta \in \mathbb{R}^d$ to the expected value of $\phi(x)$ with respect to $p_\theta(x)$. This mapping is the key that links moments (which are simple functions of the observed data) with the parameters (which are quantities that we want to estimate).

- Example (Gaussian distribution):

- * Suppose our model is a univariate Gaussian with parameters $\theta = (\mu, \sigma^2)$.
- * Then for M defined above, the moment equations are as follows:

$$M((\mu, \sigma^2)) = \mathbb{E}_{x \sim \mathcal{N}(\mu, \sigma^2)}[(x, x^2)] = (\mu, \sigma^2 + \mu^2). \quad (545)$$

- Let's see how moment mappings are useful. Suppose that someone told us some moments m^* (where $m^* = M(\theta^*)$). Then assuming M were invertible, we could solve for $\theta^* = M^{-1}(m^*)$. Existence of the inverse is known as identifiability:

- **Definition 28 (identifiability)**

- * Let $\Theta \subseteq \mathbb{R}^d$ be a set of parameters.
- * Define the confusable set for $\theta \in \Theta$ as

$$S(\theta) \stackrel{\text{def}}{=} \{\theta' \in \Theta : M(\theta) = M(\theta')\}. \quad (546)$$

This is the set of parameters that have the same moments as θ .

- * We say that Θ is *identifiable from ϕ* if for almost $\theta \in \Theta$, we have $|S(\theta)| = 1$.

- Example (Gaussian distribution): We can recover the parameters $\theta^* = (\mu^*, \sigma^{2*})$ from $m^* = (m_1^*, m_2^*)$ as follows:

- * $\mu^* = m_1^*$
- * $\sigma^{2*} = m_2^* - m_1^{2*}$

Thus, the parameters are identifiable given the first two moments.

- Plug-in

- In practice, of course, we don't have access to the true moments m^* . However, the key behind the method of moments is that we can estimate it using a sample average over the data points. These are the **empirical moments**:

$$\hat{m} \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n \phi(x^{(i)}). \quad (547)$$

- Given these empirical moments, we can simply **plug in** \hat{m} for m^* :

$$\hat{\theta} \stackrel{\text{def}}{=} M^{-1}(\hat{m}). \quad (548)$$

- How well does this procedure work? It is relatively easy to study how fast \hat{m} converges to m^* , at least in an asymptotic sense. We can show that (i) \hat{m} is close to m^* , and then use that to show that (ii) $\hat{\theta}$ is close to θ^* . The analysis reflects the conceptual simplicity of the method of moments.
- First, since \hat{m} is just an average of i.i.d. variables, we can apply the central limit theorem:

$$\sqrt{n}(\hat{m} - m^*) \xrightarrow{d} \mathcal{N}(0, \text{Var}_{x \sim p^*}(\phi(x))). \quad (549)$$

- Second, assuming that M^{-1} is continuous around m^* , we can use the delta-method to argue that $\hat{\theta}$ converges θ^* :

$$\sqrt{n}(\hat{\theta} - \theta^*) \xrightarrow{d} \mathcal{N}(0, \nabla M^{-1}(m^*) \text{Var}_{x \sim p^*}(\phi(x)) \nabla M^{-1}(m^*)^\top), \quad (550)$$

where $\nabla M^{-1}(m^*) \in \mathbb{R}^{d \times p}$ is the Jacobian matrix for the inverse moment mapping M^{-1} . Therefore, the parameter error depends on how sensitive M^{-1} is around m^* . It is also useful to note that $\nabla M^{-1}(m^*) = \nabla M(\theta^*)^\dagger$ (where \dagger denotes pseudoinverse).

- These asymptotics are rough calculations that shed some light onto when we would expect the method of moments to work well: $\nabla M(\theta) \in \mathbb{R}^{p \times d}$ must have full column rank and moreover, the first d singular values should be far away from zero. Intuitively, if we perturb θ by a little bit, we want m to move a lot, which coincides with our intuitions about the asymptotic error of parameter estimation.

- **Example 39 (exponential families)**

- Recall the connection between moments and parameters in exponential families. We will strengthen this connection by showing that maximum likelihood corresponds to method of moments for exponential families.

- Consider the standard exponential family

$$p_{\theta}(x) = \exp(\theta \cdot \phi(x) - A(\theta)), \quad (551)$$

where

- * $\phi(x) \in \mathbb{R}^d$ is the feature vector (observation function),
- * $\theta \in \mathbb{R}^d$ are the parameters, and
- * $A(\theta) = \log \int_{\mathcal{X}} \exp(\phi(x)^{\top} \theta) dx$ is the log-partition function.

- Recall that the maximum likelihood estimator $\hat{\theta}$ is defined as follows:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i=1}^n \log p_{\theta}(x^{(i)}) \quad (552)$$

$$= \arg \max_{\theta} \{\theta \cdot \hat{\mu} - A(\theta)\}, \quad (553)$$

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n \phi(x^{(i)}). \quad (554)$$

Taking derivatives and setting it to zero, we get that

$$\hat{\theta} = M(\hat{\mu}), \quad M(\mu) = (\nabla A)^{-1}(\mu). \quad (555)$$

- Maximum likelihood estimation is a convex optimization problem for exponential families, but in general, maximum likelihood is not convex, nor does it coincide with the method of moments for any feature function ϕ .
- In general, computing the inverse moment mapping M^{-1} is not easier than the maximum likelihood estimate. Method of moments is only useful if we can find an appropriate observation function ϕ such that:
 - The moment mapping M is invertible, and hopefully has singular values that are bounded below (M provides enough information about the parameters θ).
 - The inverse moment mapping M^{-1} is computationally tractable.

7.3 Method of moments for latent-variable models (Lecture 17)

- Now that we have established the principles behind method of moments, let us tackle the Naive Bayes clustering model (Example 38). Our strategy is to just start writing some moments and see what kind of equations we get (they turn out to be product of matrices).
- Preliminaries
 - Assume each document has $L \geq 3$ words (this is an unimportant technicality).

- For notational convenience, assume that each word $x_j \in \mathbb{R}^b$ is represented as an indicator vector which is equal to one in the entry of that word and zero elsewhere.

- Let's start with the first-order moments:

$$M_1 \stackrel{\text{def}}{=} \mathbb{E}[x_1] = \sum_{h=1}^k \pi_h \beta_h = B\pi. \quad (556)$$

- Note that the moments require marginalizing out the latent variables, and marginalization corresponds to matrix products.
- Interpretation: $M_1 \in \mathbb{R}^b$ is a vector of marginal word probabilities.
- Clearly this is not enough information to identify the parameters.

- We can write the second-order moments:

$$M_2 \stackrel{\text{def}}{=} \mathbb{E}[x_1 x_2^\top] = \sum_{h=1}^k \pi_h \beta_h \beta_h^\top = B \text{diag}(\pi) B^\top. \quad (557)$$

- Interpretation: $M_2 \in \mathbb{R}^{d \times d}$ is a matrix of co-occurrence word probabilities. Specifically, $M_2(u, v)$ is the probability of seeing word u with word v , again marginalizing out the latent variables.

- It turns out that the model family is non-identifiable from second-order moments. So proceed to third-order moments to get more information.

$$M_3 \stackrel{\text{def}}{=} \mathbb{E}[x_1 \otimes x_2 \otimes x_3] = \sum_{h=1}^k \pi_h \beta_h \otimes \beta_h \otimes \beta_h. \quad (558)$$

- Interpretation: $M_3 \in \mathbb{R}^{b \times b \times b}$ is a rank-3 tensor of co-occurrence word probabilities.
- We will now show that three moments is sufficient for identifying the model parameters.

- Here is one approach for using M_1, M_2, M_3 to identify $\theta = (\pi, B)$.

- Let $\eta \in \mathbb{R}^b$ be a random vector.
- Define the projection of M_3 onto η as follows:

$$M_3(\eta) \stackrel{\text{def}}{=} \sum_{h=1}^k \pi_h \beta_h \otimes \beta_h (\beta_h^\top \eta) = B \text{diag}(\pi) \text{diag}(B^\top \eta) B^\top. \quad (559)$$

- Think of $M_3(\eta)$ as M_2 (they have the same row and column space), but someone has come in and tweaked the diagonal entries. Intuitively, this gives us two different views of the latent variables.

– Now we show that given matrices M_2 and $M_3(\eta)$, we can recover the parameters $\theta = (\pi, B)$. But first, a useful lemma which captures the core computation:

– Lemma

- * Suppose we observe matrices $X = BDB^\top$ and $Y = BEB^\top$ where
 - D, E are diagonal matrices such that the ratios $\{D_{ii}/E_{ii}\}_{i=1}^k$ are all non-zero and distinct.
 - $B \in \mathbb{R}^{b \times k}$ has full column rank (this is a reasonable condition which intuitively says that all the clusters are different in a linear algebraic sense). Note this automatically implies $b \geq k$.
- * Then we can recover B (up to permutation/scaling of its columns).

– Proof:

- * Simple case
 - Assume B is invertible ($b = k$).
 - Then X and Y are also invertible.
 - Compute

$$YX^{-1} = BEB^\top B^{-\top} D^{-1} B^{-1} = B \underbrace{ED^{-1}}_{\text{diagonal}} B^{-1}. \quad (560)$$

- The RHS has the form of an eigendecomposition, so the eigenvectors of YX^{-1} are exactly the columns of B up to permutation and scaling. We actually know the scaling since the columns of B are probability distributions and must sum to 1.
- Since the diagonal elements of ED^{-1} are distinct and non-zero by assumption, the eigenvectors are distinct.
- * Now suppose X and Y are not invertible. Note that X and Y have the same column space, so we can basically project them down into that column space where they will be invertible.
- * Let $U \in \mathbb{R}^{b \times k}$ be any orthonormal basis of the column space of B (taking the SVD of M_2 suffices: $M_2 = USU^\top$). Note that U has full column rank by assumption.
- * Important: although $B \in \mathbb{R}^{b \times k}$ is not invertible, $\tilde{B} \stackrel{\text{def}}{=} U^\top B \in \mathbb{R}^{k \times k}$ is invertible.
- * So let us project X and Y onto U by both left multiplying by U^\top and right multiplying by U .
- * Then we have the following decomposition:
 - $U^\top XU = \tilde{B}D\tilde{B}^\top$
 - $U^\top YU = \tilde{B}E\tilde{B}^\top$
- * Now we are back in the simple case, which allows us to recover \tilde{B} . We can obtain $B = U\tilde{B}$.

- We apply the lemma with $X = M_2$ and $Y = M_3(\eta)$. Since η is chosen at random, $D = \text{diag}(\pi)$ and $E = \text{diag}(\pi) \text{diag}(B^\top \eta)$ will have distinct ratios with probability 1.
- Once we have recovered B , then we can recover π easily by setting $\pi = B^\dagger M_1$.
- We have demonstrated the essence of the method of moments approach [Anandkumar/Hsu/Kakade, 2012].
- We have so far that we can solve the moment equations for the Naive Bayes clustering model and recover the parameters given the moments. These techniques can be further extended to obtain consistent parameter estimates for Gaussian mixture models, hidden Markov models, Latent Dirichlet allocation, restricted PCFGs, independent component analysis, linear Bayesian networks, mixture of linear regressions, etc. An active area of research is to extend these techniques to other model families.
- We have so far shown that if we had the true moments M_2, M_3 (in the limit of infinite data), our algorithm would output the true parameters B (up to some permutation of the columns). In practice, we do not have access to the true moments, but can estimate them from data. We will still focus on the Naive Bayes clustering model (Example 38). Here is the algorithm:²⁰
- **Algorithm 8 (method of moments for Naive Bayes clustering model)**
 - Input
 - * Observations $\{(x_1^{(i)}, x_2^{(i)}, x_3^{(i)})\}_{i=1}^n$
 - * Unit vector $\eta \in \mathbb{R}^b$ (remains fixed for the algorithm, but can be chosen uniformly at random)
 - Step 1: Compute empirical moments $\hat{M}_2, \hat{M}_3(\eta) \in \mathbb{R}^{b \times b}$:

$$\hat{M}_2 \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n x_1^{(i)} \otimes x_2^{(i)\top} \quad (561)$$

$$\hat{M}_3(\eta) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n x_1^{(i)} x_2^{(i)\top} (x_3^{(i)\top} \eta). \quad (562)$$

Note that without loss of generality, we can assume \hat{M}_2 is symmetric, since we can replace \hat{M}_2 with $\frac{1}{2}(\hat{M}_2 + \hat{M}_2^\top)$, which is also a consistent estimator of M_2 . Same goes for $\hat{M}_3(\eta)$.

- Step 2: Compute a whitening matrix $\hat{W} \in \mathbb{R}^{b \times k}$: Compute the SVD (remember, \hat{M}_2 is symmetric)

$$\hat{M}_2 = \hat{U} \hat{S} \hat{U}^\top \quad (563)$$

²⁰This is an improved version, where rather than using eigendecomposition, we use SVD, which is algorithmically simpler and more stable.

and take the k largest singular values/vectors:

$$\hat{W} = \hat{U}_{1:k} \hat{S}_{1:k}^{-\frac{1}{2}}. \quad (564)$$

- Step 3: Apply the whitening matrix to $\hat{M}_3(\eta)$ and compute the SVD: Let $\hat{v}_1, \dots, \hat{v}_k$ be the left singular vectors of:

$$\hat{W}^\top \hat{M}_3(\eta) \hat{W}. \quad (565)$$

- Step 4: Return parameters $\hat{B} \in \mathbb{R}^{b \times k}$:

$$\hat{B}_i \propto (\hat{W}^\top)^\dagger \hat{v}_i. \quad (566)$$

Note that we know each column \hat{B}_i is a probability distribution, so we only need it up to normalization constants.

- Our analysis of Algorithm 8 consists of two parts:
 - Given infinite number of examples, does the algorithm output the correct parameters? We saw a preliminary analysis of this last time.
 - Given finite number of examples, how does the error in the moment estimates affect the error in the resulting parameters? This require some matrix perturbation theory.
- **Theorem 36 (correctness of Algorithm 8)**
 - Assume that B has full column rank and $\pi \succ 0$.
 - In the limit of infinite data, Algorithm 8 outputs \hat{B} such that $\hat{B} \xrightarrow{P} B$ (up to permutation of the columns).
- Proof of Theorem 36
 - In the limit of infinite data, we assume all the quantities are exact (so we replace \hat{M}_2 with M_2).
 - To warmup, let's start with first-order moments:

$$\mathbb{E}[x_1 \mid h] = \beta_h, \quad (567)$$

$$\mathbb{E}[x_1] = \sum_{h=1}^k \pi_h \beta_h = B\pi. \quad (568)$$

In computing second-order moments, we rely on the fact that x_1 and x_2 are uncorrelated given h :

$$M_2 = \mathbb{E}[x_1 x_2^\top] = \sum_{h=1}^k \pi_h \beta_h \beta_h^\top \quad (569)$$

$$= B \operatorname{diag}(\pi) B^\top \quad (570)$$

$$= CC^\top. \quad (571)$$

For convenience, let's define $C = B \operatorname{diag}(\pi)^{\frac{1}{2}}$ so that we have $M_2 = CC^\top$. Now, the third-order moments include a random scalar variable $x_3^\top \eta$, which just contributes to the diagonal (η is some fixed random vector.)

$$M_3(\eta) = \mathbb{E}[x_1 x_2^\top (x_3^\top \eta)] = \sum_{h=1}^k \pi_h (\beta_h^\top \eta) \beta_h \beta_h^\top \quad (572)$$

$$= B \operatorname{diag}(\pi) \operatorname{diag}(B^\top \eta) B^\top \quad (573)$$

$$= C \operatorname{diag}(B^\top \eta) C^\top. \quad (574)$$

- Since M_2 is rank k , W actually whitens M_2 in that $W^\top M_2 W = I$. This means that $(W^\top C)(W^\top C)^\top = I$, so that $W^\top C$ is an orthogonal matrix. This matrix will play a significant role later, so we name it

$$F \stackrel{\text{def}}{=} W^\top C. \quad (575)$$

- Now, whitening $M_3(\eta)$ using the same W results in

$$W^\top M_3(\eta) W = V \operatorname{diag}(B^\top \eta) V^\top. \quad (576)$$

This reveals that the columns of V are the singular vectors of $W^\top M_3(\eta) W$. Recall, that we also had

$$W^\top M_3(\eta) W = F \operatorname{diag}(B^\top \eta) F^\top. \quad (577)$$

Since singular vectors are unique up to scaling z (and permutation, which is not important), we can conclude that $W^\top C = F = V \operatorname{diag}(z)$.

- The final step is to unwhiten:

$$B \operatorname{diag}(\pi)^{\frac{1}{2}} = C = (W^\top)^\dagger V \operatorname{diag}(z). \quad (578)$$

The normalization constant is actually unnecessary since we know that the columns of B must sum to 1, so we can write

$$B_i \propto (W^\top)^\dagger V_i. \quad (579)$$

- For the singular vectors to be unique at each point, we required that B have rank k and that $\text{diag}(B^\top \eta)$ have distinct elements, which will happen with probability 1 if η is chosen at random.
- Towards sample complexity
 - Now if we apply our algorithm on the estimated empirical moments \hat{M}_2, \hat{M}_3 , how close is the resulting parameters \hat{B} to the true parameters B ?
 - The way we should think about this issue is in terms of **perturbation analysis**. If the algorithm is a relatively smooth function of its input, then we expect that introducing noise to the inputs (replacing M_2, M_3 with \hat{M}_2, \hat{M}_3) will not alter the parameter estimates \hat{B} too much.
 - Now, we just need to study the error propagation through each step. The only place where we deal with randomness is in the first step. After that, it's mostly bounding the spectral properties of matrices based on operations that are performed on them. We will mostly be using the **operator norm** $\sigma_1(\cdot)$ for matrices, which is equal to the largest singular value.
- **Theorem 37 (sample complexity of Algorithm 8)**
 - For sufficiently large n , we have with probability at least $1 - \delta$ (over randomness in data) that the following holds: with probability at least $3/4$ (over the random choice of η), for each cluster $i = 1, \dots, k$:

$$\|\hat{B}_i - B_i\|_2 = O_p \left(\frac{k^3 \sqrt{\log(1/\delta)}}{\pi_{\min}^2 \sigma_k(B)^3 \sqrt{n}} \right). \quad (580)$$

- Legend
 - * n : number of examples
 - * k : number of clusters
 - * $p_{\min} = \min_i \pi_i$: smallest prior probability of a cluster (hard to estimate parameters if we don't get samples from that cluster very often)
 - * $\sigma_k(B)$: k -th largest singular value of the conditional probabilities (if two clusters are close, this value is small, meaning it's hard to tell the clusters apart)
- The probability of $3/4$ isn't a serious issue since we can repeat the procedure q times with independently random choices of η to drive up this probability exponentially to $1 - (1/4)^q$.
- Proof sketch of Theorem 37
 - The full proof is quite long and involved, so we refer the reader to the “Two SVDs Suffice” paper. Instead, we will just highlight the structure of the proof and the key ideas.

- Recall that the plan is to study the propagation of error from the empirical moments all the way to the parameter estimates. In the process, we are multiplying, adding, inverting matrices, taking SVDs. We need to understand the impact on error each of these operations has.
- Here are a few basic identities regarding largest and smallest singular values to get us started:

- * Addition: given matrices A and B ,

$$\sigma_1(A + B) \leq \sigma_1(A) + \sigma_1(B). \quad (581)$$

- * Multiplication: given matrices A and B ,

$$\sigma_1(AB) \leq \sigma_1(A)\sigma_1(B). \quad (582)$$

- * Inverse: if $A \in \mathbb{R}^{b \times k}$ have full column rank k

$$\sigma_1(A) = \sigma_k(A^\dagger)^{-1}, \quad (583)$$

where $A^\dagger = (A^\top A)^{-1}A^\top$ is the pseudo-inverse. Note that here the fact that we are dealing with matrices requires some special care, as we go from first singular value to the k -th.

- The next two are important theorems in matrix perturbation theory that allow us to control the singular values and vectors of perturbed matrices, respectively:

- * Weyl's theorem:

- Let $A, E \in \mathbb{R}^{b \times k}$ be matrices.
- Then for all $i = 1, \dots, k$:

$$|\sigma_i(A + E) - \sigma_i(A)| \leq \sigma_1(E). \quad (584)$$

- Interpretation: the amount by which the singular values of the perturbed matrix $A + E$ change is bounded by the largest singular value of the perturbation E . We will invoke this theorem for σ_1 and σ_k .

- * Wedin's theorem:

- Suppose A has the following SVD:

$$A = \begin{pmatrix} U_1 \\ U_2 \\ U_3 \end{pmatrix} \begin{pmatrix} \Sigma_1 & 0 \\ 0 & \Sigma_2 \\ 0 & 0 \end{pmatrix} \begin{pmatrix} V_1 \\ V_2 \end{pmatrix}^\top. \quad (585)$$

- Let $\hat{A} = A + E$ and let $\hat{U}_1, \hat{U}_2, \hat{U}_3, \hat{\Sigma}_1, \hat{\Sigma}_2, \hat{V}_1, \hat{V}_2$ be the corresponding SVD of \hat{A} .
- Let Φ be the matrix of **canonical angles** between U_1 and \hat{U}_1 :

$$\Phi_{i,j} = \arccos(U_{1;i}^\top \hat{U}_{1;j}). \quad (586)$$

- Suppose that for each i , $\sigma_i(\hat{\Sigma}_1) \geq \alpha + \delta$ and $\sigma_i(\Sigma_i) \leq \alpha$.
- Then

$$\sigma_1(\sin \Phi) \leq \frac{\sigma_1(E)}{\delta}, \quad (587)$$

where $\sin(\Phi)$ is applied elementwise.

- Interpretation: this theorem allows us to control the error of the singular vectors. If the first k singular values are significantly larger than the rest, then the angles between the first k are bounded.

– Step 1: Bounding errors in moments $M_2, M_3(\eta)$

- * Even though $M_2, M_3(\eta)$ are matrices, we can bound their vectorized forms.
- * Generically, let π be a probability vector and $\hat{\pi}$ be an empirical average over n multinomial draws from π . Then by the central limit theorem, we have

$$\sqrt{n}(\hat{\pi} - \pi) \xrightarrow{d} \mathcal{N}(0, \text{Cov}(\pi)). \quad (588)$$

For a multinomial, we have

$$\text{Cov}(\pi) = \text{diag}(\pi) - \pi\pi^\top. \quad (589)$$

Upper bounding the variance with the second moment and rewriting, we have:

$$\hat{\pi} - \pi = O_p\left(\frac{\pi}{\sqrt{n}}\right). \quad (590)$$

- * Now, applying this result to M_2 and $M_3(\eta)$ to get a handle on the errors on the moments:

$$\epsilon_{M_2} = \sigma_1(\hat{M}_2 - M_2) = O_p\left(\frac{\sigma_1(M_2)}{\sqrt{n}}\right) \quad (591)$$

$$\epsilon_{M_3(\eta)} = \sigma_1(\hat{M}_3(\eta) - M_3(\eta)) = O_p\left(\frac{\sigma_1(M_3(\eta))}{\sqrt{n}}\right). \quad (592)$$

The constants are tiny: for example, M_2 is a distribution, so its singular values are all bounded by 1.

- * For the remainder of the proof, assume n is large enough so that

$$\boxed{\epsilon_{M_2} \leq \frac{\sigma_k(M_2)}{2}}. \quad (593)$$

– Step 2: Bounding errors in the whitening matrix W

- * Recall $C = B \text{diag}(\pi)^{\frac{1}{2}}$ and that $M_2 = CC^\top$.
- * Recall $F = W^\top C$ and that $FF^\top = W^\top M_2 W = I$.

- * Let $\hat{F} = \hat{W}^\top C$.
- * Our goal is to bound

$$\sigma_1(F - \hat{F}). \quad (594)$$

- * Note that \hat{W} doesn't quite whiten M_2 , but in general, suppose we have $\hat{W}^\top M_2 \hat{W} = ADA^\top$ (some PSD matrix).
- * Given this, we have:
 - $\hat{W} = WAD^{\frac{1}{2}}A^\top$
 - $\hat{F} = AD^{\frac{1}{2}}A^\top F$
- * Then

$$\sigma_1(\hat{W})^2 = \frac{1}{\sigma_k(\hat{M}_2)} \quad (595)$$

$$\leq \frac{1}{\sigma_k(M_2) - \epsilon_{M_2}} \quad [\text{Weyl's theorem}] \quad (596)$$

$$\leq \frac{2}{\sigma_k(M_2)} \quad [\text{assumption}] \quad (597)$$

$$\leq \frac{2}{\sigma_k(C)^2}. \quad (598)$$

- * Then

$$\sigma_1(\hat{F} - F) \leq \sigma_1(F)\sigma_1(I - AD^{\frac{1}{2}}A^\top) \quad (599)$$

$$= \sigma_1(I - AD^{\frac{1}{2}}A^\top) \quad [\sigma_1(F) = 1] \quad (600)$$

$$\leq \sigma_1(I - ADA^\top) \quad (601)$$

$$\leq \sigma_1(\hat{W}^\top(\hat{M}_{2;1:k} - M_2)\hat{W}) \quad (602)$$

$$\leq \sigma_1(\hat{W})^2[\sigma_1(\hat{M}_{2;1:k} - \hat{M}_2) + \sigma_1(\hat{M}_2 - M_2)] \quad (603)$$

$$= \sigma_1(\hat{W})^2[\sigma_{k+1}(\hat{M}_2) + \sigma_1(\hat{M}_2 - M_2)] \quad (604)$$

$$\leq 2\sigma_1(\hat{W})^2\epsilon_{M_2} \quad [\text{since } M_2 \text{ is rank } k] \quad (605)$$

$$\leq \frac{4}{\sigma_k(C)^2}\epsilon_{M_2} \quad [\text{substitute previous result}]. \quad (606)$$

- * Similarly, we can show:

$$\sigma_1(\hat{W}^\dagger - W^\dagger) \leq \frac{6\sigma_1(C)}{\sigma_k(C)^2}\epsilon_{M_2}. \quad (607)$$

and

$$\sigma_1(\Pi - \Pi_W) \leq \frac{4}{\sigma_k(C)^2}\epsilon_{M_2}. \quad (608)$$

* For the higher-order moments, define the whitened versions:

$$\begin{aligned} \cdot T &= W^\top M_3(\eta) W \\ \cdot \hat{T} &= \hat{W}^\top \hat{M}_3(\eta) \hat{W} \\ \cdot \epsilon_T &= \sigma_1(\hat{T} - T) \end{aligned}$$

* Similarly, we can also bound the third-order moments:

$$\epsilon_T = O_p \left(\frac{\epsilon_{M_2}}{\sqrt{\pi_{\min}} \sigma_k(C)^2} + \frac{\epsilon_{M_3}}{\sigma_k(C)^3} \right) \quad (609)$$

– Step 3: Bounding errors in the singular values and vectors

- * Let σ_i and v_i be the singular values and vectors of T .
- * Let $\hat{\sigma}_i$ and \hat{v}_i be the singular values and vectors of \hat{T} .
- * First, a lemma that shows if the singular values are bounded from below and separated, then the error in the singular values is small:

$$\sigma_i \geq \Delta \quad (610)$$

$$\sigma_i - \sigma_{i+1} \geq \Delta. \quad (611)$$

Then

$$\|v_i - \hat{v}_i\|_2 \leq \frac{2\sqrt{k}}{\Delta - \epsilon_T} \epsilon_T. \quad (612)$$

* To show this, let $\theta = \arccos(v_i \cdot \hat{v}_i)$ be the angle between the singular vectors.

$$\|v_i - \hat{v}_i\|_2^2 = 2(1 - \cos(\theta)) \leq 2(1 - \cos^2(\theta)) = 2\sin^2(\theta). \quad (613)$$

By Weyl's theorem,

$$|\hat{\sigma}_i - \sigma_j| \geq |\sigma_i - \sigma_j| - \epsilon_T \geq \Delta - \epsilon_T. \quad (614)$$

By Wedin's theorem,

$$|\sin(\theta)| \leq \frac{\sqrt{2k}}{\Delta - \epsilon_T} \epsilon_T. \quad (615)$$

* Second, we show that if we choose η uniformly at random, then Δ will be sufficiently large. Specifically, with probability at least $1 - \delta$, we have $\Delta \geq \Omega(\delta/k^{2.5})$. For each pair of clusters, we can ensure the separation to be $\frac{1}{\sqrt{k}}$ (sum of k independent variables), and we pay k^2 more for a union bound over all pairs of clusters.

* The final result is that with probability at least $1 - \delta$,

$$\|v_i - \hat{v}_i\|_2 \leq O \left(\frac{2k^3}{\delta} \epsilon_T \right). \quad (616)$$

– Step 4: Bounding errors in the parameters

* Application of the same ideas (see paper for more details).

7.4 Summary (Lecture 17)

- Computationally, the method of moments estimator is very efficient in that it requires making one pass over the data to estimate the moments, whereas EM requires making as many passes as necessary to get convergence.
- Statistically, the method of moments estimator is not efficient: that is, the asymptotic variance of this estimator (550) will be higher than that of the maximum likelihood estimator (379). Intuitively this is because we are only using the first three moments and higher-order moments provide more information. However, the statistical deficiency is made up for by the computational gains. These methods are perfect for the large data regime, where we really need to be able to process large amounts of data quickly. For unsupervised learning using latent-variable models, this is exactly the regime we're in. In practice, lower statistical efficiency can be problematic. It is beneficial to use the method of moments as initialization and run local optimization (EM), which will improve the (empirical) log-likelihood.
- The method of moments provides a general framework for exploring new procedures. By selecting different observation functions ϕ , we can reveal more or less information about the data, making it harder or easier to invert the moment mapping to obtain parameter estimates.
- The key idea of method of moments for these latent-variable models is to have three views of h which are conditionally independent. Using two views allows us to identify the column space, but not the parameters, while three views also pins down the exact basis.
- In the context of latent-variable models where parameter estimation has traditionally been intractable due to non-convexity, we see that method of moments can provide a fresh perspective. By using just the first-, second-, and third-order moments, we can obtain consistent parameter estimates efficiently by simply performing a spectral decomposition.

7.5 References

- Anandkumar/Hsu/Kakade, 2012: A Method of Moments for Mixture Models and Hidden Markov Models
- Anandkumar/Foster/Hsu/Kakade/Liu, 2012: Two SVDs Suffice: Spectral decompositions for probabilistic topic modeling and latent Dirichlet allocation
- Anandkumar/Ge/Hsu/Kakade/Telgarsky, 2012: Spectral learning of latent-variable models

8 Conclusions and outlook

8.1 Review (Lecture 18)

- The goal of this course is to provide a theoretical understanding of why machine learning algorithms work. To undertake this endeavor, we have developed many powerful tools (e.g., convex optimization, uniform convergence) and applied them to the classic machine learning setting: learning a predictor given a training set and applying to unseen test examples.
- To obtain more clarity, it is useful to decompose the error into approximation, estimation, and optimization error:




Figure 8: Cartoon showing error decomposition into approximation, estimation, and optimization errors.

- Approximation error has to do with how much potential the hypothesis class has (e.g., Gaussian kernel versus linear kernel, large norm versus smaller norm)
- Estimation error has to do with how well you're able to learn, which depends on the complexity of the hypothesis and the amount of data you have.
- Optimization error is how well your algorithm is able to approximate the perfect optimizer (empirical risk minimizer).

Sometimes approximation and optimization error are hard to distinguish. For example, in kernel methods, random Fourier features can be viewed as defining a smaller hypothesis space or doing an approximate job optimizing the original kernel-based objective. An algorithm implicitly defines a class of hypotheses which are accessible by the algorithm.

- In the batch setting, we wish to study the generalization ability of learning algorithms. The key quantity was excess risk

$$L(\hat{h}) - L(h^*). \quad (617)$$

- We could use uniform convergence, which studies

$$\sup_{h \in \mathcal{H}} |L(h) - \hat{L}(h)|. \quad (618)$$

* FIGURE: [empirical and expected risk, convergence]

* Key ideas:

- Concentration: moment generating function of sub-Gaussian random variables
- Measures of complexity of hypothesis class: Rademacher complexity, covering numbers, VC dimension

- We could also use asymptotic statistics to get

$$L(\hat{\theta}) = L(\theta^*) + (\hat{\theta} - \theta^*)^\top \nabla L^2(\theta^*)(\hat{\theta} - \theta^*) + \dots \quad (619)$$

$$\hat{\theta} - \theta^* = -\nabla^2 \hat{L}(\theta^*) \nabla \hat{L}(\theta^*) + \dots, \quad (620)$$

which vastly simplifies the calculations needed.

- In online learning, nature adversarially chooses convex loss functions f_1, \dots, f_T . The online mirror descent learner produces

$$w_t = \arg \min_{w \in S} \left\{ \psi(w) + \sum_{i=1}^{t-1} w \cdot z_i \right\}, \quad (621)$$

We analyze the regret:

$$\text{Regret}(u) = \sum_{t=1}^T [f_t(w_t) - f_t(u)]. \quad (622)$$

The decomposition into approximation, estimation, and optimization errors is not perfect here.

- Key ideas:

- * Convexity allows us to linearize losses to produce an upper bound on regret
- * Strongly convex regularizer: tradeoff stability with fitting the data
- * ψ allows us to adapt to the geometry of the problem and arrive at algorithm such as EG

- Results

- Our results for estimation error depend on both the hypothesis class \mathcal{H} and the number of examples. The exact dependence depends on the structure of the problem, which we summarize here (there are analogues both in the online and the batch settings):

- * Classic parametric rate: $1/\sqrt{n}$
 - * If we have realizability or strong convexity: $1/n$
 - * Default dependence on dimension: d
 - * If we assume sparsity (k nonzeros): $k \log d$
 - * In infinite dimensions (RKHS), we depend on the norm of the weight vector: $\|w\|_2$
- But machine learning isn't all about prediction where training and test examples are drawn i.i.d. from the same distribution. In the following, we will point out a few other directions.

8.2 Changes at test time (Lecture 18)

- Suppose we train a predictor, and deploy it in real-life. Theory says that as long as the test distribution doesn't deviate from training, we have guarantees on the predictor's behavior. But we do not live in a stationary world.
- Example: Google Flu Trends
 - Trained on 2003–2008 data, predict flu based on (50M) common queries, got 97% accuracy.
 - In 2009, vastly underestimated (due to swine flu, people searched differently).
 - In 2011–2013, overpredicted.
- Online learning doesn't solve this problem
 - One might think that online learning doesn't place distribution over the data, so it's robust.
 - But the analyses are with respect to an expert which is static, and if the world is changing under you, being static is a pretty lame thing to do. So the bounds do hold, but the statements are quite weak.
- Covariate shift
 - Training distribution: $x \sim p^*$
 - Test distribution: $x \sim q^*$
 - * Example: object recognition in different lighting conditions / camera
 - Assume $p^*(y | x)$ is fixed across both training and test
 - Assume lots of unlabeled examples drawn from both p^* and q^* which can be used to estimate the marginal distributions.
 - Instance weighting (simplest approach)
 - * Idea: upweight examples that are underrepresented at training time.

* Estimator:

$$\hat{\theta} = \arg \min_{\theta} \hat{L}(\theta), \quad \hat{L}(\theta) = \frac{1}{n} \sum_{i=1}^n \hat{w}(x) \ell((x, y); \theta). \quad (623)$$

- * If we had infinite unlabeled data, we could use the weights $\hat{w}(x) = p^*(x)/q^*(x)$, from which it's easy to check that $\hat{\theta}$ is an unbiased estimator of the generalization error at test time.
- * There are two problems with this:
 - First, in practice, we don't have infinite unlabeled data, so we would need to estimate $\hat{w}(x)$.
 - Second, we need to assume that q^* is absolutely continuous with respect to p^* (must see all test examples at training time). Otherwise, we have no hope of doing well.
- * There are several procedures that address these problems as well as analyses. For example, one can use asymptotic statistics:
 - [Shimodaira: Improving predictive inference under covariate shift by weighting the log-likelihood function](#)

- Domain adaptation / multi-task learning

- In domain adaptation, even $p^*(y | x)$ might be different between train and test.
 - * Example: my email and your email
- Techniques
 - * The general idea is to solve joint ERM problem where assume that weight vectors are close. Let $W = [w_1, \dots, w_T] \in \mathbb{R}^{d \times T}$ be the matrix of weights for T tasks.
 - * We can think about performing joint learning where we regularize W using one of the following:
 - Assume weight vectors are similar in Euclidean distance: $\sum_{i \neq j} \|w_i - w_j\|^2$
 - Assume weight vectors lie in the same low-dimensional subspace: use trace norm $\|W\|_*$
 - Assume there exists a sparse set of features that is shared by all tasks: use block norm $\|W\|_{2,1}$, which is the sum of the L2 norms of the rows.
 - * Neural networks provide a natural and compelling solution: just have all the tasks share the same hidden layer.
- As far as theoretical analyses, the key intuition is that the regularizer reduces the effective hypothesis space from T independent weight vectors to T highly-related weight vectors.
 - * [Maurer, 2006: Bounds for Linear Multi-Task Learning](#)

- Example: deep neural networks are non-robust
 - Szegedy/Zaremba/Sutskever/Bruna/Erhan/Goodfellow/Fergus, 2014: Intriguing properties of neural networks
 - This paper shows that one can take a high-accuracy neural network for object recognition, perturb the input by a very small adversarial amount to make the predictor incorrect. The perturbation is imperceptible to the naked eye and is obtained via an optimization problem like this:

$$\min_{r \in \mathbb{R}^d} (f(x + r) - y_{\text{wrong}})^2 + \lambda \|r\|^2, \quad (624)$$
 where r is the perturbation, f is the trained neural network, and x is the input.
 - Note that if we had a classifier based on a Gaussian kernel, we couldn't do this, because the Gaussian kernel is smooth.
- Robustness at test time
 - At test time, suppose that up to K features can be zeroed out adversarially.
 - We can optimize for this using robust optimization. The following paper shows that for classification, the robust optimization version results in a quadratic program.
 - Globerson/Roweis, 2006: Nightmare at Test Time: Robust Learning by Feature Deletion

8.3 Alternative forms of supervision (Lecture 18)

- Active learning (learner chooses examples non i.i.d.)
 - Example: learning binary classifiers in 1D
 - * Suppose data is generated according to $y = \mathbb{I}[x \geq \theta]$ for some unknown θ .
 - * If we sample i.i.d. from some distribution over x , our expected risk falls as $O(1/n)$.
 - * However, if we actively choose points (using binary search), our expected risk falls as $O(e^{-cn})$, which is exponentially faster.
 - Intuition: query examples that we are most uncertain about. In reality, we have noise, and sampling random i.i.d. is not terrible in the beginning.
 - Dasgupta, 2006: Coarse sample complexity bounds for active learning
 - Bach, 2007: Active learning for misspecified generalized linear models
- Semi-supervised learning
 - Motivation: labeled data is expensive, unlabeled data is cheap

- But how does knowing $p^*(x)$ (from unlabeled data) help you with $p^*(y | x)$ (prediction problem). There is no free lunch, so we need to make some assumptions about the relationship between $p^*(x)$ and $p^*(y | x)$. For example, in classification, we can assume that the decision boundary defined by $[p^*(y | x) \geq \frac{1}{2}]$ doesn't cut through high-density regions of p^* .
- Theory
 - * $p^*(x)$ + assumptions about relationship defines a compatibility function $\chi(h)$
 - * Reduced hypothesis class \mathcal{H} to ones which are compatible $\{h \in \mathcal{H} : \chi(h) = 1\}$
 - * Unlabeled data helps because we're reducing the hypothesis class
- Balcan/Blum, 2009: A Discriminative Model for Semi-Supervised Learning

8.4 Interaction between computation and statistics (Lecture 18)

- This class is mostly about the statistical properties of learning algorithms, for which there is by now quite solid theory for. One of the most interesting avenues for future research is how computation plays a role in this story. The situation is rather complex.
- Favorable relationship: large optimization error results in smaller effective hypothesis space and actually controls estimation error.
 - Early stopping: relate to L2 regularization
 - Low-rank kernel approximations
- Unfavorable relationship: good estimators are hard to compute
 - Graphical models (partition function): pseudolikelihood
 - * Suppose $x \in \{0, 1\}^d$.
 - * $p_\theta(x) = W_\theta(x)/Z(\theta)$
 - * The partition function $Z(\theta) = \sum_x W_\theta(x)$ is computationally expensive, taking in the worst case $O(2^d)$ time.
 - * However, we can maximize the pseudolikelihood $\prod_{i=1}^d p_\theta(x_i | x_{-i})$, which takes $O(d)$ time.
 - * Pseudolikelihood is computationally more efficient, but statistically less efficient.
 - Learning latent variable models (e.g., Gaussian mixture models):
 - * Maximum likelihood is non-convex and hard to compute.
 - * The method of moments estimator (which works under some assumptions) is easy to compute but statistically less efficient.
 - Other constraints on learning: communication, memory, privacy
- Neural networks

- Here, the story is a bit more complex.
- On one hand, neural networks are difficult to train. People often talk about the vanishing gradient problem in the context of gradient-based optimization, where we are nowhere near a good solution, and yet the gradient is near zero. This happens when the weights w_j are too large, which saturates the sigmoid. To avoid this, careful initialization and step size selection are important.
- On the other hand, if we fully optimized existing neural network models, there is some chance they would just overfit, and the fact that we're not fully optimizing is actually central to their ability to generalize.

A Appendix

A.1 Notation

In general, we will not be religious about using uppercase letters to denote random variables or bold letters to denote vectors or matrices. The type should hopefully be clear from context.

- General definitions

- $[n] = \{1, \dots, n\}$
- For a sequence v_1, \dots, v_n :
 - * Let $v_{i:j} = (v_i, v_{i+1}, \dots, v_{j-1}, v_j)$ be the subsequence from i to j inclusive.
 - * Let $v_{<i} = v_{1:i-1}$.
- ∇f : gradient of a differentiable function f
- $\partial f(v)$: set of subgradients of a convex function f
- Indicator (one-zero) function:

$$\mathbb{I}[\text{condition}] \stackrel{\text{def}}{=} \begin{cases} 1 & \text{if condition is true} \\ 0 & \text{otherwise.} \end{cases} \quad (625)$$

- Probability simplex:

$$\Delta_d \stackrel{\text{def}}{=} \left\{ w \in \mathbb{R}^d : w \succeq 0 \text{ and } \sum_{i=1}^d w_i = 1 \right\}. \quad (626)$$

- Euclidean projection:

$$\Pi_S(w) \stackrel{\text{def}}{=} \arg \min_{u \in S} \|u - w\|_2 \quad (627)$$

is the closest point (measured using Euclidean distance) to w that's in S .

- Linear algebra

- Inner (dot) product: given two vectors $u, v \in \mathbb{R}^n$, the dot product is $u^\top v = u \cdot v = \langle u, v \rangle = \sum_{i=1}^n u_i v_i$.
- Outer product: for a vector $v \in \mathbb{R}^n$, let $v^\otimes = vv^\top$
- Positive semidefinite (PSD) matrix: a square matrix $A \in \mathbb{R}^{n \times n}$ is PSD iff A is symmetric and $v^\top A v \geq 0$ for all vectors $v \in \mathbb{R}^n$.
- Eigenvalues: for a PSD matrix $A \in \mathbb{R}^{n \times n}$, let $\lambda_1(A), \dots, \lambda_n(A)$ be the eigenvalues of A , sorted in decreasing order.

- $\text{tr}(A)$: for a square matrix $A \in \mathbb{R}^{n \times n}$, $\text{tr}(A)$ denotes the trace of A , the sum of the diagonal entries ($\text{tr}(A) = \sum_{i=1}^n A_{ii}$).
 - * $\text{tr}(ABC) = \text{tr}(BCA)$, but $\text{tr}(ABC) \neq \text{tr}(BAC)$ in general
 - * $\text{tr}(A) = \sum_{i=1}^n \lambda_i(A)$
- $\text{diag}(v)$: for a vector $v \in \mathbb{R}^d$, a matrix whose diagonal entries are the components of v and off-diagonal entries are zero.
- Probability and statistics
 - Probability measure: \mathbb{P}
 - Expectation:
 - * $\mathbb{E}[m(x)] = \int_{\mathcal{X}} m(x)p(x)dx$, where $p(x)$ is a density function
 - * $\mathbb{E}[m(x)] = \sum_{x \in \mathcal{X}} m(x)p(x)$, where $p(x)$ is a probability mass function
- Functional analysis
 - **Definition 29 (Cauchy sequence)**

A Cauchy sequence in a metric space (X, ρ) is $(x_i)_{i \geq 1}$ such that for any $\epsilon > 0$, there exists an integer n such that all $i, j > n$, $\rho(x_i, x_j) < \epsilon$.
 - **Definition 30 (complete metric space)**

A complete metric space (X, ρ) is one where every Cauchy sequence $(x_i)_{i \geq 1}$ converges to a limit point $x^* \in X$. Intuition: if the elements of the sequence are getting arbitrarily close to each other, they are getting close to some particular point in the space.
 - $L^p(\mathcal{X})$ is the space of all measurable functions f with finite p -norm:

$$\|f\|_p \stackrel{\text{def}}{=} \left(\int |f(x)|^p dx \right)^{1/p} < \infty. \quad (628)$$
 - Example: if $\mathcal{X} = \mathbb{R}$, $f(x) = 1$ is not in L^p for any $p < \infty$.
 - $L^2(\mathcal{X})$ is the set of all square integrable functions.
 - $L^\infty(\mathcal{X})$ is the set of all bounded functions.
- We will try to stick with the following conventions:
 - x : input
 - y : output
 - d : dimensionality
 - n : number of examples
 - t : iteration number

- T : total number of iterations
- f : (convex) function
- w : weight vector
- θ : parameters
- L : Lipschitz constant
- λ : amount of regularization
- η : step size
- $p^*(x, y)$: true distribution of data
- In general:
 - * v^* denotes the optimal value of some variable v .
 - * \hat{v} denotes an empirical estimate of some variable v based on training data.

A.2 Basic theorems

- Jensen's inequality

- For any convex function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ and random variable $X \in \mathbb{R}^d$,

$$f(\mathbb{E}[X]) \leq \mathbb{E}[f(X)]. \quad (629)$$

- Hölder's inequality

- For any real vectors $u, v \in \mathbb{R}^d$,

$$|u \cdot v| \leq \|u\|_p \|v\|_q, \quad (630)$$

for $1/p + 1/q = 1$, where $\|u\|_p = (\sum_{j=1}^d |u_j|^p)^{1/p}$ is the p -norm.

- For $p = q = 2$, this is the Cauchy-Schwartz inequality.
- Another important case is $p = 1$ and $q = \infty$.
- Note that this result holds in greater generality (for L_p spaces).